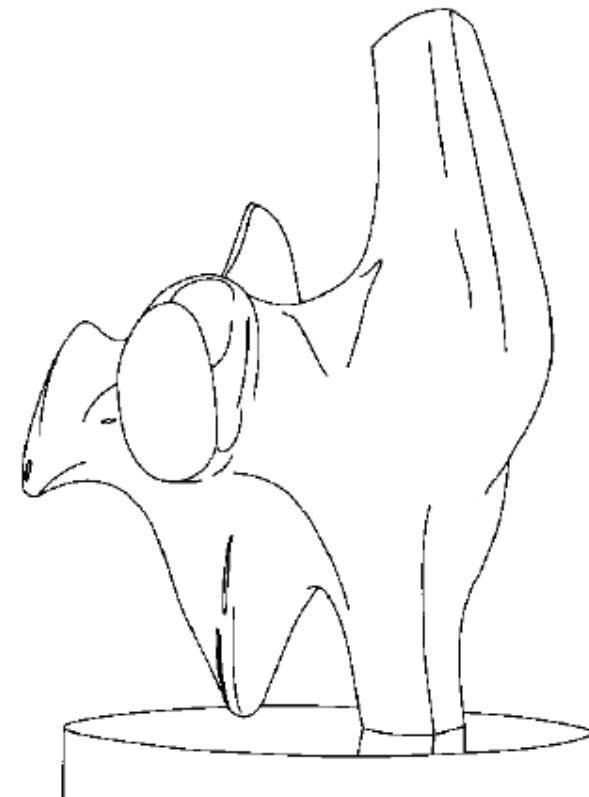
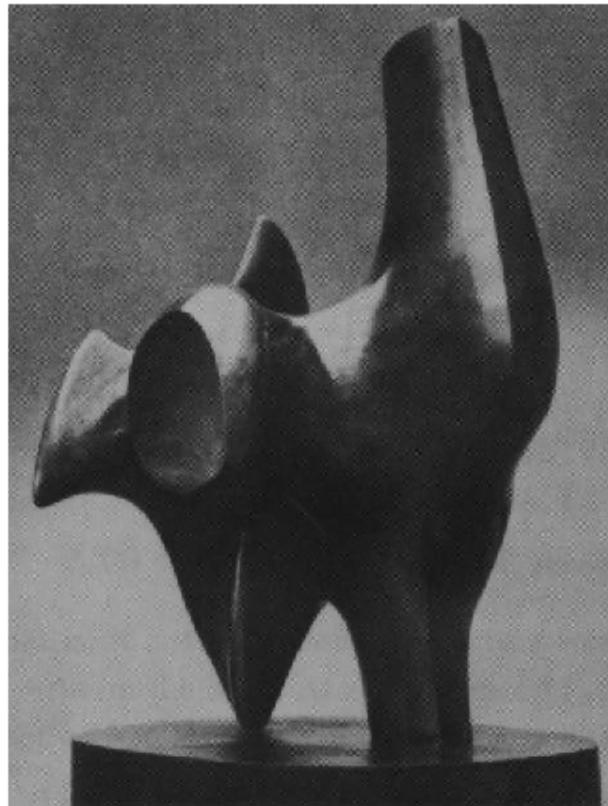


Edge Detection

- Slide content from Derek Hoiem, Cornelia Fermüller, Steve Seitz

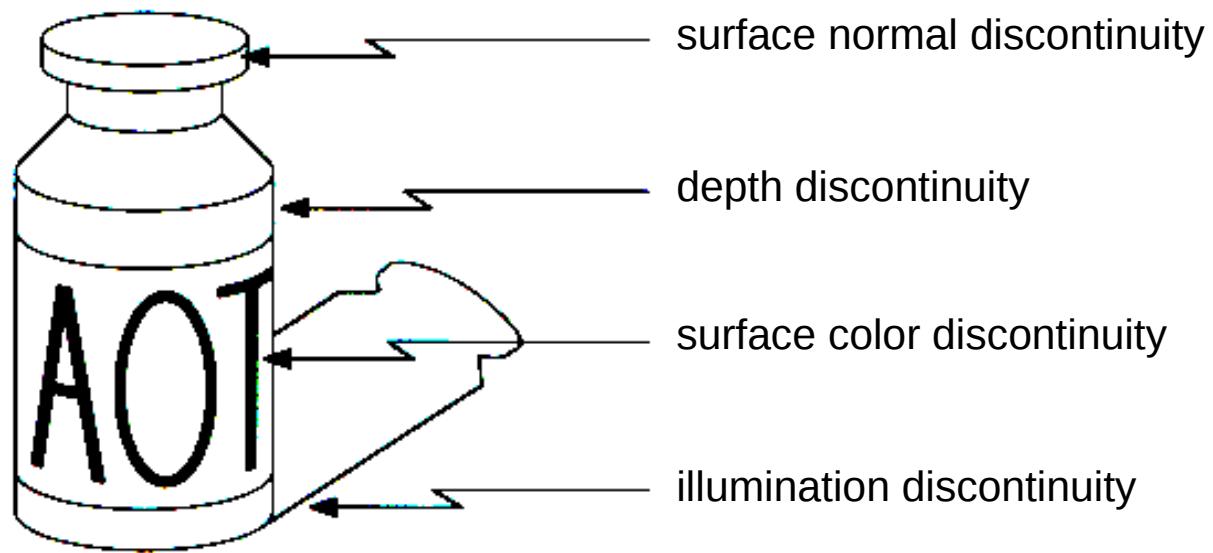
Edge detection



Convert a 2D image into a set of curves

- Extracts salient features of the scene
- More compact than pixels

Origin of Edges

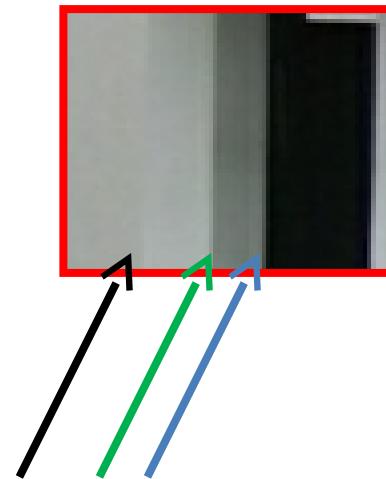


Edges are caused by a variety of factors

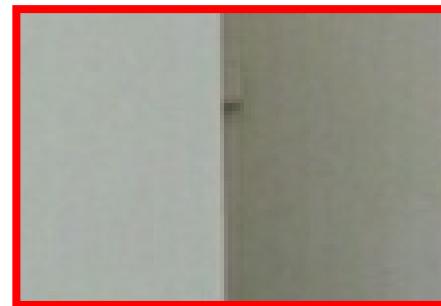
Closeup of edges



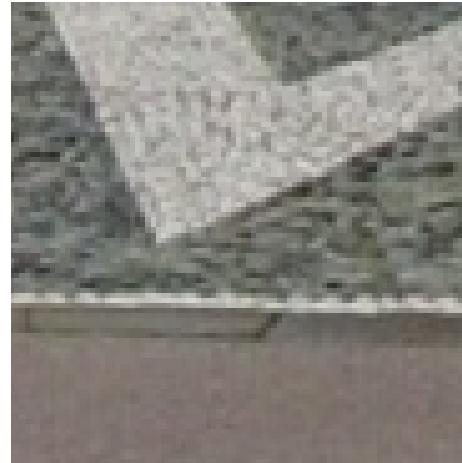
Closeup of edges



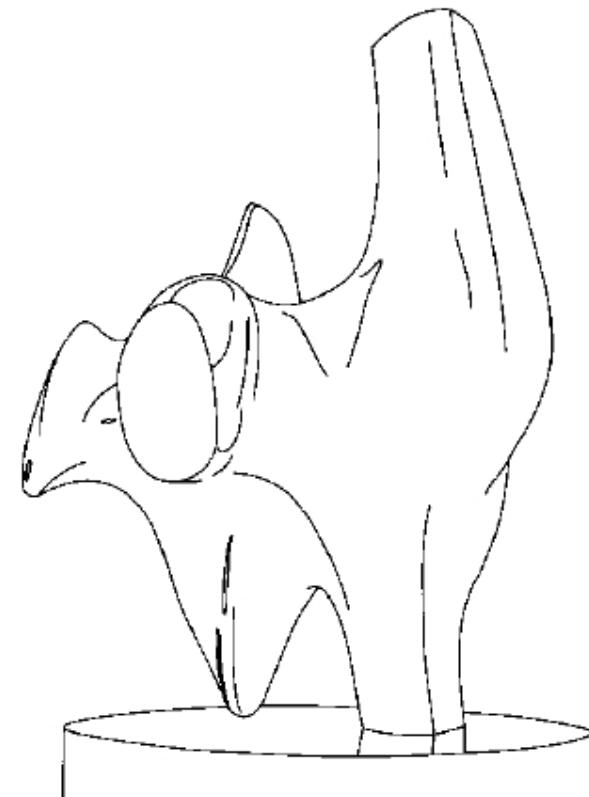
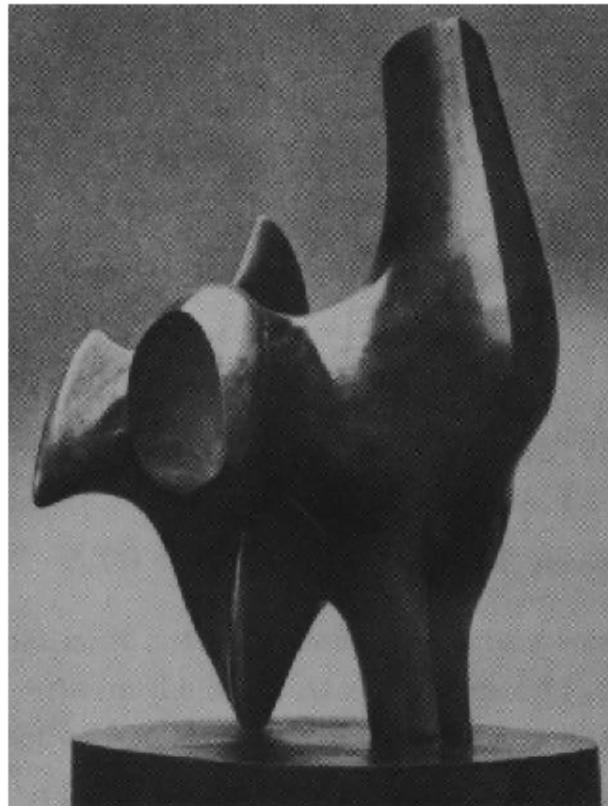
Closeup of edges



Closeup of edges

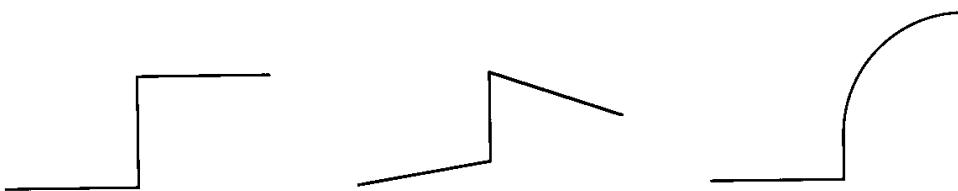


Edge detection

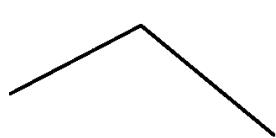


How can you tell that a pixel is on an edge?

Profiles of image intensity edges



Step Edges



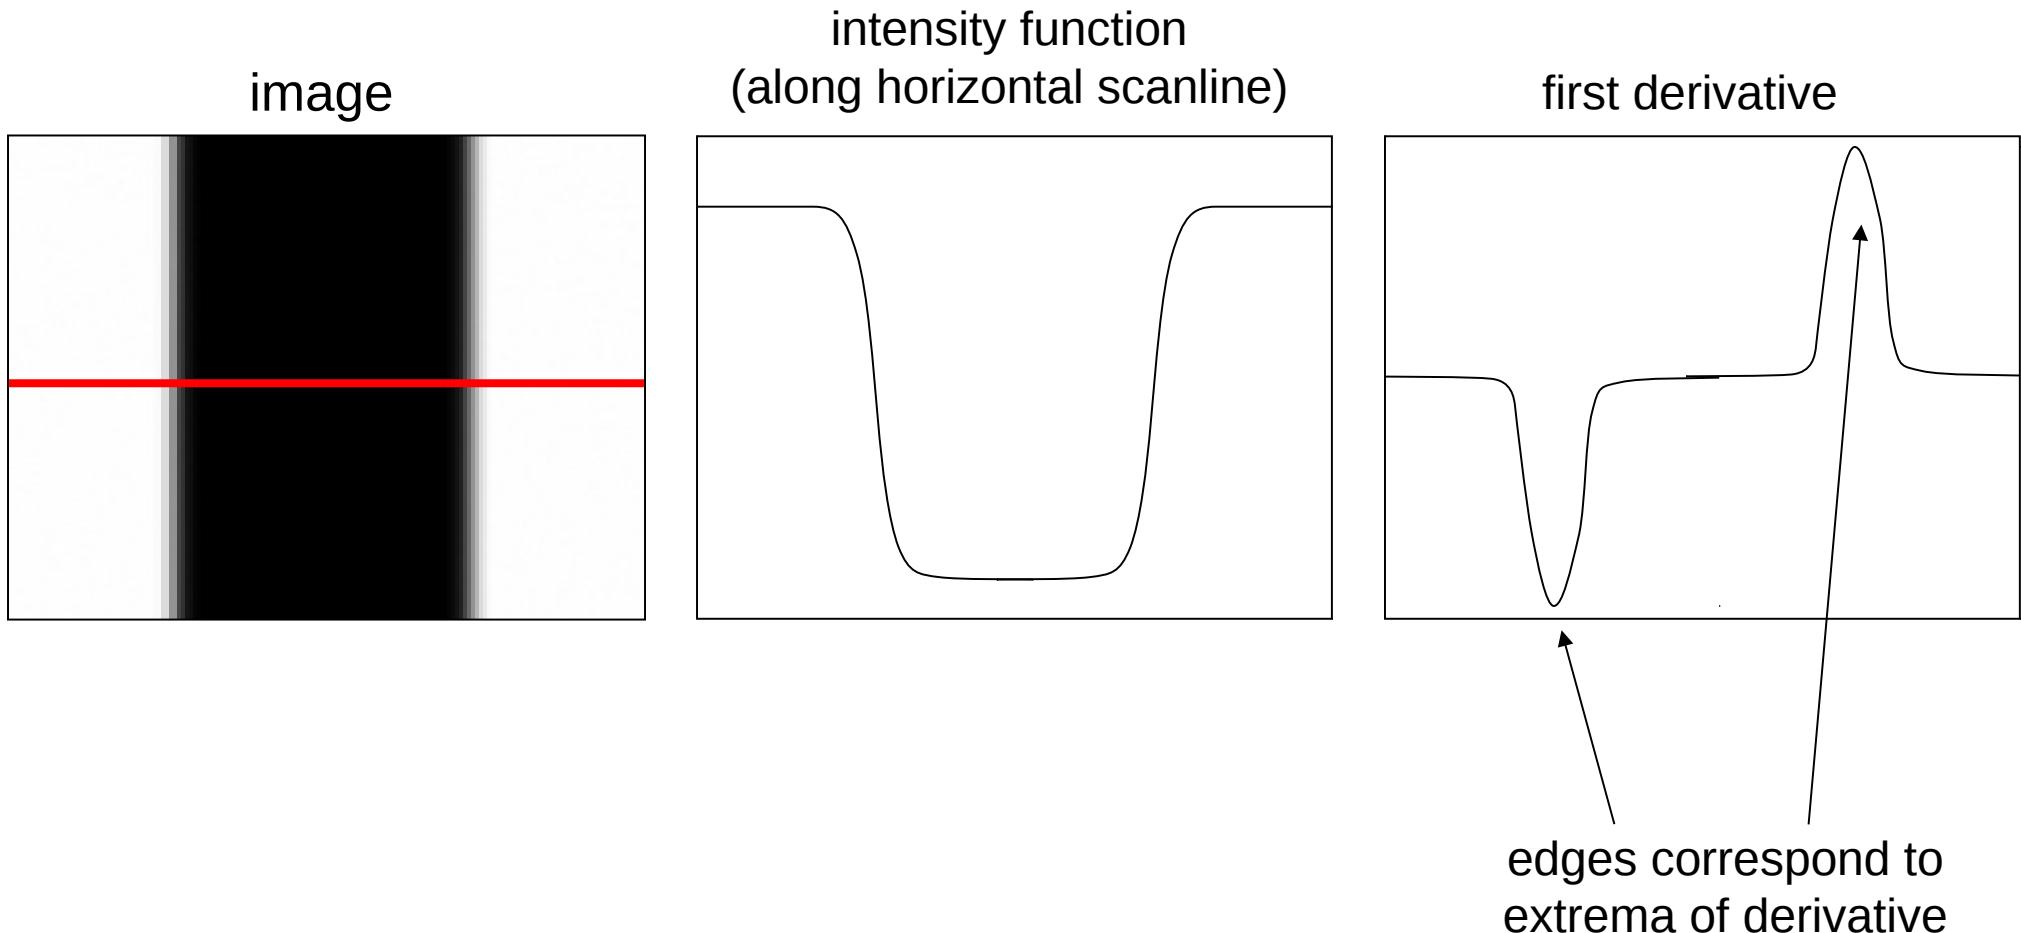
Roof Edge



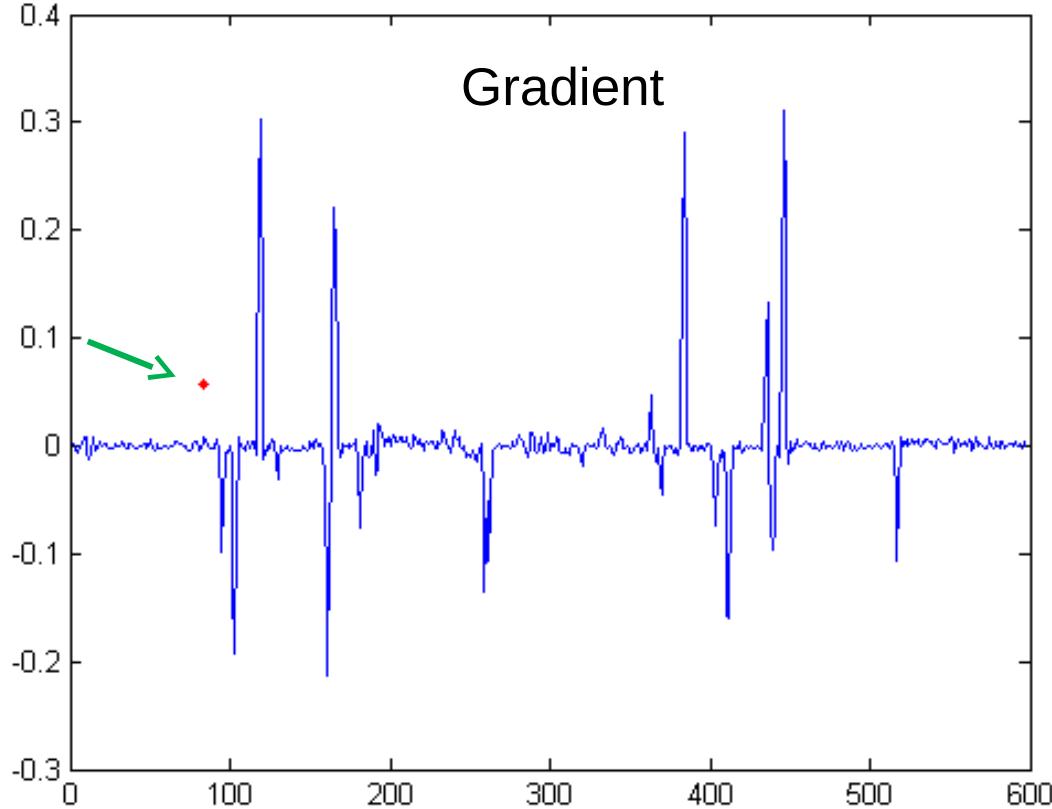
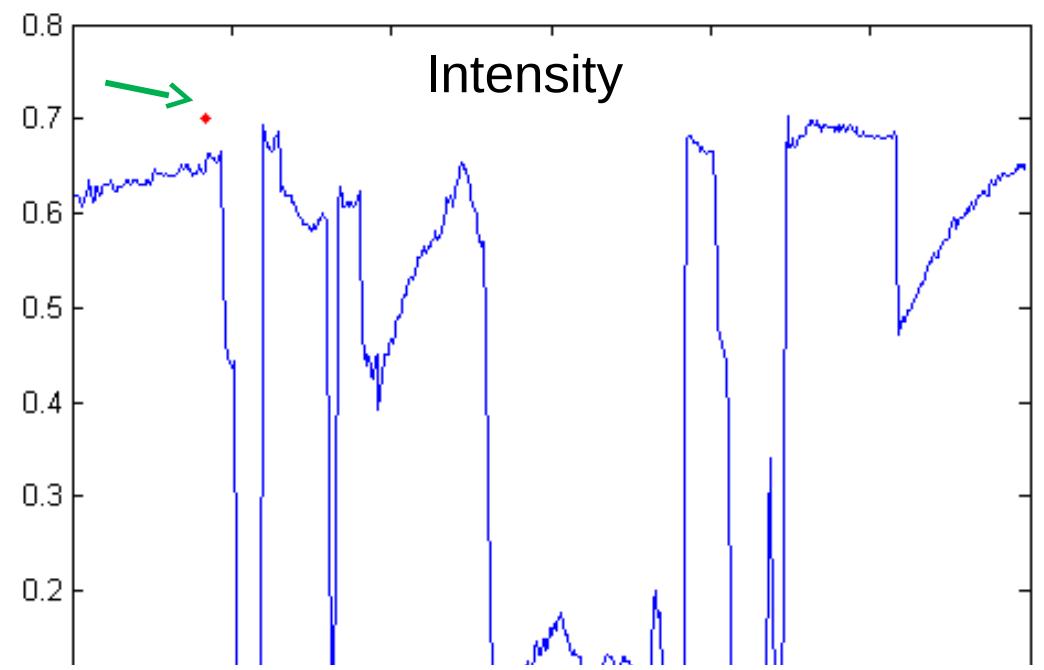
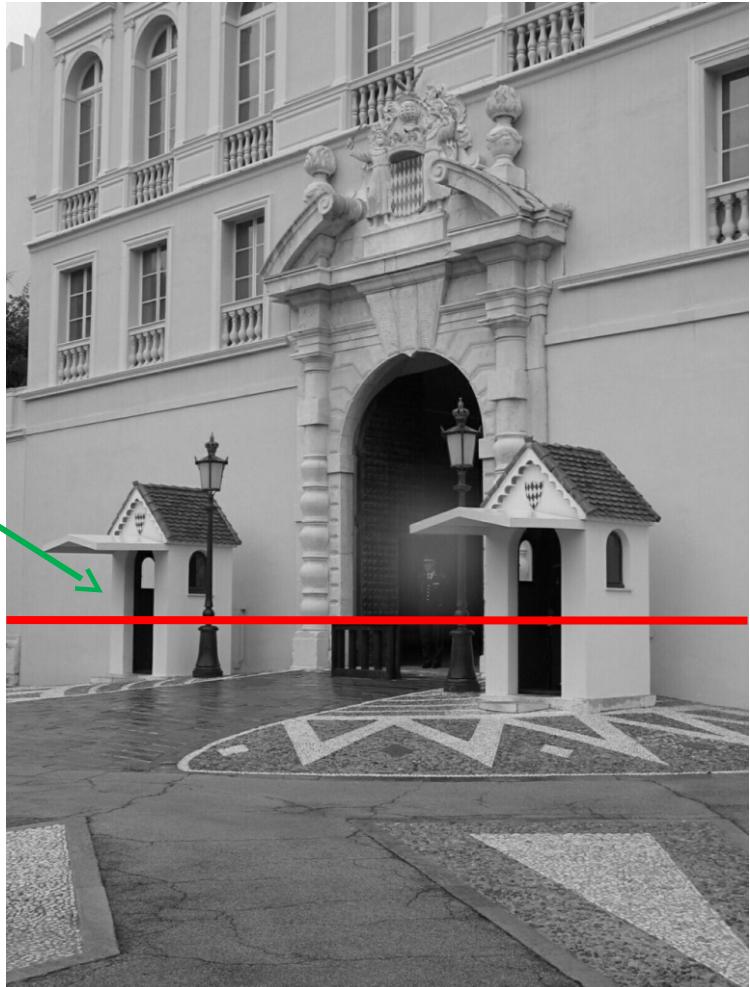
Line Edges

Characterizing edges

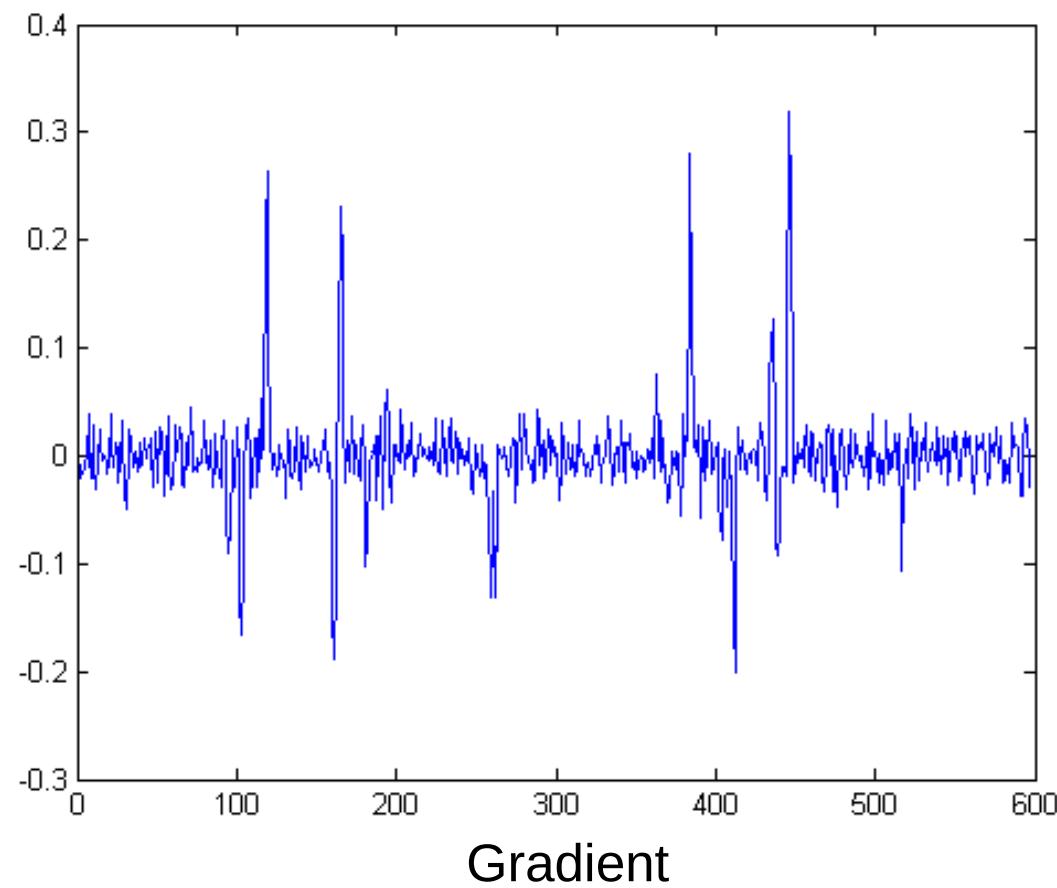
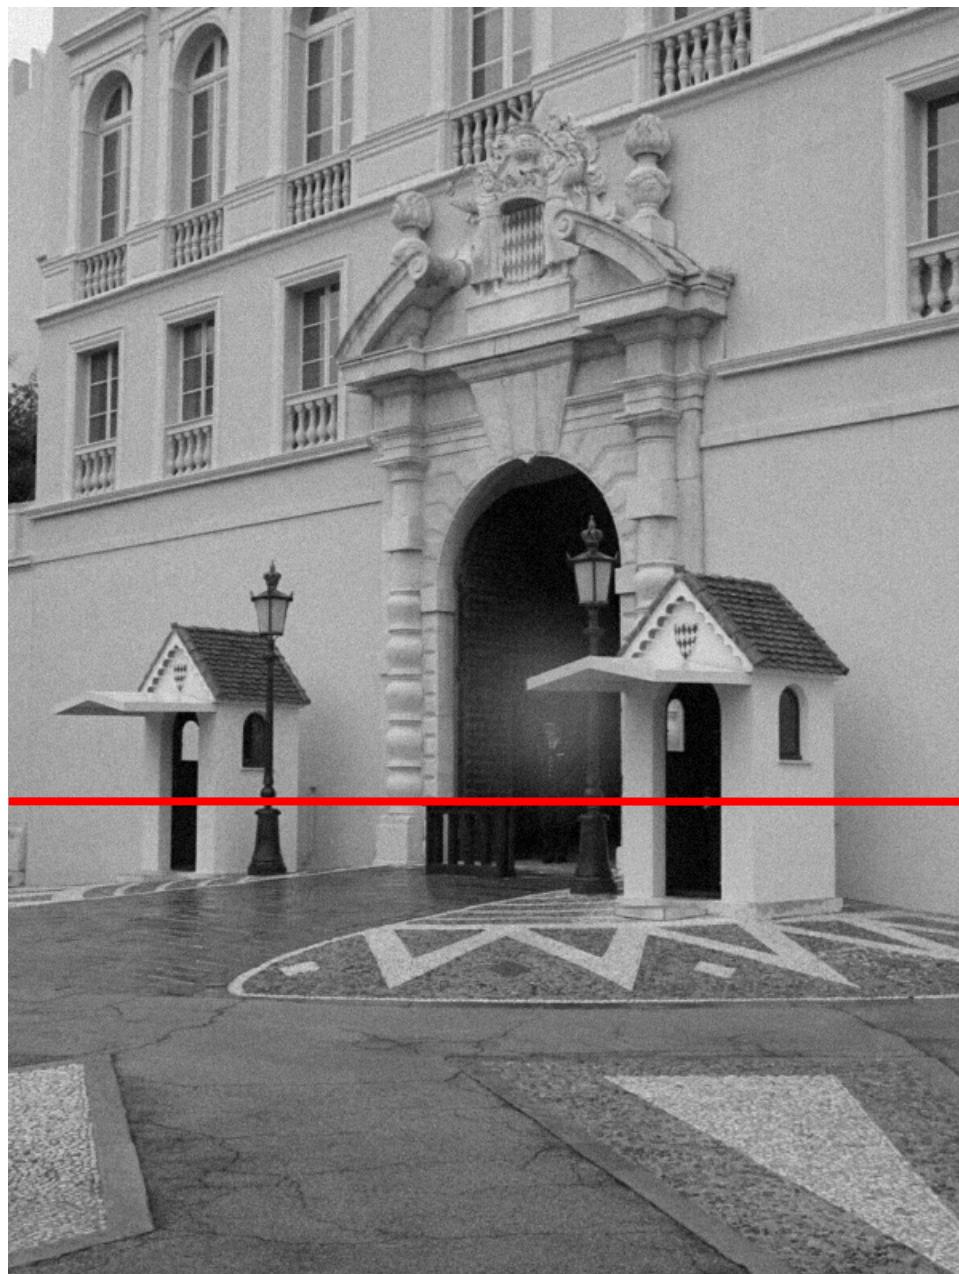
- An edge is a place of rapid change in the image intensity function



Intensity profile

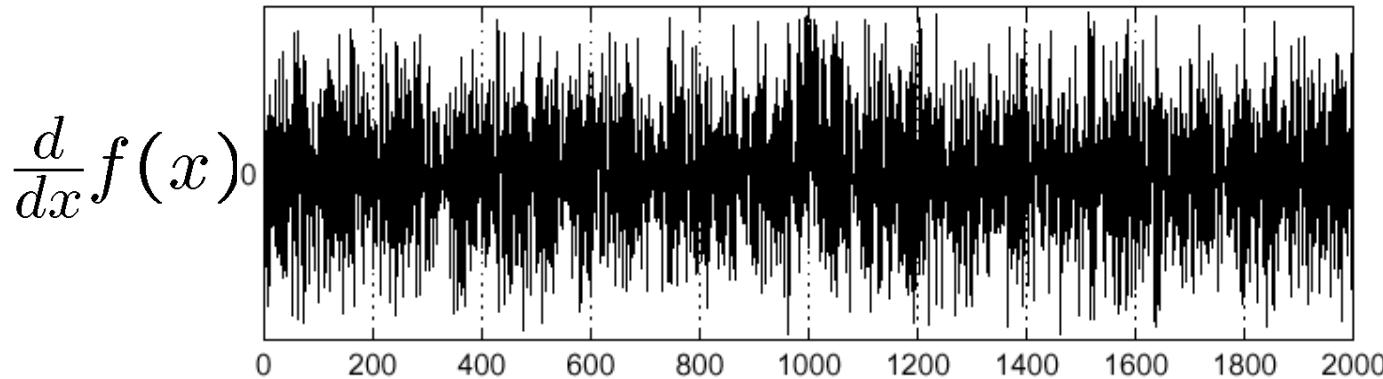
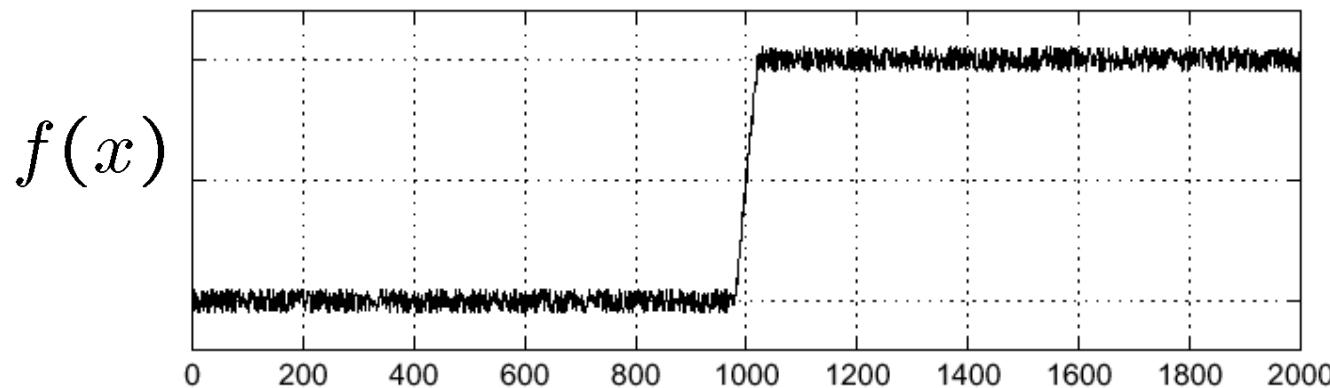


With a little Gaussian noise



Effects of noise

- Consider a single row or column of the image
 - Plotting intensity as a function of position gives a signal



Where is the edge?

Edge is Where Change Occurs

Change is measured by derivative in 1D

Biggest change, derivative has maximum magnitude

Or 2nd derivative is zero.

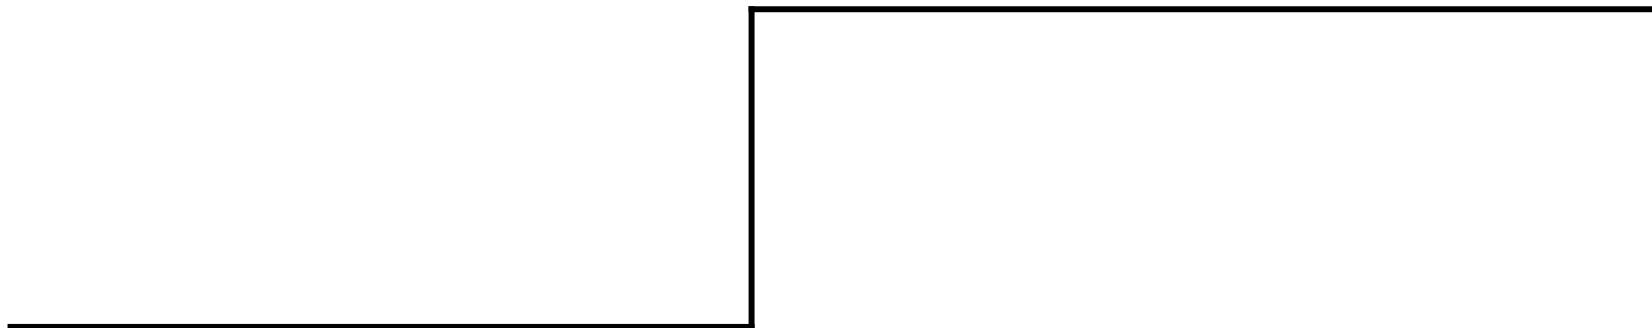
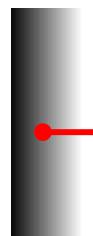


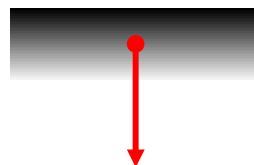
Image gradient

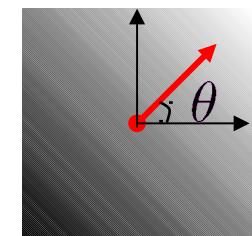
The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity


$$\nabla f = \left[\frac{\partial f}{\partial x}, 0 \right]$$


$$\nabla f = \left[0, \frac{\partial f}{\partial y} \right]$$


$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient direction is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

•

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x}\right)^2 + \left(\frac{\partial f}{\partial y}\right)^2}$$

The discrete gradient

How can we differentiate a *digital* image $f[x,y]$?

- Option 1: reconstruct a continuous image, then take gradient
- Option 2: take discrete derivative (finite difference)

$$\frac{\partial f}{\partial x}[x, y] \approx f[x + 1, y] - f[x, y]$$

The Sobel operator

Better approximations of the derivatives exist

- The *Sobel* operators below are very commonly used

$$\frac{1}{8} \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$

s_x

$$\frac{1}{8} \begin{array}{|c|c|c|} \hline 1 & 2 & 1 \\ \hline 0 & 0 & 0 \\ \hline -1 & -2 & -1 \\ \hline \end{array}$$

s_y

- The standard defn. of the Sobel operator omits the 1/8 term
 - doesn't make a difference for edge detection
 - the 1/8 term **is** needed to get the right gradient value, however

Gradient operators

 Δ_1 $\begin{matrix} 0 & 1 \\ -1 & 0 \end{matrix}$ Δ_2 $\begin{matrix} 1 & 0 \\ 0 & -1 \end{matrix}$

(a)

 Δ_1 $\begin{matrix} -1 & 0 & 1 \\ -1 & 0 & 1 \\ -1 & 0 & 1 \end{matrix}$ Δ_2 $\begin{matrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{matrix}$

(b)

 Δ_1 $\begin{matrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{matrix}$ Δ_2 $\begin{matrix} 1 & 2 & 1 \\ 0 & 0 & 0 \\ -1 & -2 & -1 \end{matrix}$

(c)

 Δ_1 $\begin{matrix} -3 & -1 & 1 & 3 \\ -3 & -1 & 1 & 3 \\ -3 & -1 & 1 & 3 \\ -3 & -1 & 1 & 3 \end{matrix}$ Δ_2 $\begin{matrix} 3 & 3 & 3 & 3 \\ 1 & 1 & 1 & 1 \\ -1 & -1 & -1 & -1 \\ -3 & -3 & -3 & -3 \end{matrix}$

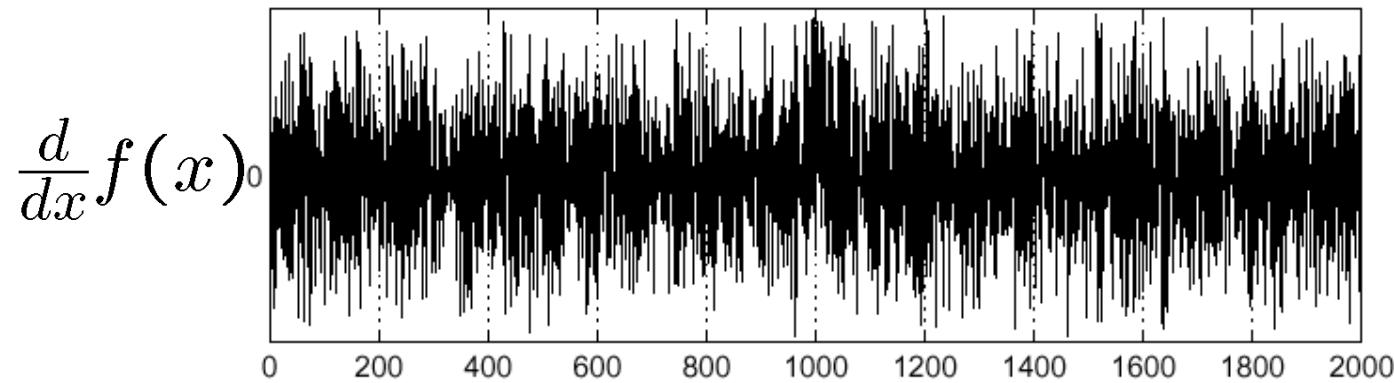
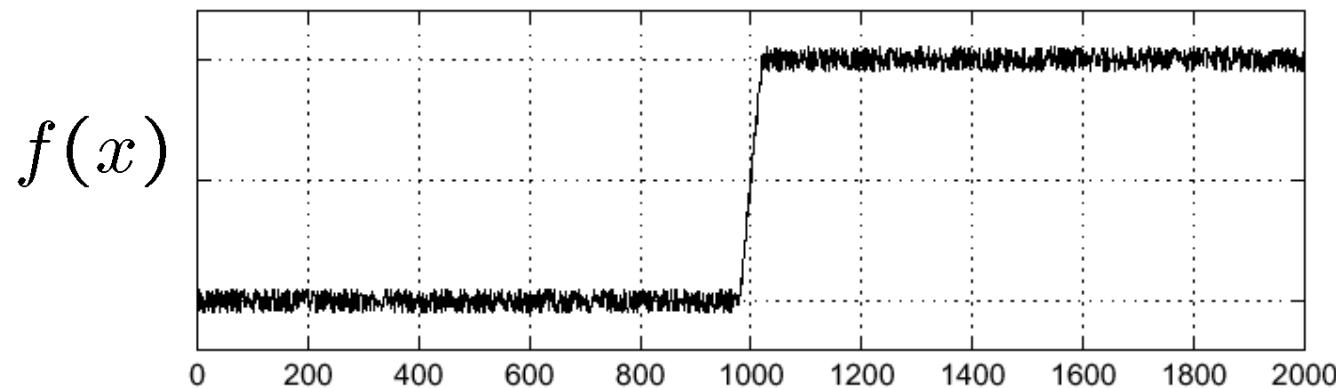
(d)

- (a): Roberts' cross operator (b): 3x3 Prewitt operator
(c): Sobel operator (d) 4x4 Prewitt operator

Effects of noise

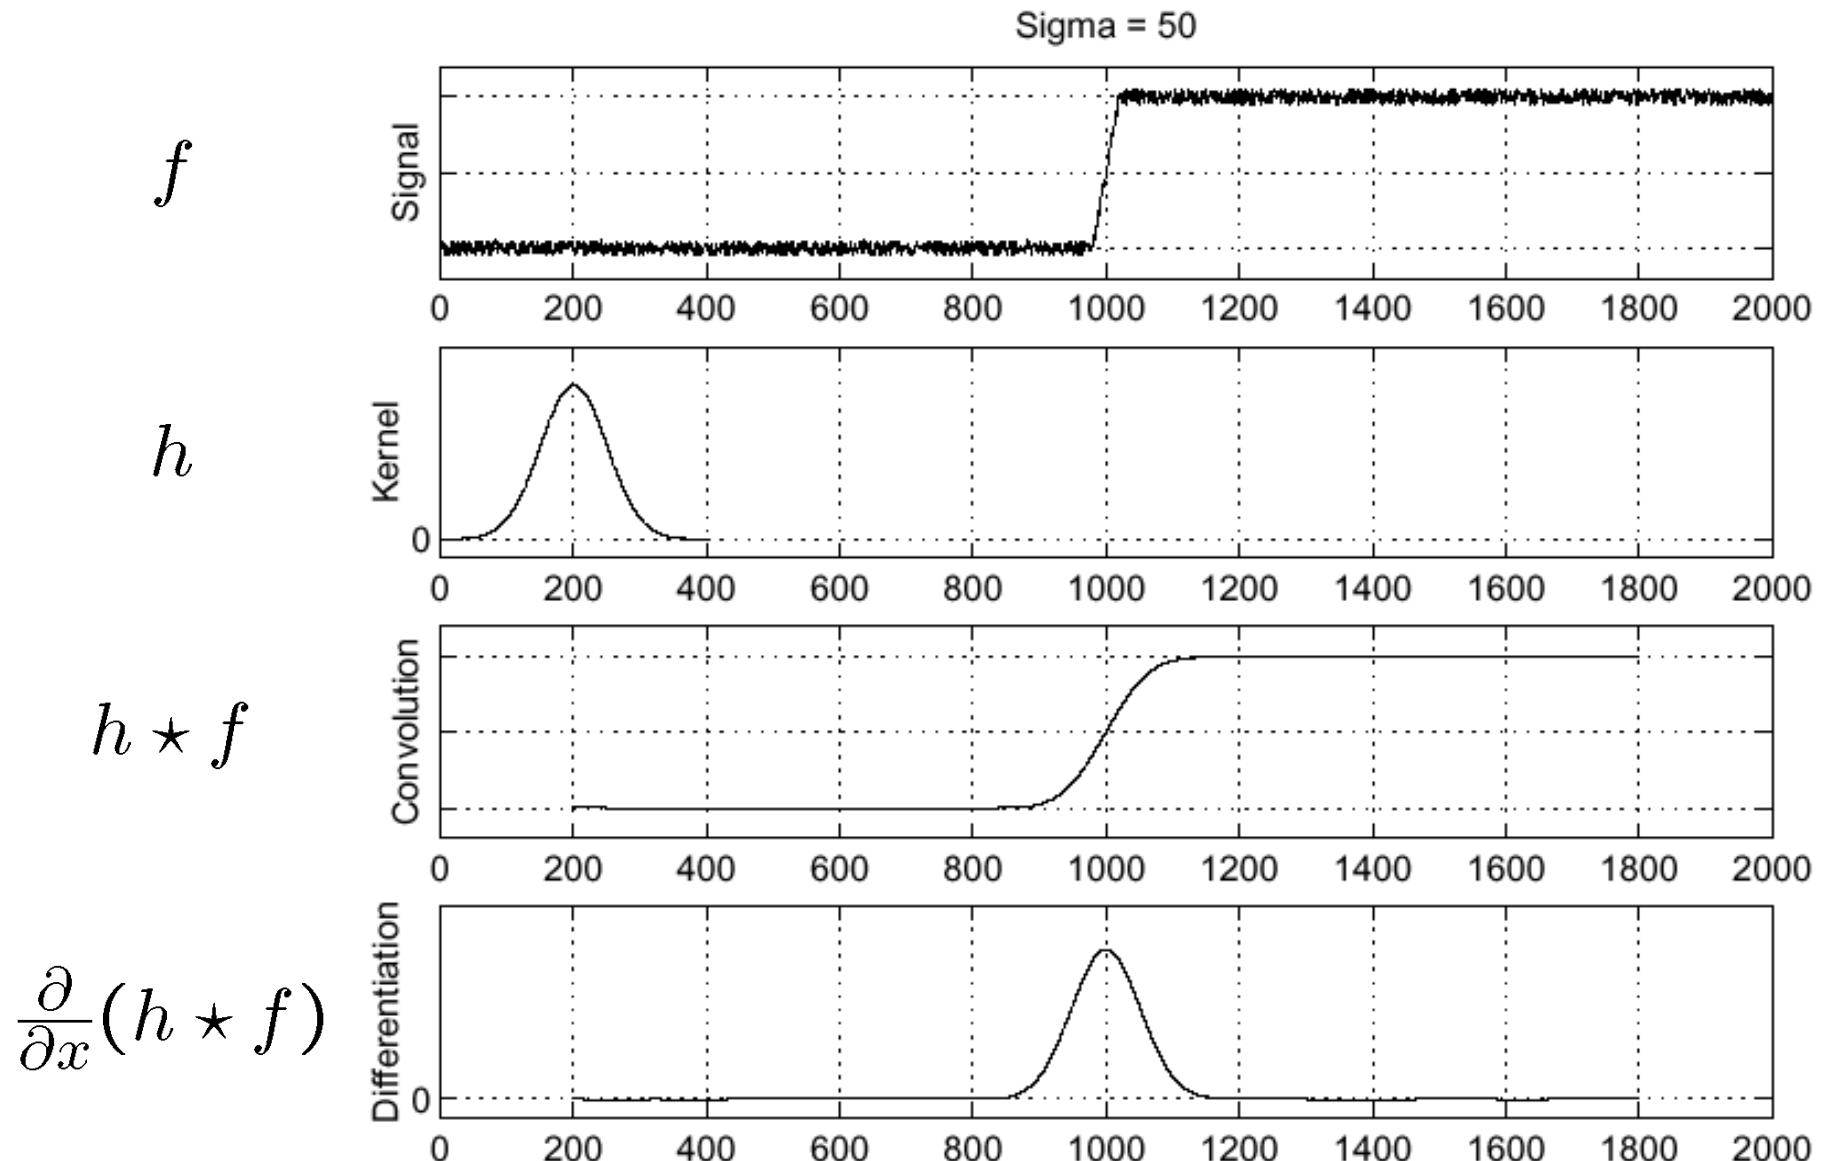
Consider a single row or column of the image

- Plotting intensity as a function of position gives a signal



Where is the edge?

Solution: smooth first



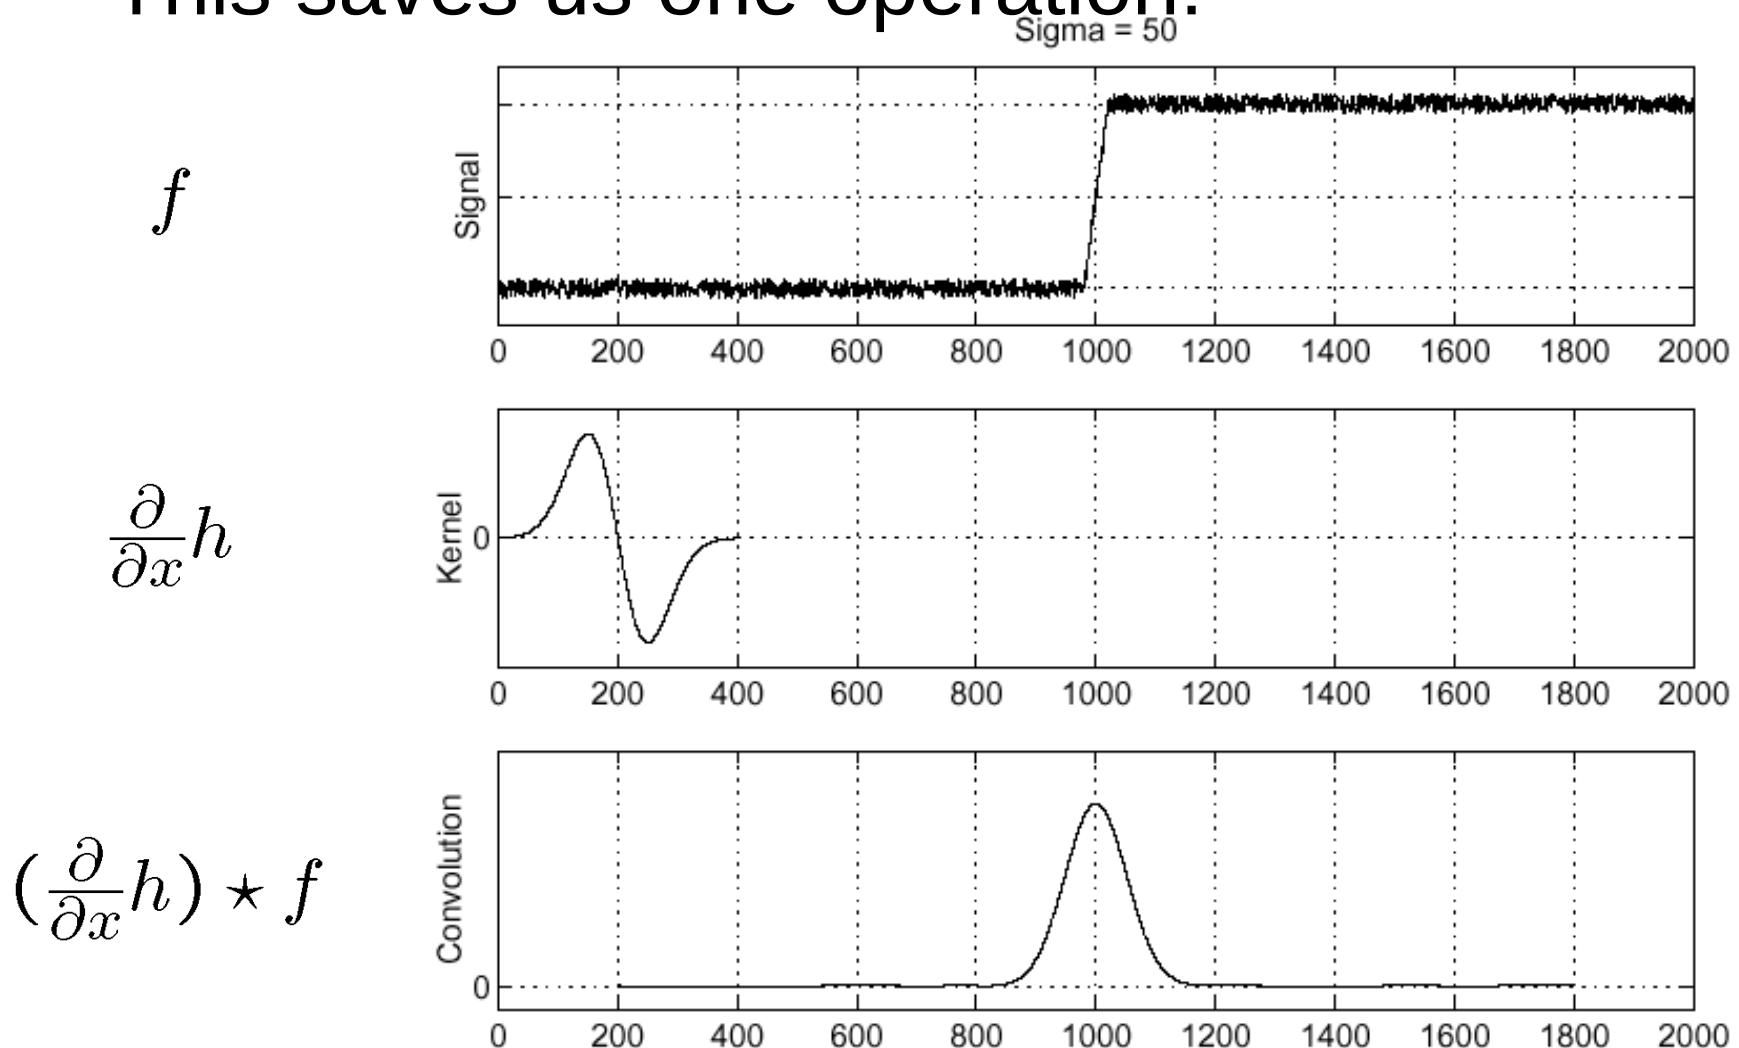
Where is the edge? Look for peaks in

$$\frac{\partial}{\partial x}(h \star f)$$

Derivative theorem of convolution

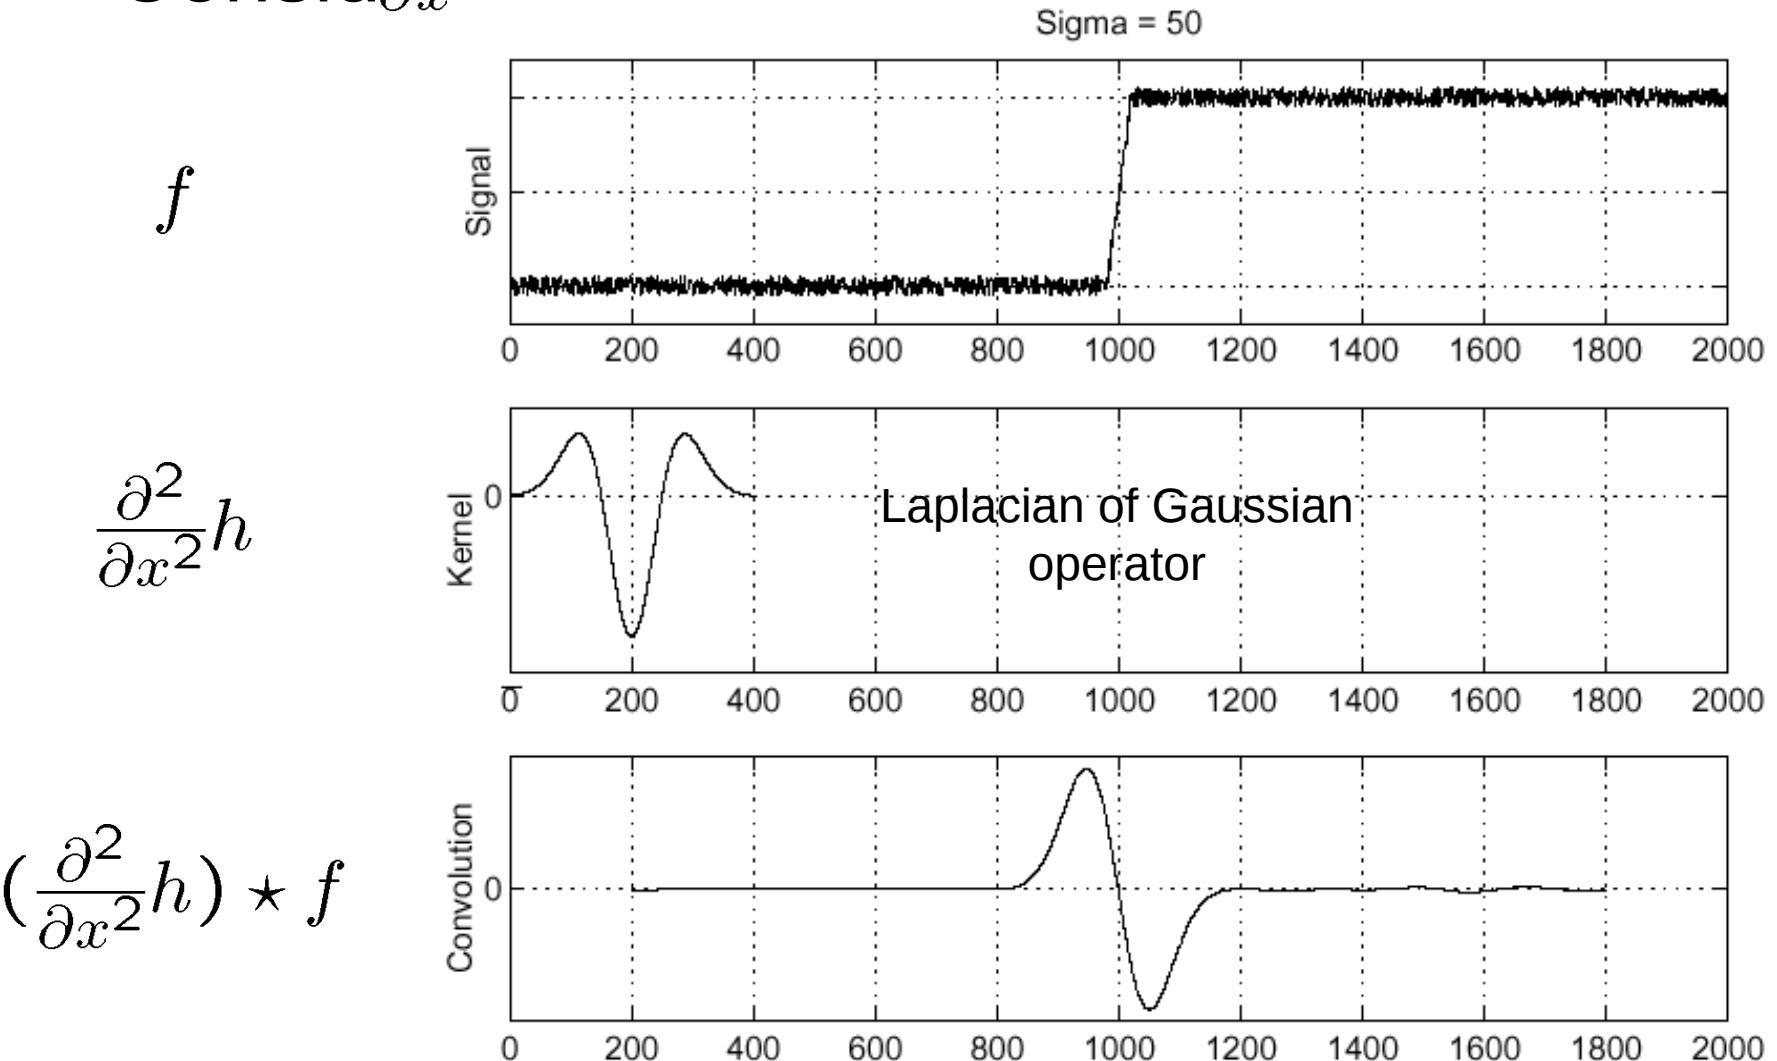
$$\frac{\partial}{\partial x}(h \star f) = (\frac{\partial}{\partial x}h) \star f$$

This saves us one operation:



Laplacian of Gaussian

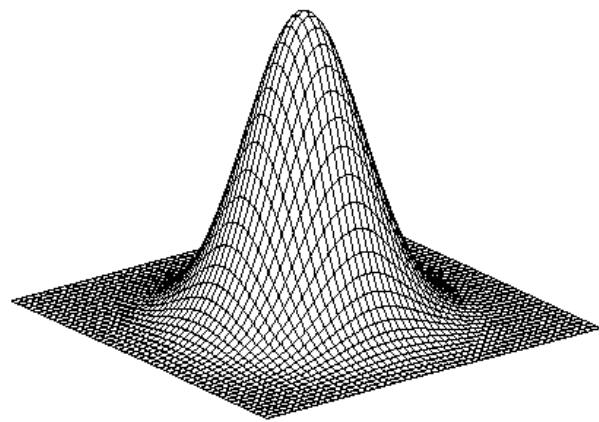
Consider $\frac{\partial^2}{\partial x^2}(h \star f)$



Where is the edge?

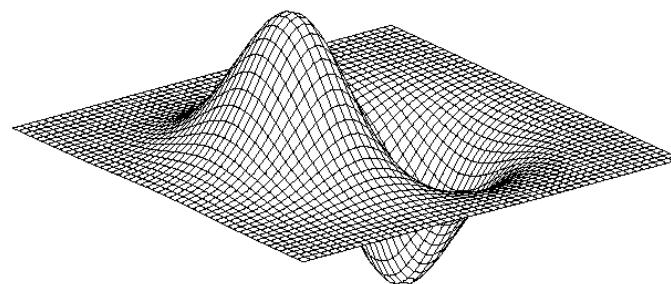
Zero-crossings of bottom graph

2D edge detection filters



Gaussian

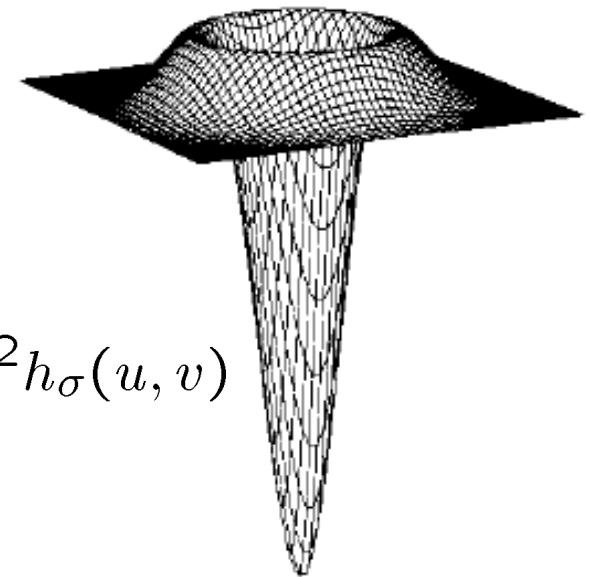
$$h_\sigma(u, v) = \frac{1}{2\pi\sigma^2} e^{-\frac{u^2+v^2}{2\sigma^2}}$$



derivative of Gaussian

$$\frac{\partial}{\partial x} h_\sigma(u, v)$$

Laplacian of Gaussian



$$\nabla^2 h_\sigma(u, v)$$

∇^2 is the **Laplacian** operator:

$$\nabla^2 f = \frac{\partial^2 f}{\partial x^2} + \frac{\partial^2 f}{\partial y^2}$$

Edge detection by subtraction



original

Edge detection by subtraction



smoothed (5x5 Gaussian)

Edge detection by subtraction

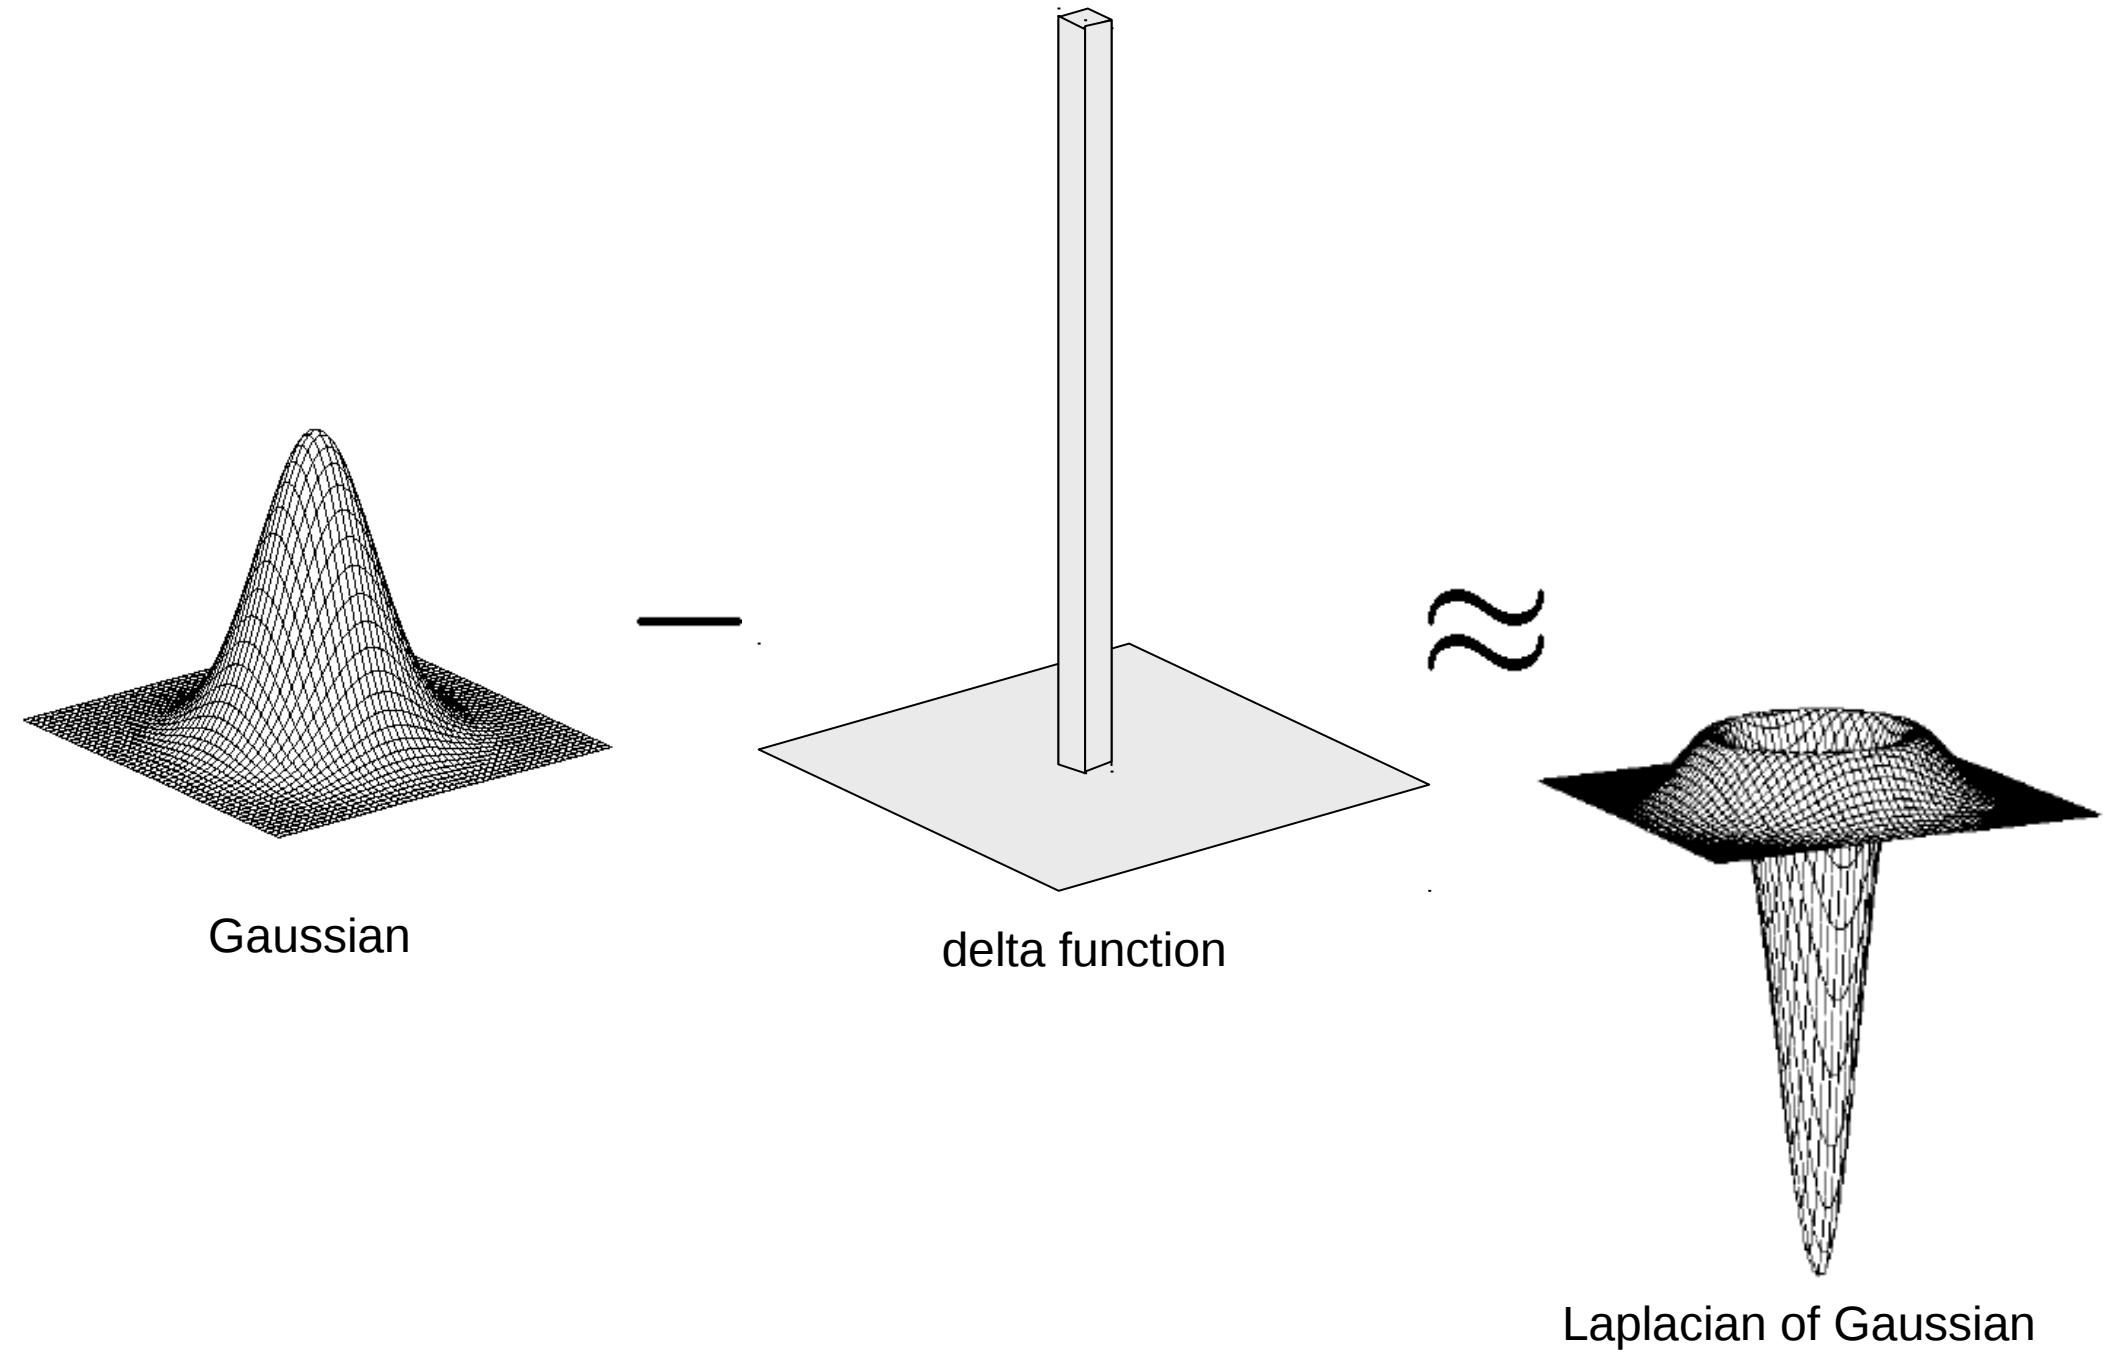


smoothed – original
(scaled by 4, offset +128)

Why does
this work?

filter demo

Gaussian - image filter



Optimal Edge Detection: Canny

Assume:

- Linear filtering
- Additive iid Gaussian noise

Edge detector should have:

- Good Detection. Filter responds to edge, not noise.
- Good Localization: detected edge near true edge.
- Single Response: one per edge.

Optimal Edge Detection: Canny (continued)

Optimal Detector is approximately
Derivative of Gaussian.

Detection/Localization trade-off

- More smoothing improves detection
- And hurts localization.

This is what you might guess from (detect
change) + (remove noise)

The Canny edge detector



original image (Lena)

The Canny edge detector



norm of the gradient

The Canny edge detector



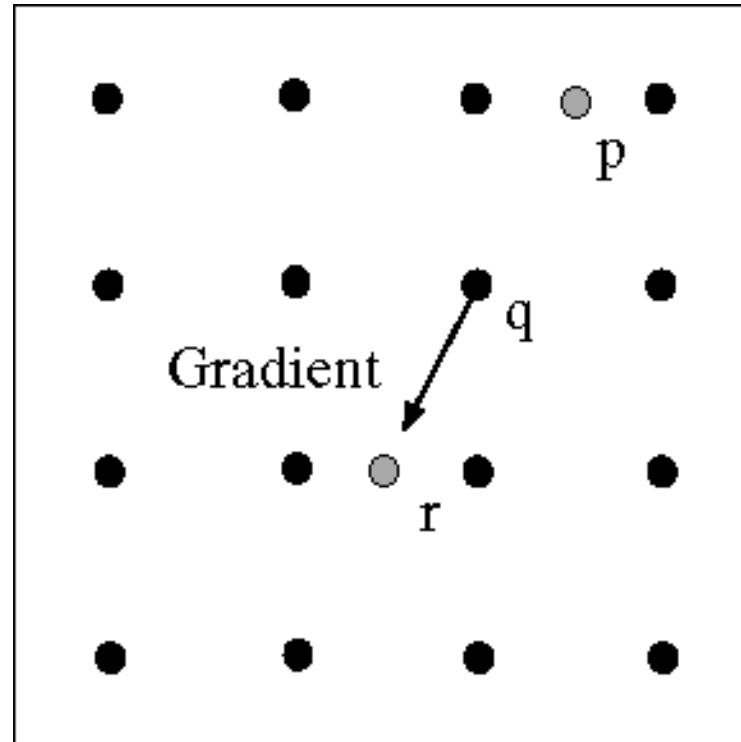
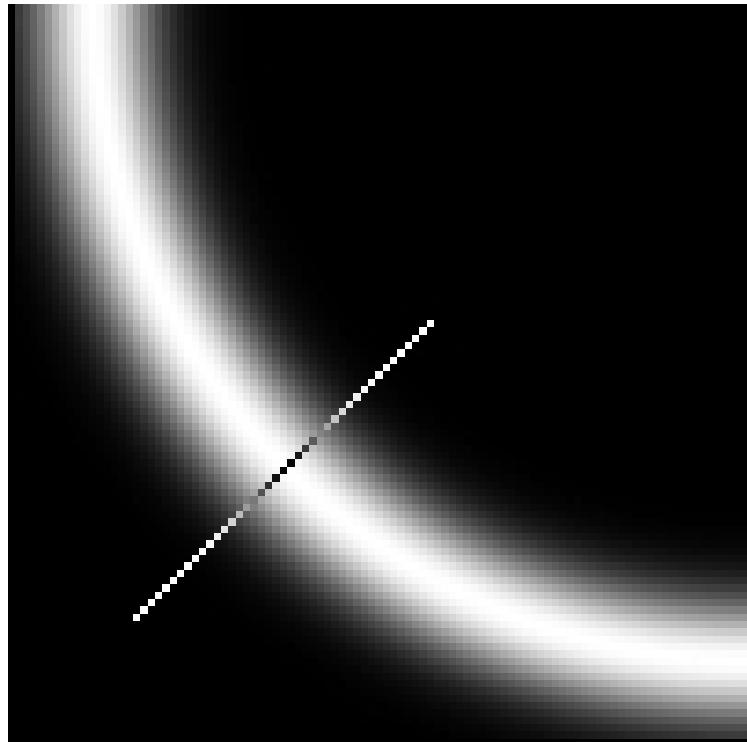
thresholding

The Canny edge detector



thinning
(non-maximum suppression)

Non-maximum suppression

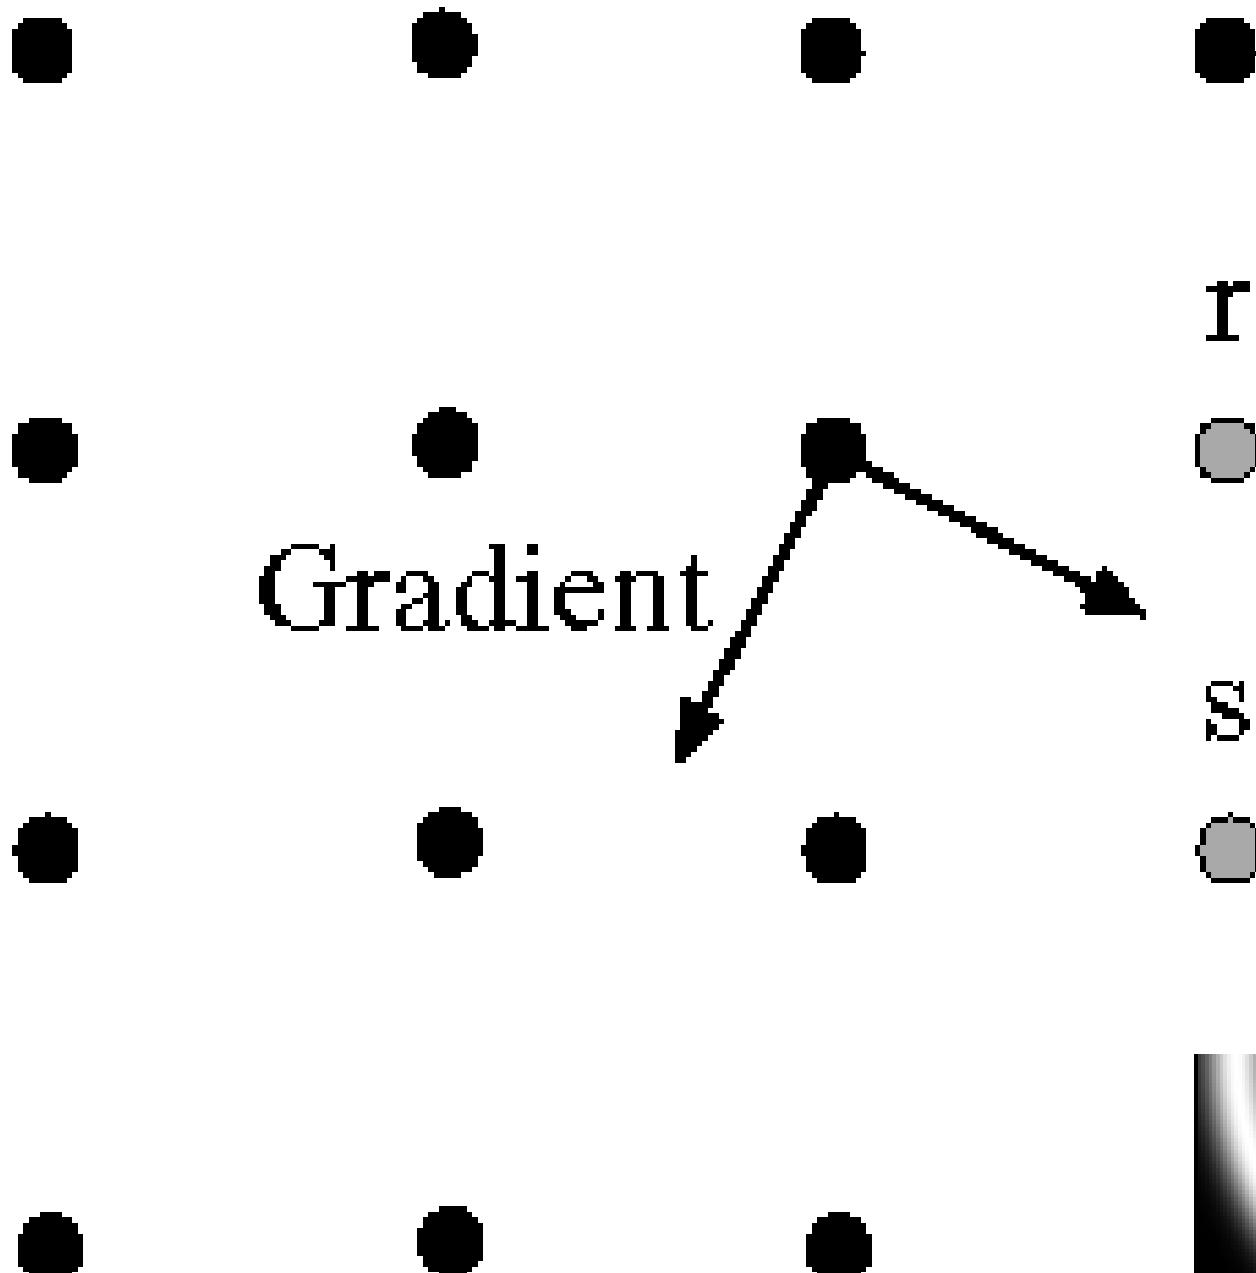


Check if pixel is local maximum along gradient direction
– requires checking interpolated pixels p and r

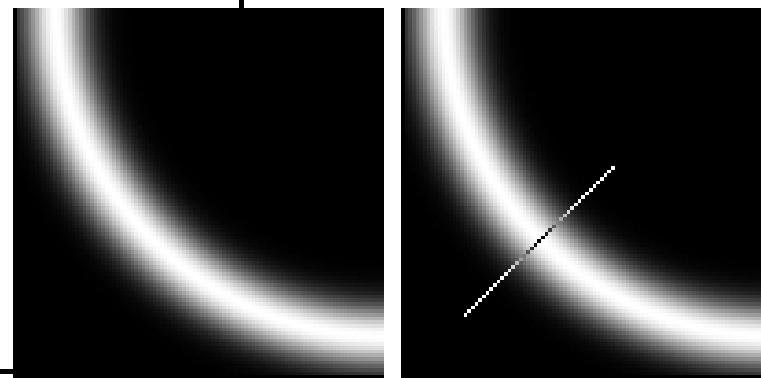
Predicting the next edge point

Assume the marked point is an edge point. Then we construct the tangent to the edge curve (which is normal to the gradient at that point) and use this to predict the next points (here either r or s).

Gradient



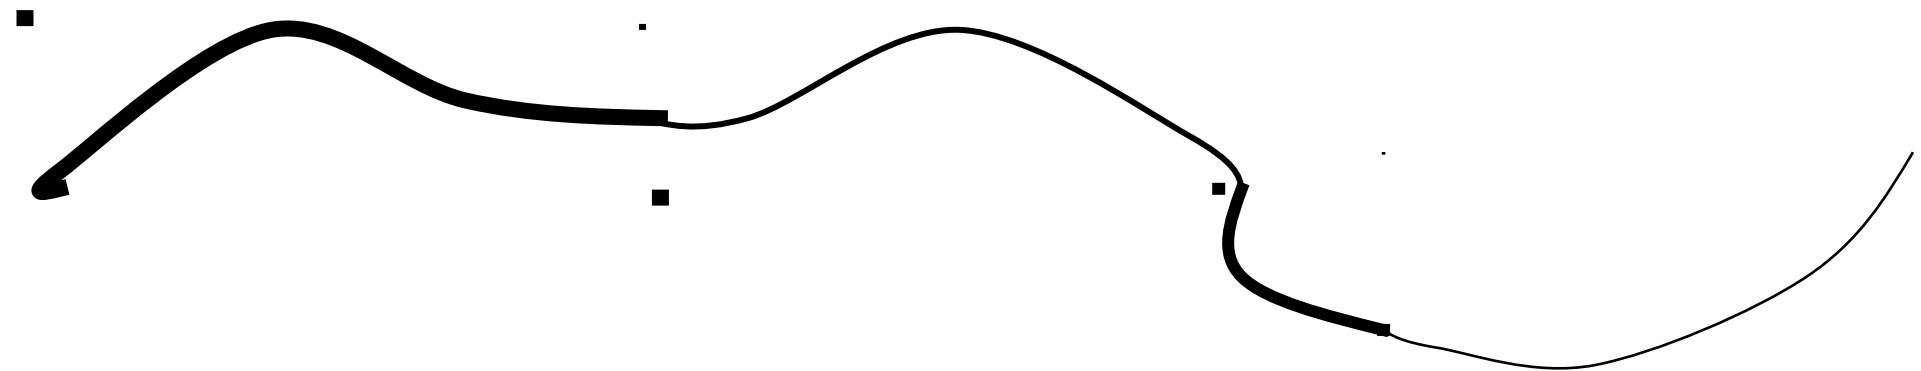
(Forsyth & Ponce)



Hysteresis

Check that maximum value of gradient
value is sufficiently large

- drop-outs? use **hysteresis**
 - use a high threshold to start edge curves and a low threshold to continue them.



Effect of σ (Gaussian kernel size)



original



Canny with $\sigma = 1$

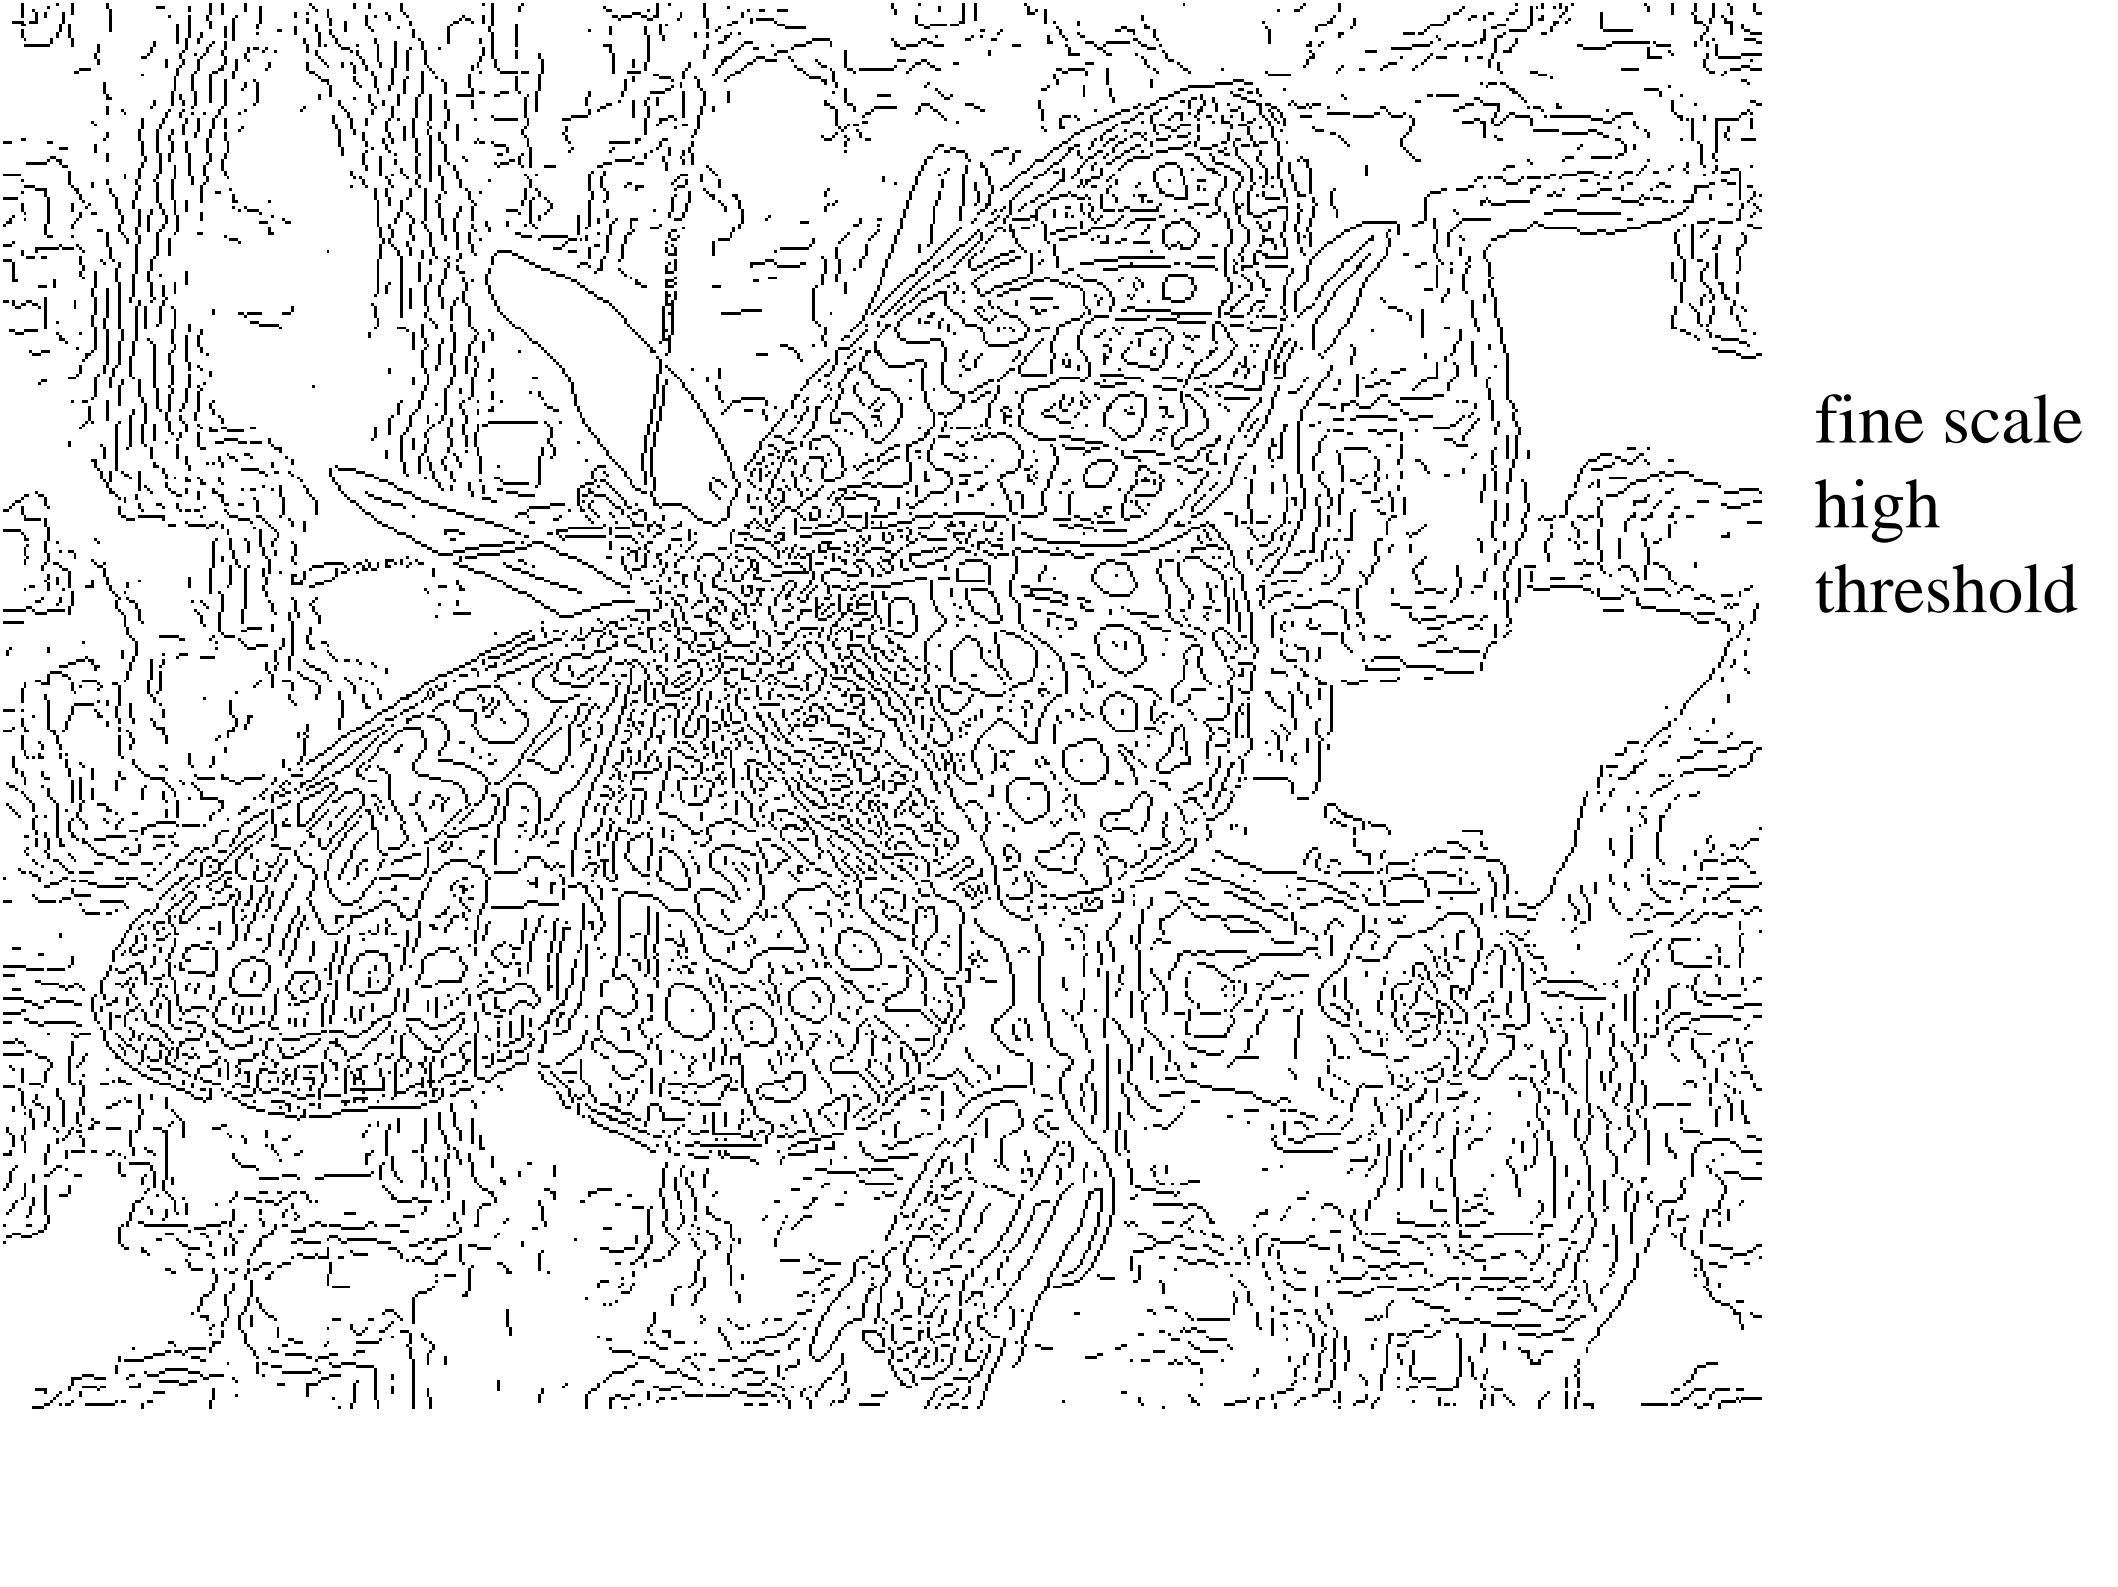


Canny with $\sigma = 2$

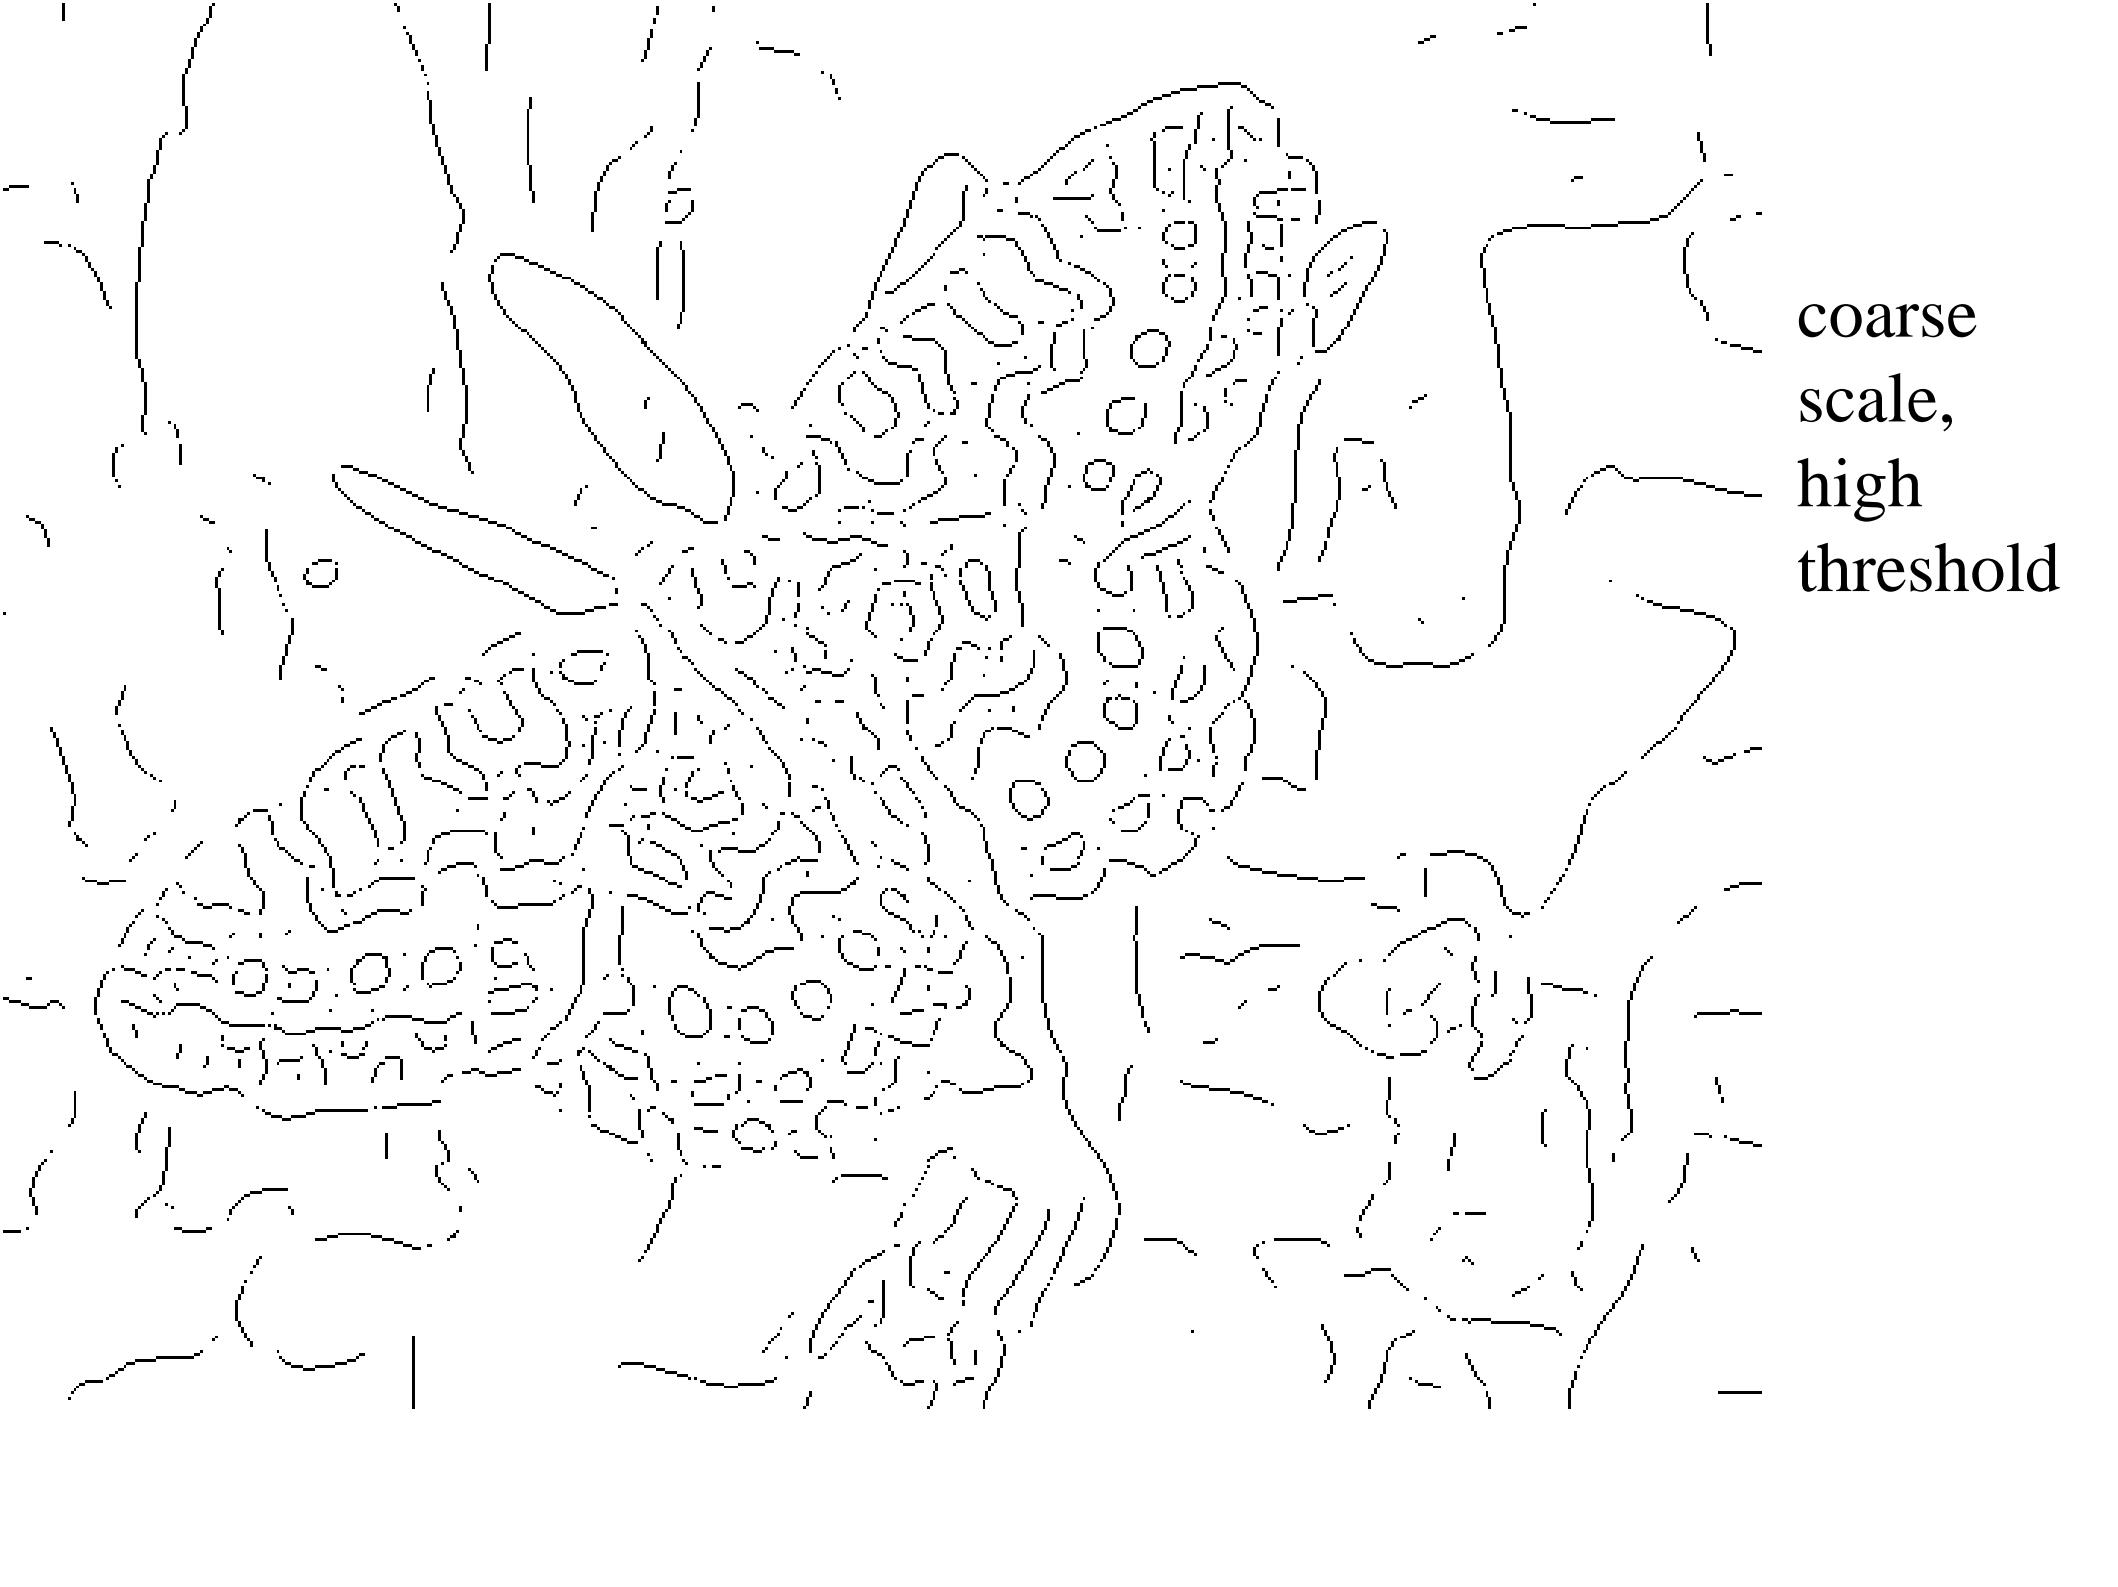
The choice of σ depends on desired behavior

- large σ detects large scale edges
- small σ detects fine features



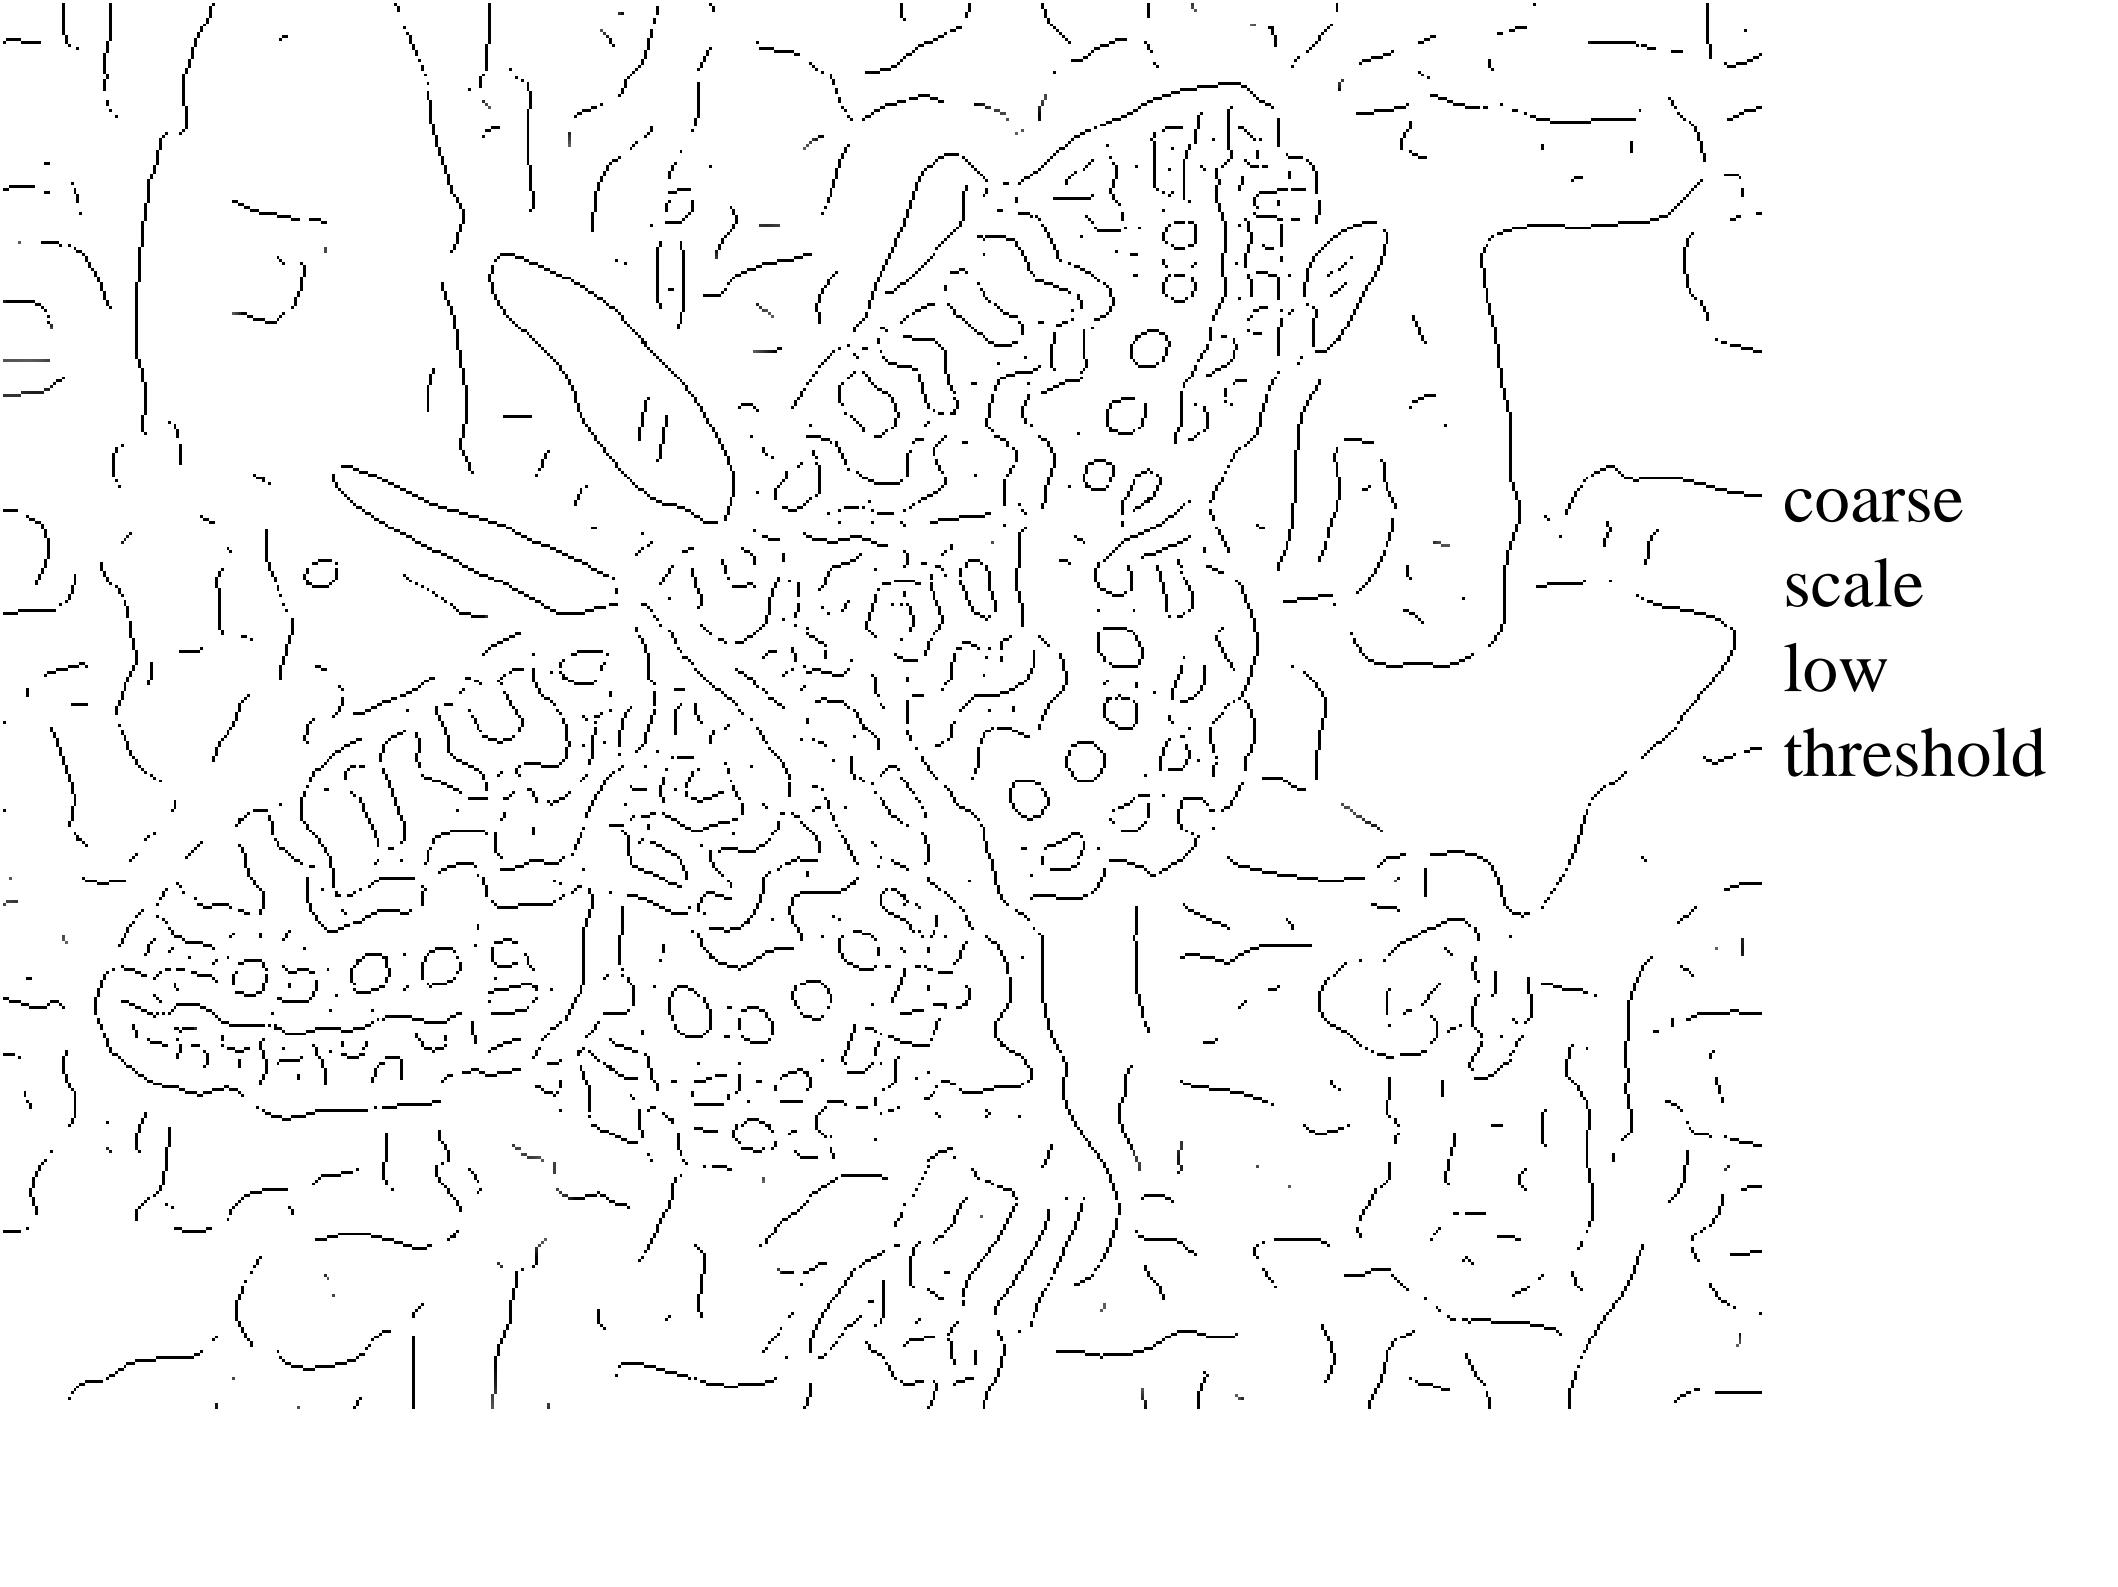


**fine scale
high
threshold**



A high-resolution grayscale image of a brain scan. Overlaid on the image is a coarse segmentation map, represented by thick black outlines that divide the brain into large, irregular regions. These regions represent different tissue types or structures at a larger scale than a fine-grained segmentation. The overall appearance is more abstract and less detailed than a standard anatomical brain scan.

coarse
scale,
high
threshold



coarse
scale
low
threshold

So, what scale to choose?

It depends what we're looking for.

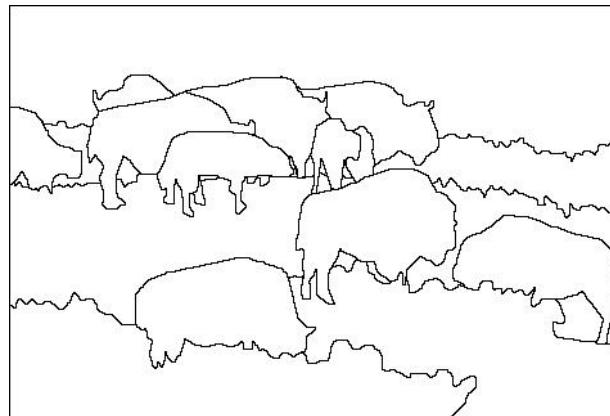


Learning to detect boundaries

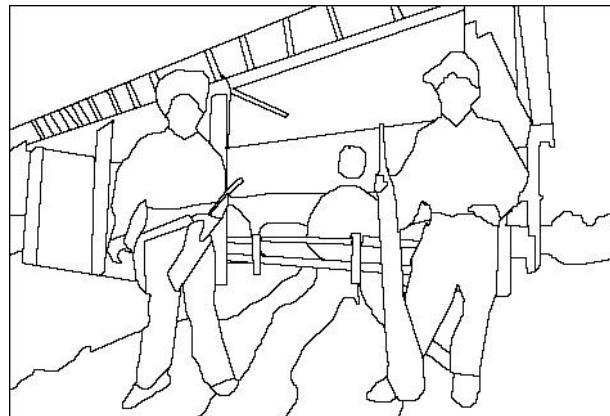
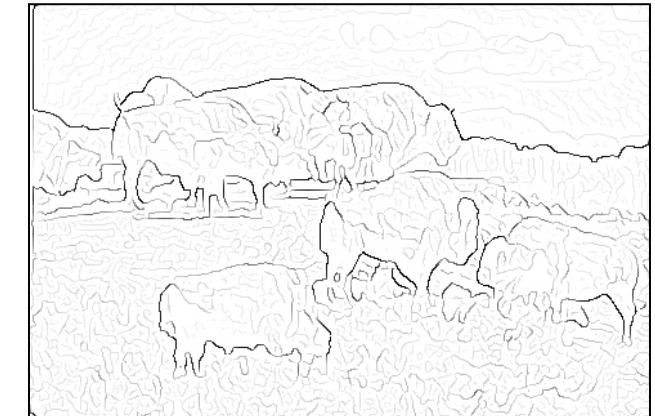
image



human segmentation



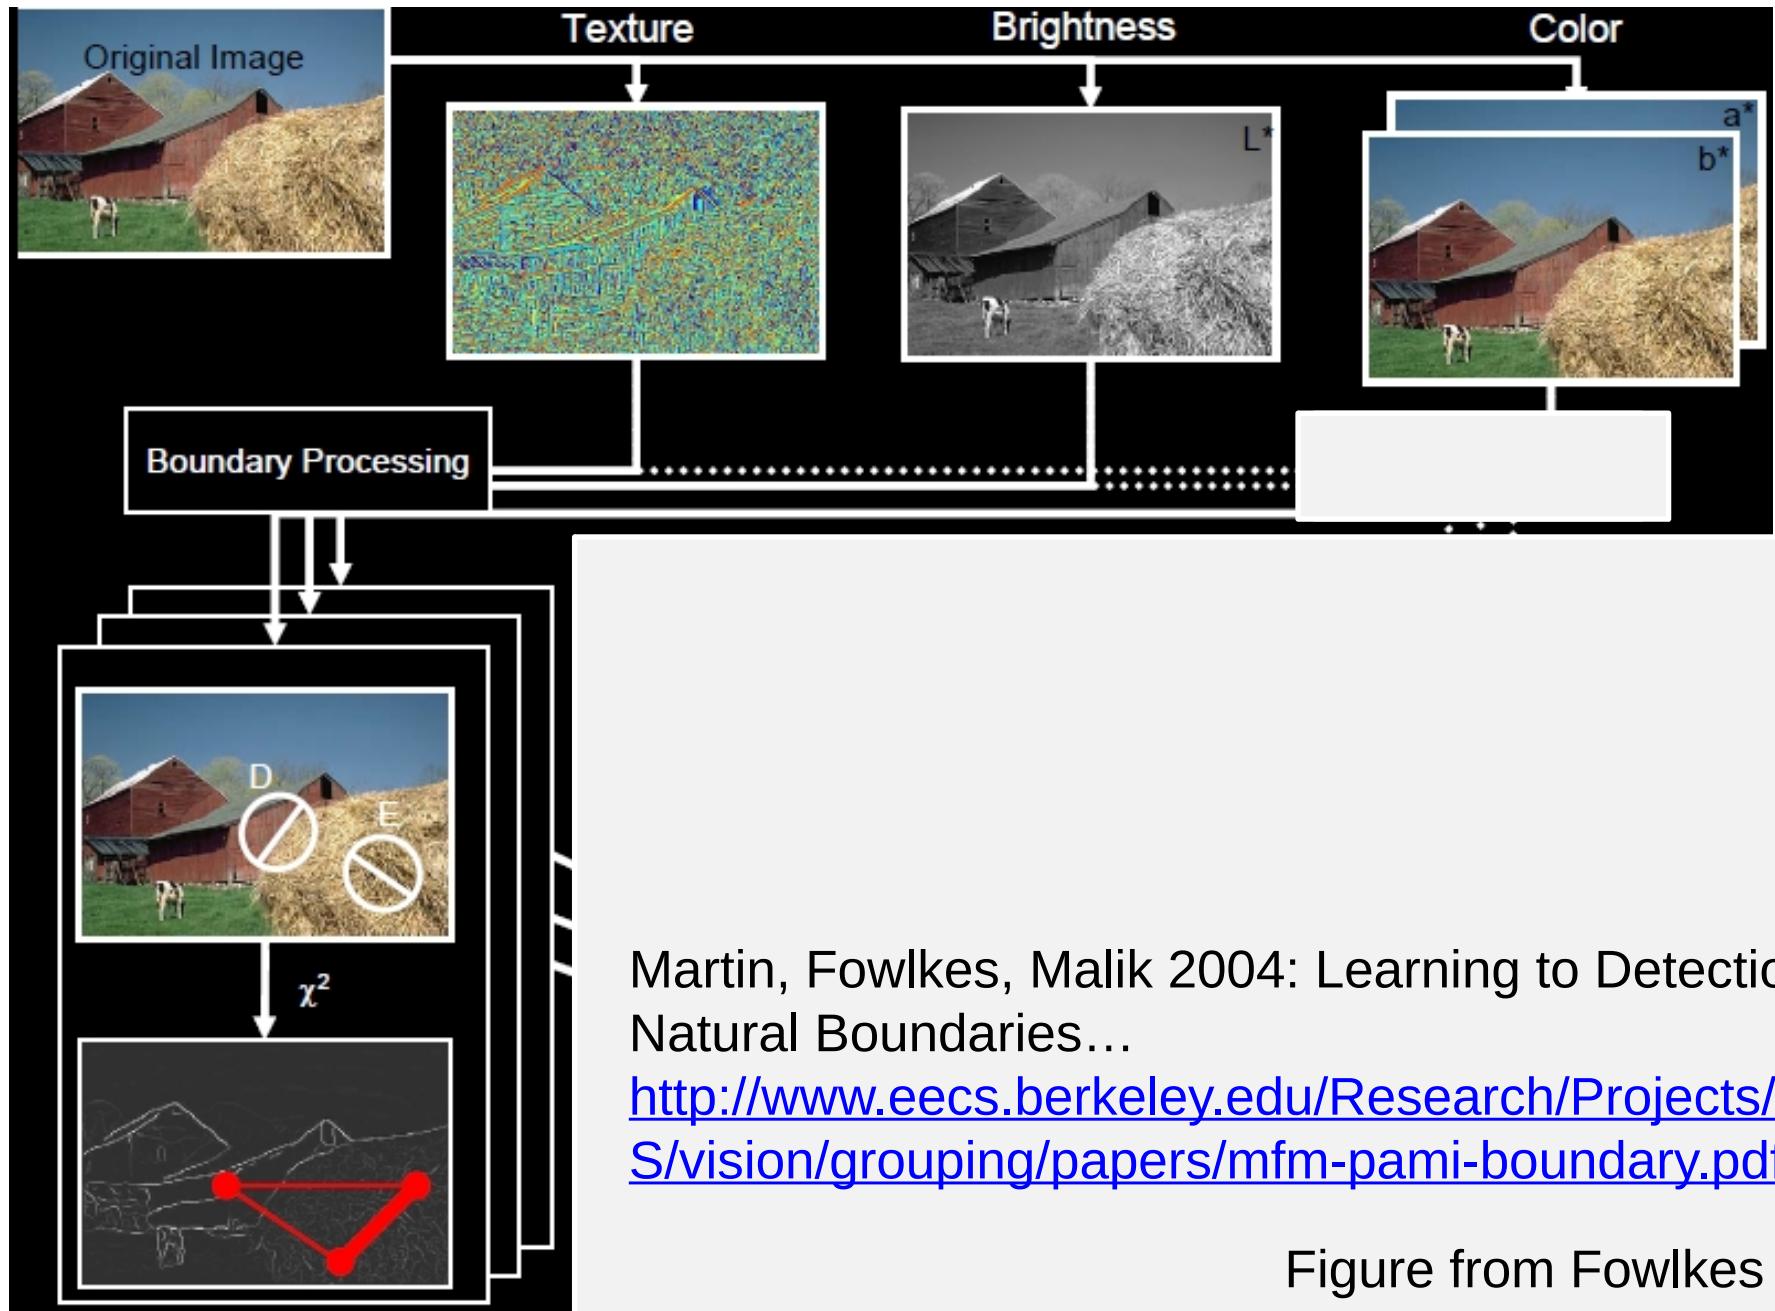
gradient magnitude



- Berkeley segmentation database:

<http://www.eecs.berkeley.edu/Research/Projects/CS/vision/grouping/segbench/>

pB boundary detector



pB Boundary Detector

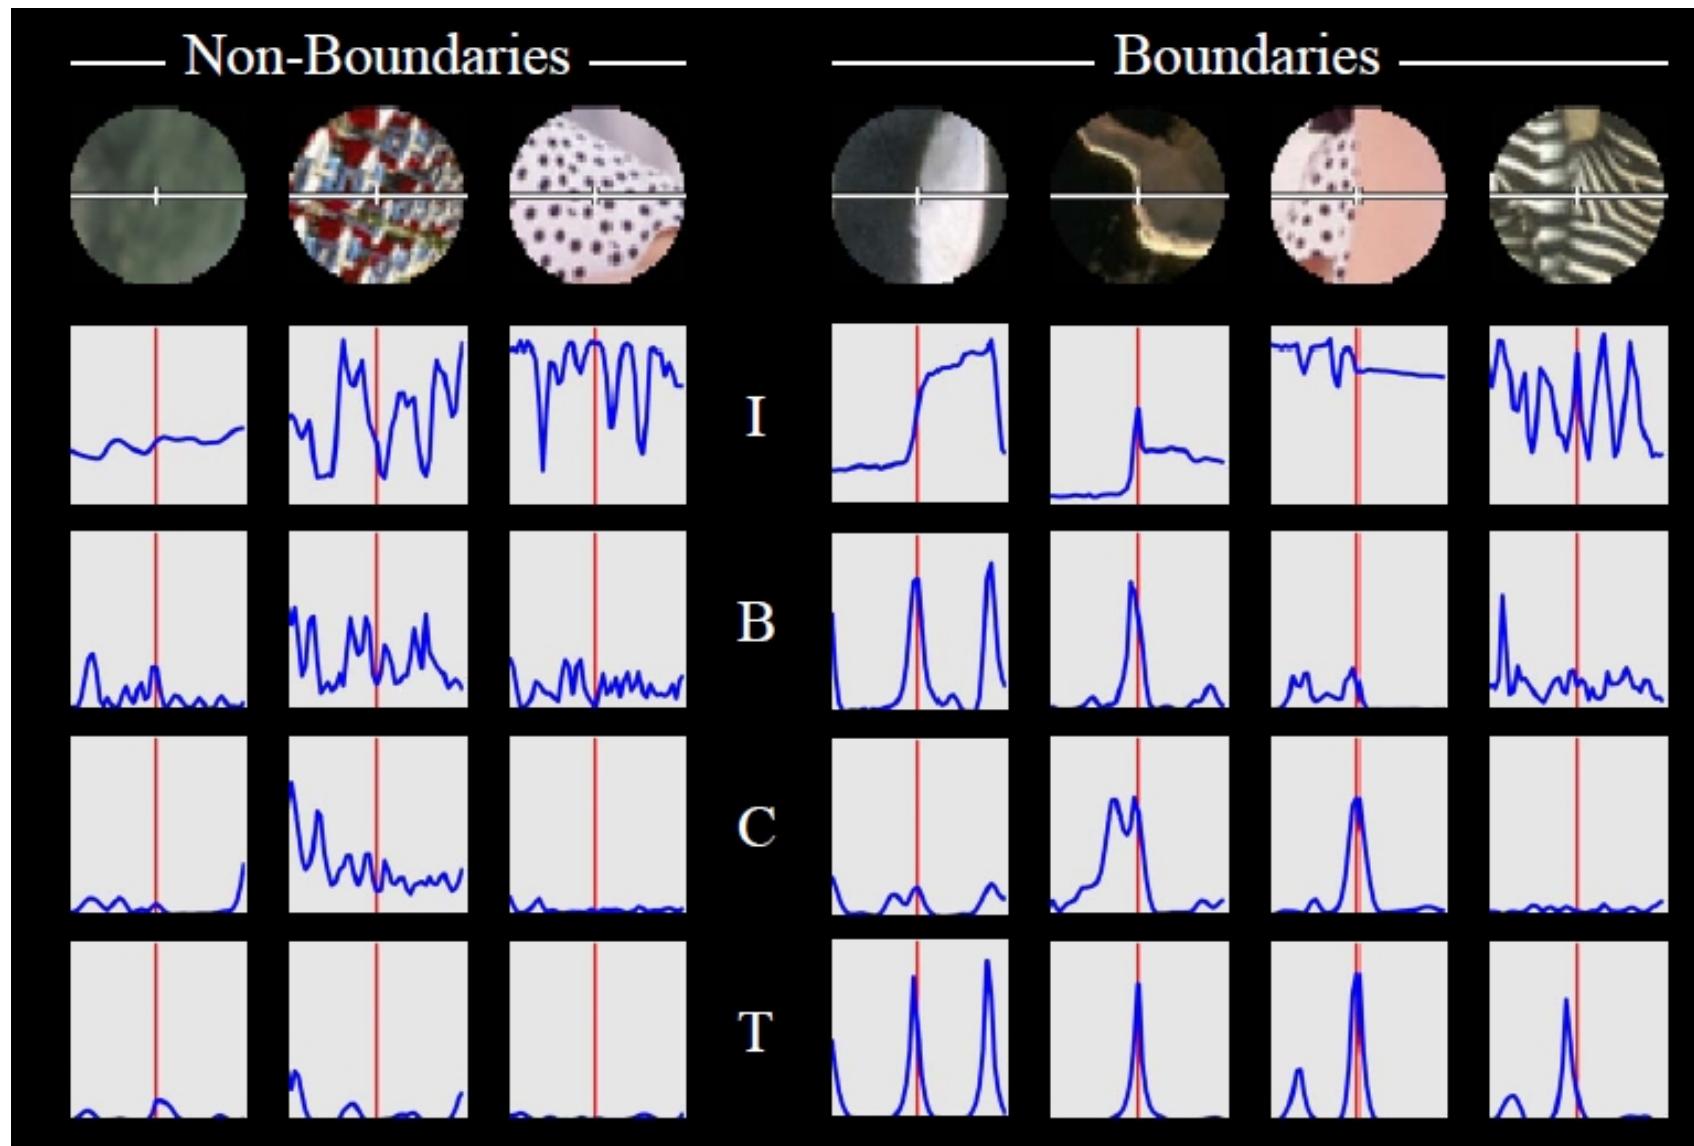
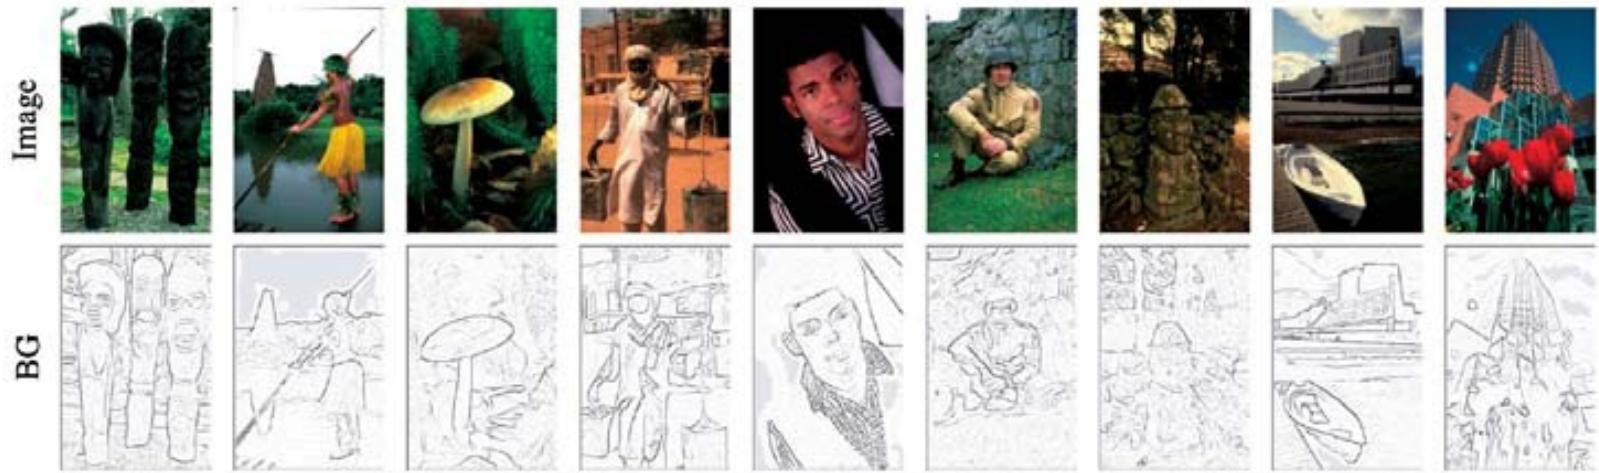
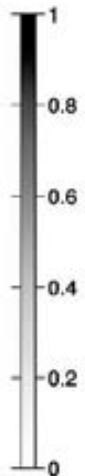


Figure from Fowlkes

Brightness



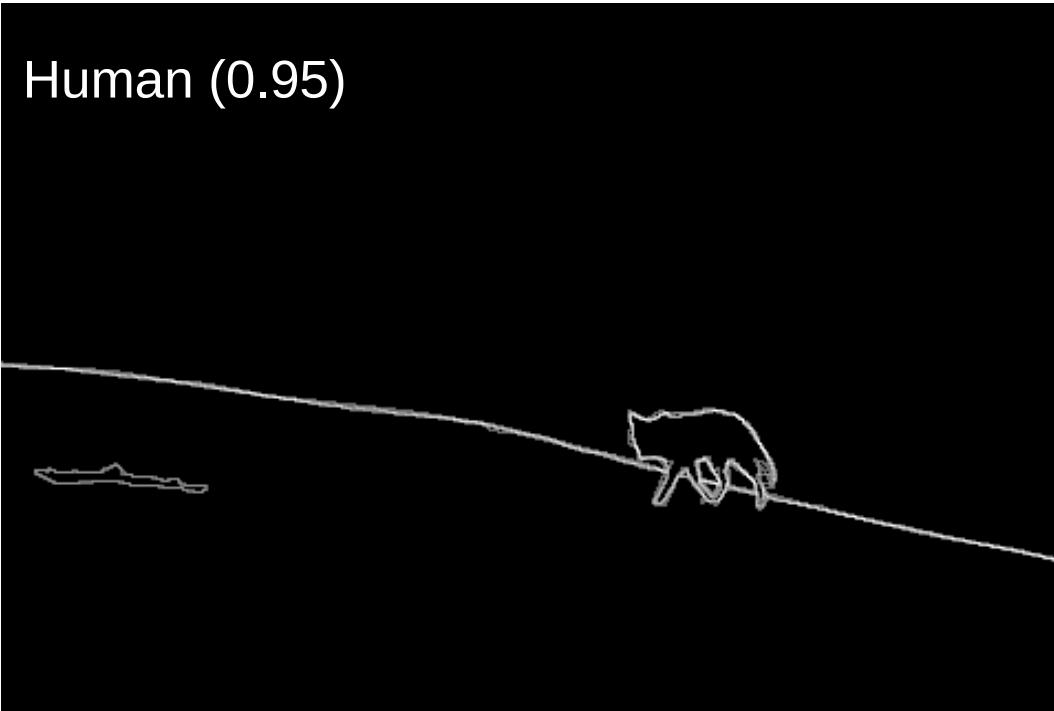
Color



Results



Human (0.95)

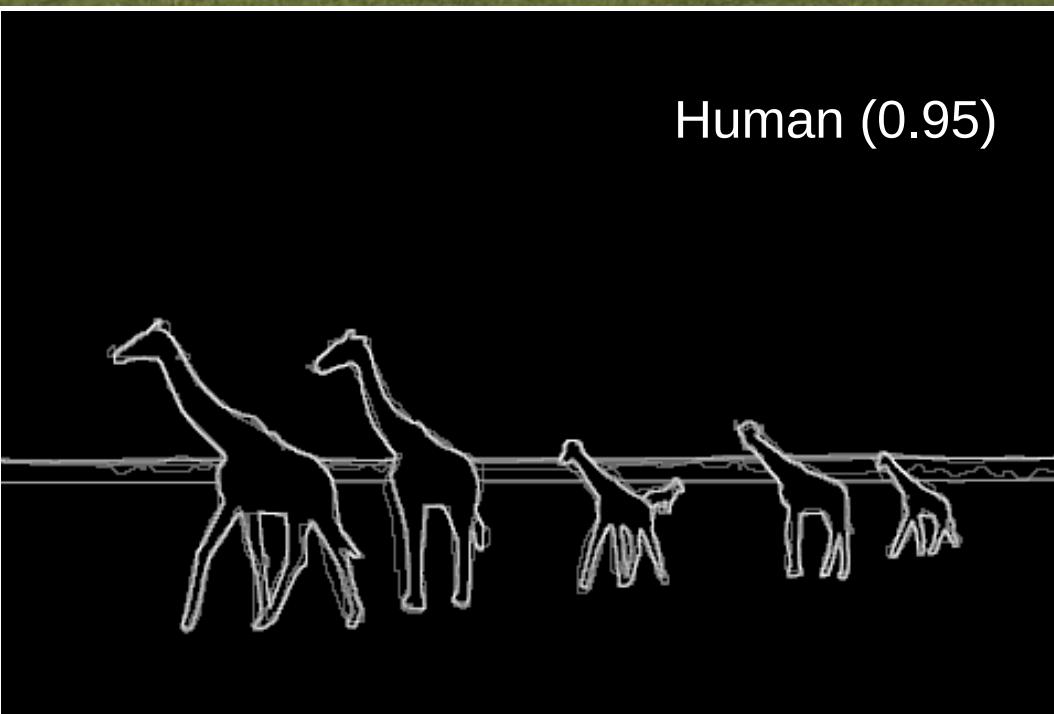


Pb (0.88)

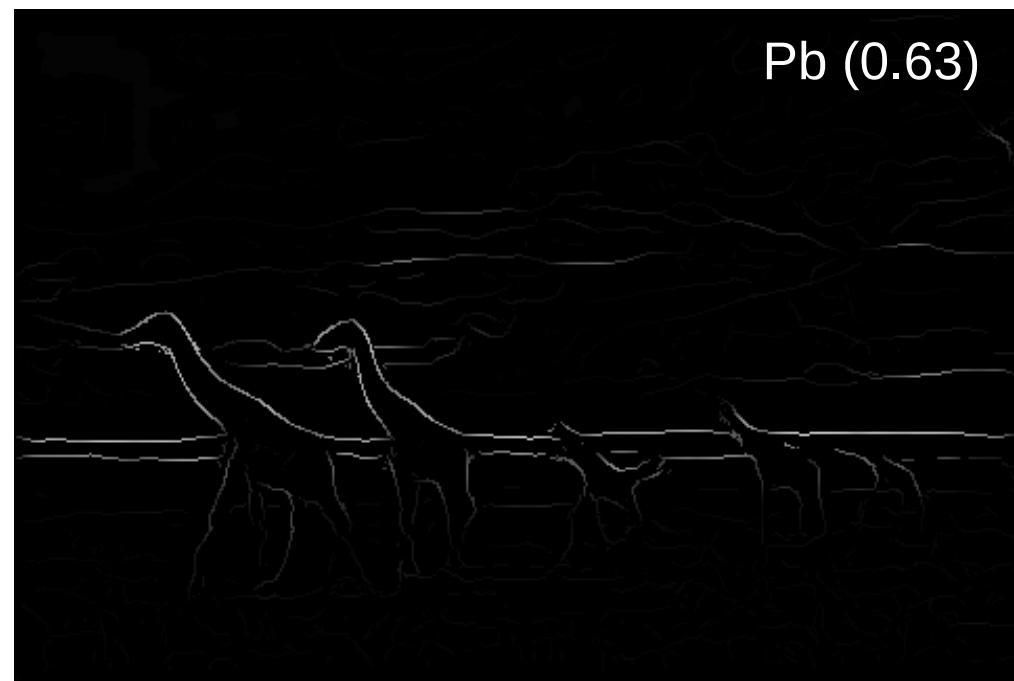


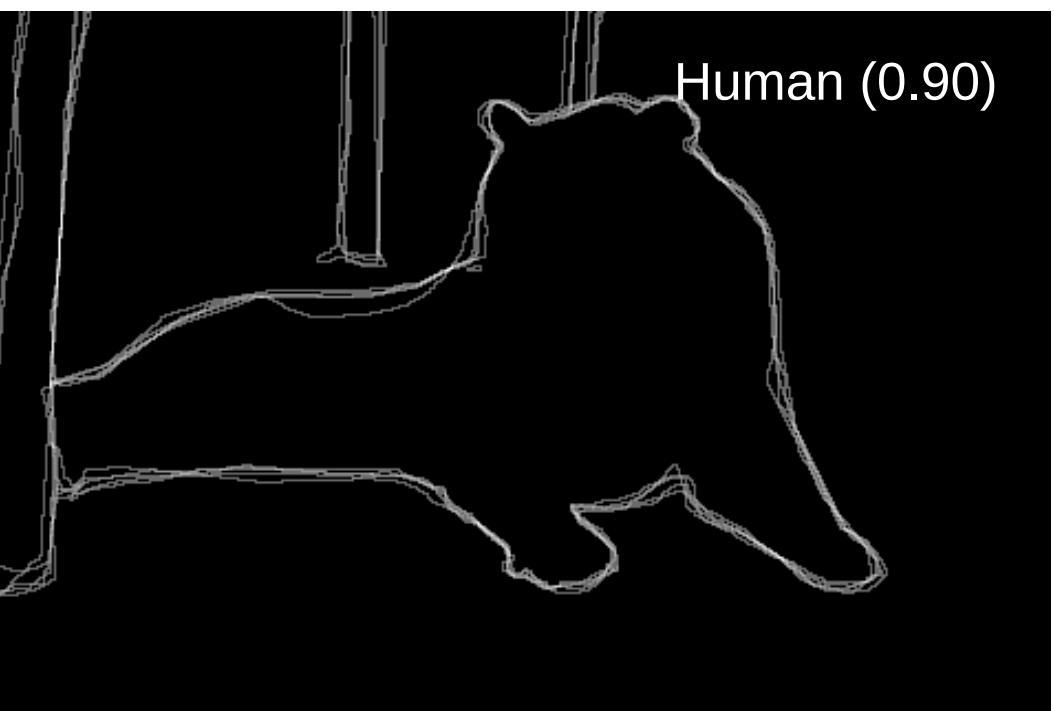


Human (0.95)



Pb (0.63)





For more: <http://www.eecs.berkeley.edu/Research/Projects/CS/vision/bsds/bench/html/108082-color.html>

State of edge detection

- Local edge detection is mostly solved
 - Intensity **gradient**, color, **texture**
- Some methods to take into account longer contours, but could probably do better
- Poor use of object and high-level information