

# Texture

Vinay P. Namboodiri

- Slide credit: Kristen Grauman

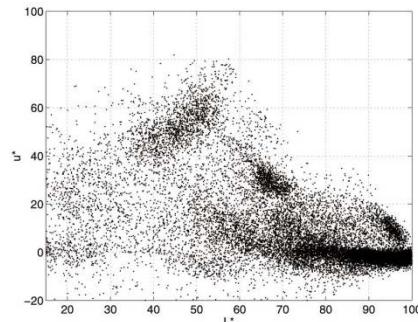
# Review

# Mean shift clustering

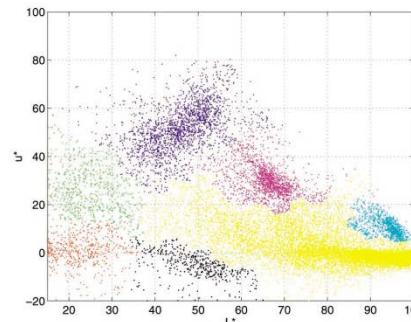
- The mean shift algorithm seeks *modes* of the given set of points
  1. Choose kernel and bandwidth
  2. For each point:
    - a) Center a window on that point
    - b) Compute the mean of the data in the search window
    - c) Center the search window at the new mean location
    - d) Repeat (b,c) until convergence
  3. Assign points that lead to nearby modes to the same cluster

# Segmentation by Mean Shift

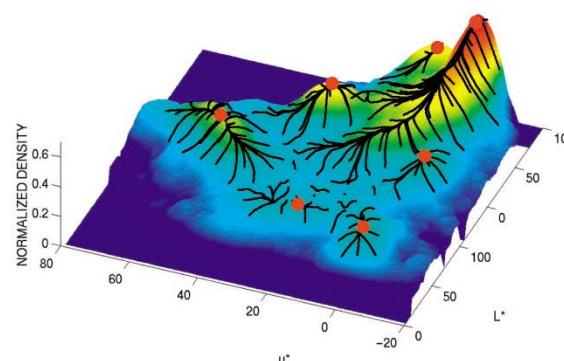
- Compute features for each pixel (color, gradients, texture, etc); also store each pixel's position
- Set kernel size for features  $K_f$  and position  $K_s$
- Initialize windows at individual pixel locations
- Perform mean shift for each window until convergence
- Merge modes that are within width of  $K_f$  and  $K_s$



(a)



(b)



# Mean shift segmentation results



# Meyer's watershed segmentation

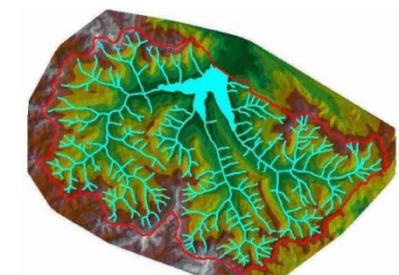
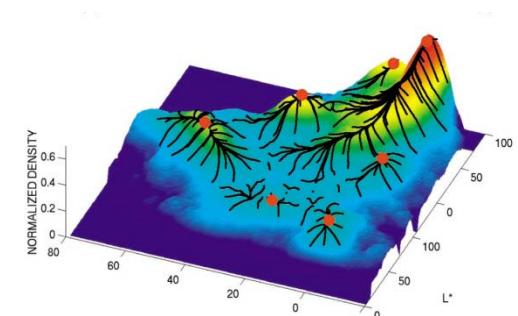
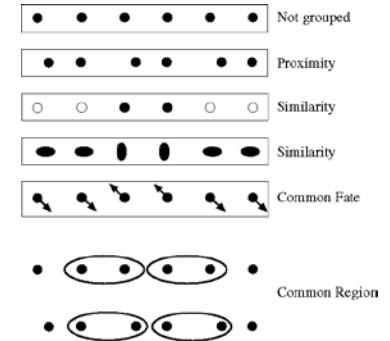
1. Choose local minima as region seeds
2. Add neighbors to priority queue, sorted by value
3. Take top priority pixel from queue
  1. If all labeled neighbors have same label, assign that label to pixel
  2. Add all non-marked neighbors to queue
4. Repeat step 3 until finished (all remaining pixels in queue are on the boundary)

Matlab: `seg = watershed(bnd_im)`

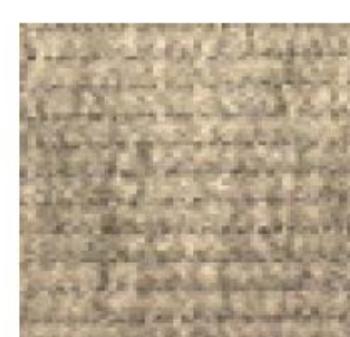
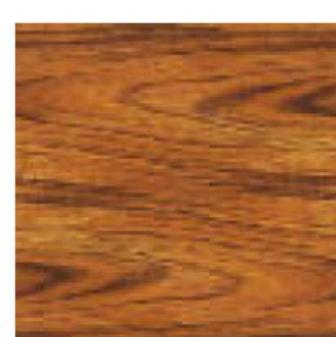
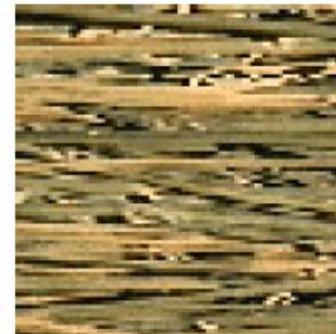
Meyer 1991

# Things to remember

- Gestalt cues and principles of organization
- Uses of segmentation
  - Efficiency
  - Better features
  - Want the segmented object
- Mean-shift segmentation
  - Good general-purpose segmentation method
  - Generally useful clustering, tracking technique
- Watershed segmentation
  - Good for hierarchical segmentation
  - Use in combination with boundary prediction

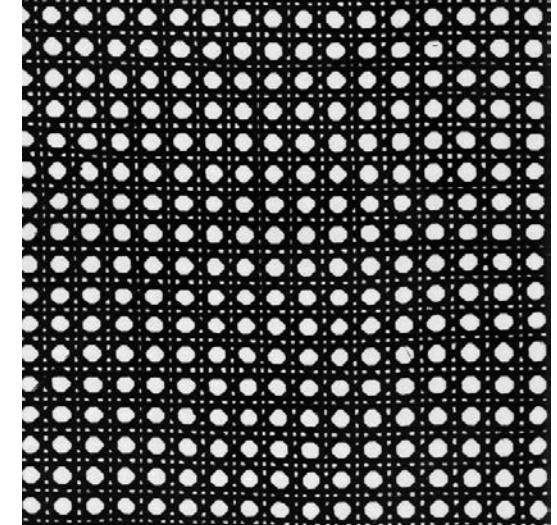
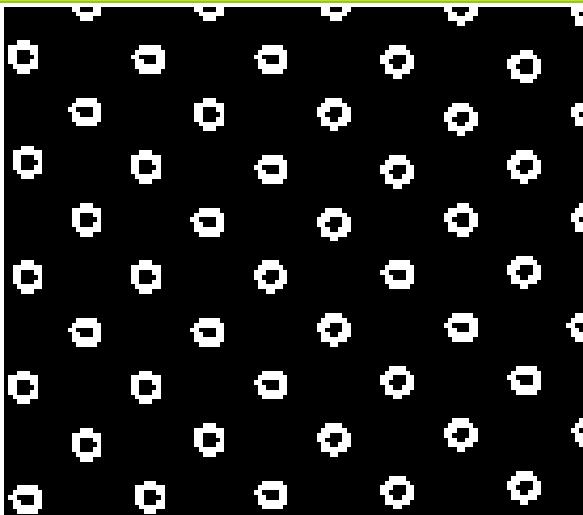
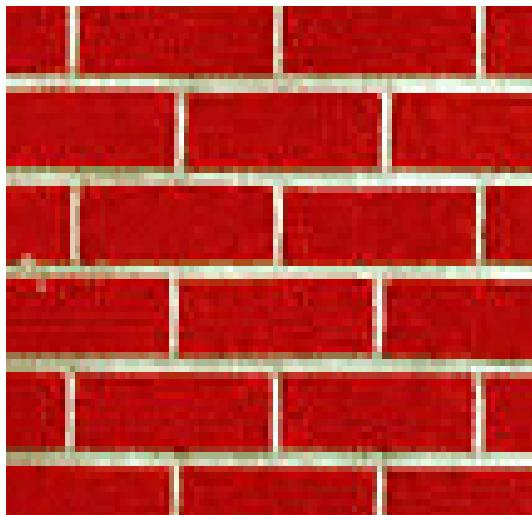


# Today: Texture

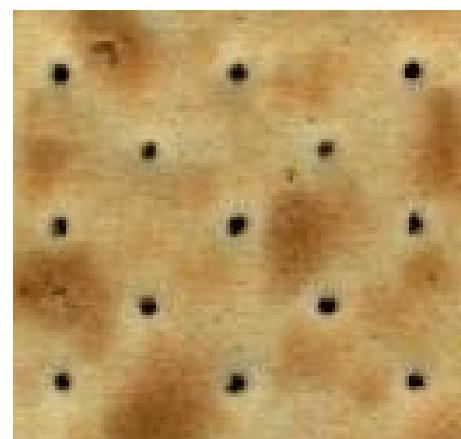
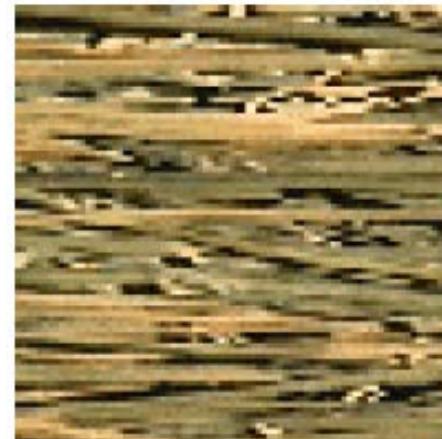


What defines a texture?

# Includes: more regular patterns



# Includes: more random patterns

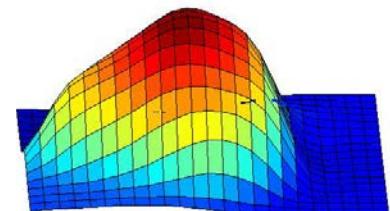
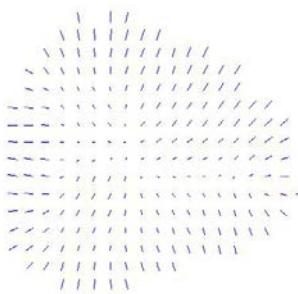
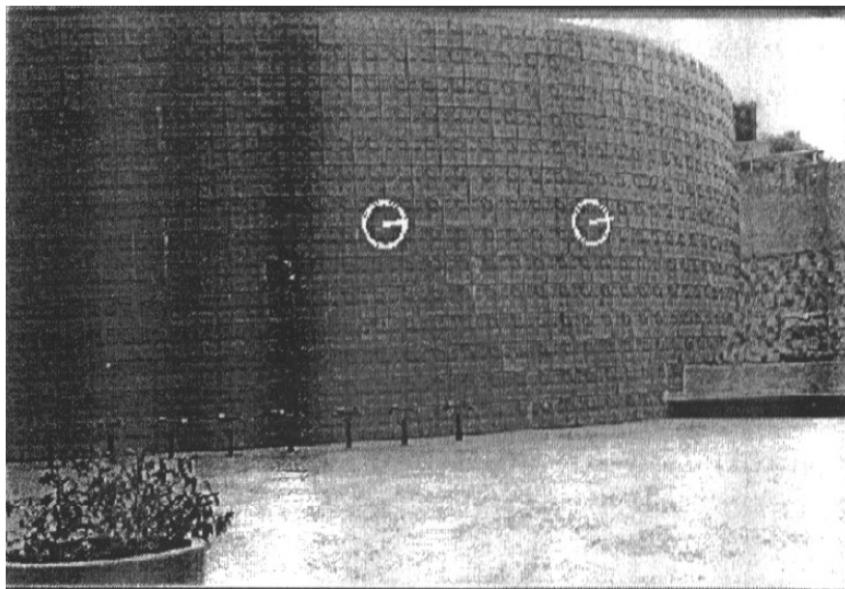


# Texture-related tasks

- **Shape from texture**
  - Estimate surface orientation or shape from image texture

# Shape from texture

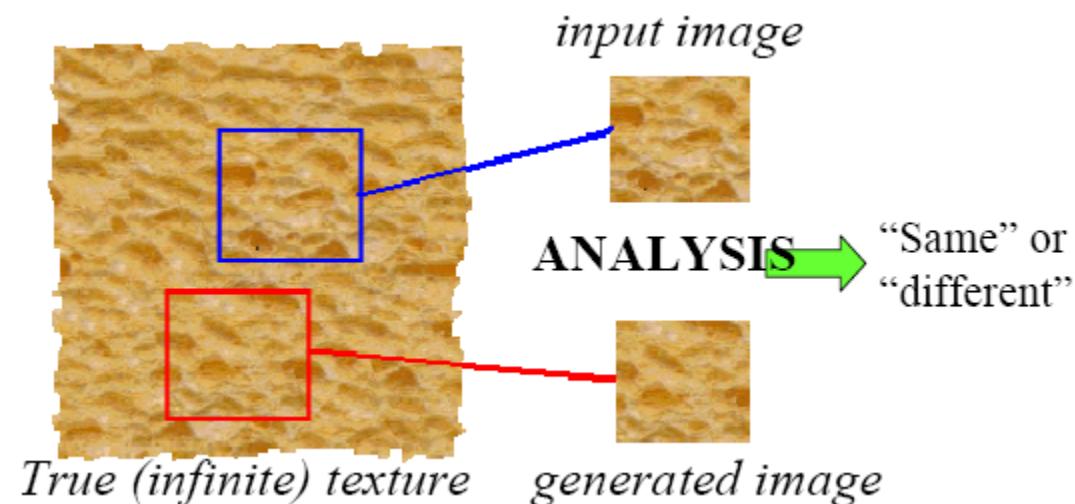
- Use deformation of texture from point to point to estimate surface shape



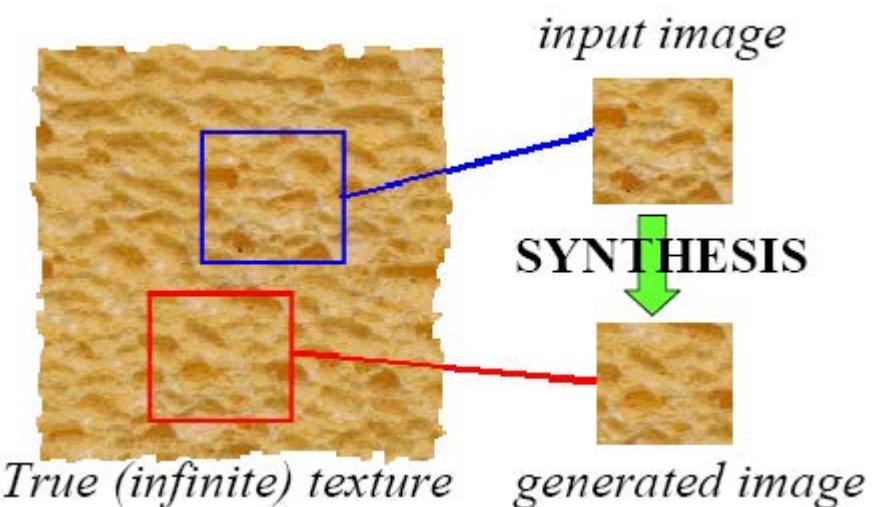
# Texture-related tasks

- **Shape from texture**
  - Estimate surface orientation or shape from image texture
- **Segmentation/classification** from texture cues
  - Analyze, represent texture
  - Group image regions with consistent texture
- **Synthesis**
  - Generate new texture patches/images given some examples

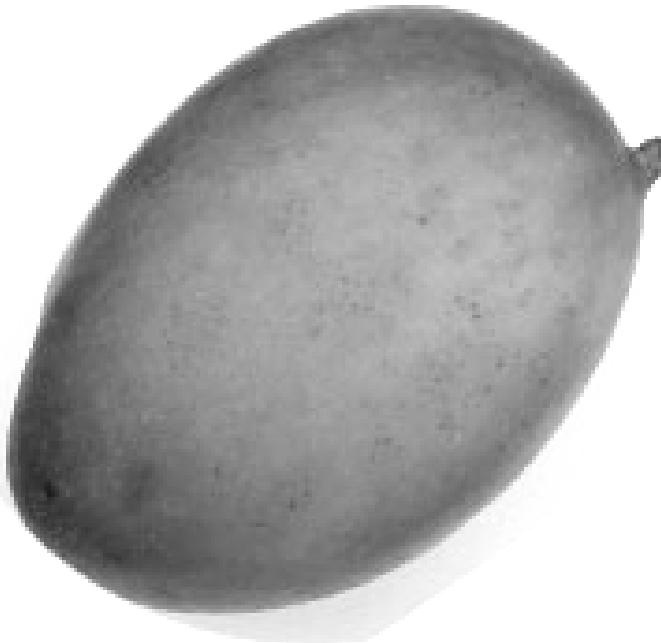
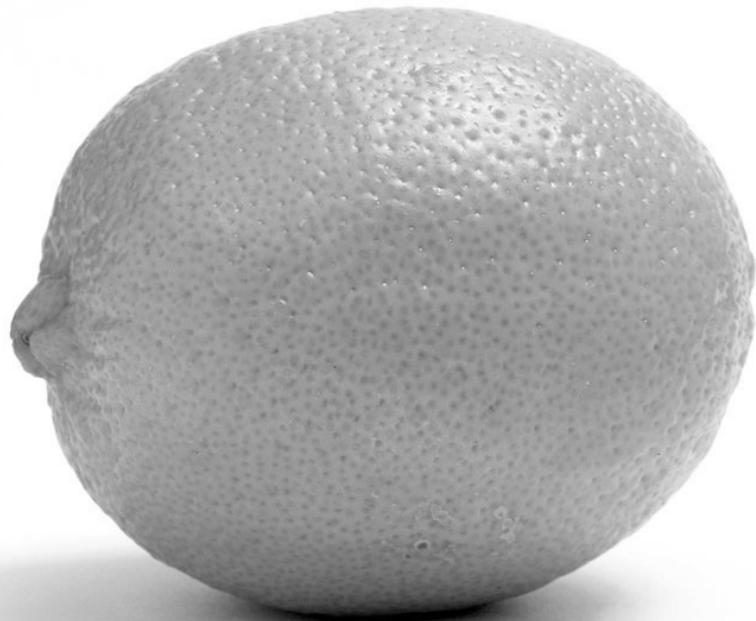
# Analysis vs. Synthesis



Why analyze  
texture?











What kind of response will we get with an edge detector for these images?



...and for this image?

# Why analyze texture?

Importance to perception:

- Often indicative of a material's properties
- Can be important appearance cue, especially if shape is similar across objects
- Aim to distinguish between shape, boundaries, and texture

Technically:

- Representation-wise, we want a feature one step above “building blocks” of filters, edges.

# Psychophysics of texture

- Some textures distinguishable with *preattentive* perception— without scrutiny, eye movements [Julesz 1975]

Same or different?

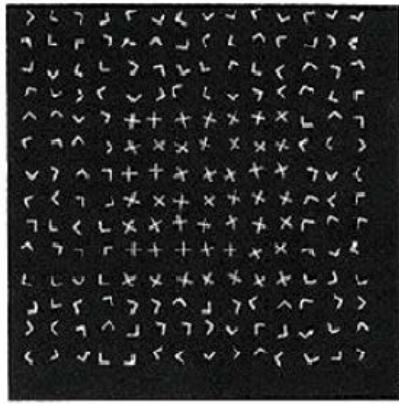
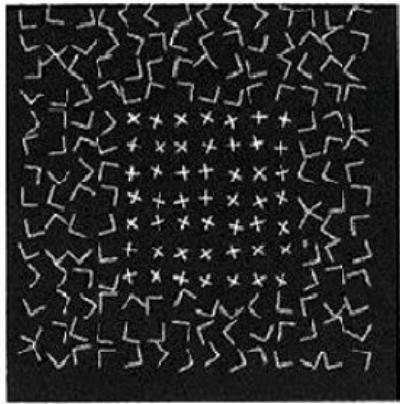
A 10x10 grid of black L-shaped blocks on a white background. The blocks are arranged in a staggered pattern, where each block is positioned such that its vertical stem is aligned with the horizontal bar of the block directly above it. This creates a continuous, winding path across the entire grid.

A 10x10 grid of black L-shaped blocks on a white background. The blocks are arranged in a staggered pattern, creating a tessellated effect. Each block is composed of three squares: one vertical column of two squares and one horizontal row of two squares, meeting at the top-right square of the column.

A 7x7 grid of black L-shaped blocks on a white background. Each block is composed of three squares: one vertical column of two squares and one horizontal square extending from the middle of the left side of the column. The blocks are arranged in a staggered pattern, where each row is offset by one square relative to the row above it. The grid starts with a block in the top-left corner and continues across seven rows and seven columns.

A 10x10 grid of black L-shaped blocks on a white background. The blocks are arranged in a staggered pattern, creating a tessellated effect. Each block is composed of three squares: one vertical column of two squares and one horizontal row of two squares extending from the middle of the column. The blocks are oriented such that they do not overlap, leaving a consistent gap between them.

# Capturing the local patterns with image measurements



[Bergen &  
Adelson,  
*Nature* 1988]

Scale of  
patterns  
influences  
discriminability

Size-tuned  
linear filters

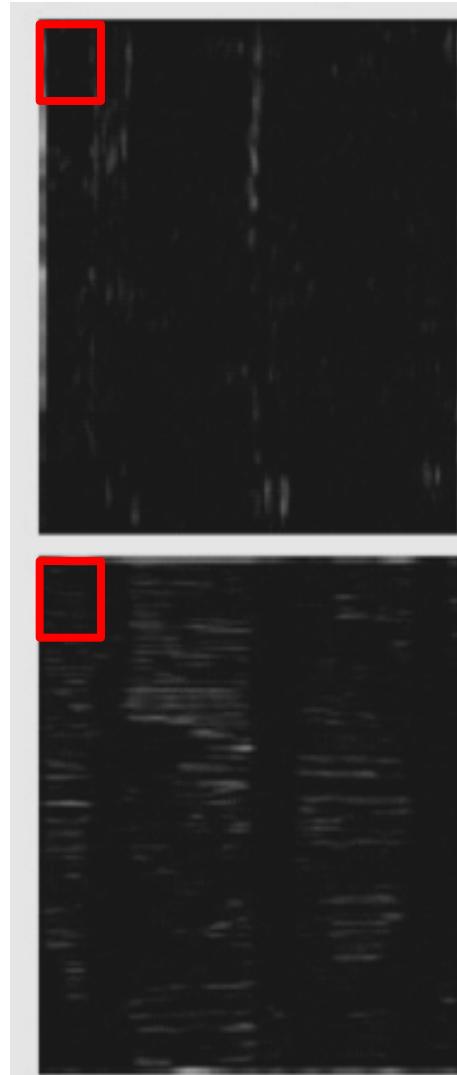
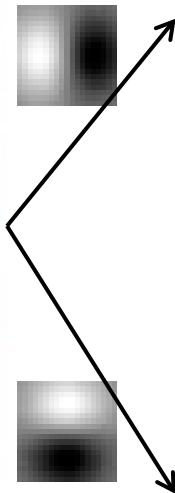
# Texture representation

- Textures are made up of repeated local patterns, so:
  - Find the patterns
    - Use filters that look like patterns (spots, bars, raw patches...)
    - Consider magnitude of response
  - Describe their statistics within each local window
    - Mean, standard deviation
    - Histogram
    - Histogram of “prototypical” feature occurrences

# Texture representation: example



original image



derivative filter  
responses, squared

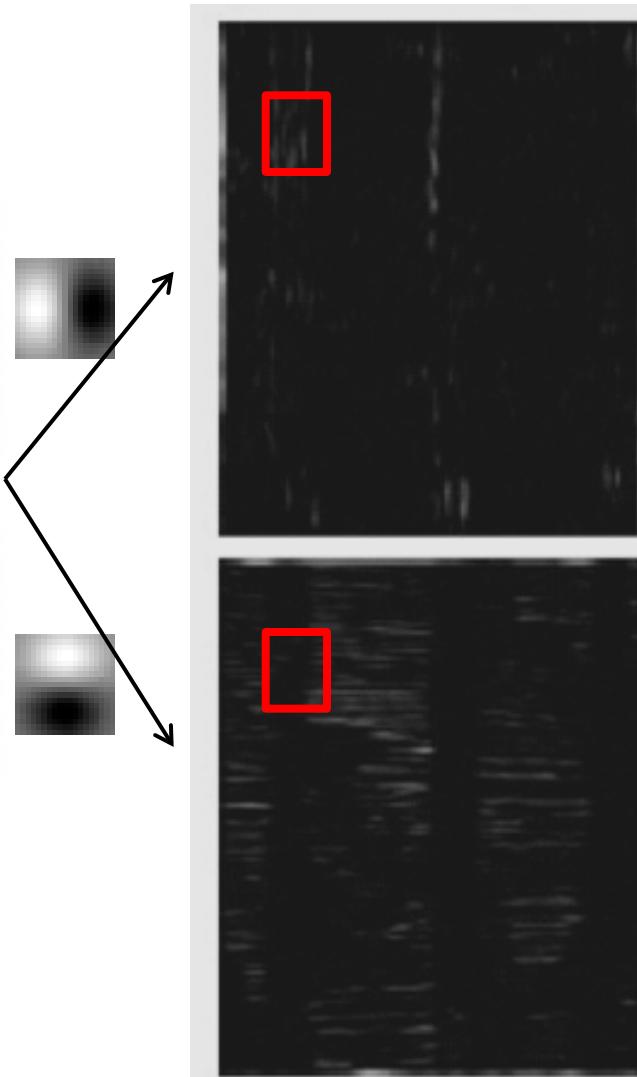
|         | <u>mean<br/><math>d/dx</math><br/>value</u> | <u>mean<br/><math>d/dy</math><br/>value</u> |
|---------|---|---|
| Win. #1 | 4   | 10  |
| ⋮       | ⋮   | ⋮   |
| ⋮       | ⋮   | ⋮   |
| ⋮       | ⋮   | ⋮   |

statistics to  
summarize patterns  
in small windows

# Texture representation: example



# original image



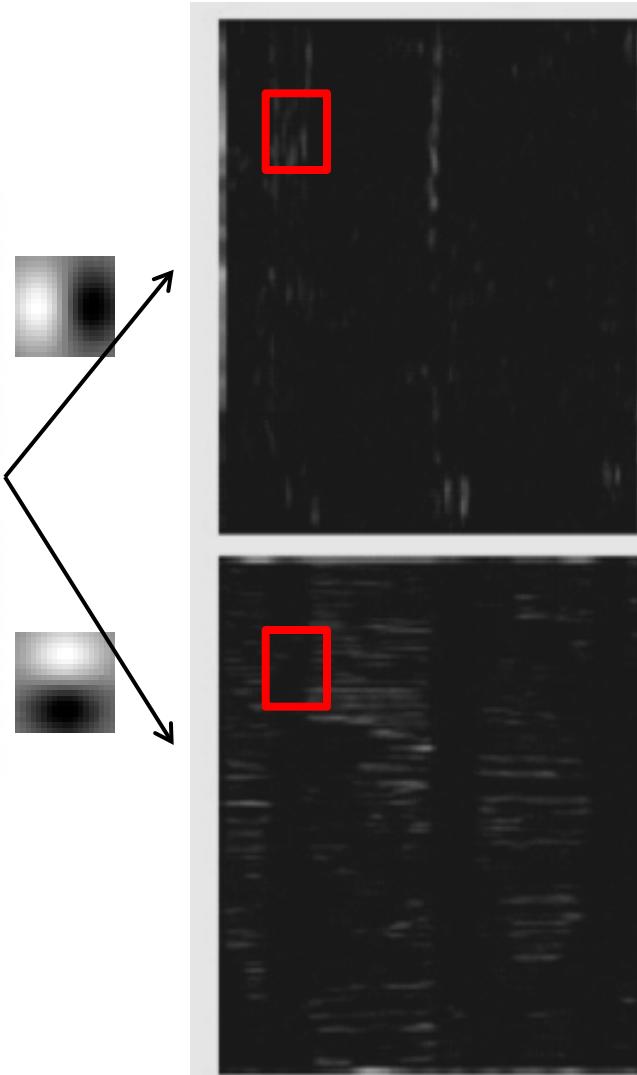
## derivative filter responses, squared

**statistics to  
summarize patterns  
in small windows**

# Texture representation: example



# original image



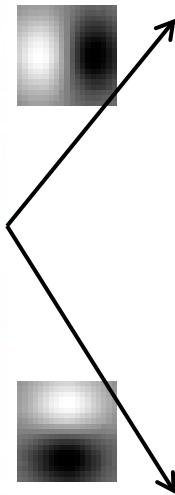
## derivative filter responses, squared

**statistics to  
summarize patterns  
in small windows**

# Texture representation: example



original image

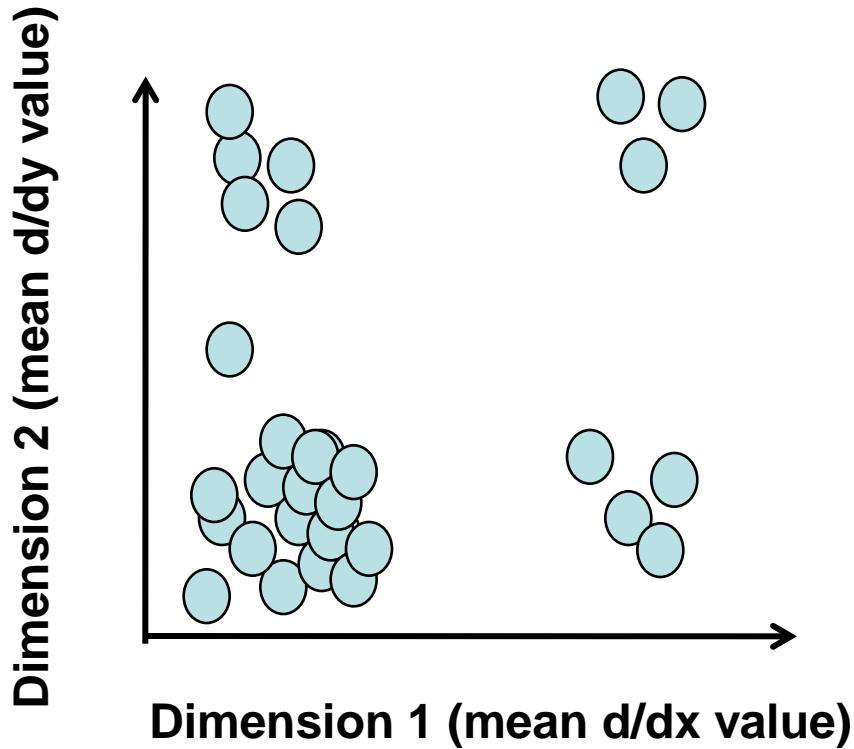


derivative filter  
responses, squared

|         | <u>mean<br/><math>d/dx</math><br/>value</u> | <u>mean<br/><math>d/dy</math><br/>value</u> |
|---------|---|---|
| Win. #1 | 4   | 10  |
| Win.#2  | 18  | 7   |
| :       |   |   |
| Win.#9  | 20  | 20  |
|         |   |   |
|         | ⋮   |   |

statistics to  
summarize patterns  
in small windows

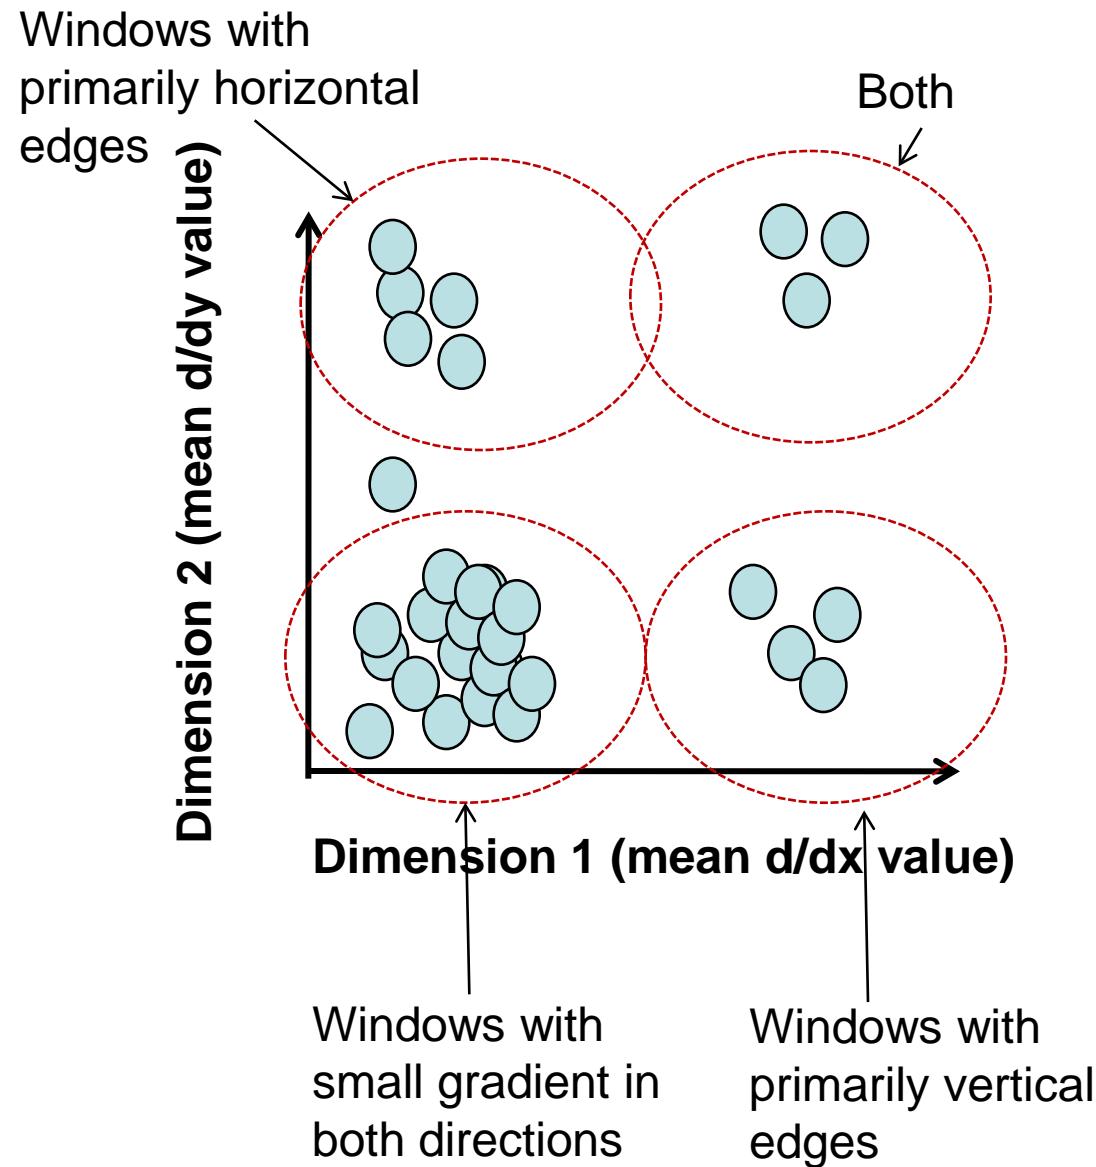
# Texture representation: example



|         | <u>mean<br/><math>d/dx</math><br/>value</u> | <u>mean<br/><math>d/dy</math><br/>value</u> |
|---------|---|---|
| Win. #1 | 4   | 10  |
| Win. #2 | 18  | 7   |
| ⋮       |   |   |
| Win. #9 | 20  | 20  |
| ⋮       |   |   |

**statistics to  
summarize patterns  
in small windows**

# Texture representation: example



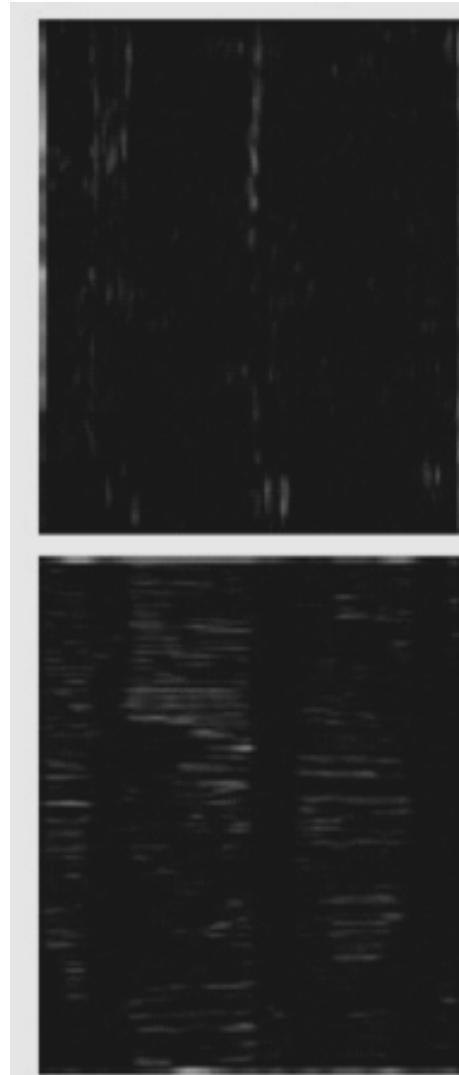
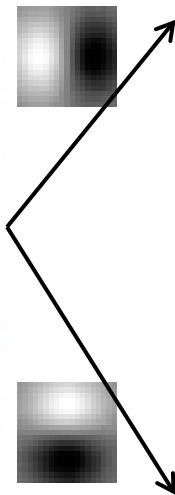
|         | <u>mean</u><br><u><math>d/dx</math></u><br><u>value</u> | <u>mean</u><br><u><math>d/dy</math></u><br><u>value</u> |
|---------|---|---|
| Win. #1 | 4   | 10  |
| Win.#2  | 18  | 7   |
| ⋮       | ⋮   | ⋮   |
| Win.#9  | 20  | 20  |
| ⋮       | ⋮   | ⋮   |

**statistics to summarize patterns in small windows**

# Texture representation: example



original image

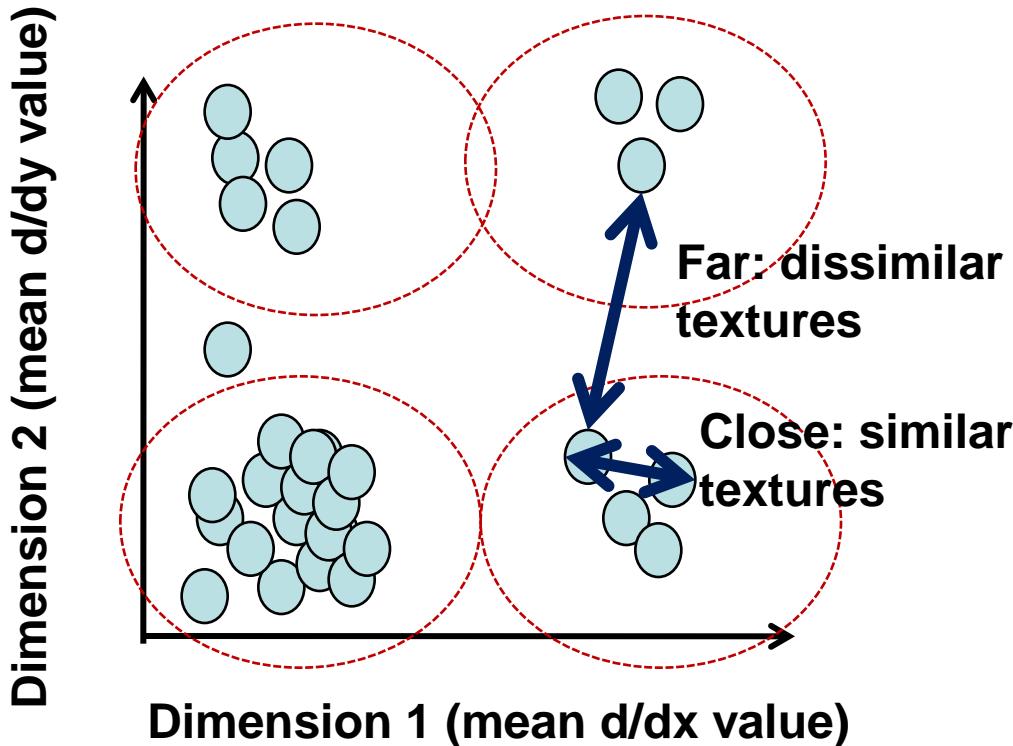


derivative filter  
responses, squared



visualization of the  
assignment to  
texture “types”

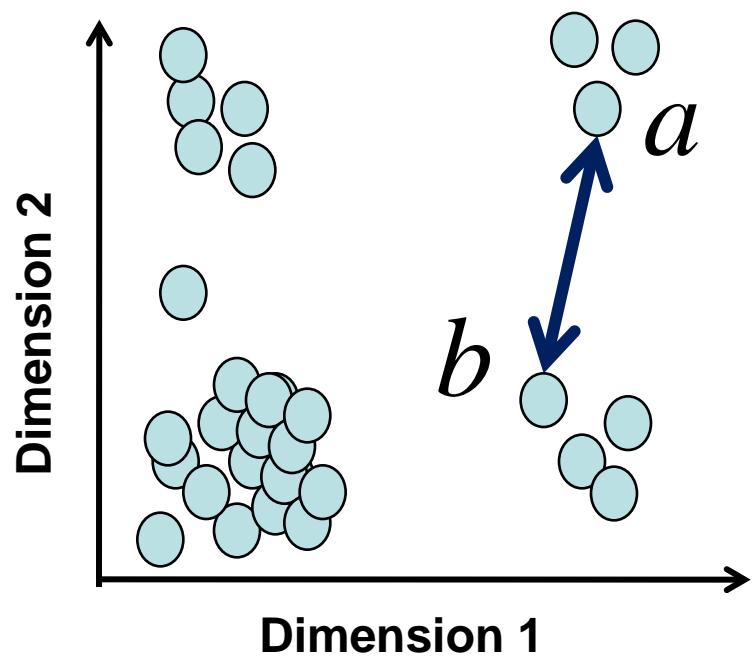
# Texture representation: example



|         | <u>mean<br/><math>d/dx</math><br/>value</u> | <u>mean<br/><math>d/dy</math><br/>value</u> |
|---------|---|---|
| Win. #1 | 4   | 10  |
| Win.#2  | 18  | 7   |
| :       |   |   |
| Win.#9  | 20  | 20  |
| ⋮       |   |   |

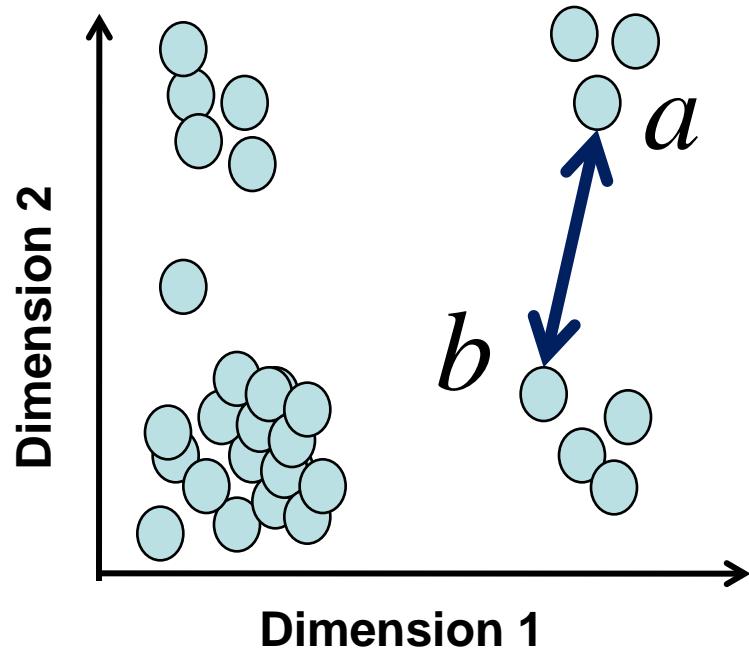
statistics to  
summarize patterns  
in small windows

# Texture representation: example

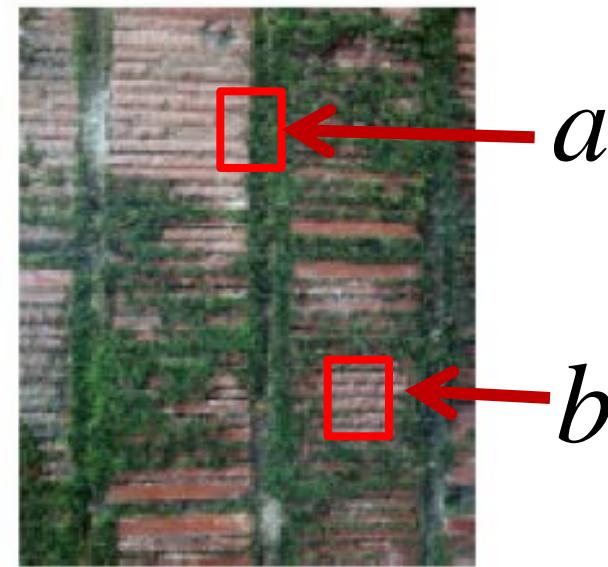


$$D(a,b) = \sqrt{(a_1 - b_1)^2 + (a_2 - b_2)^2}$$

# Texture representation: example

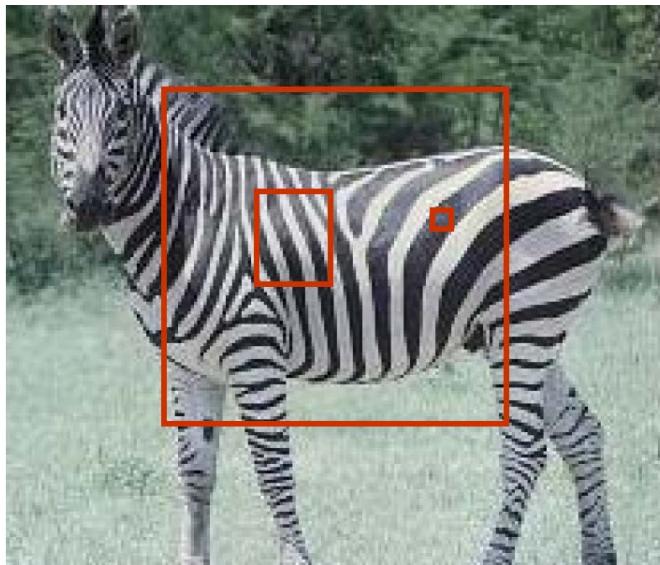


Distance reveals how dissimilar texture from window *a* is from texture in window *b*.



# Texture representation: window scale

- We're assuming we know the relevant window size for which we collect these statistics.

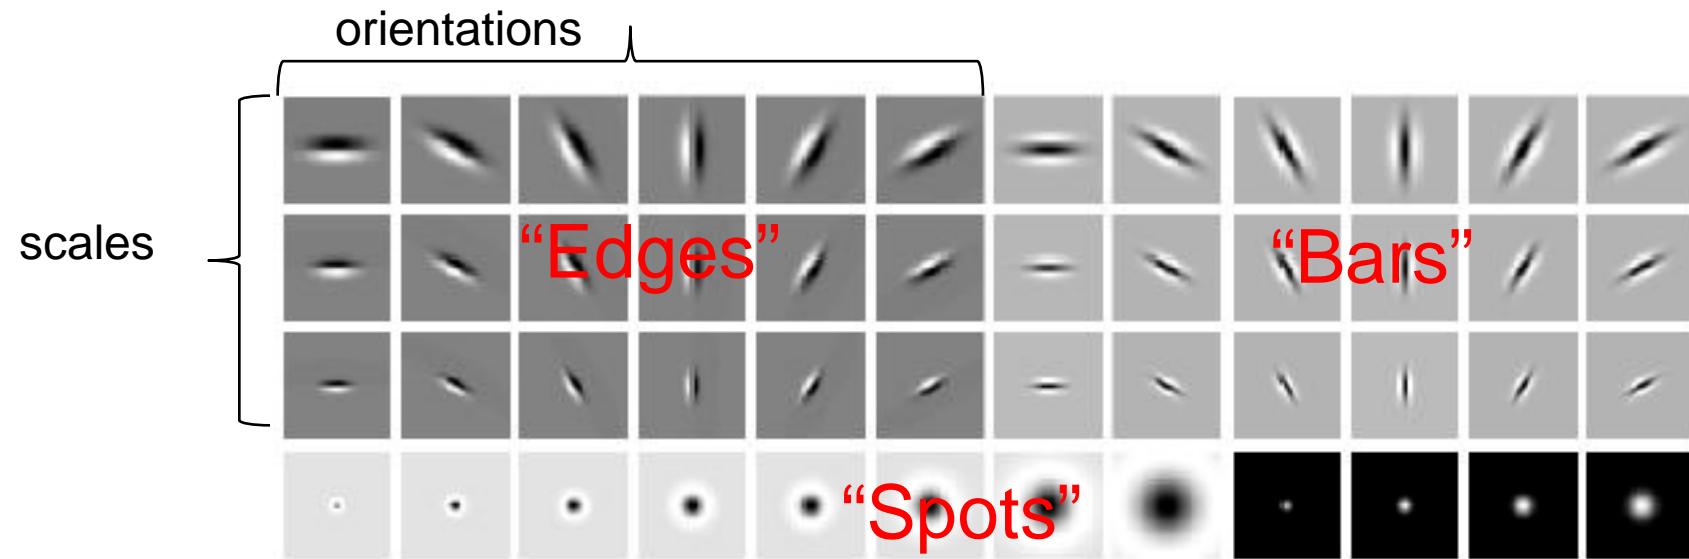


Possible to perform scale selection by looking for window scale where texture description not changing.

# Filter banks

- Our previous example used two filters, and resulted in a 2-dimensional feature vector to describe texture in a window.
  - x and y derivatives revealed something about local structure.
- We can generalize to apply a collection of multiple ( $d$ ) filters: a “filter bank”
- Then our feature vectors will be  $d$ -dimensional.
  - still can think of nearness, farness in feature space

# Filter banks



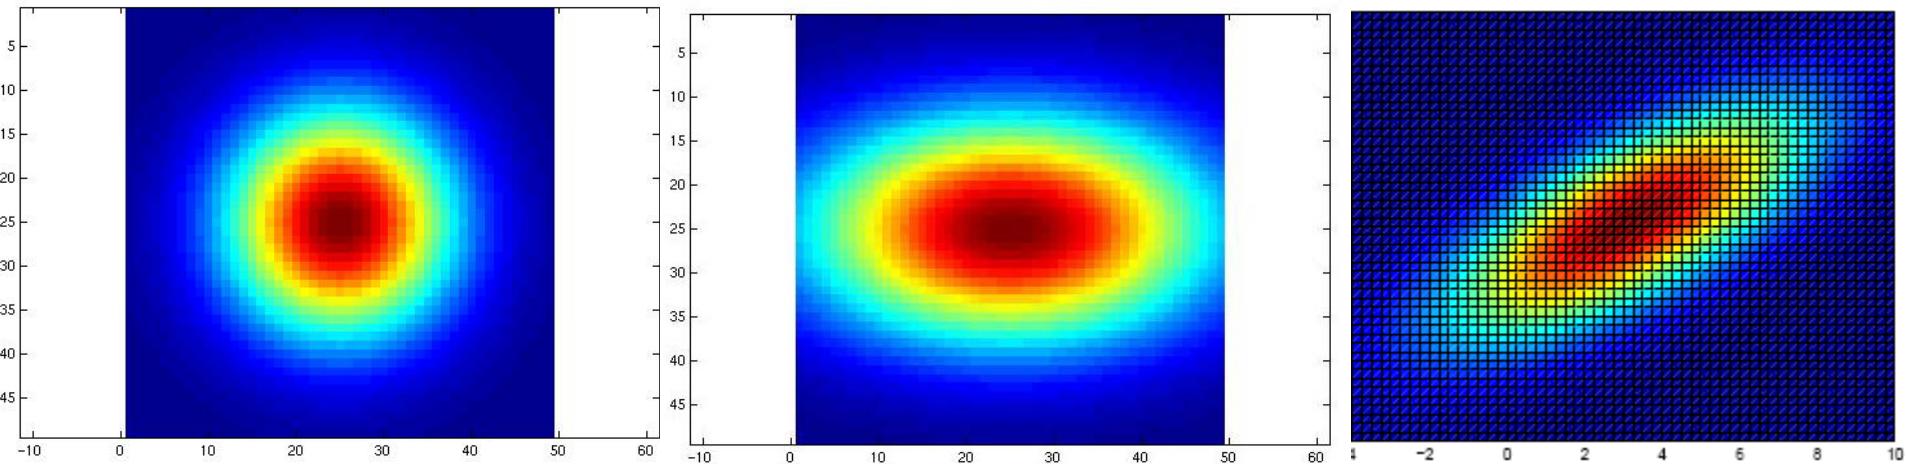
- What filters to put in the bank?
  - Typically we want a combination of scales and orientations, different types of patterns.

Matlab code available for these examples:

<http://www.robots.ox.ac.uk/~vgg/research/texclass/filters.html>

# Multivariate Gaussian

$$p(x; \mu, \Sigma) = \frac{1}{(2\pi)^{n/2} |\Sigma|^{1/2}} \exp \left( -\frac{1}{2} (x - \mu)^T \Sigma^{-1} (x - \mu) \right)$$



$$\Sigma = \begin{bmatrix} 9 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 16 & 0 \\ 0 & 9 \end{bmatrix}$$

$$\Sigma = \begin{bmatrix} 10 & 5 \\ 5 & 5 \end{bmatrix}$$

# Filter bank

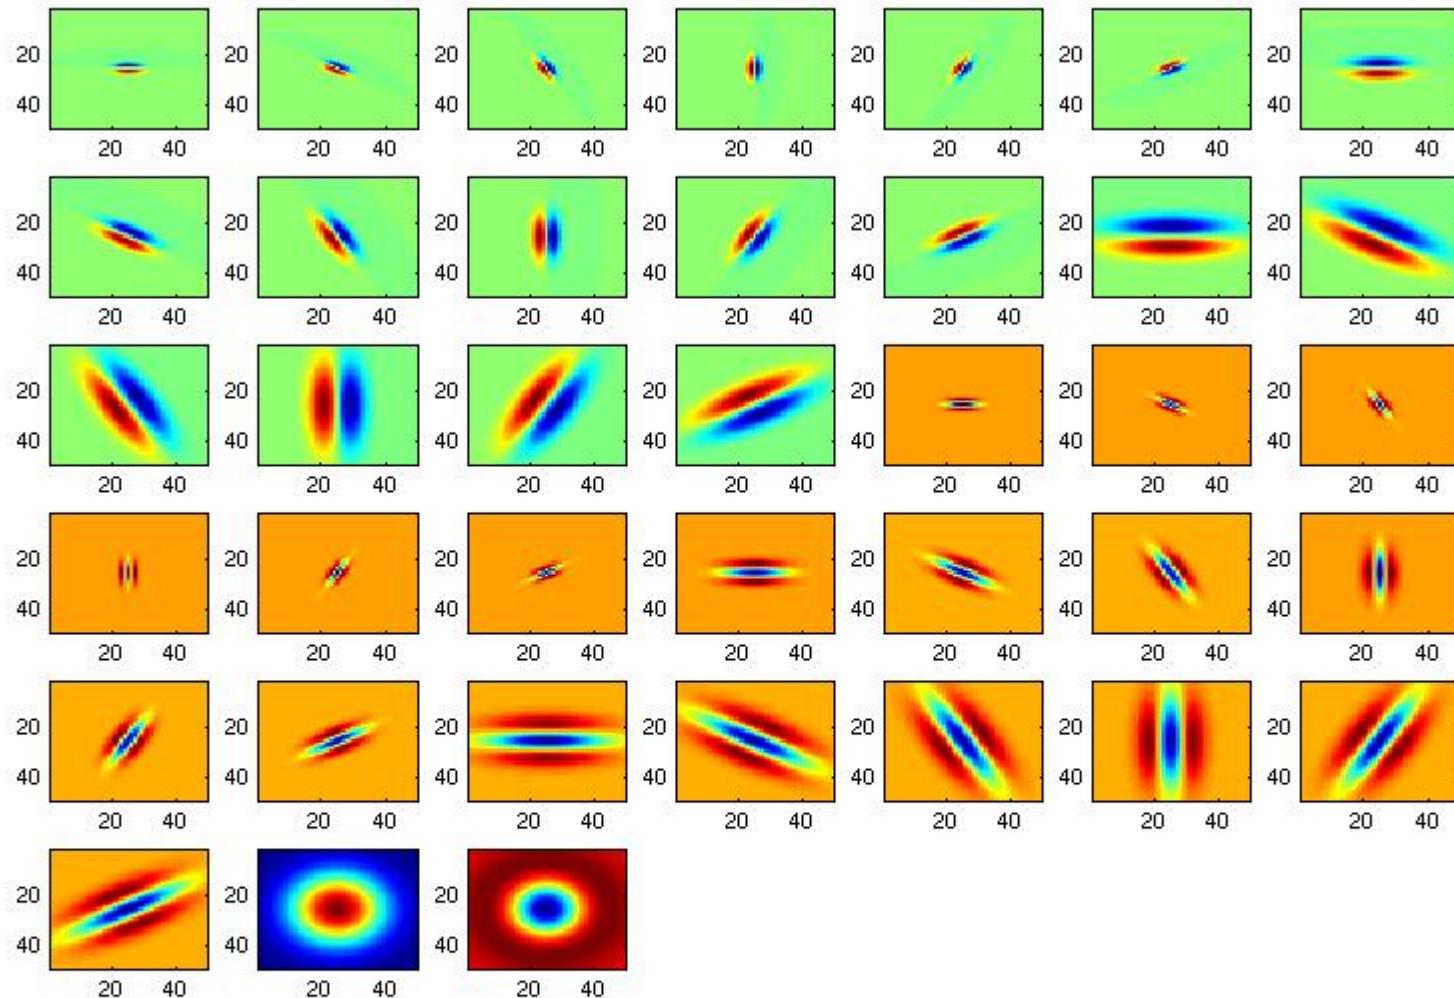
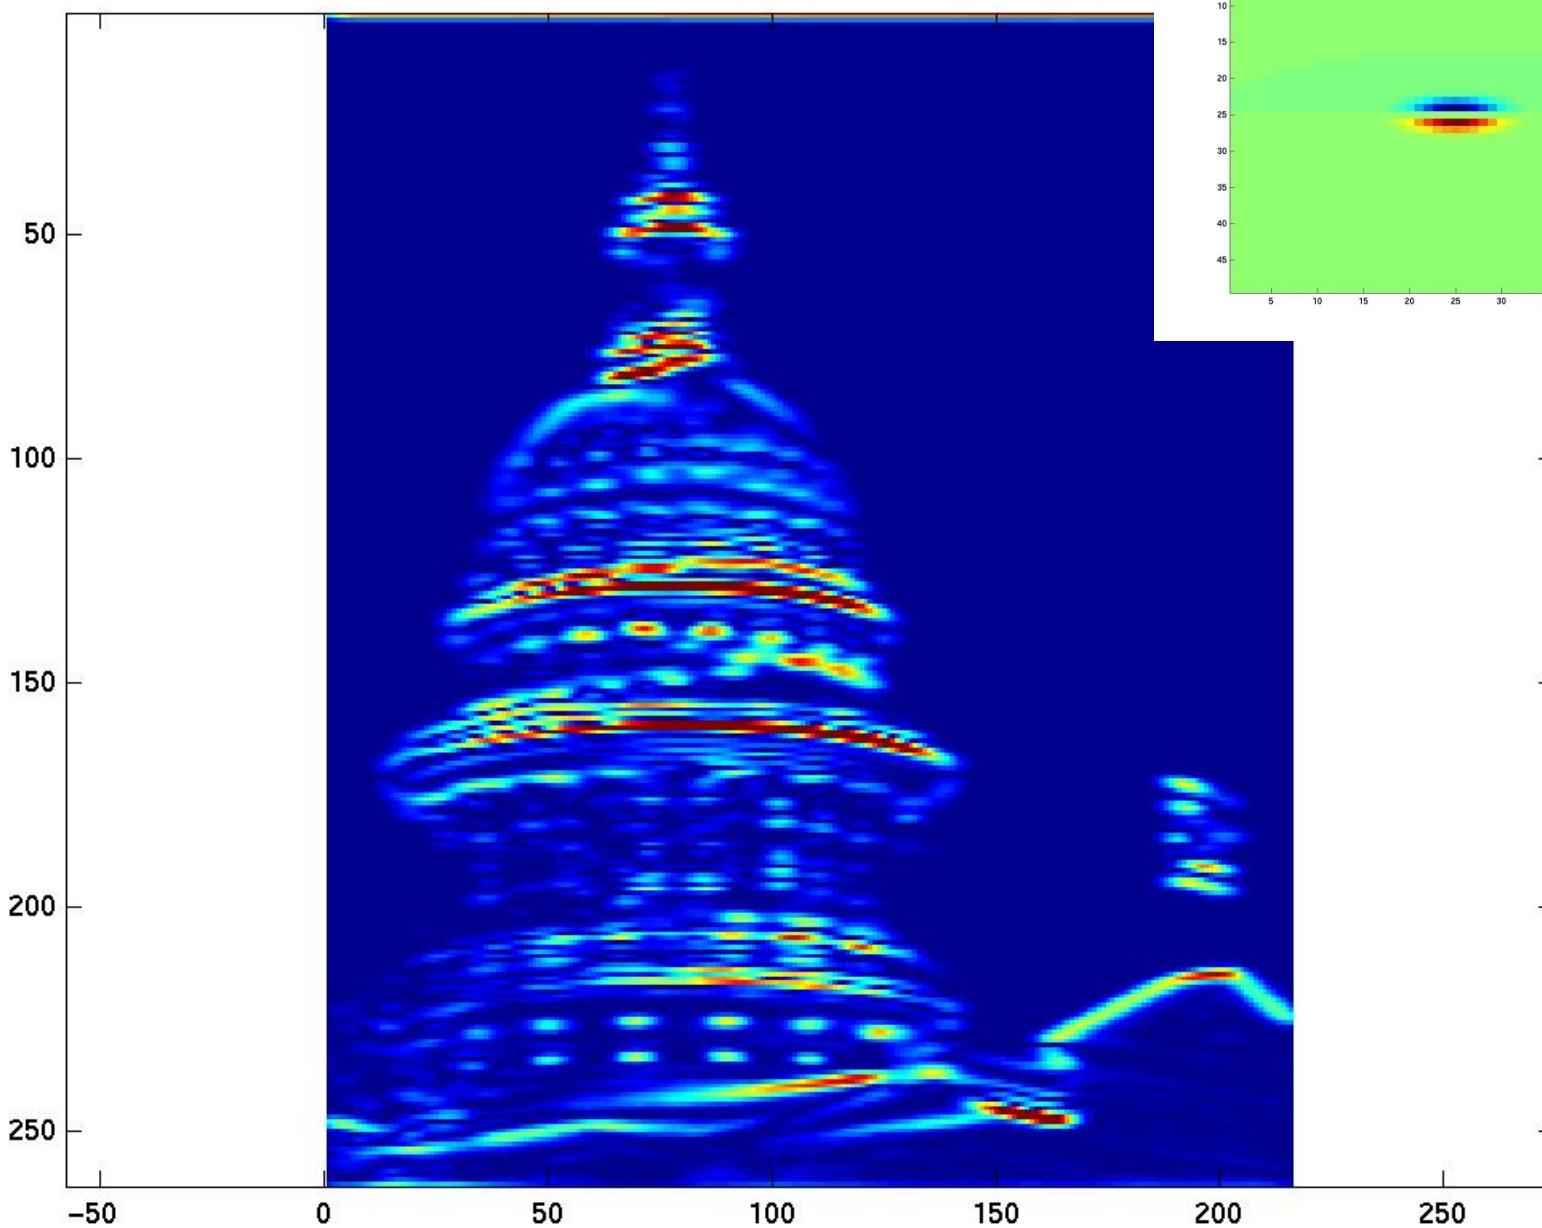
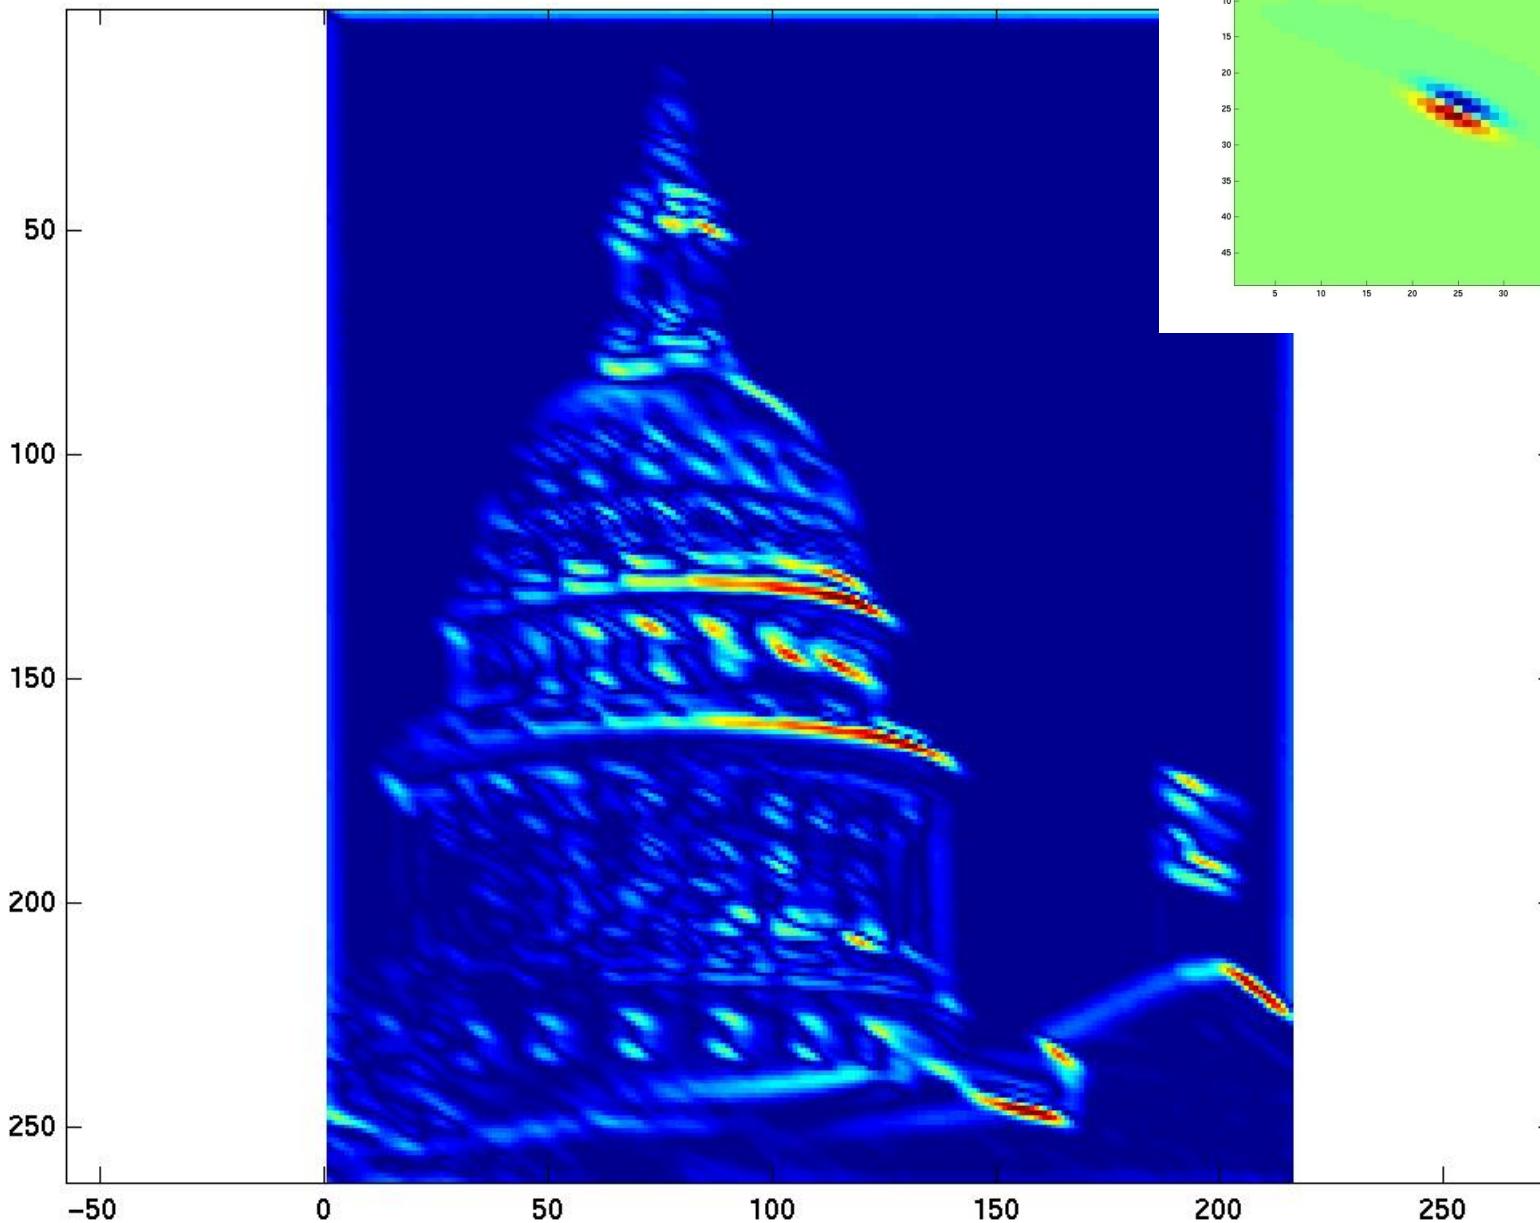


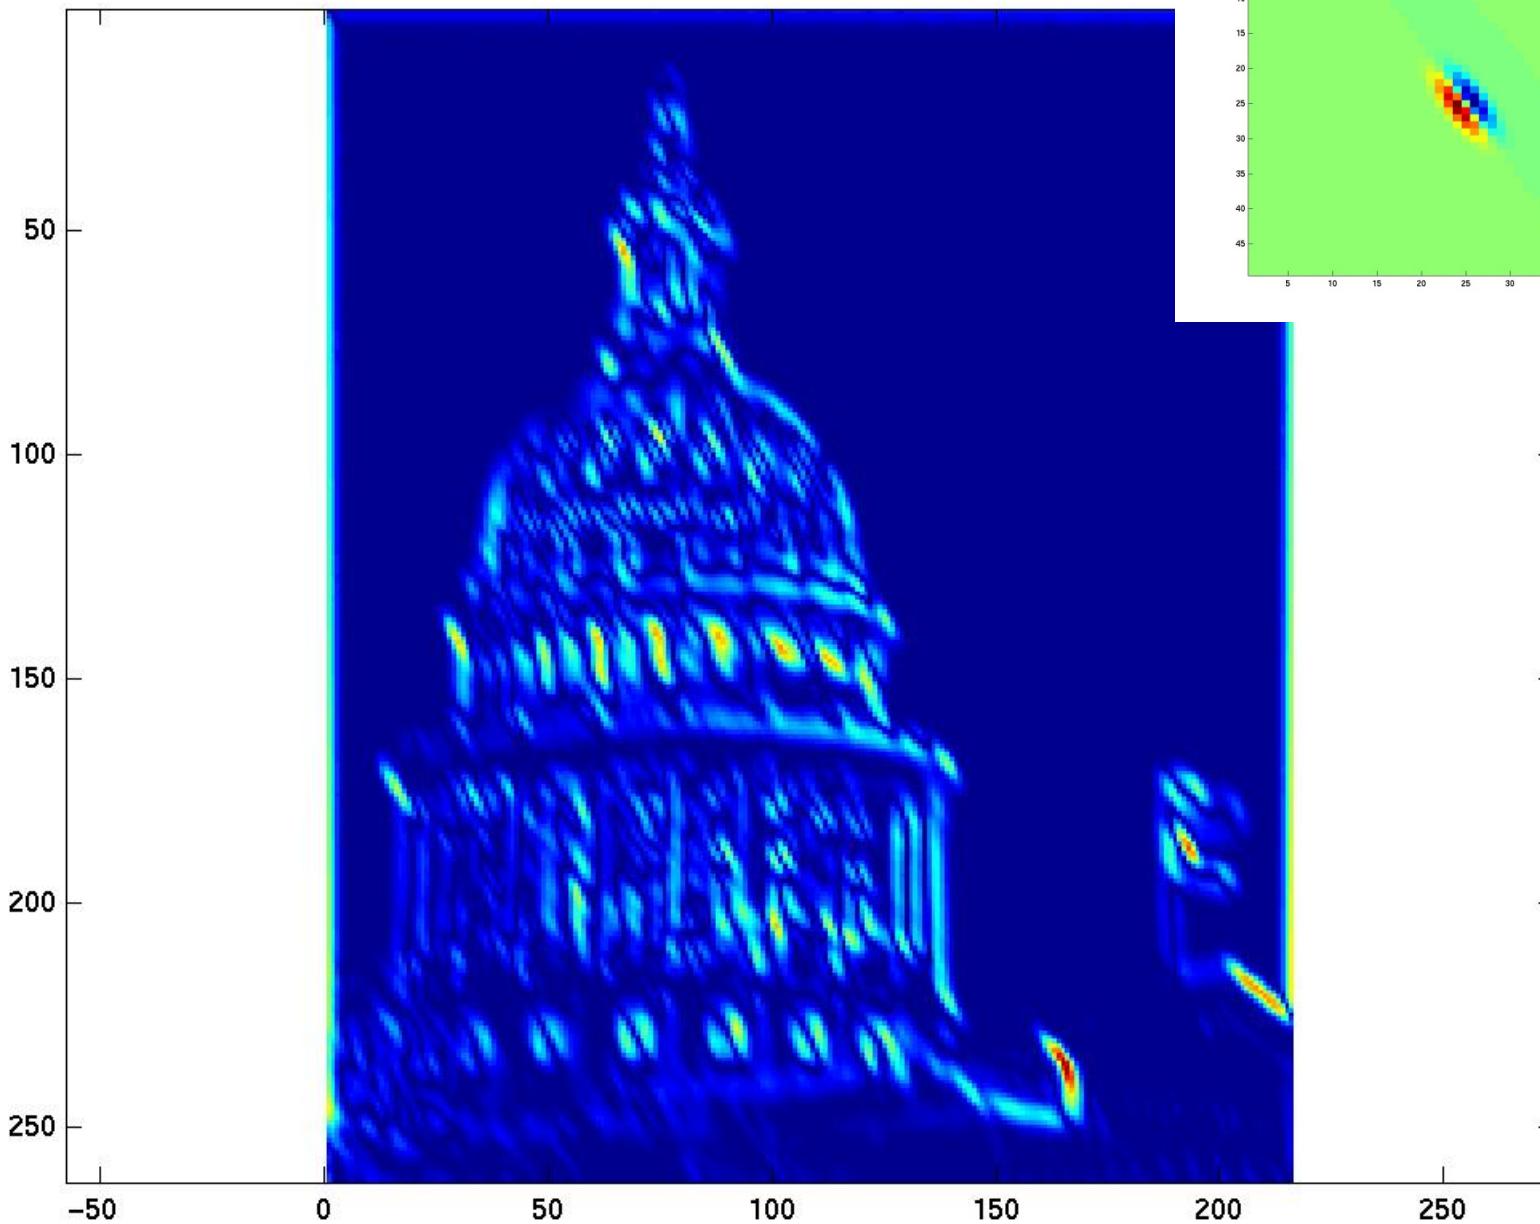
 Image from <http://www.texasexplorer.com/austincap2.jpg>

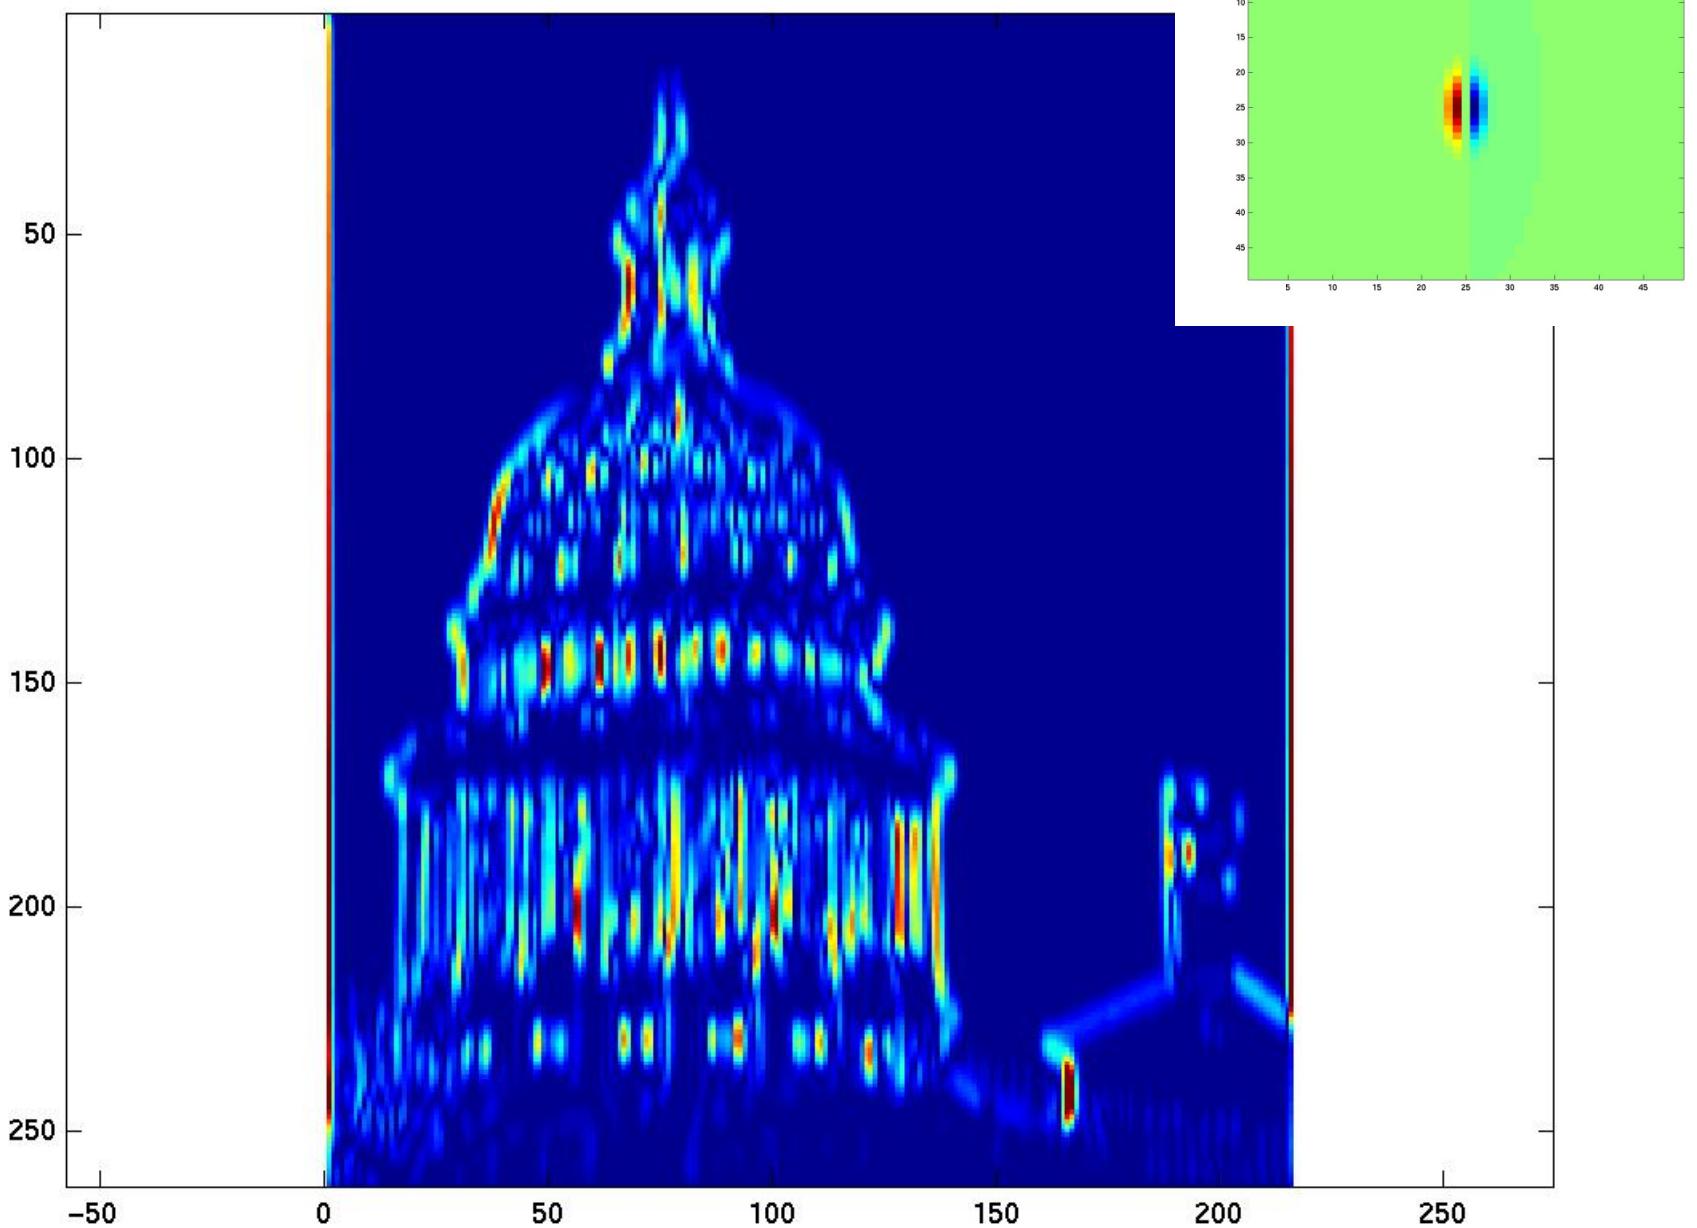


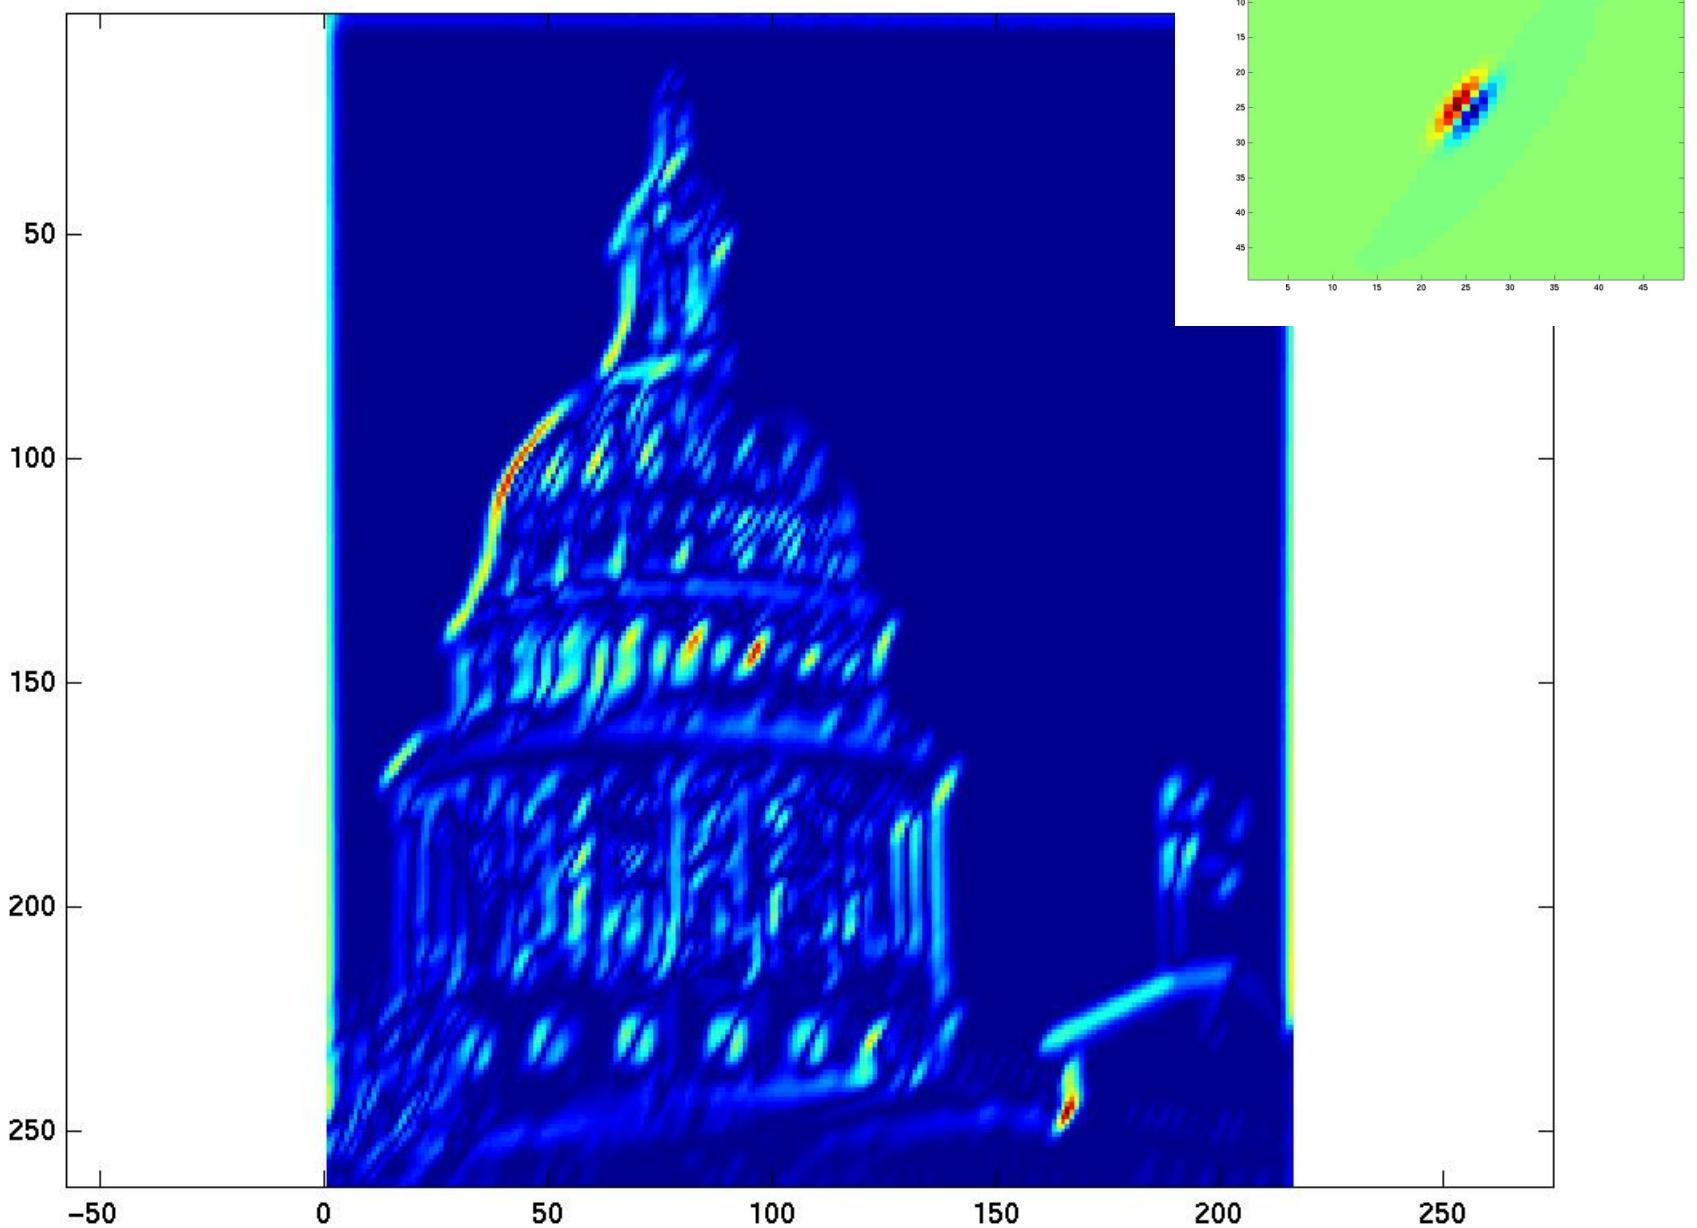
 Kristen Grauman

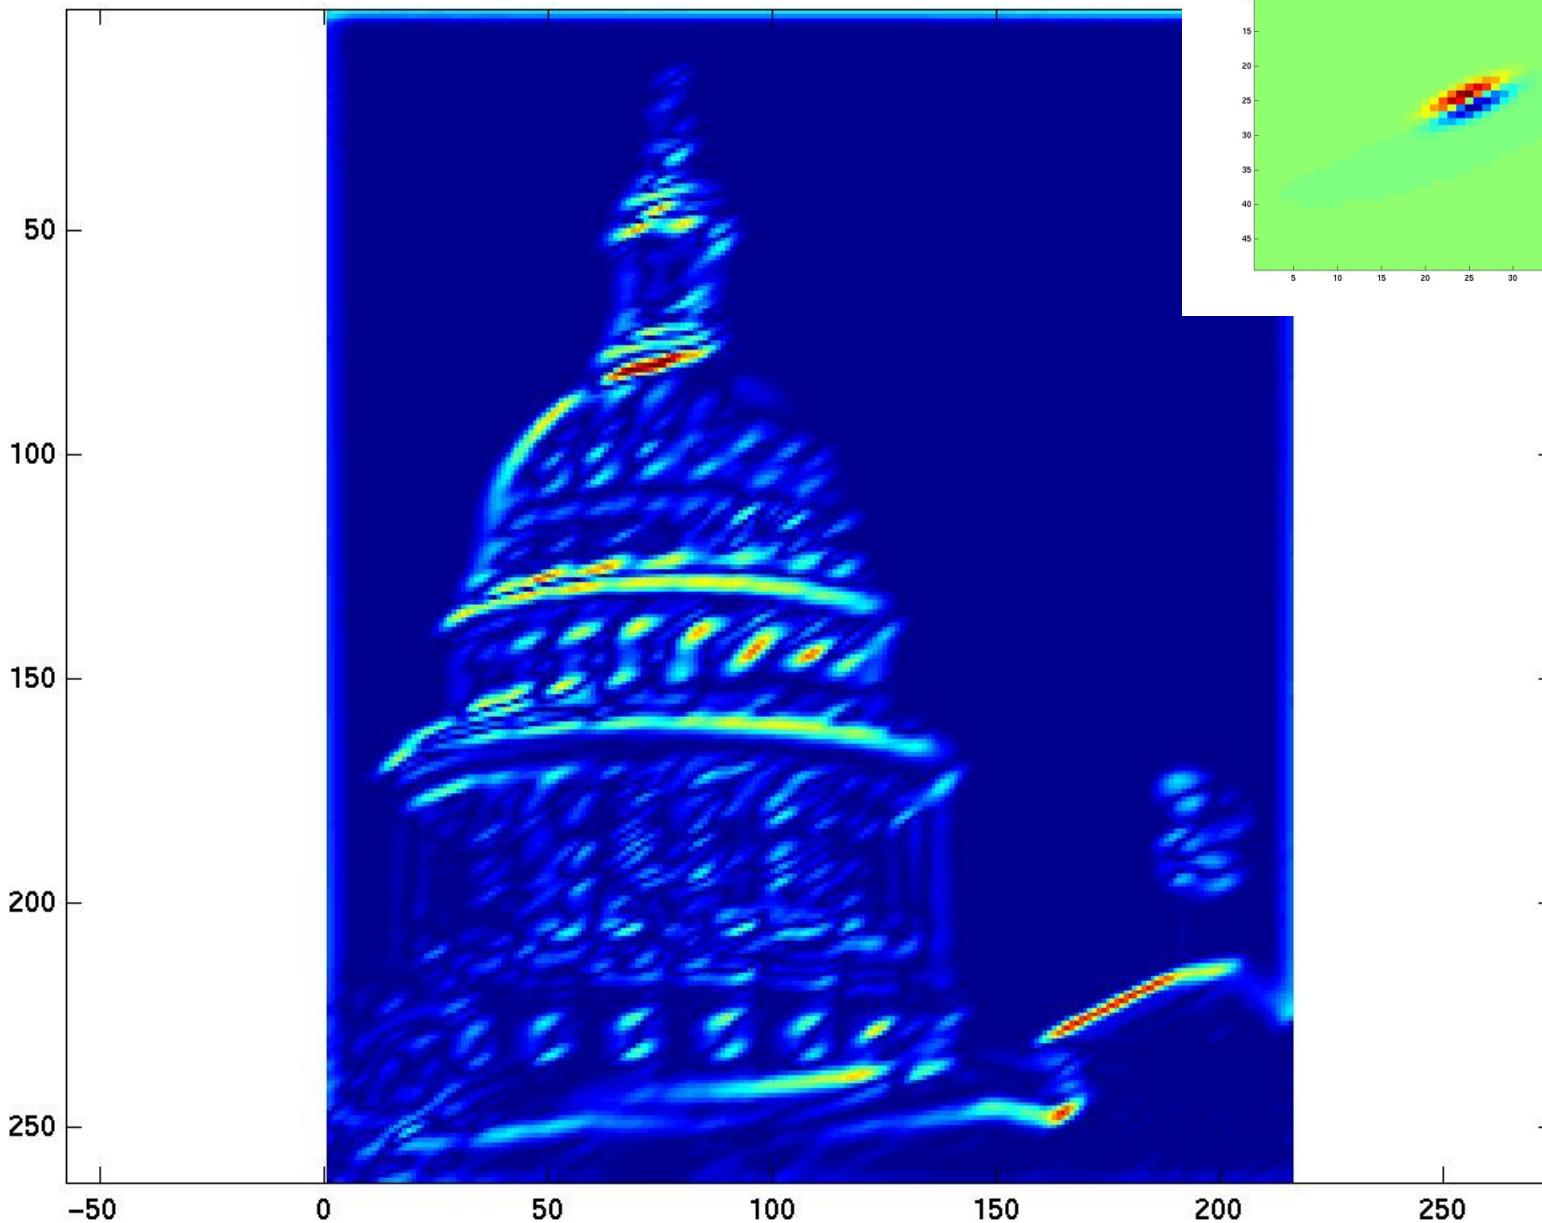


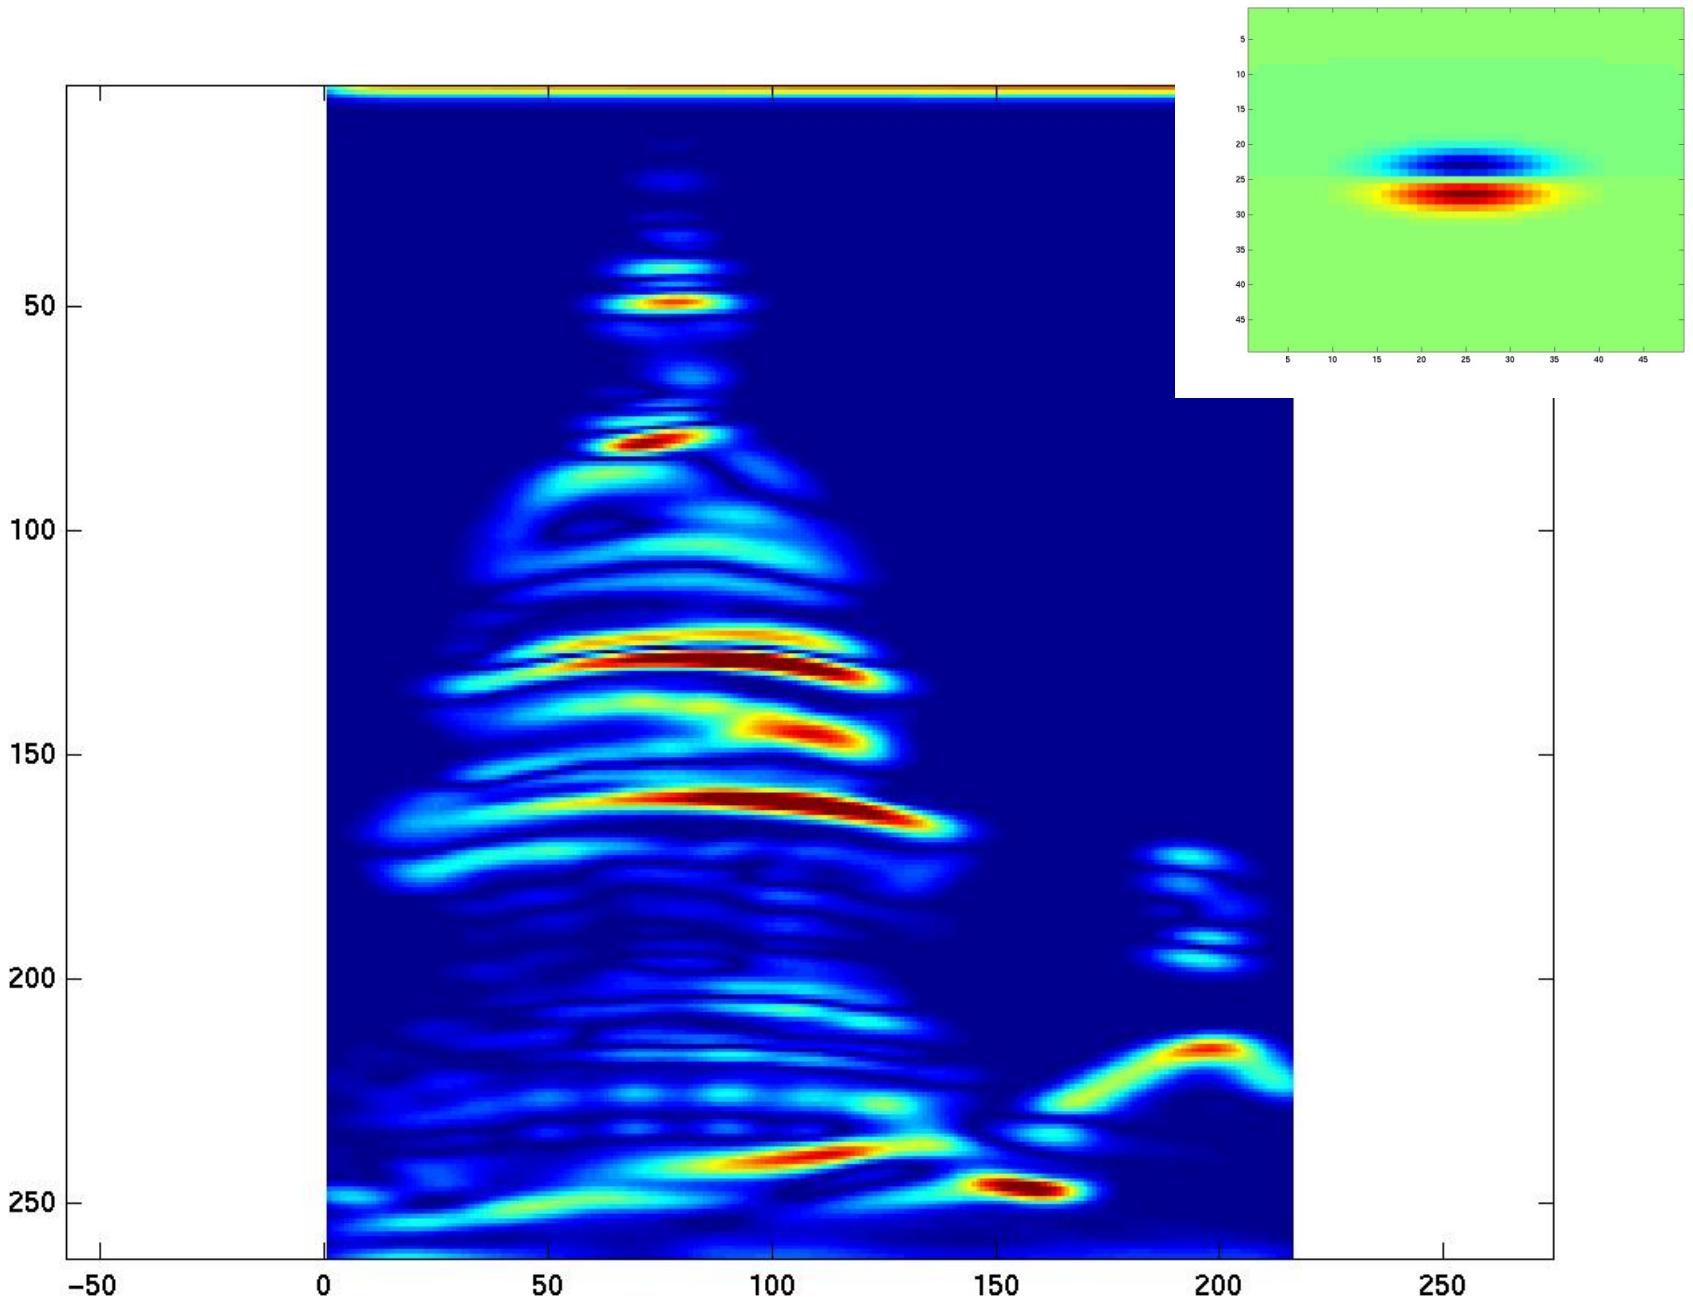


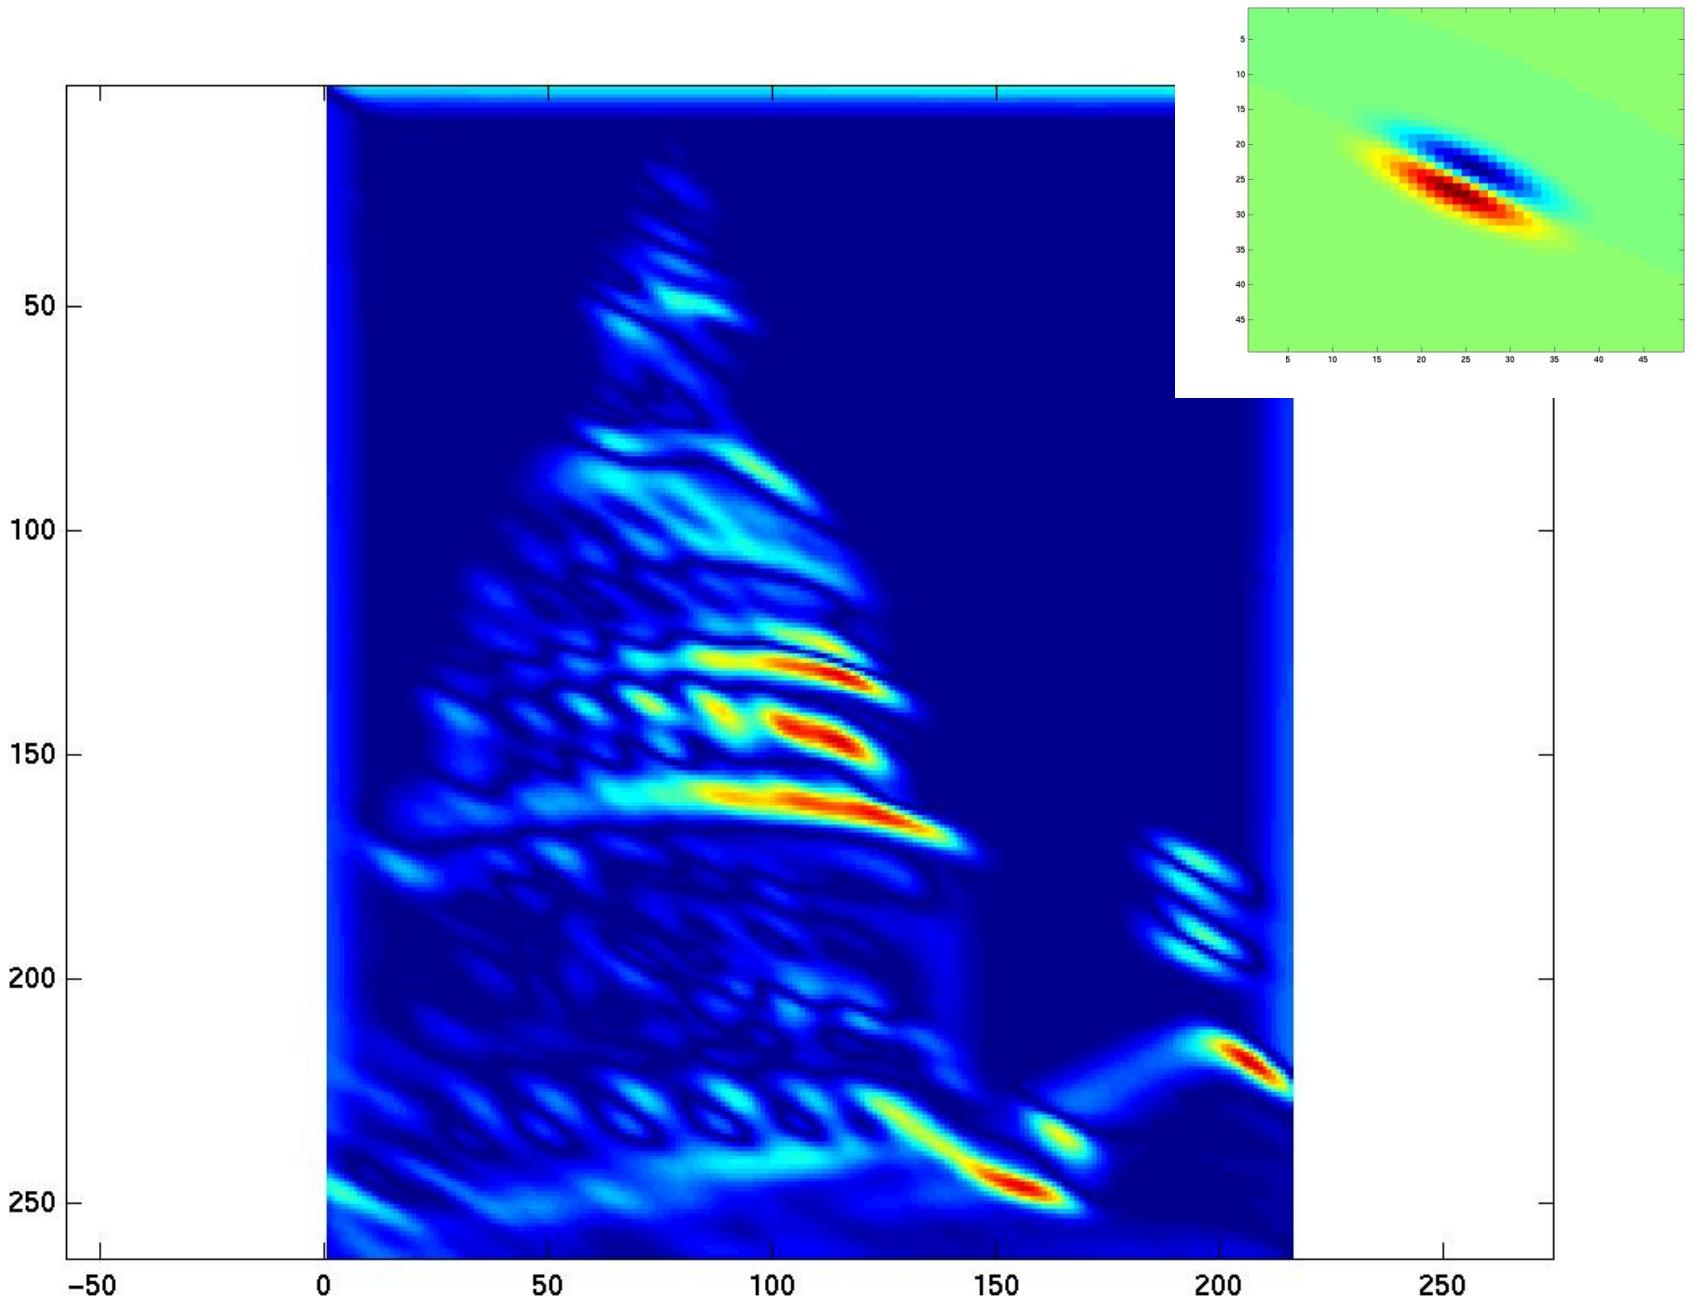


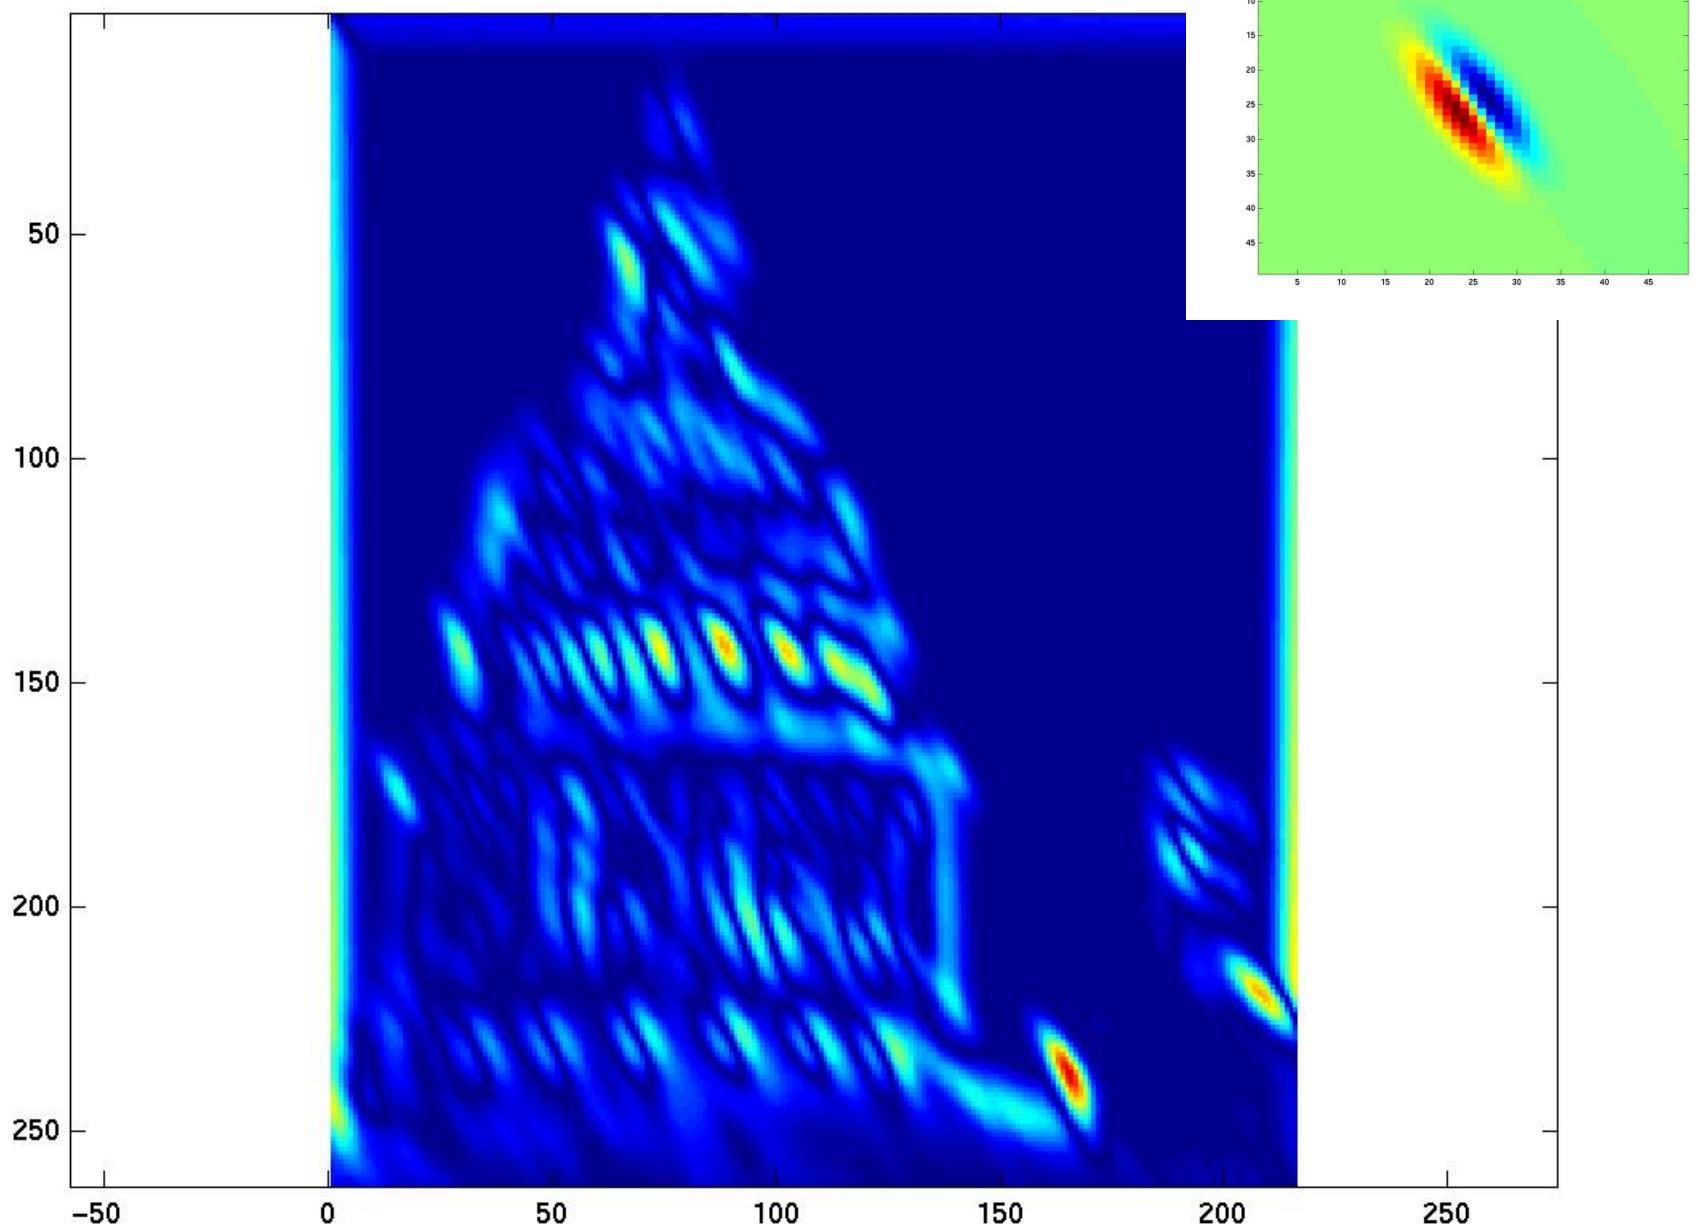


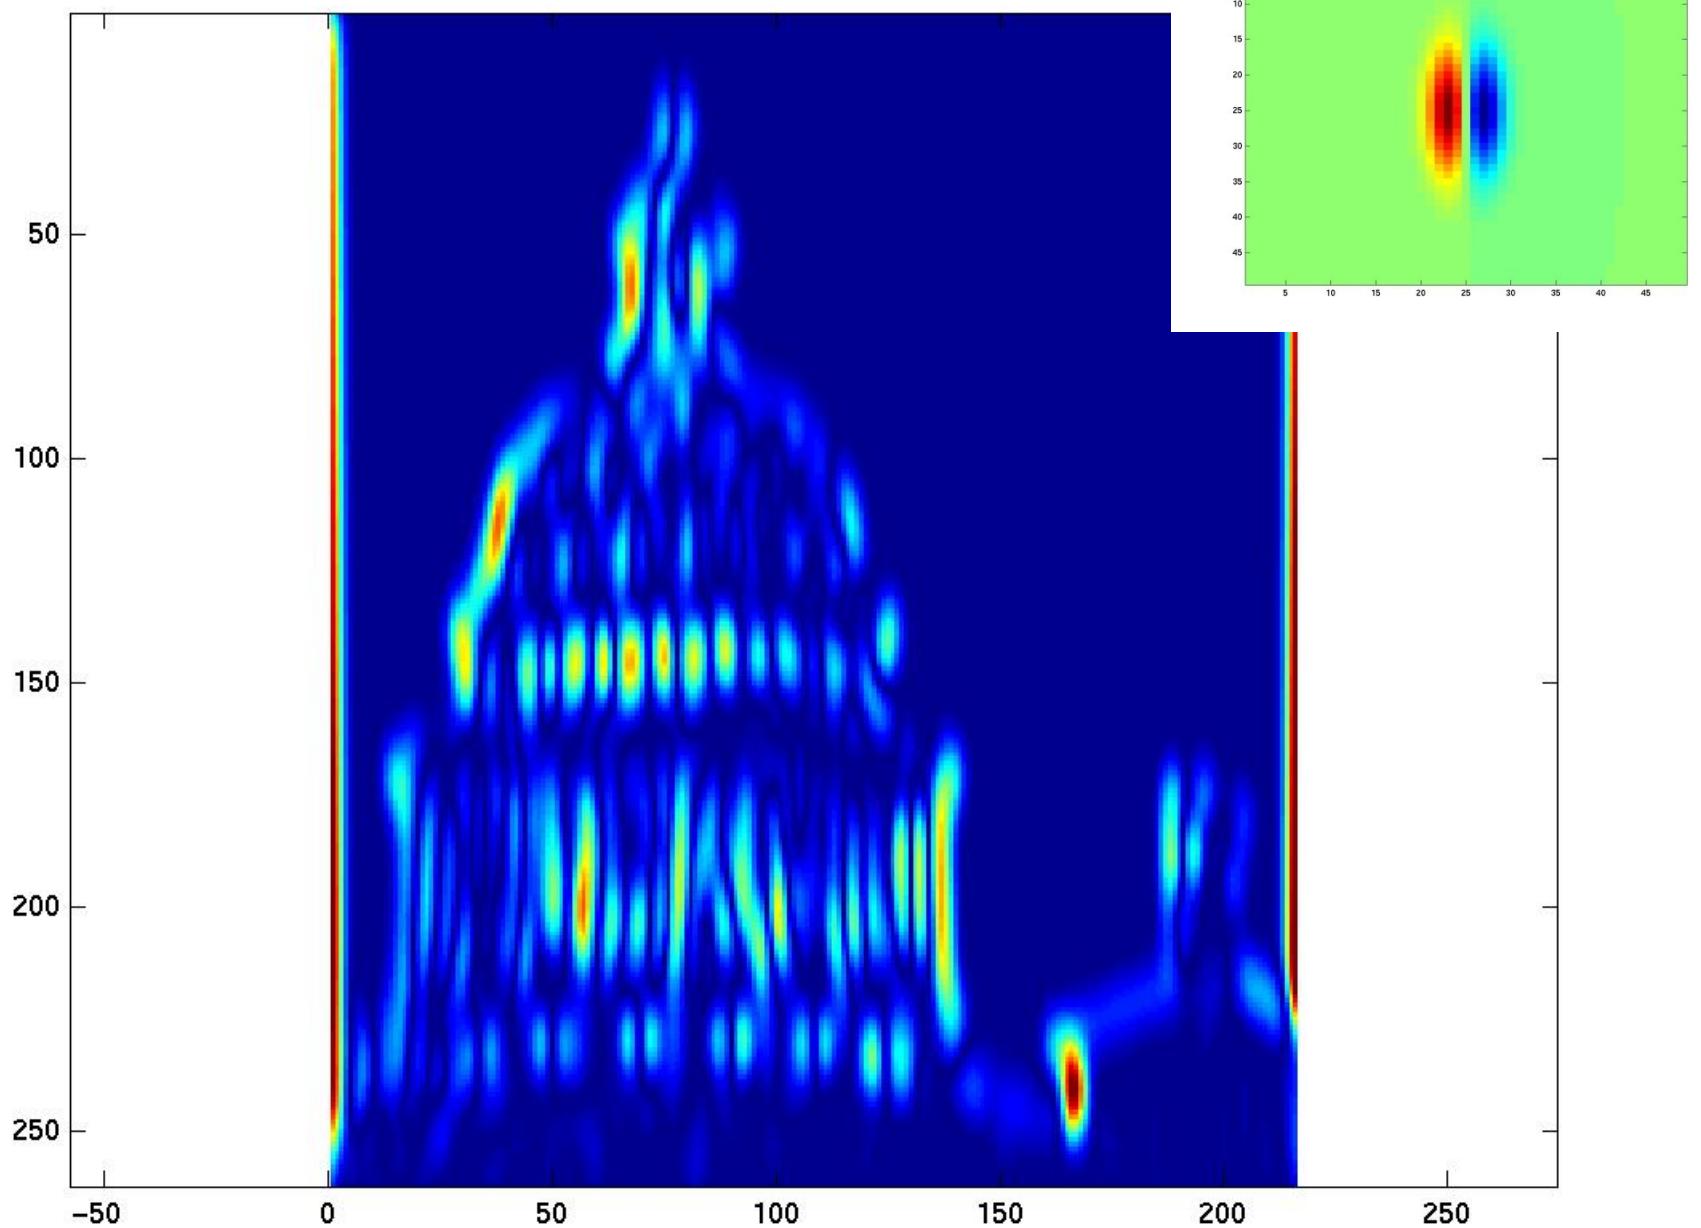


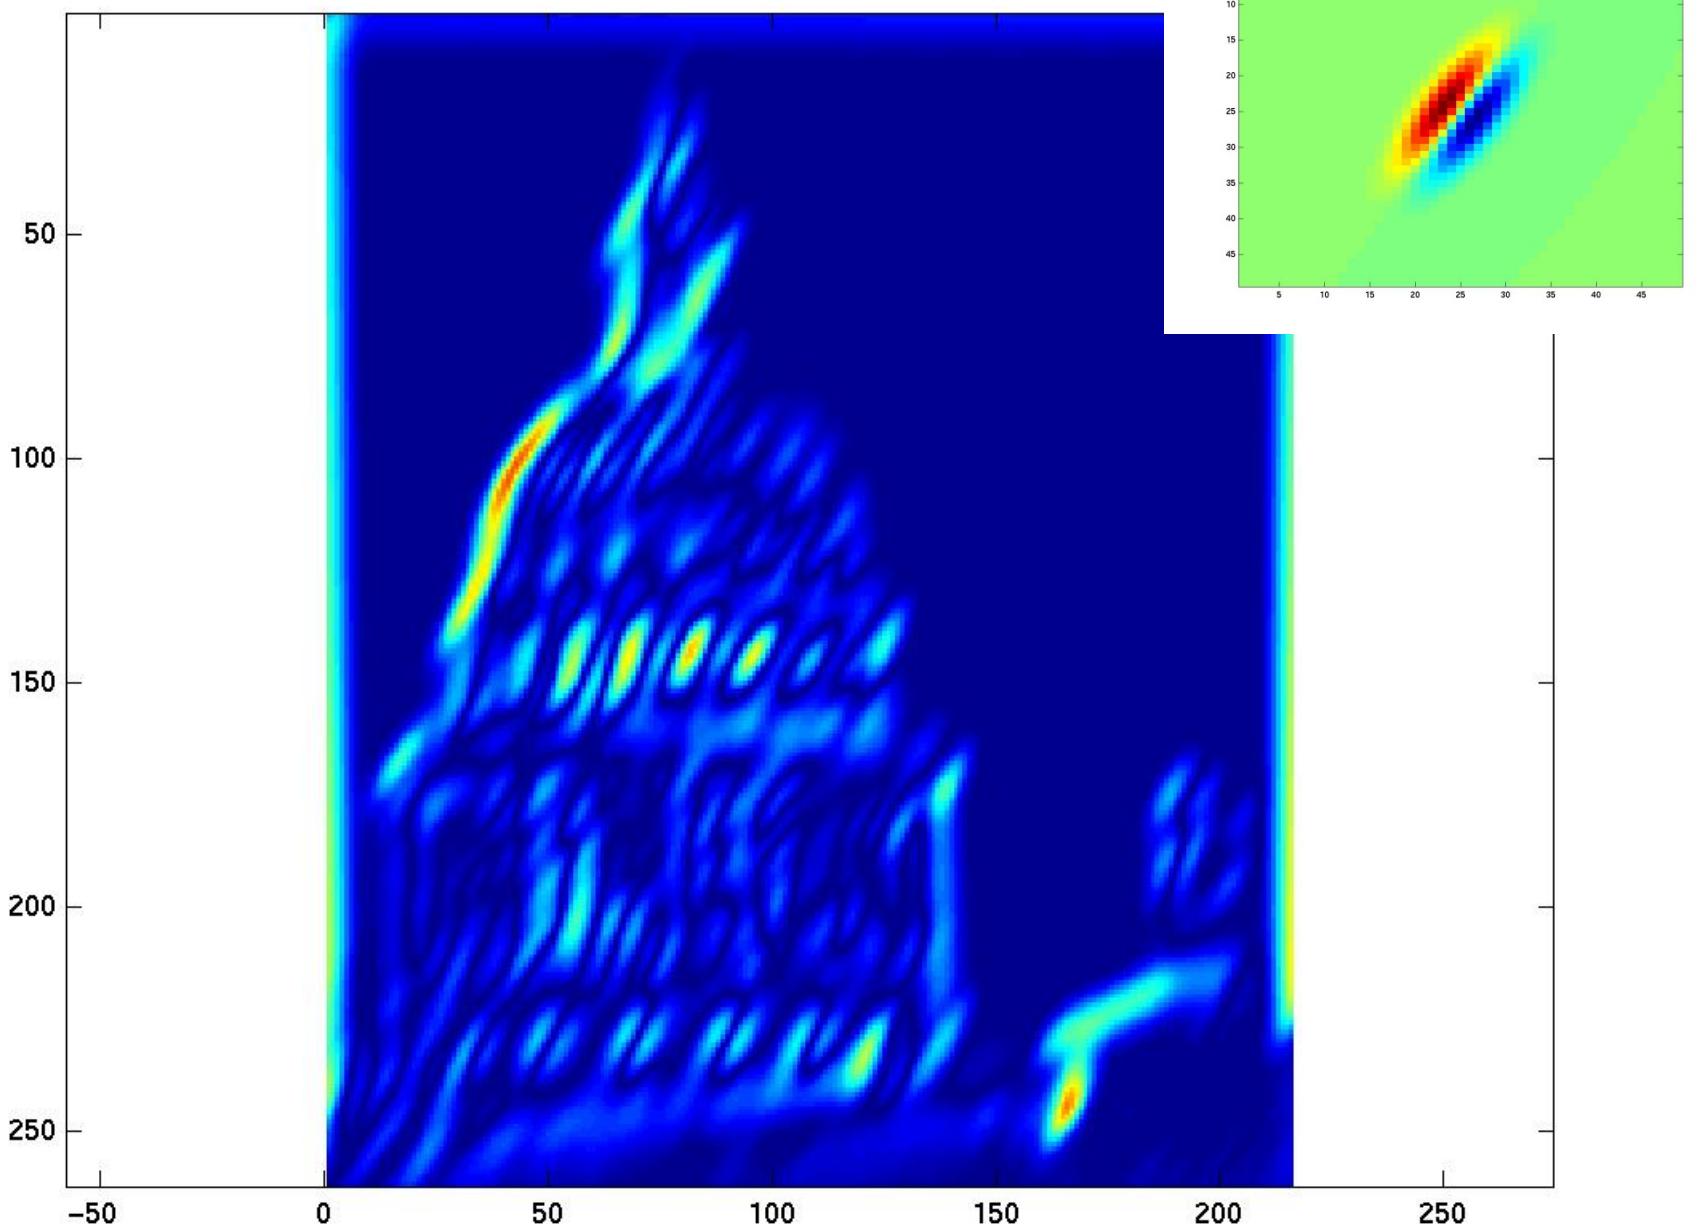


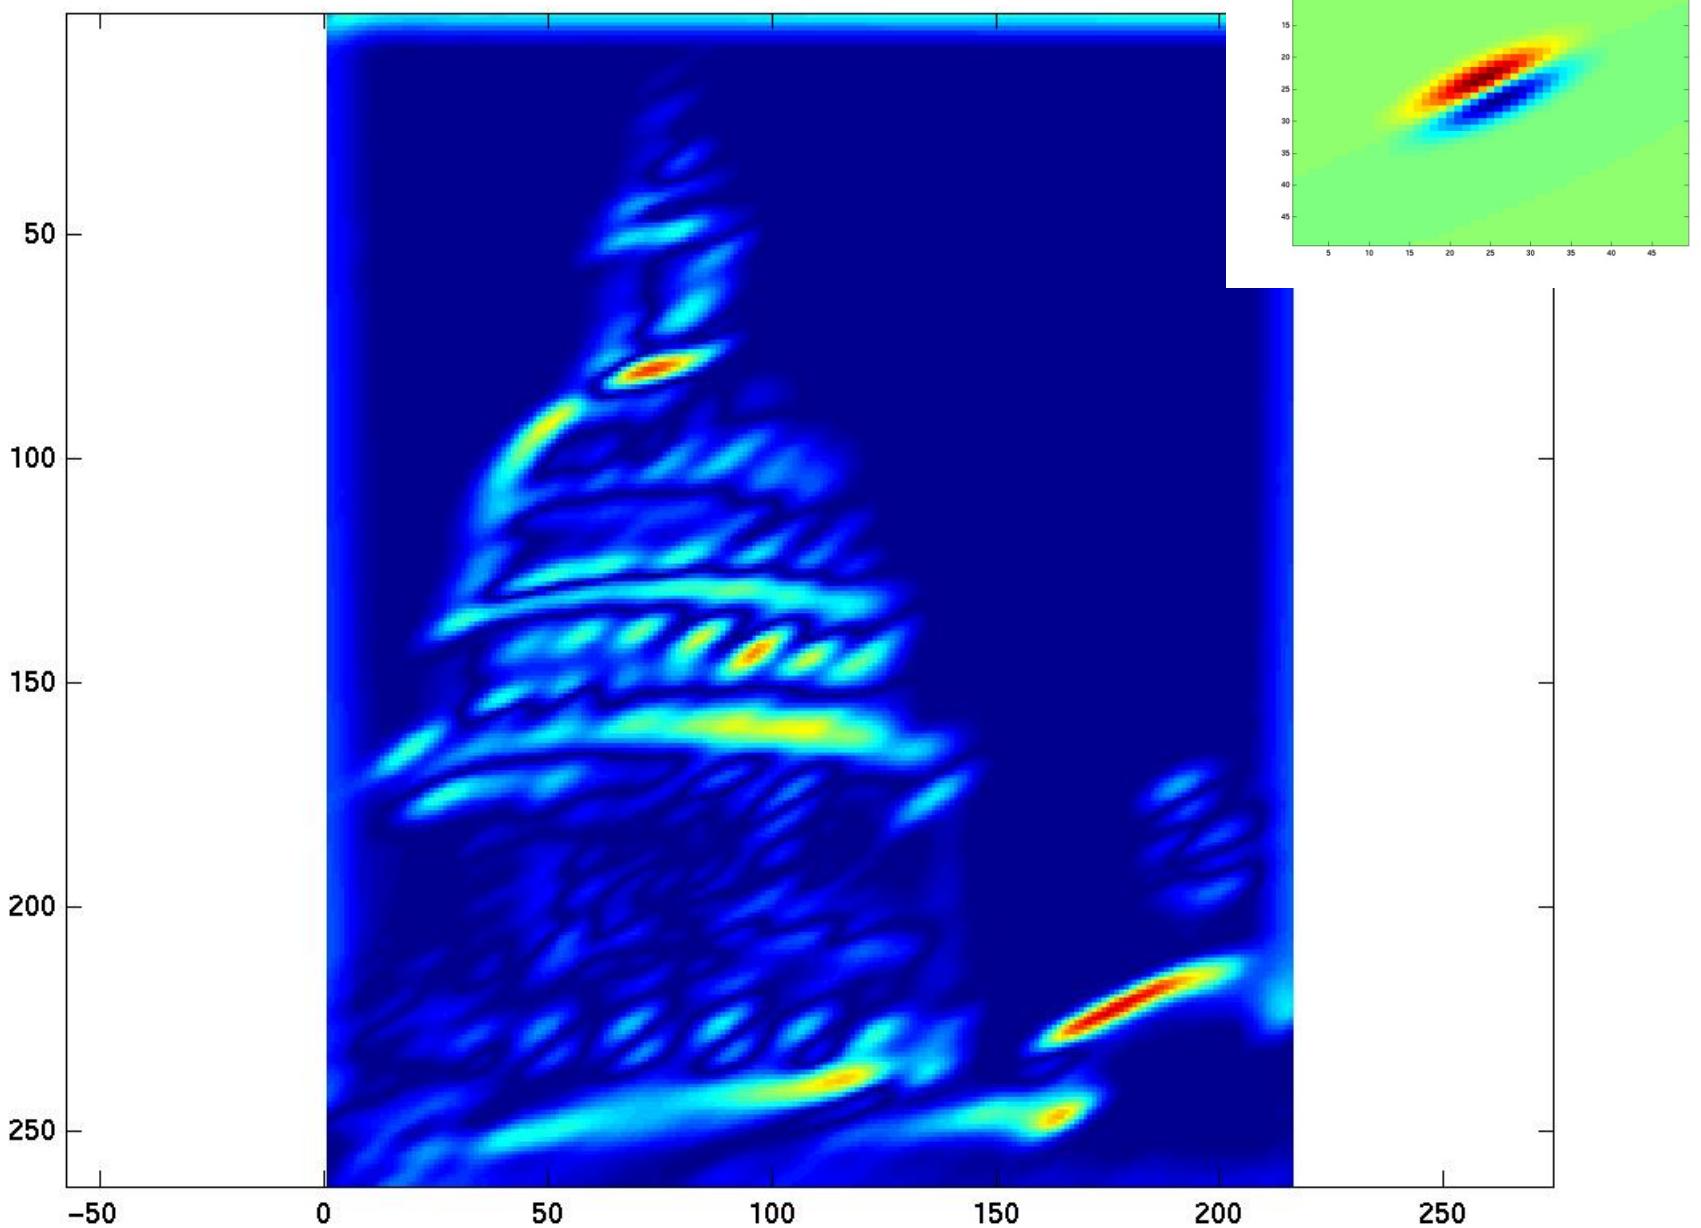


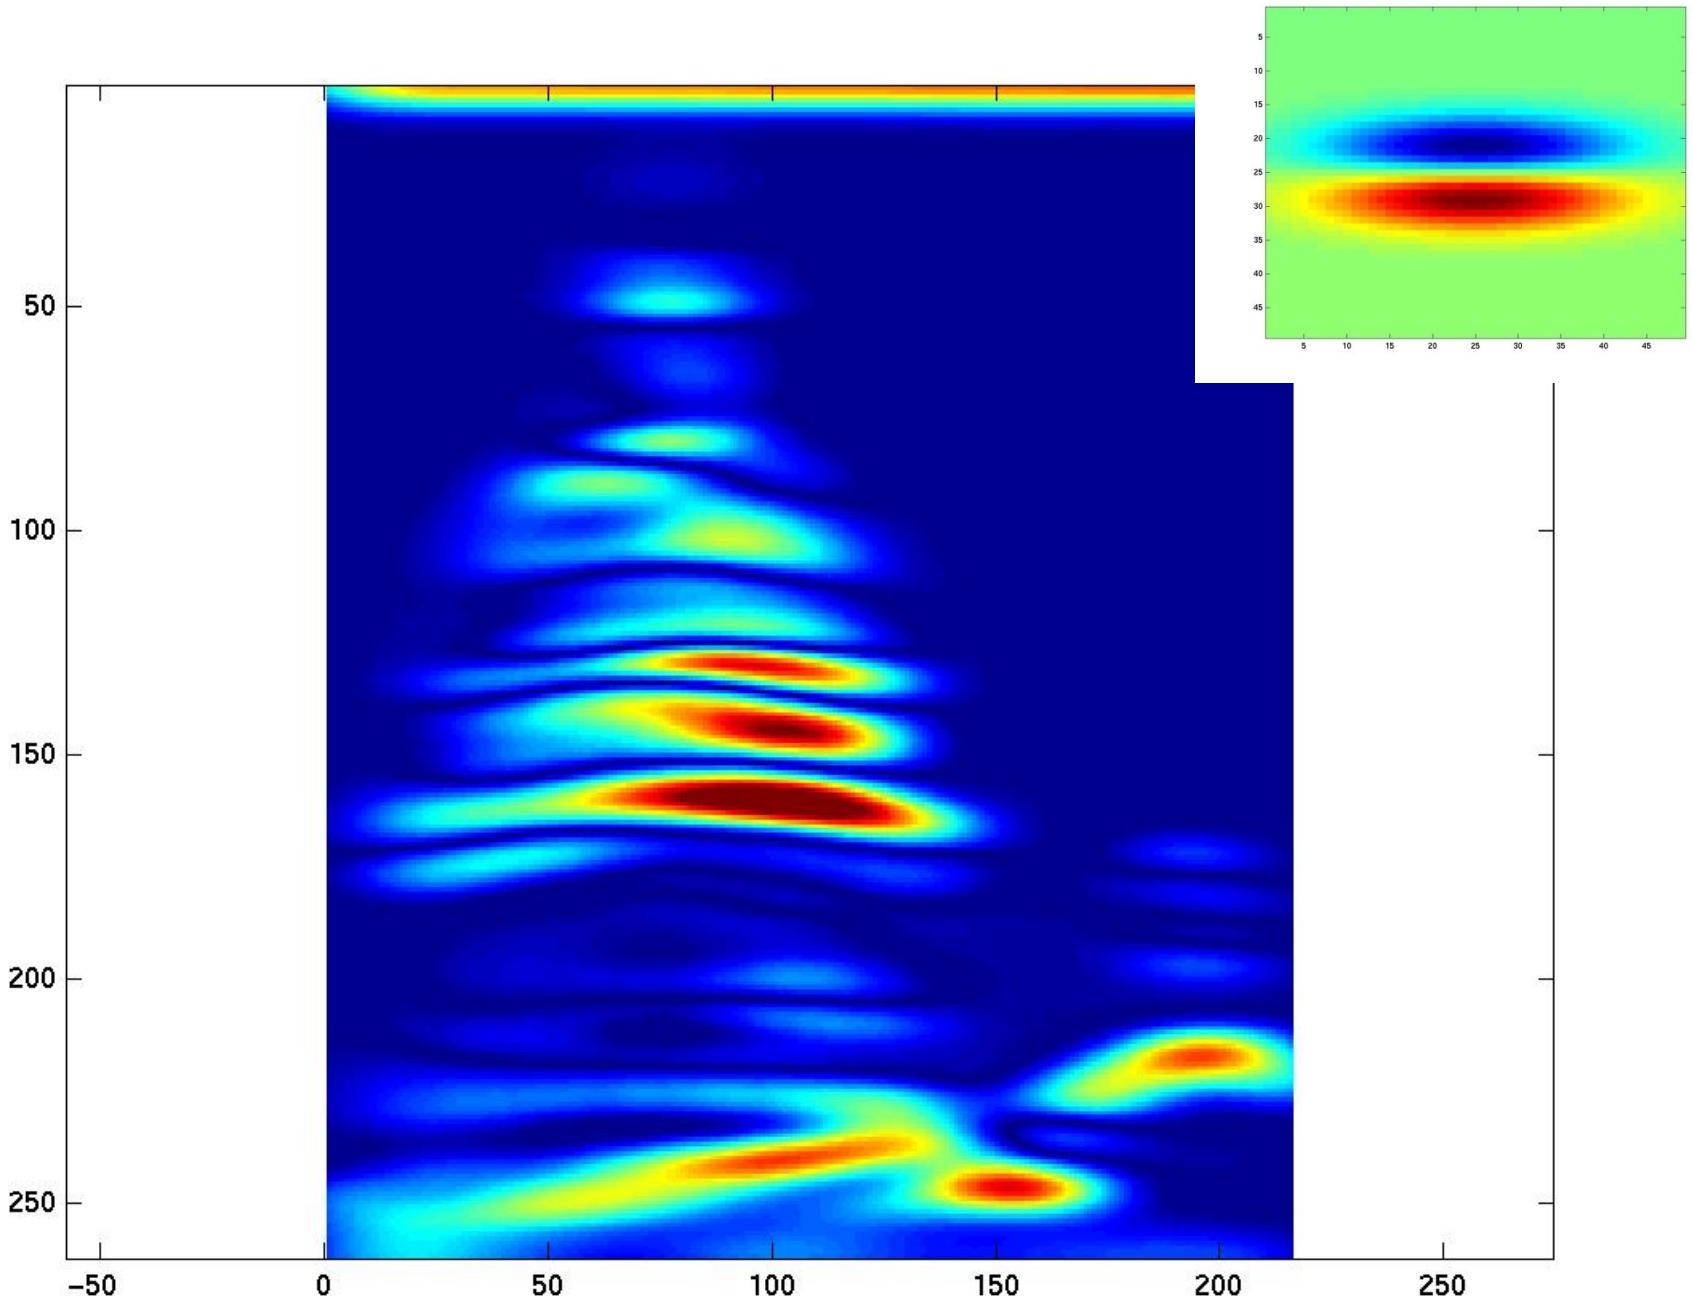


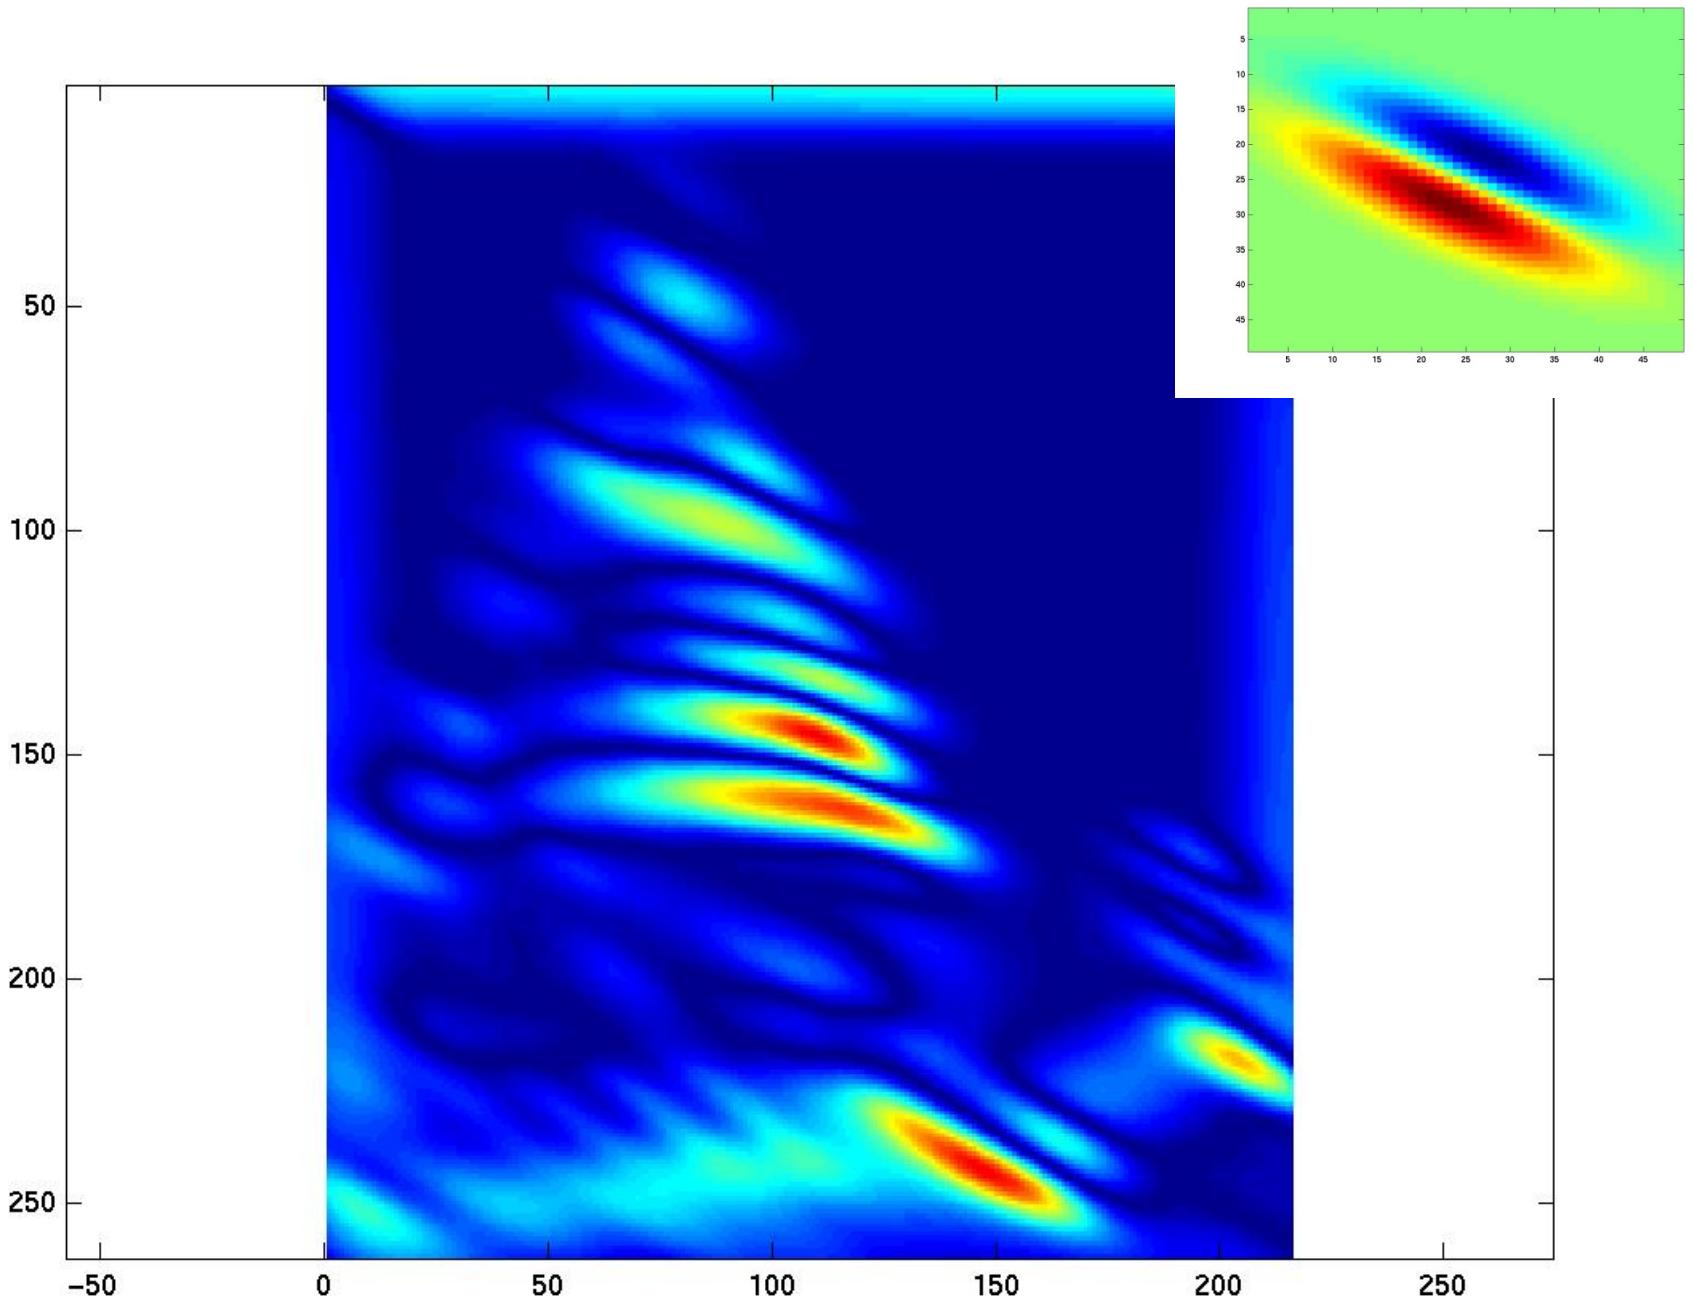


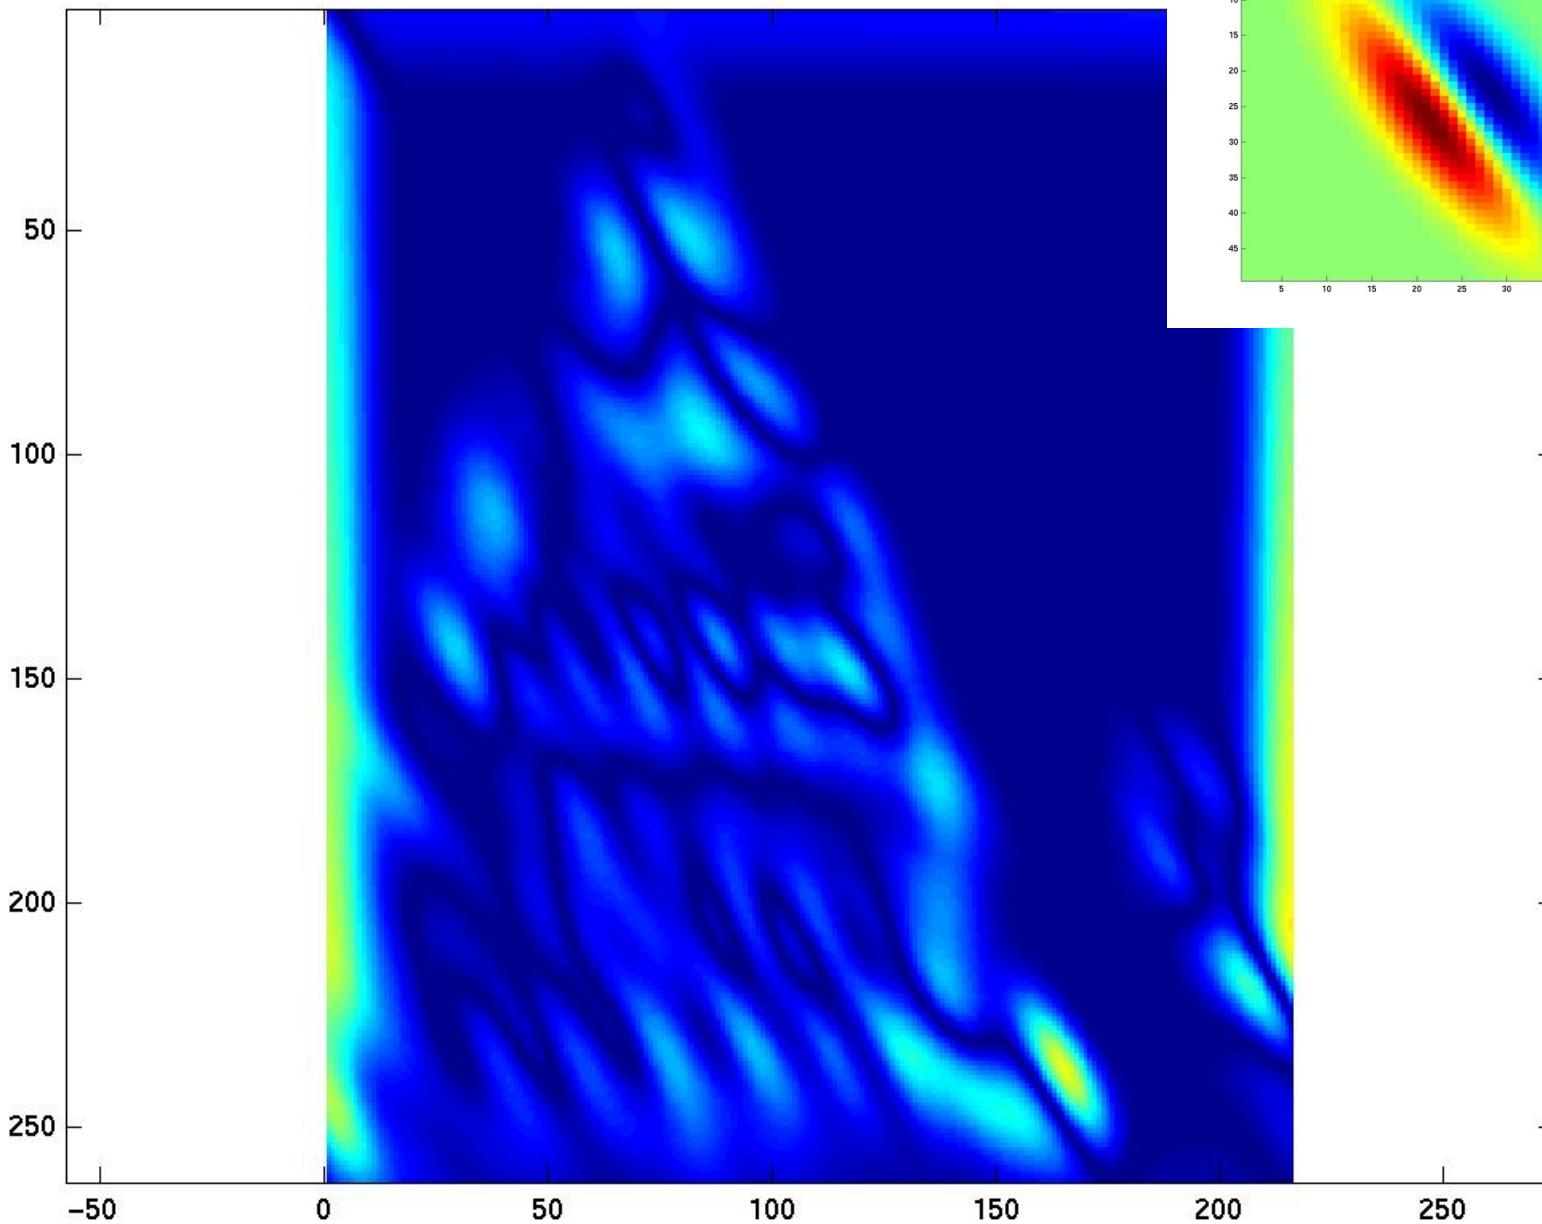


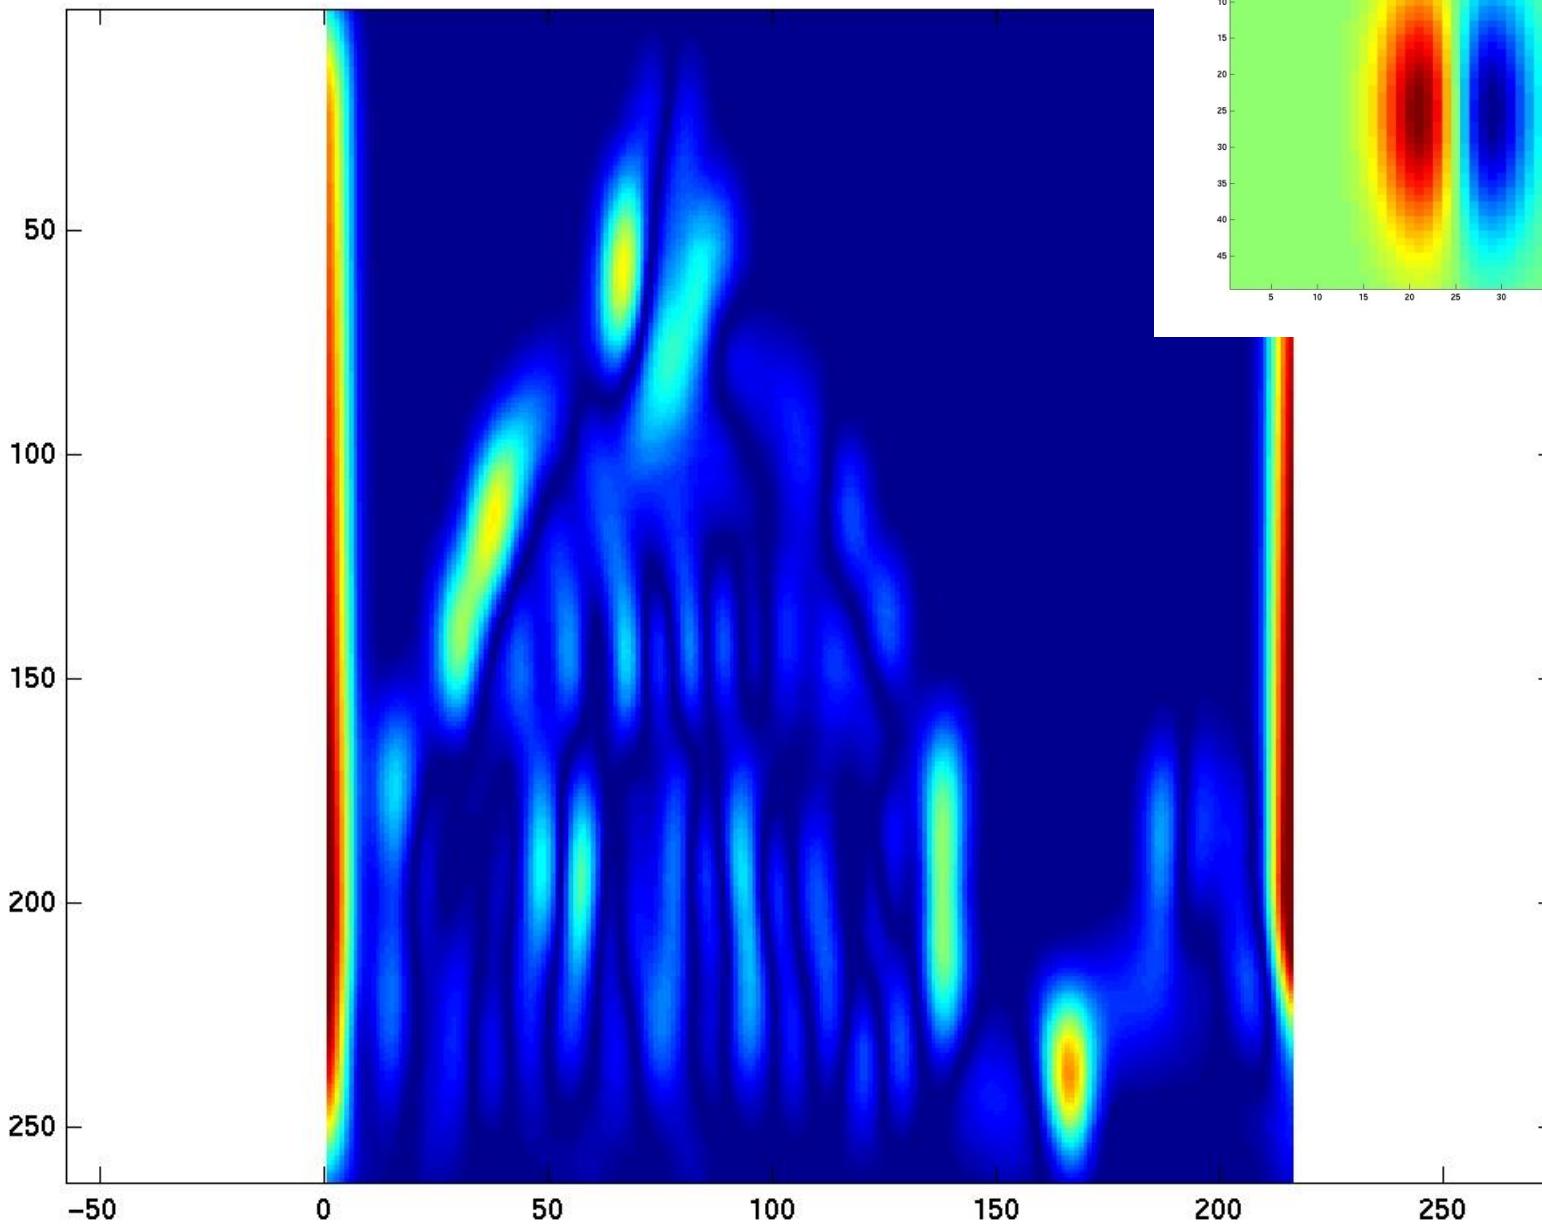


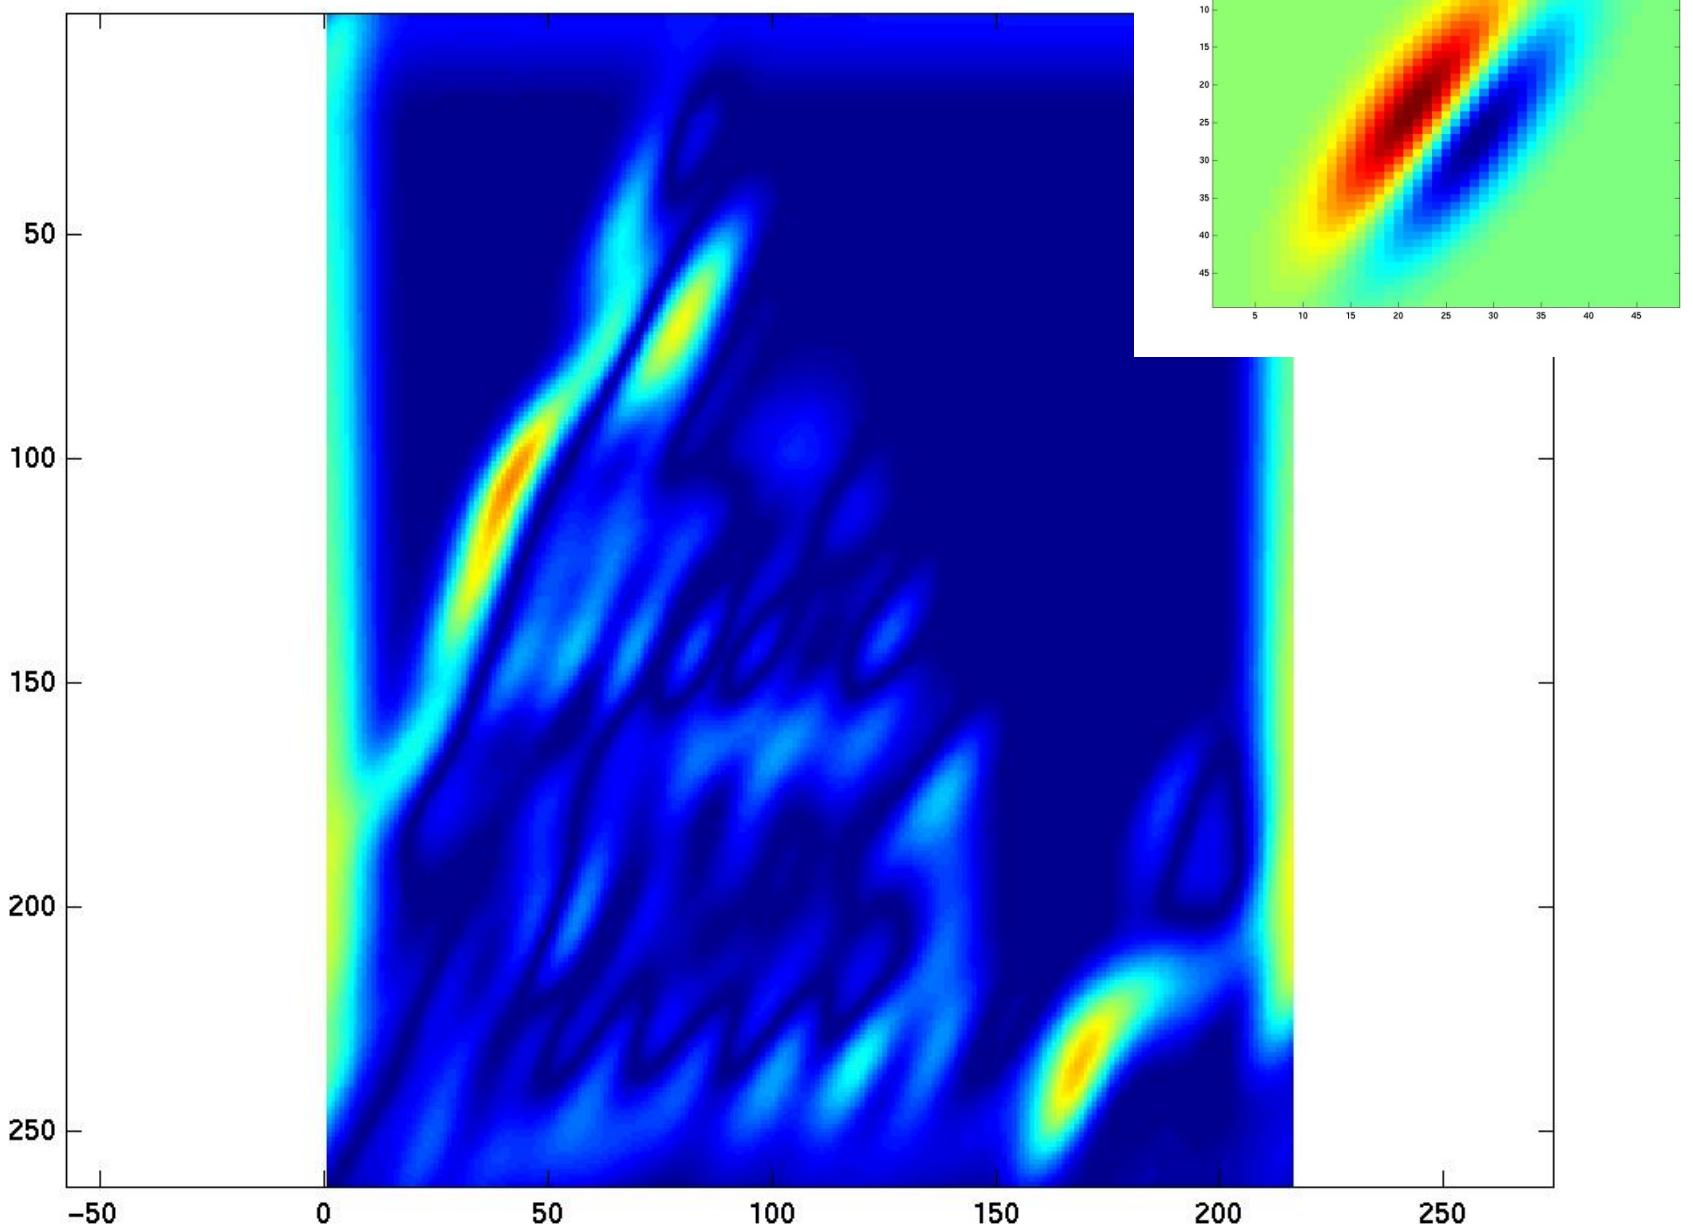


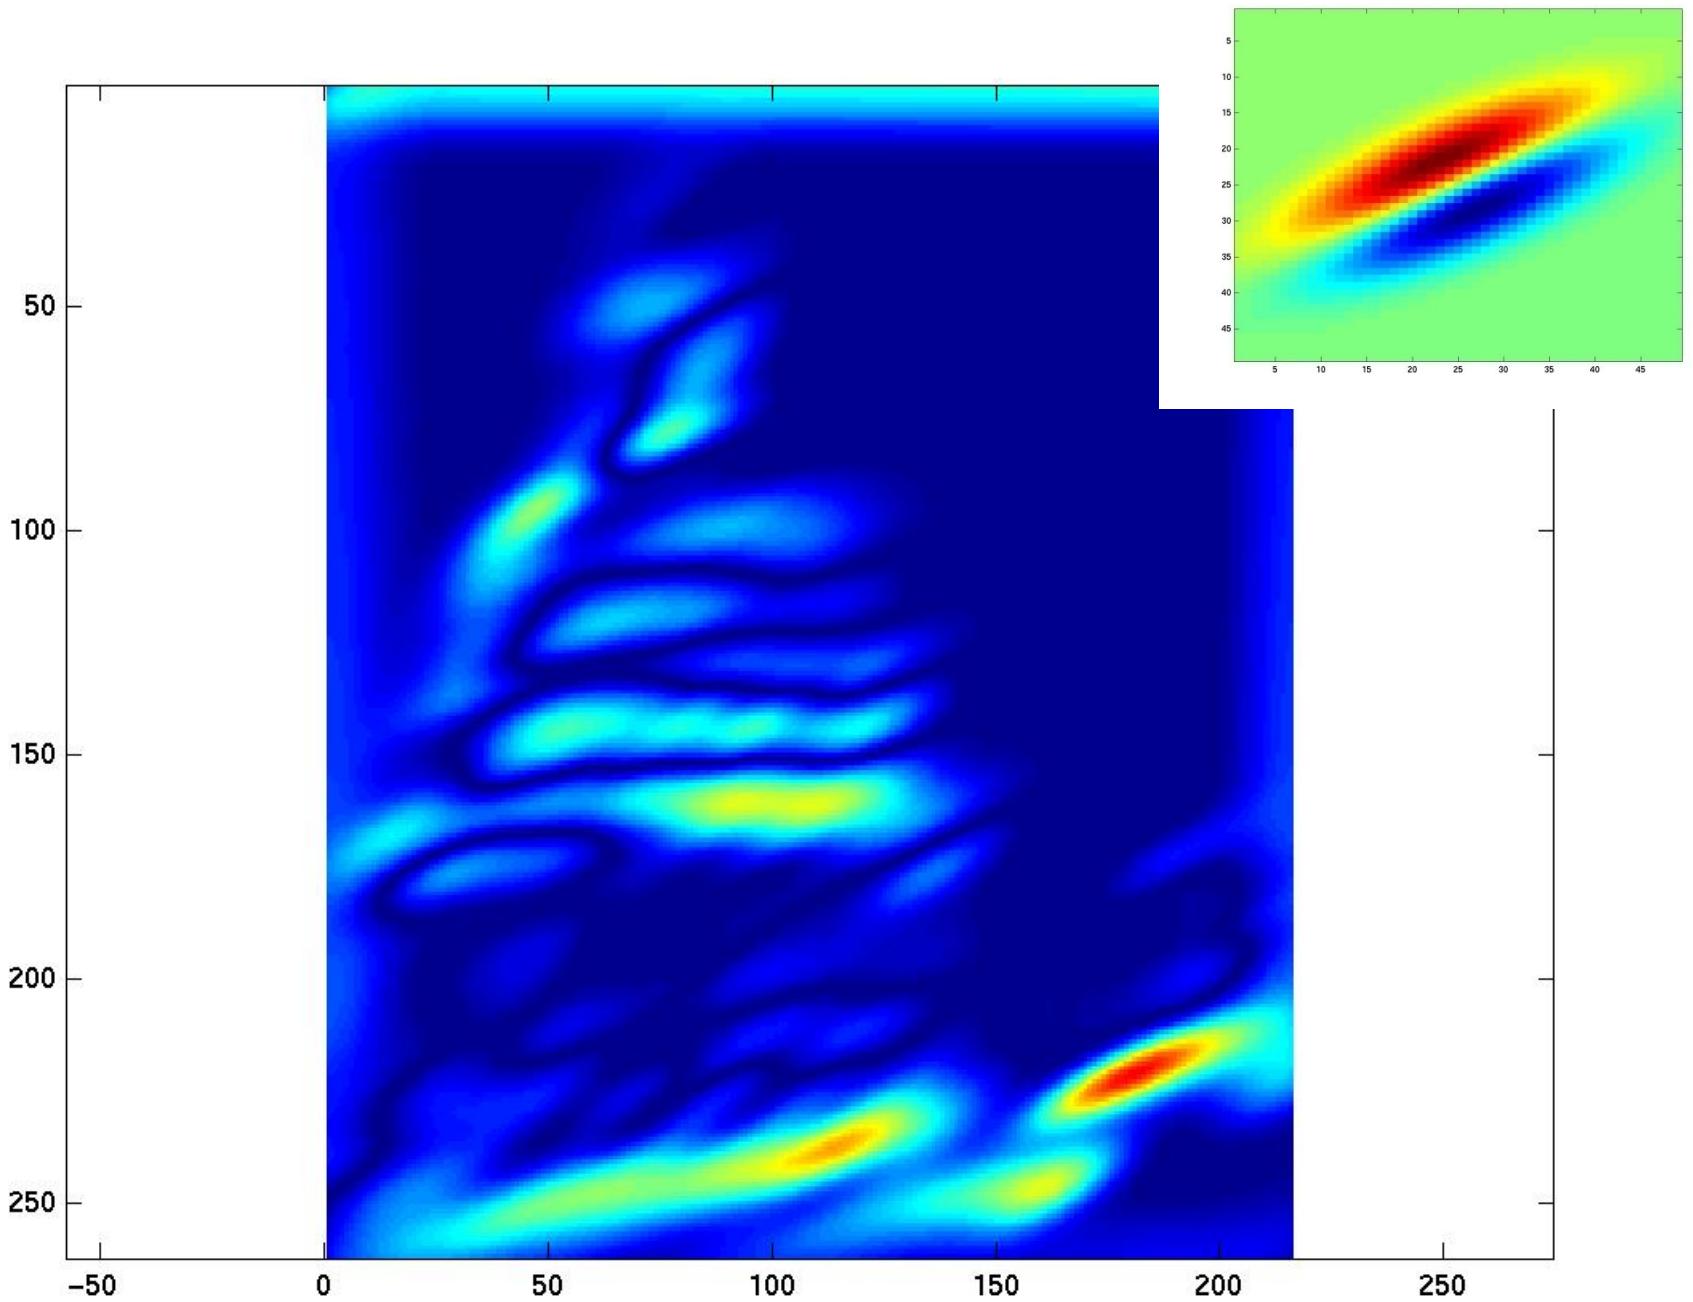


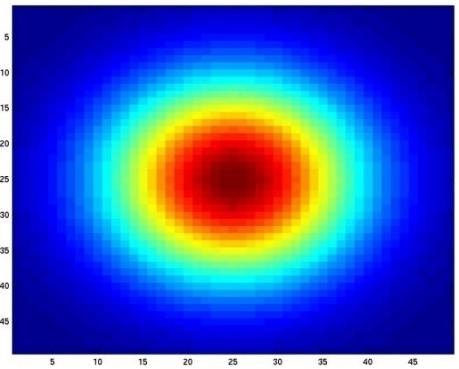
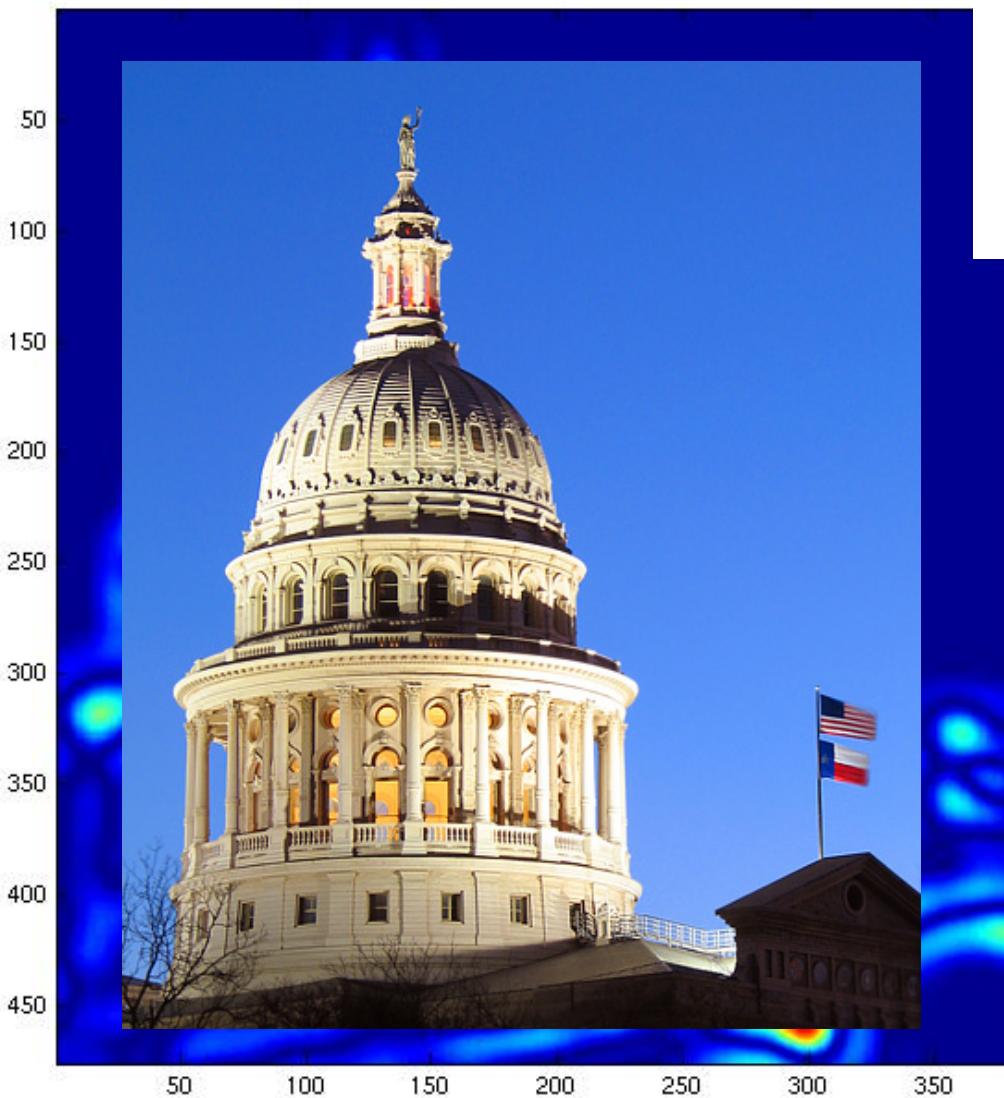




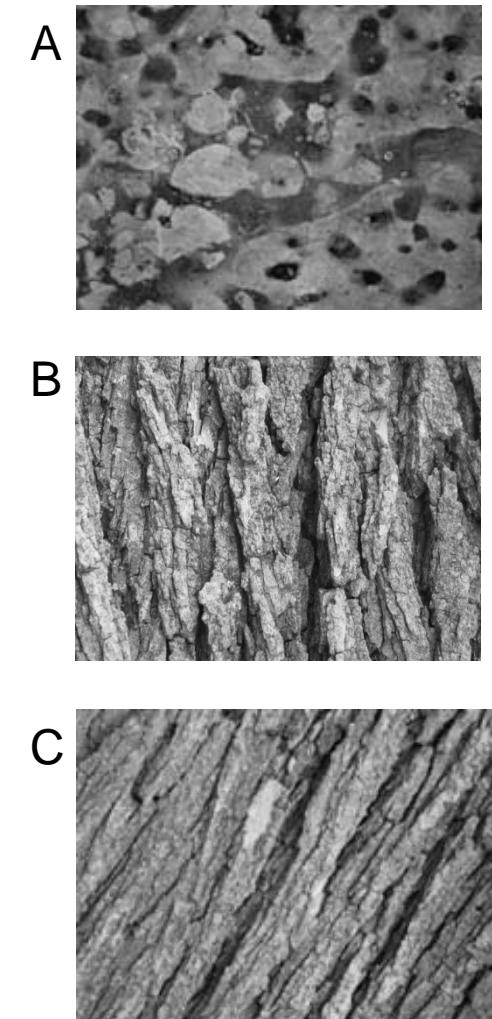
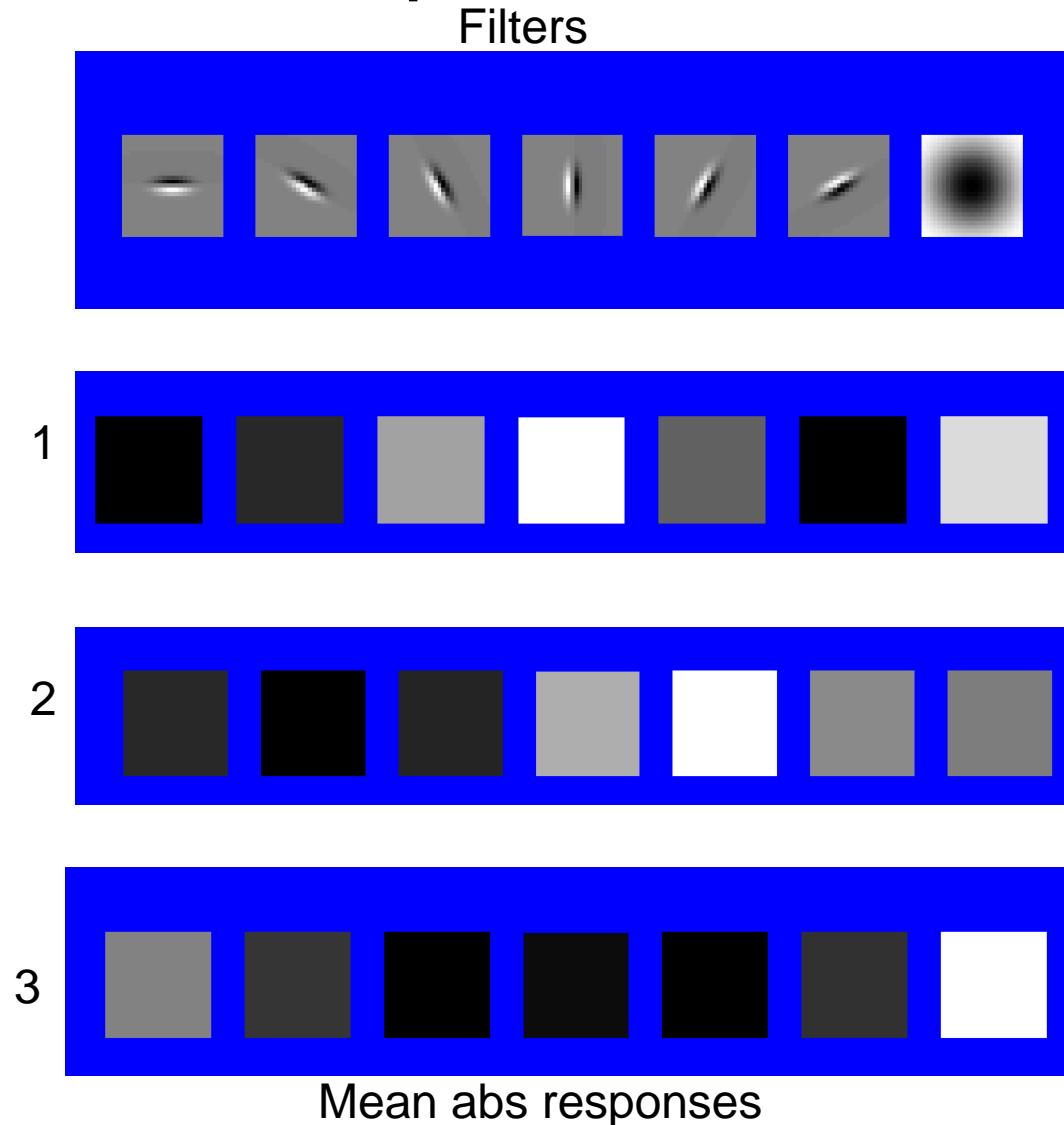




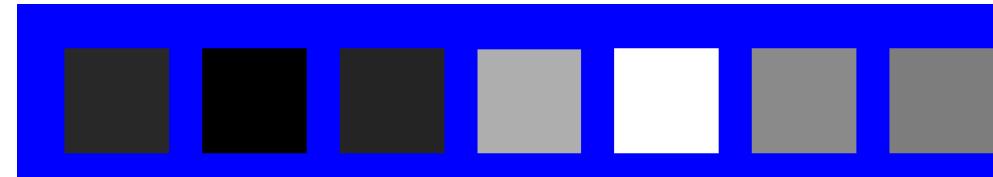
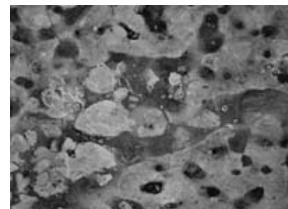
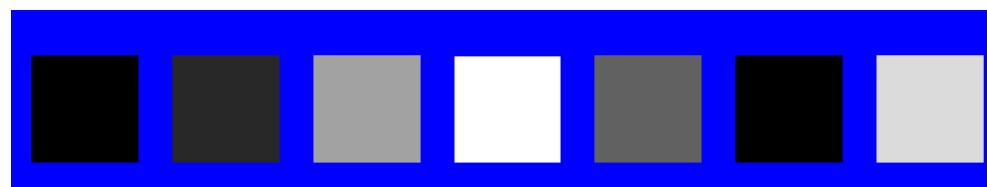
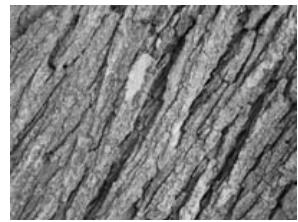
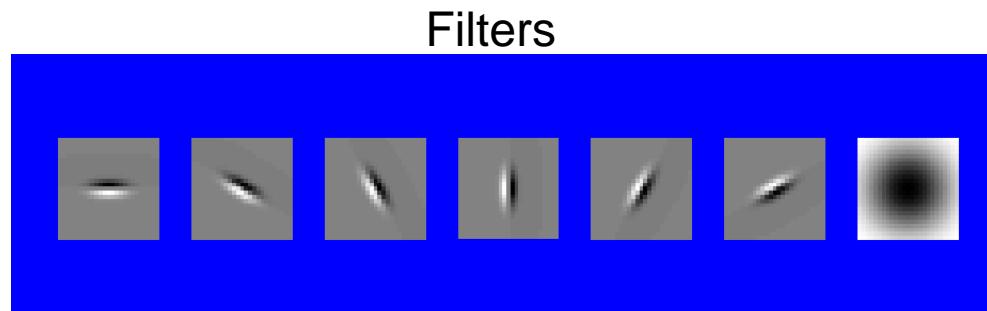




# You try: Can you match the texture to the response?

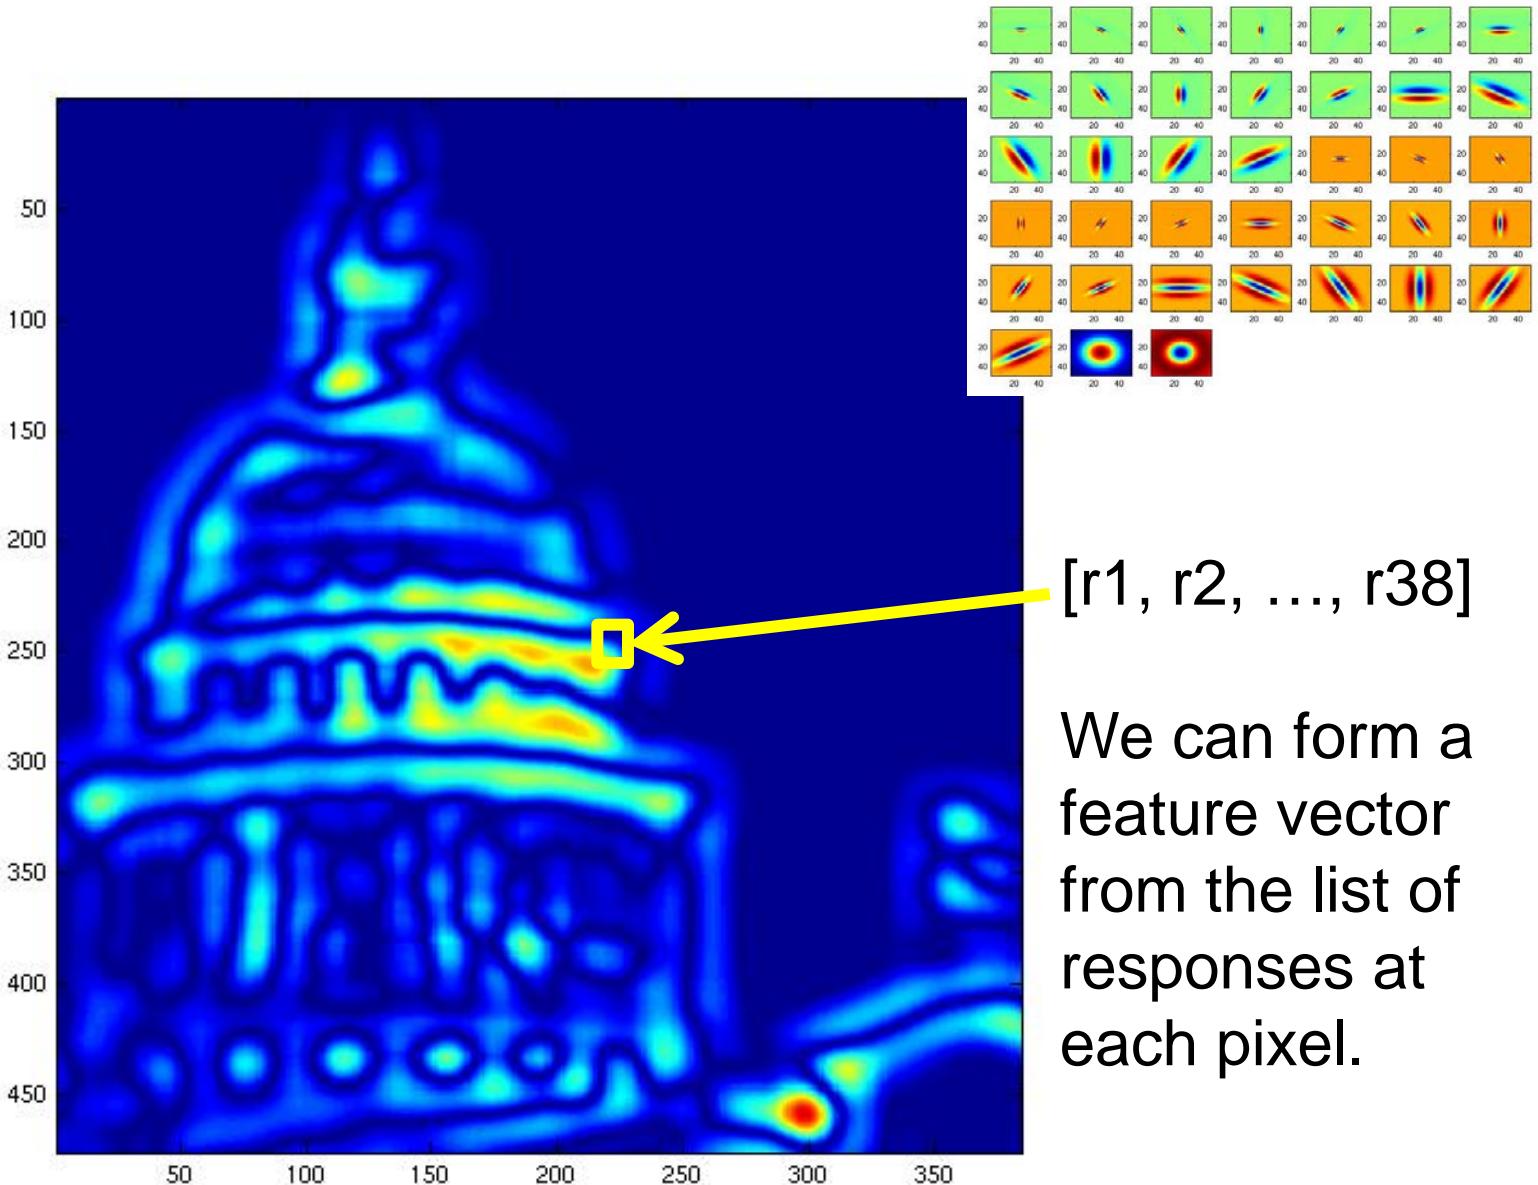


# Representing texture by mean abs response



Mean abs responses

Derek Hoiem

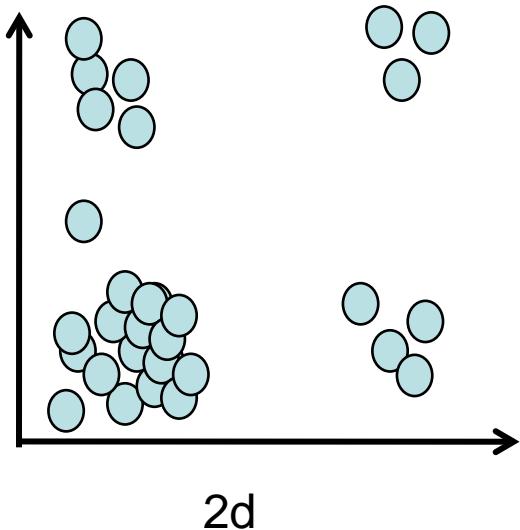


We can form a feature vector from the list of responses at each pixel.

# $d$ -dimensional features

$$D(a, b) = \sqrt{\sum_{i=1}^d (a_i - b_i)^2}$$

Euclidean distance ( $L_2$ )

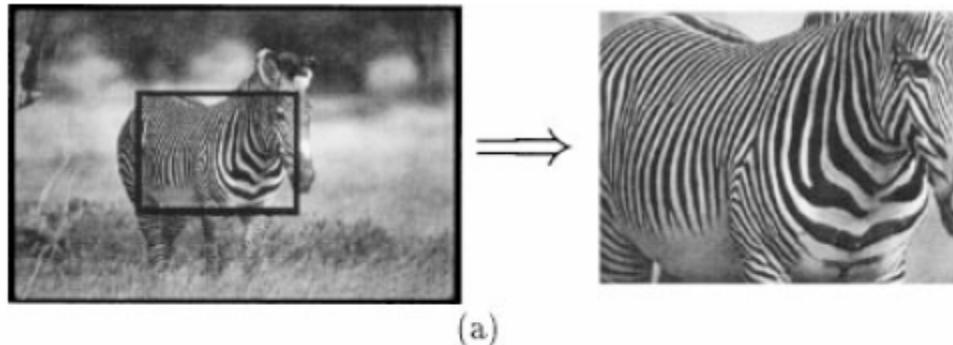


Example uses of  
texture in vision:  
analysis

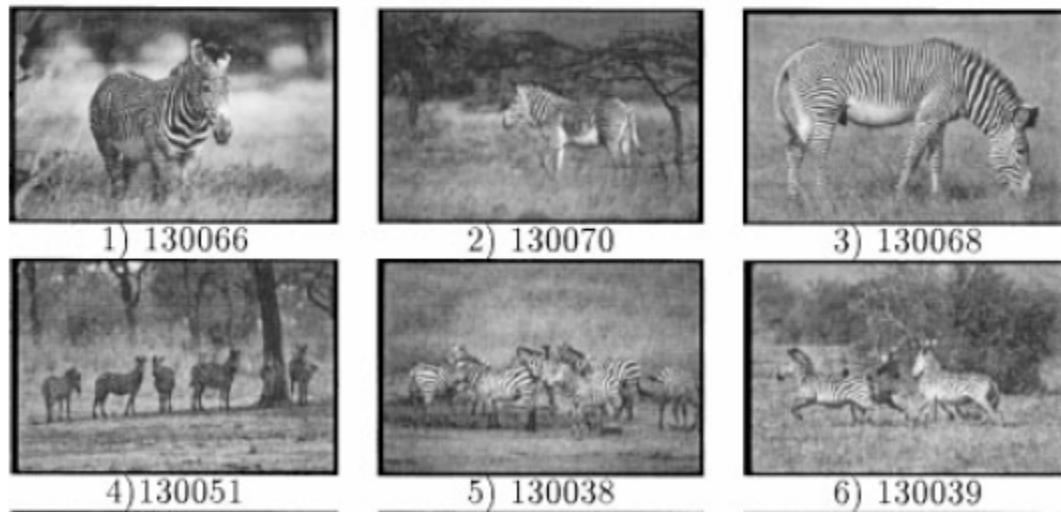
# Classifying materials, “stuff”



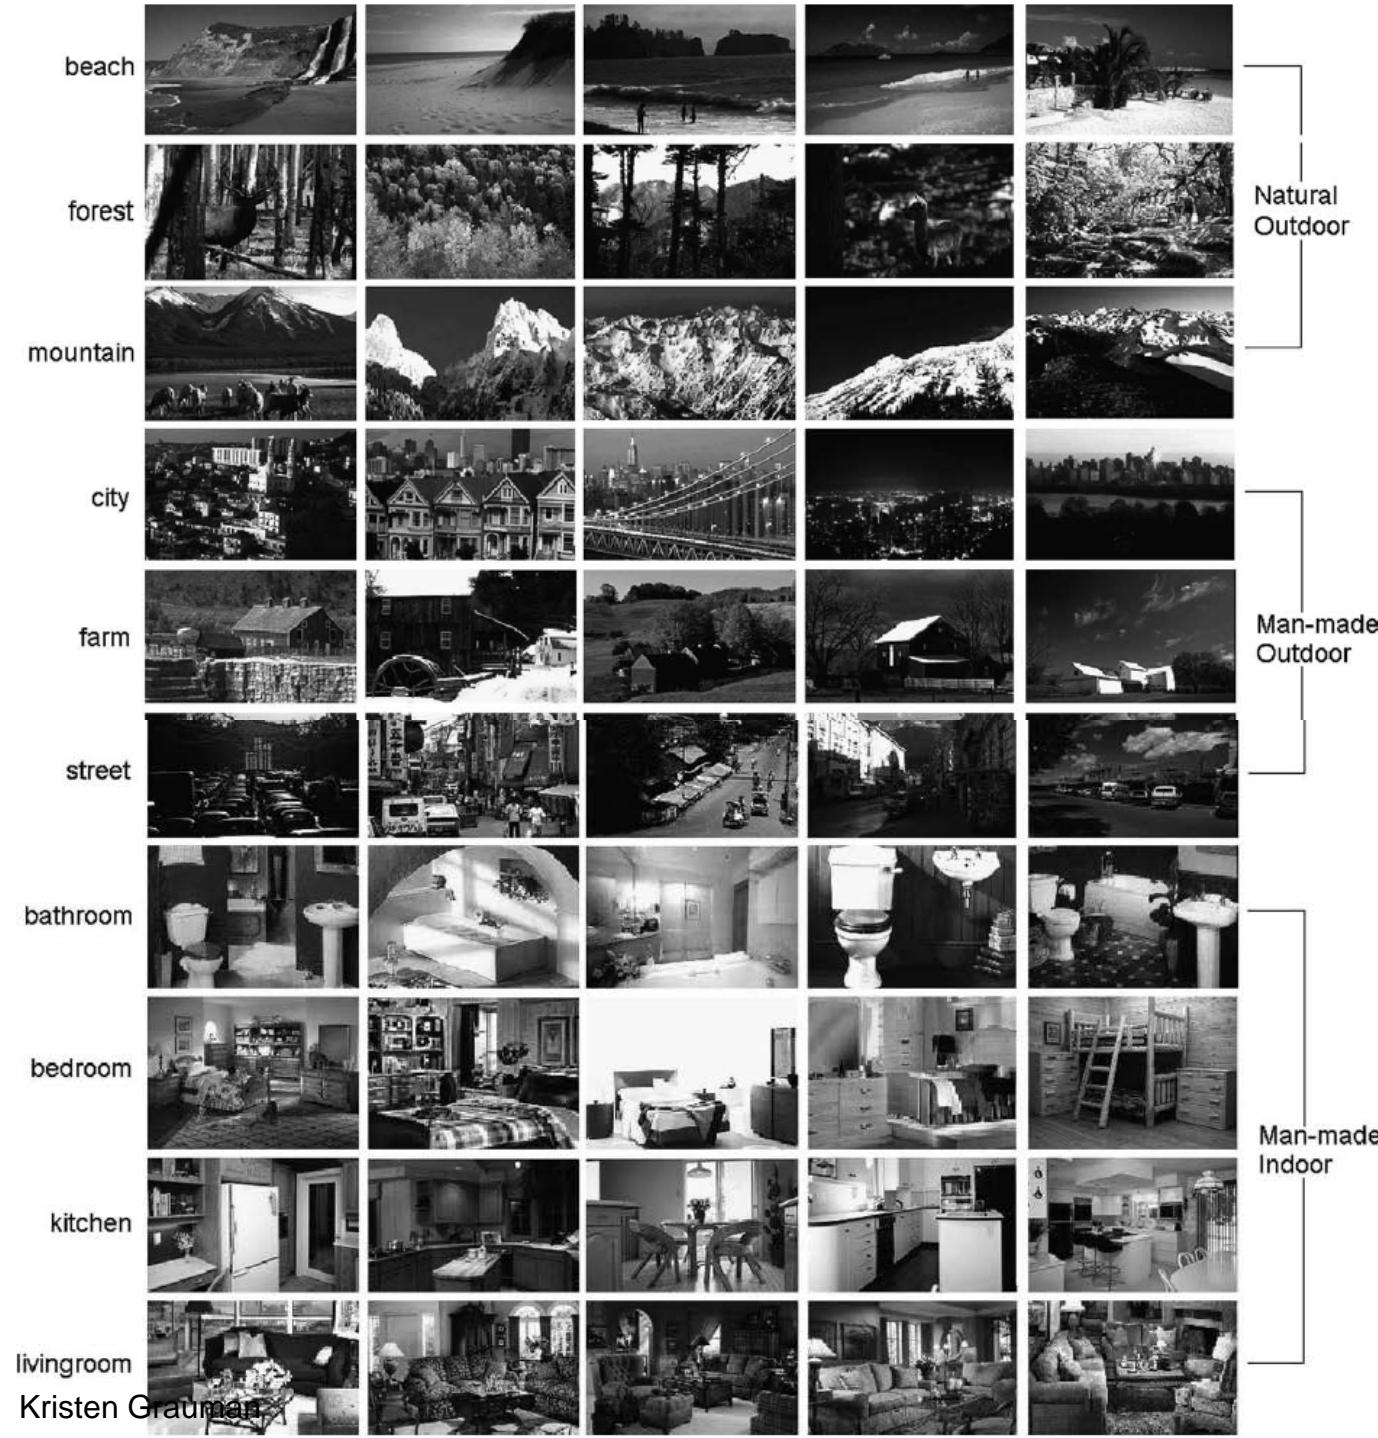
Figure by Varma  
& Zisserman



## Texture features for image retrieval



Y. Rubner, C. Tomasi, and L. J. Guibas. The earth mover's distance as a metric for image retrieval. *International Journal of Computer Vision*, 40(2):99-121, November 2000,



# Characterizing scene categories by texture

L. W. Renninger and J. Malik. When is scene identification just texture recognition? Vision Research 44 (2004) 2301–2311



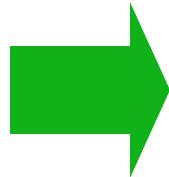
## Segmenting aerial imagery by textures

# Texture-related tasks

- **Shape from texture**
  - Estimate surface orientation or shape from image texture
- **Segmentation/classification** from texture cues
  - Analyze, represent texture
  - Group image regions with consistent texture
- **Synthesis**
  - Generate new texture patches/images given some examples

# Texture synthesis

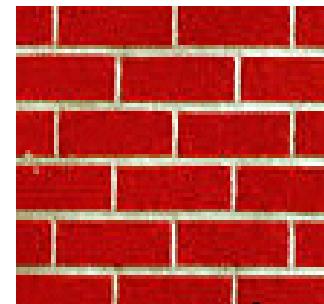
- Goal: create new samples of a given texture
- Many applications: virtual environments, hole-filling, texturing surfaces



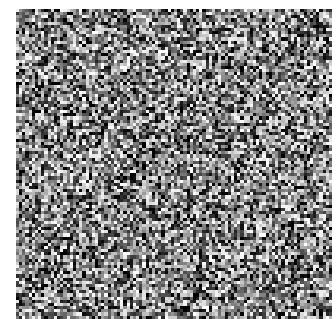
# The Challenge

- Need to model the whole spectrum: from repeated to stochastic texture

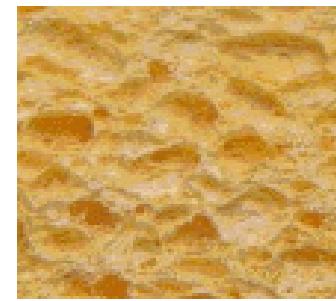
Alexei A. Efros and Thomas K. Leung, “Texture Synthesis by Non-parametric Sampling,” Proc. International Conference on Computer Vision (ICCV), 1999.



repeated



stochastic



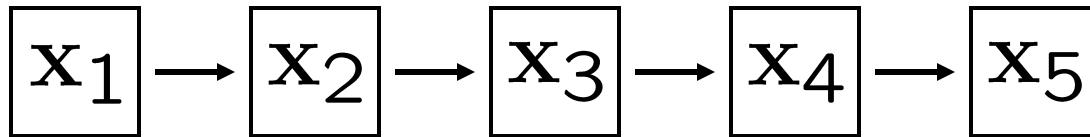
Both?

# Markov Chains

---

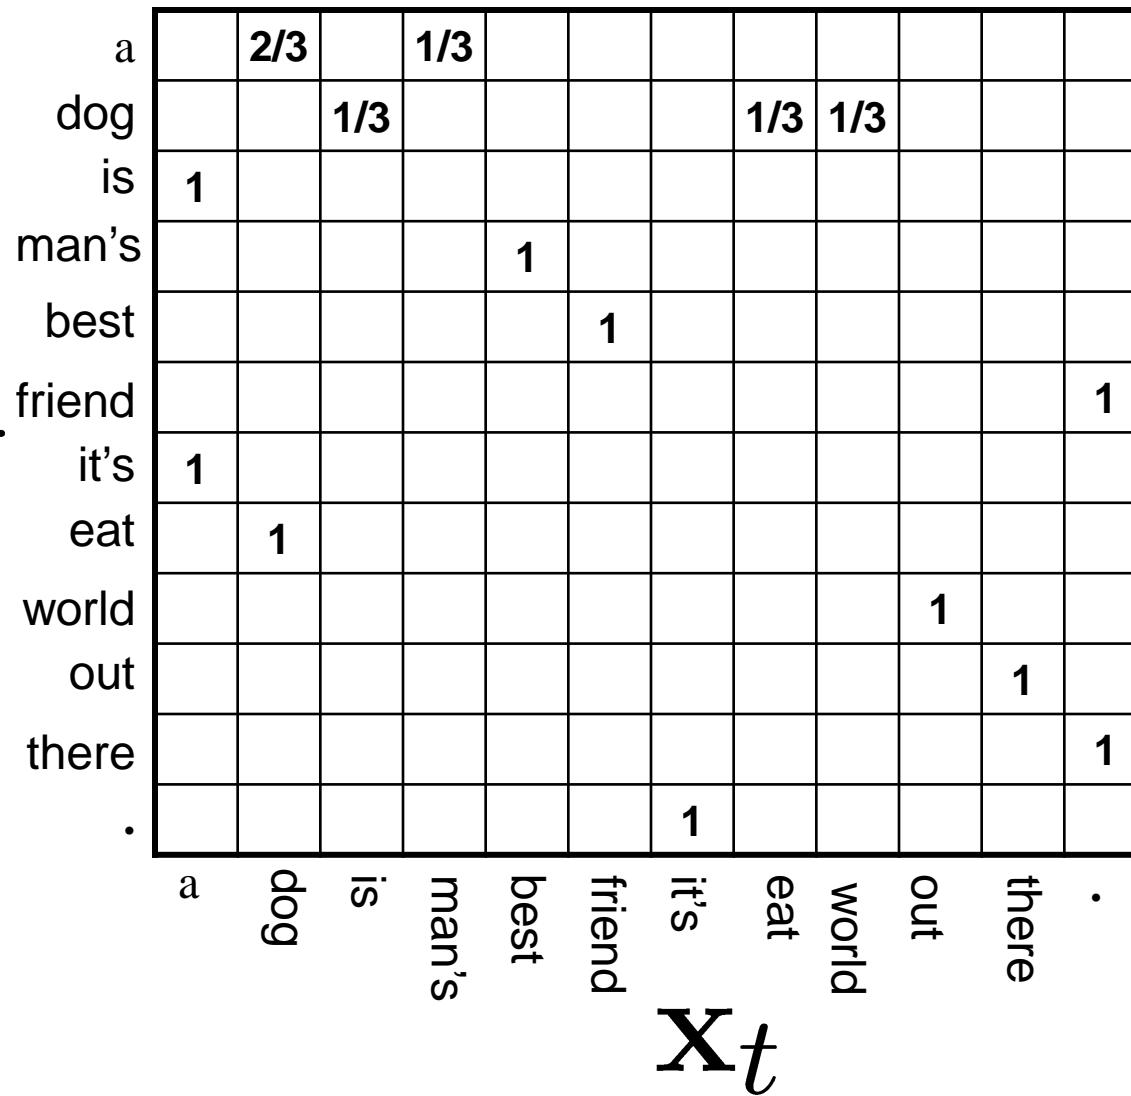
## Markov Chain

- a sequence of random variables  $\mathbf{x}_1, \mathbf{x}_2, \dots, \mathbf{x}_n$
- $\mathbf{x}_t$  is the **state** of the model at time t



# Markov Chain Example: Text

**“A dog is a man’s best friend. It’s a dog eat dog world out there.”**



# Text synthesis

---

Create plausible looking poetry, love letters, term papers, etc.

## Most basic algorithm

1. Build probability histogram
  - find all blocks of N consecutive words/letters in training documents
  - compute probability of occurrence  $p(\mathbf{x}_t | \mathbf{x}_{t-1}, \dots, \mathbf{x}_{t-(n-1)})$

**WE NEED TO EAT CAKE**

# Text synthesis

- Results:
  - “*As I've commented before, really relating to someone involves standing next to impossible.*”
  - “*One morning I shot an elephant in my arms and kissed him.*”
  - “*I spent an interesting evening recently with a grain of salt*”

Dewdney, “A potpourri of programmed prose and prosody” *Scientific American*, 1989.

# Markov Random Field

---

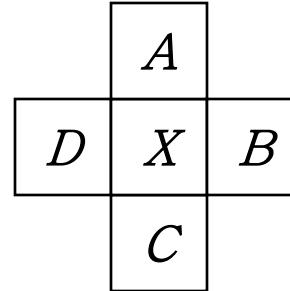
## A Markov random field (MRF)

- generalization of Markov chains to two or more dimensions.

### First-order MRF:

- probability that pixel  $X$  takes a certain value given the values of neighbors  $A$ ,  $B$ ,  $C$ , and  $D$ :

$$P(X|A, B, C, D)$$

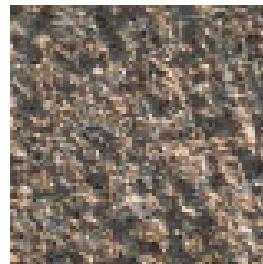


# Texture Synthesis [\[Efros & Leung, ICCV 99\]](#)

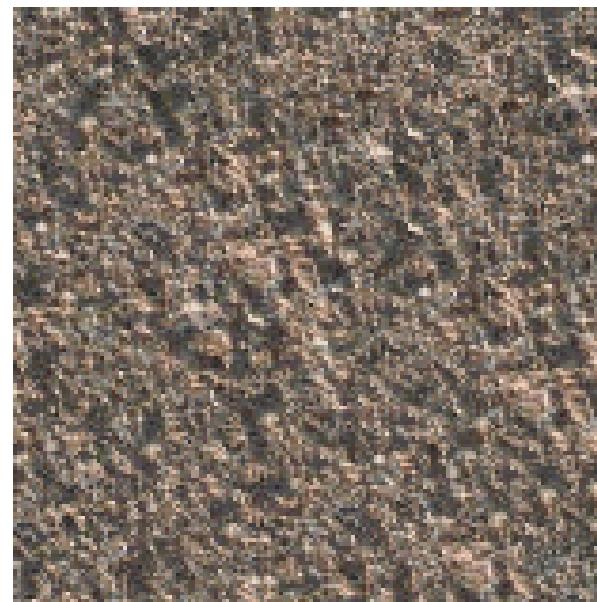
---

Can apply 2D version of text synthesis

Texture corpus  
(sample)



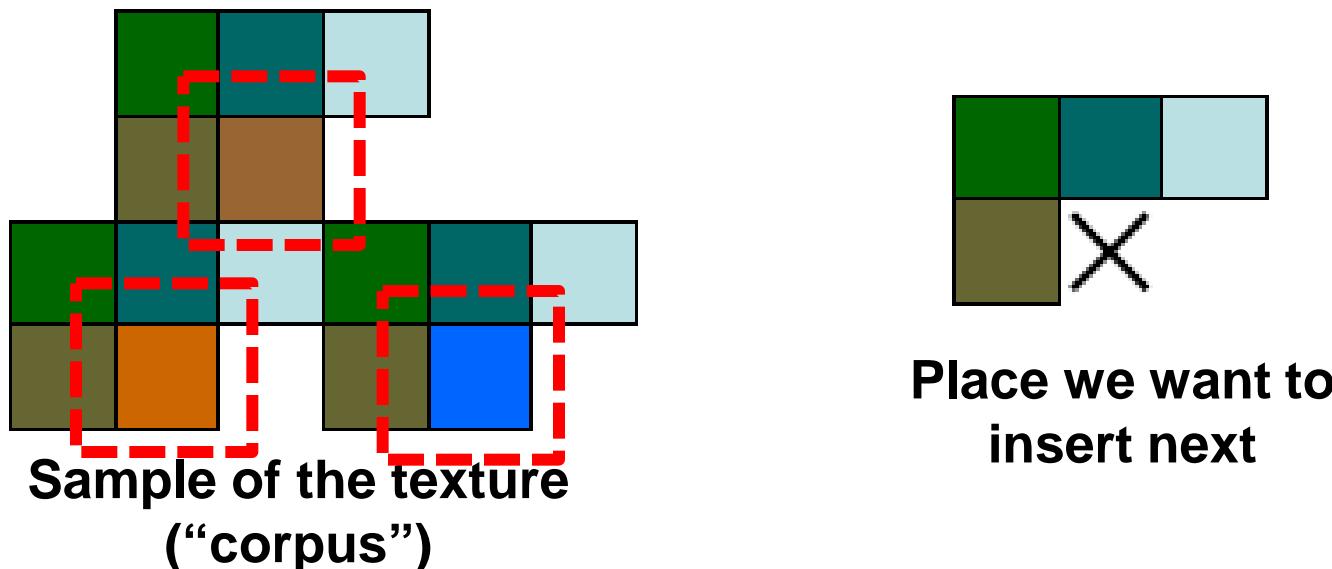
Output



# Texture synthesis: intuition

Before, we inserted the next word based on existing nearby words...

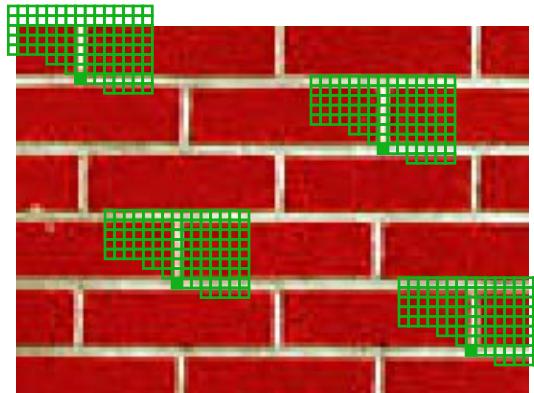
Now we want to insert **pixel intensities** based on existing nearby pixel values.



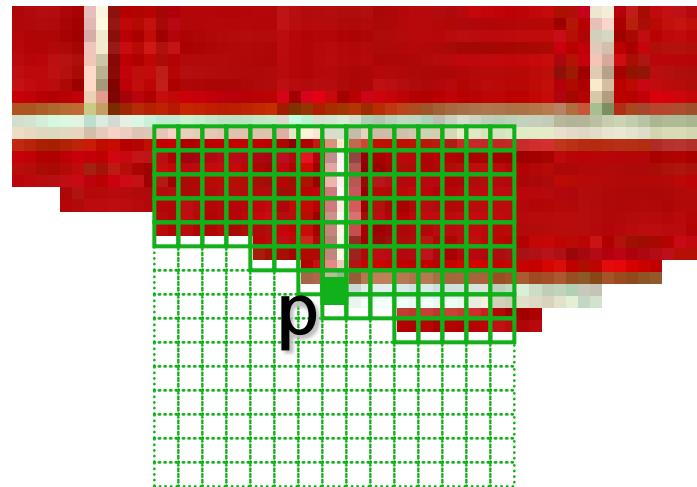
Distribution of a value of a pixel is conditioned on its neighbors alone.

# Synthesizing One Pixel

---



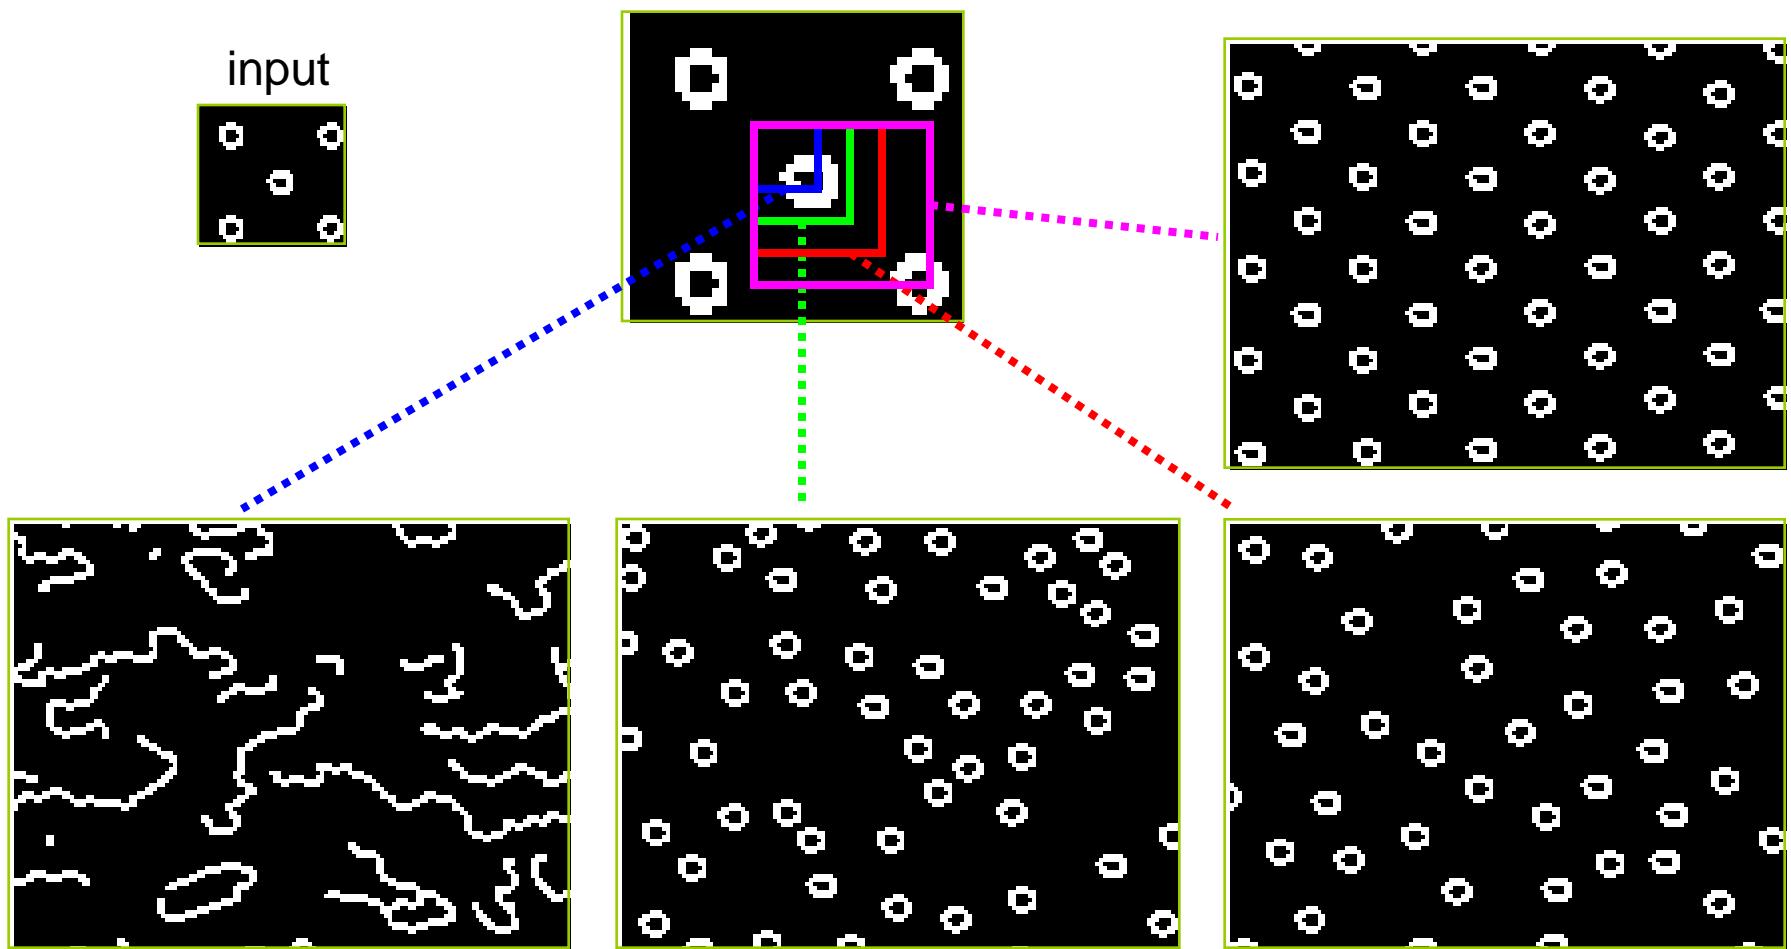
input image



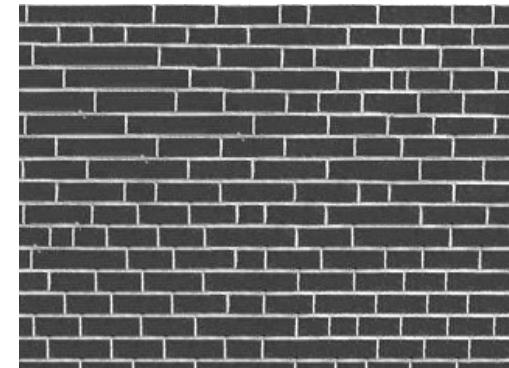
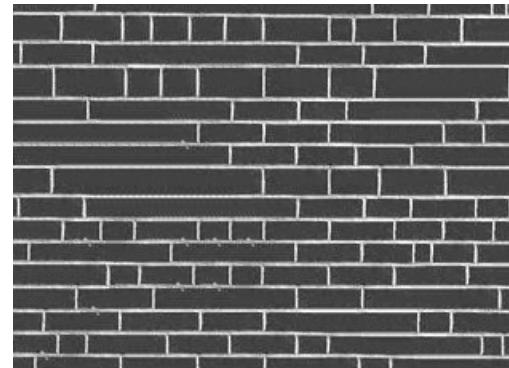
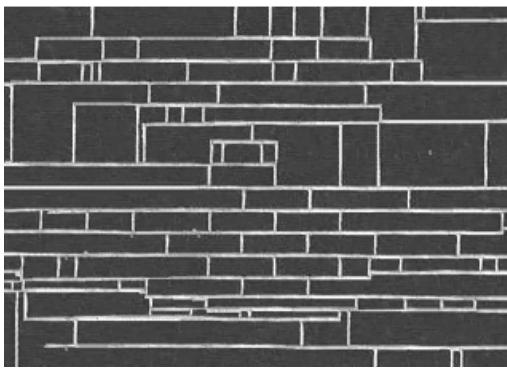
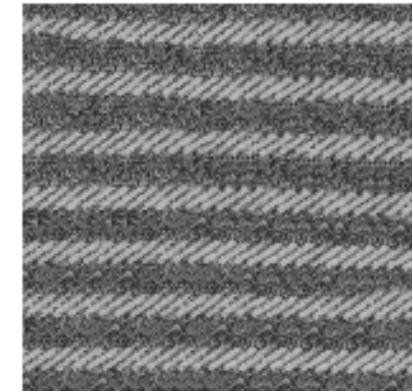
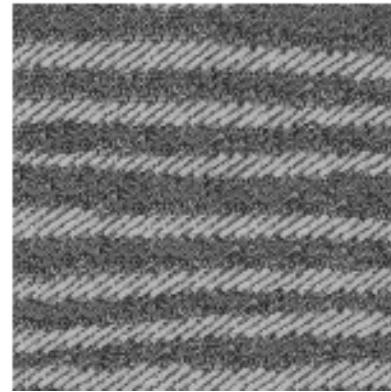
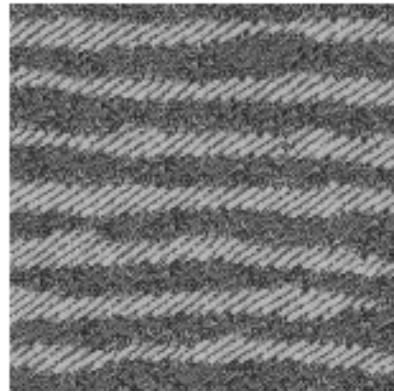
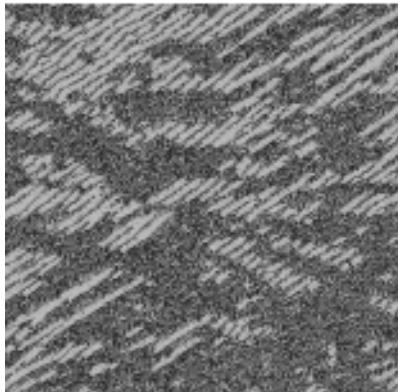
synthesized image

- What is  $P(x|\text{neighborhood of pixels around } x)$  ?
- Find all the windows in the image that match the neighborhood
- To synthesize  $x$ 
  - pick one matching window at random
  - assign  $x$  to be the center pixel of that window
- An **exact** neighbourhood match might not be present, so find the **best** matches using **SSD error** and randomly choose between them, preferring better matches with higher probability

# Neighborhood Window



# Varying Window Size



Increasing window size



# Growing Texture

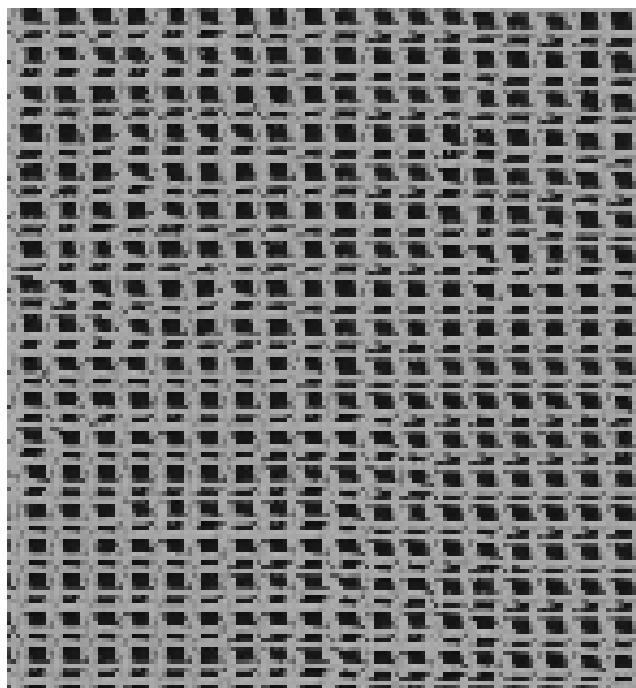
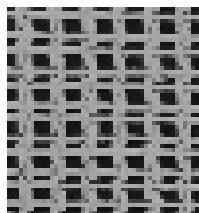
---



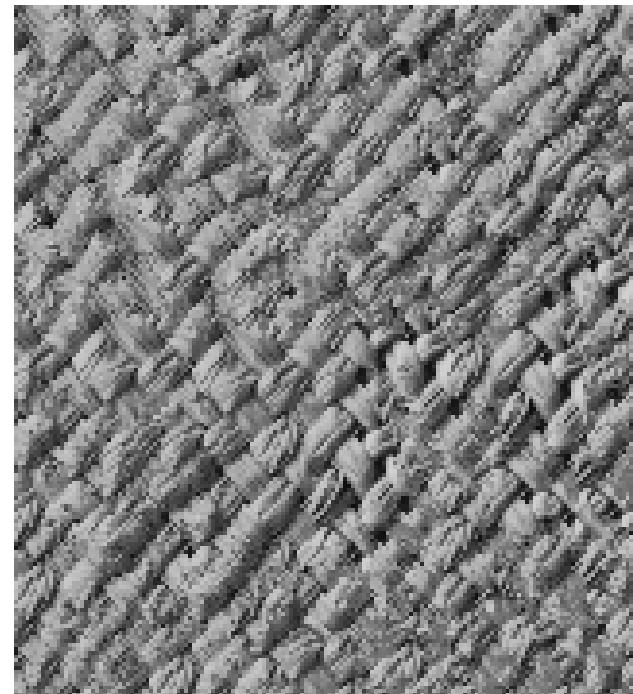
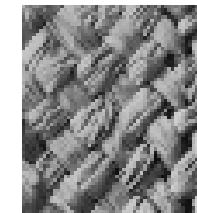
- Starting from the initial image, “grow” the texture one pixel at a time

# Synthesis results

french canvas



rafia weave

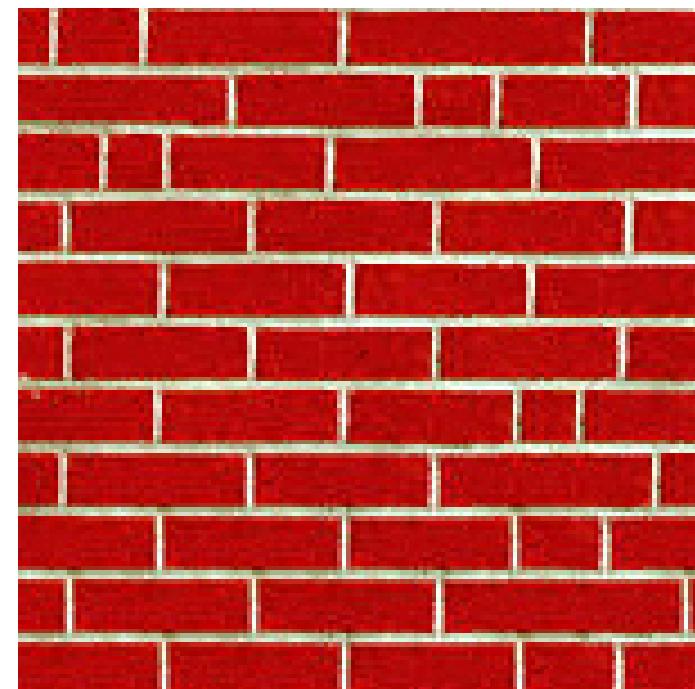
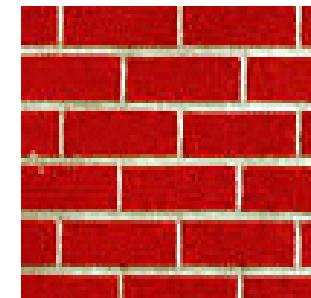


# Synthesis results

white bread



brick wall

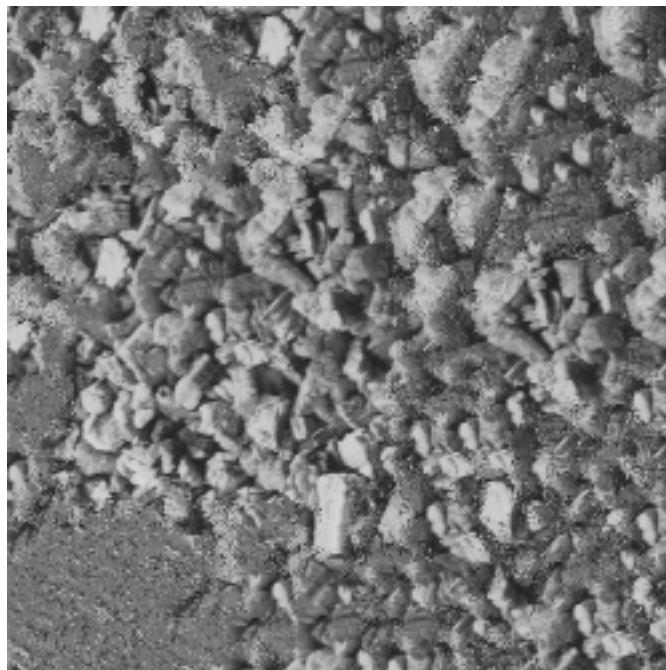


# Synthesis results

uring in the unsensato  
r Dick Gephardt was fai  
rful riff on the looming  
nly asked, "What's your  
tions?" A heartfelt sigh  
story about the emergenc  
es against Clinton. "Boy  
g people about continuin  
ardt began, patiently obs  
s, that the legal system h  
e with this latest tanger

ithairm, them . "Whnephartfe lartifelintomimen  
el ck Clirtioout ormain thartfelins. f aut s anetc  
the ry onst wartfe lck Gephntoomimeationl sigab  
Cliooufit Clinut Cll riff on, hat's yordn, parut tly  
ons yoontonsteht wasked, paim t sahe loo' riff on l  
nskoneploourtfeas leil A nst Clit, "Wleontongal s  
k Cirtioouirtfepe óng pme abegal fartfenstemem  
itiensteneltorydt telemeaphinsverdt was agemer  
ff ons artientont Cling peme as urfe atih, "Boui s  
hal s fartfelt sig pedrl̄dt ske abounutie aboutioo  
itfaone newwas yous aboonthardt thatins fain, ped,  
ains, them, pabout wasy arfut cōuitly d, ln A h  
ole emthrdingbooneme agas fa bontinsyst Clinut  
ory about continst Clipeouinst Cloke agatiff out C  
stome zinemen tly ardt beoraboul n, thenly as t C  
cons faimeme Diontont wat coutlyohgans as fan  
ien, phrtfaul, "Wbaut cout congagal cōmininga  
mifmst Cliiy abon al coountha.ermungairt tf oun  
The looorystan loontieph. intly on, theoplegatick  
ul tatteeontly atie Diontiomi wal s f tbegae ener  
mthahsat's enenhhmas fan, "intchthorw ahons v

# Failure Cases

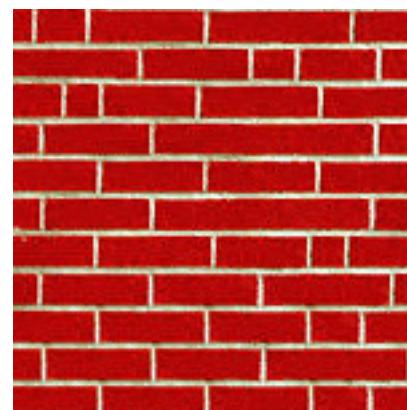
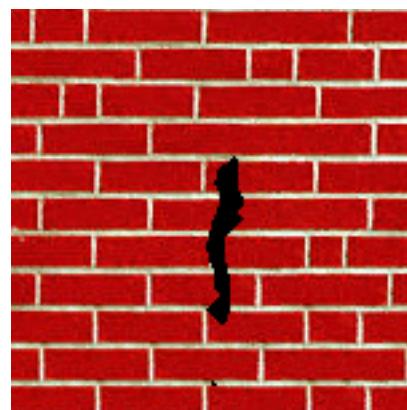
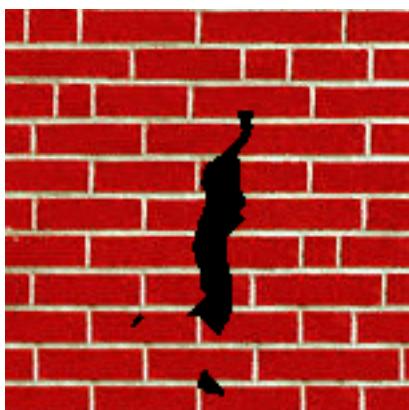


**Growing garbage**

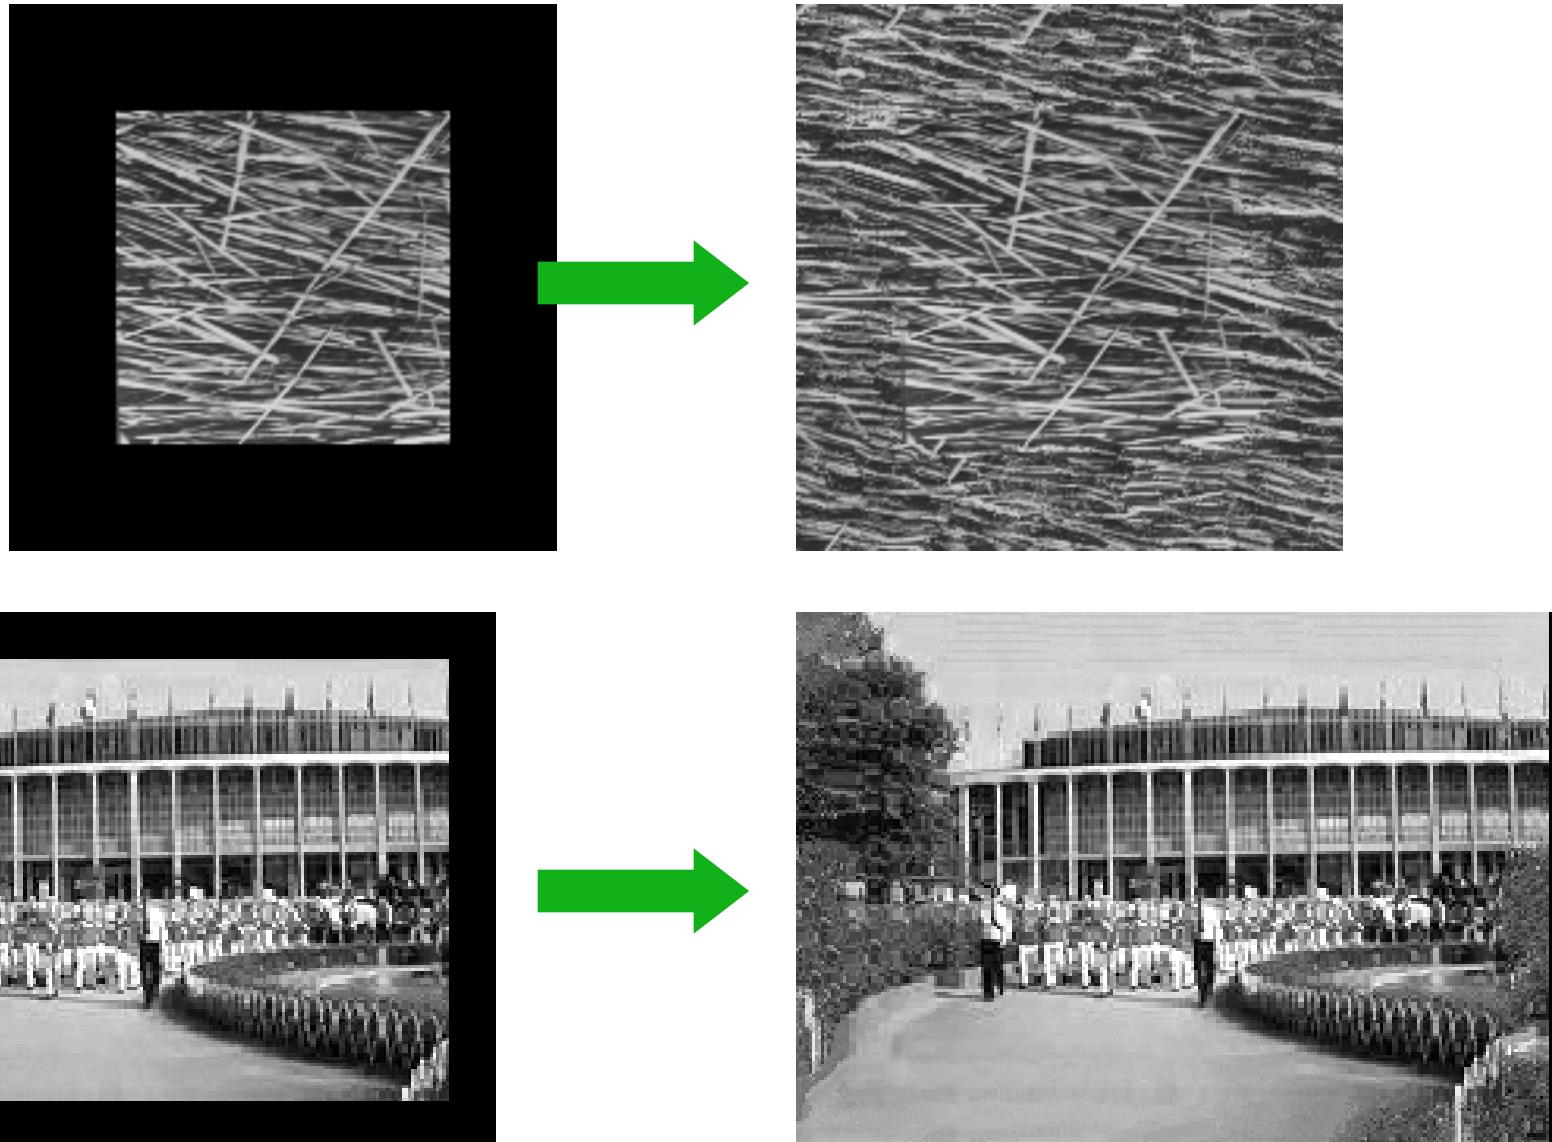


**Verbatim copying**

# Hole Filling



# Extrapolation

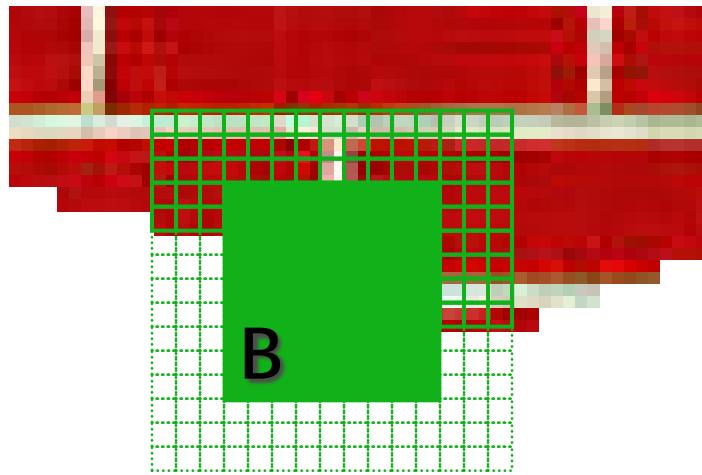


# Efros and Leung algorithm

```
function GrowImage(SampleImage, Image, WindowSize)
    while Image not filled do
        progress = 0
        PixelList = GetUnfilledNeighbors(Image)
        foreach Pixel in PixelList do
            Template = GetNeighborhoodWindow(Pixel)
            BestMatches = FindMatches(Template, SampleImage)
            BestMatch = RandomPick(BestMatches)
            if (BestMatch.error < MaxErrThreshold) then
                Pixel.value = BestMatch.value
                progress = 1
            end
        end
        if progress == 0
            then MaxErrThreshold = MaxErrThreshold * 1.1
    end
    return Image
end
```

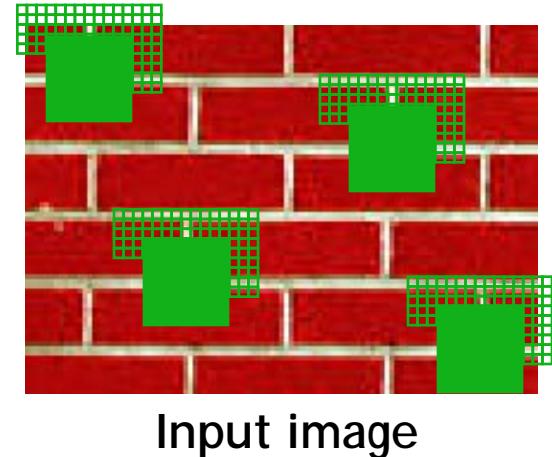
- The Efros & Leung algorithm
  - Simple
  - Surprisingly good results
  - Synthesis is easier than analysis!
  - ...but very slow

# Image Quilting [Efros & Freeman 2001]



Synthesizing a block

non-parametric sampling

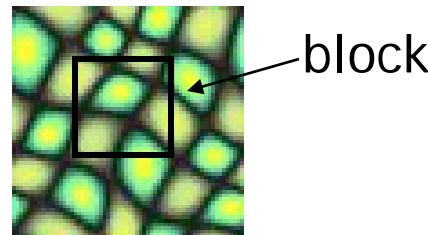


Input image

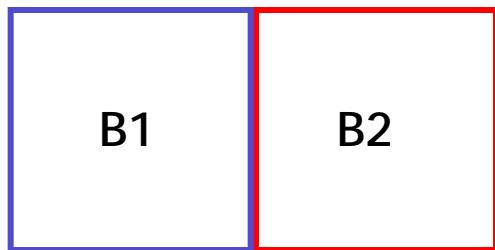
- Observation: neighbor pixels are highly correlated

Idea: unit of synthesis = block

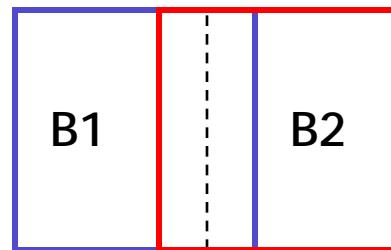
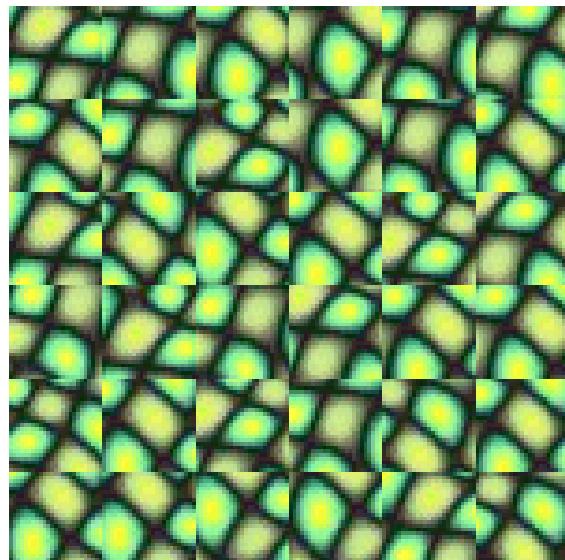
- Exactly the same but now we want  $P(B | N(B))$
- Much faster: synthesize all pixels in a block at once



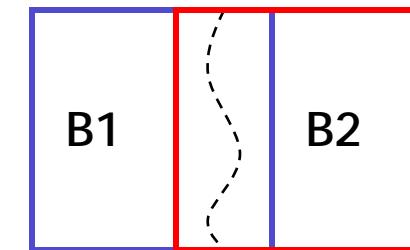
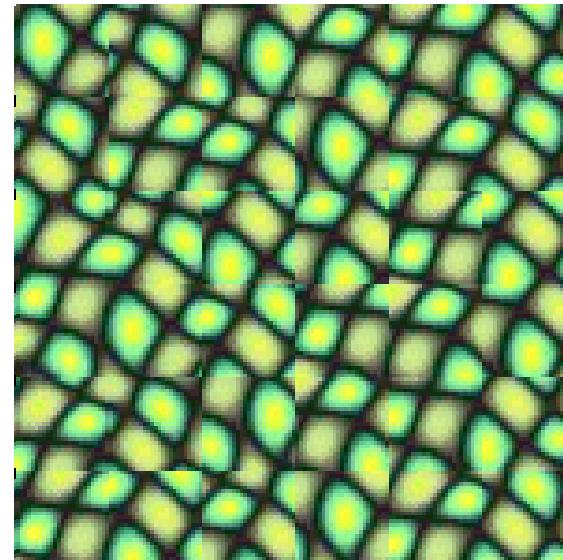
Input texture



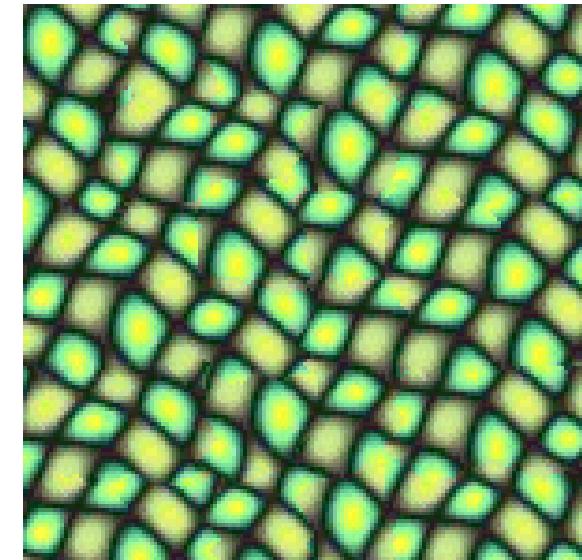
Random placement  
of blocks



Neighboring blocks  
constrained by overlap

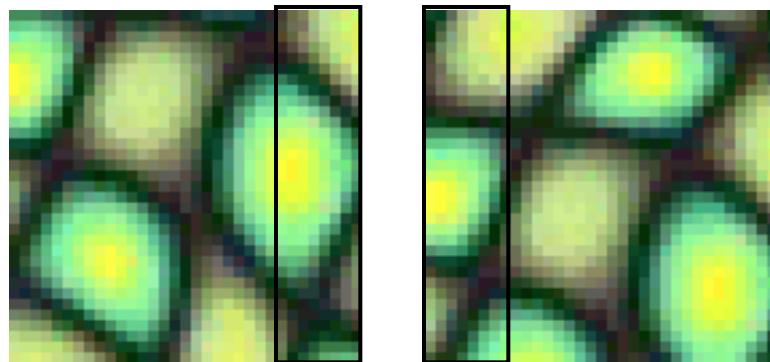


Minimal error  
boundary cut

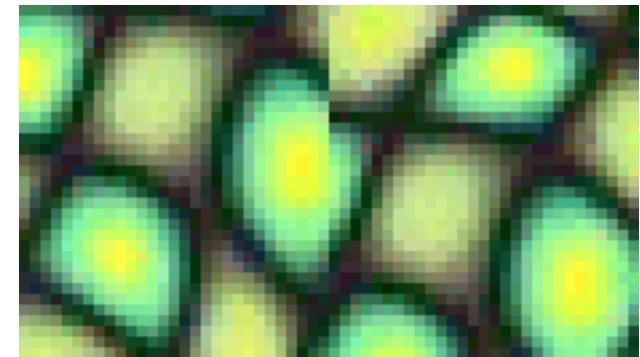


# Minimal error boundary

overlapping blocks

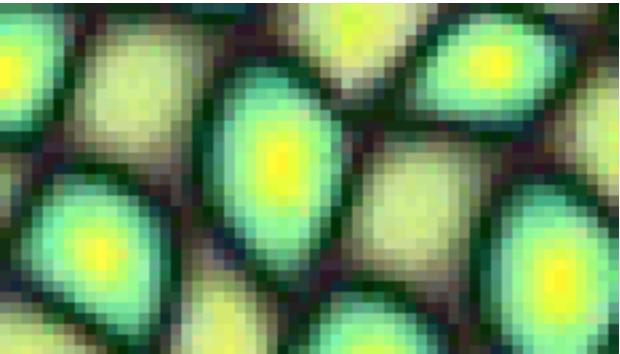


vertical boundary

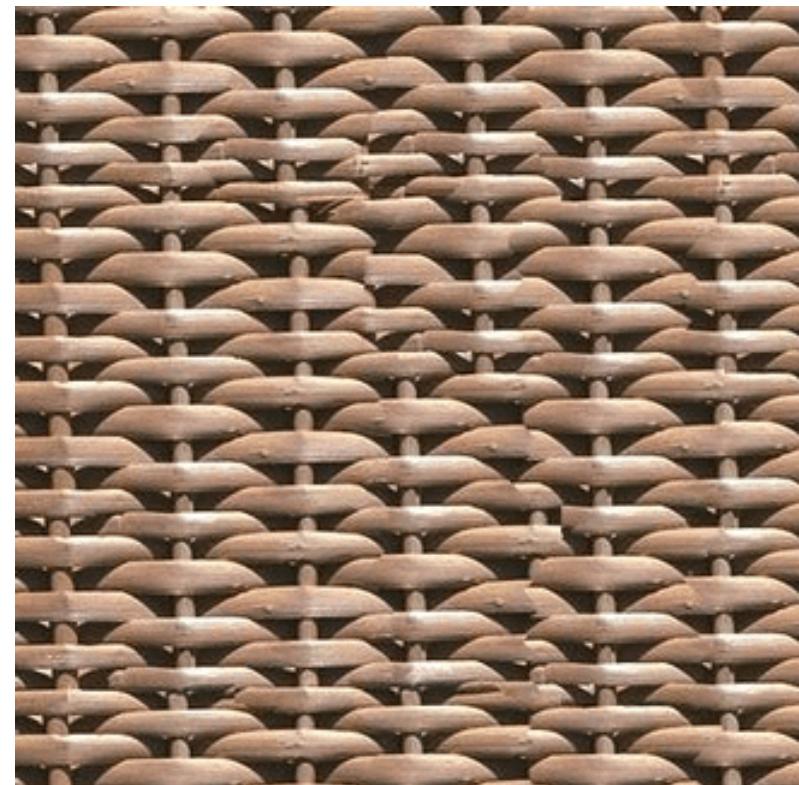
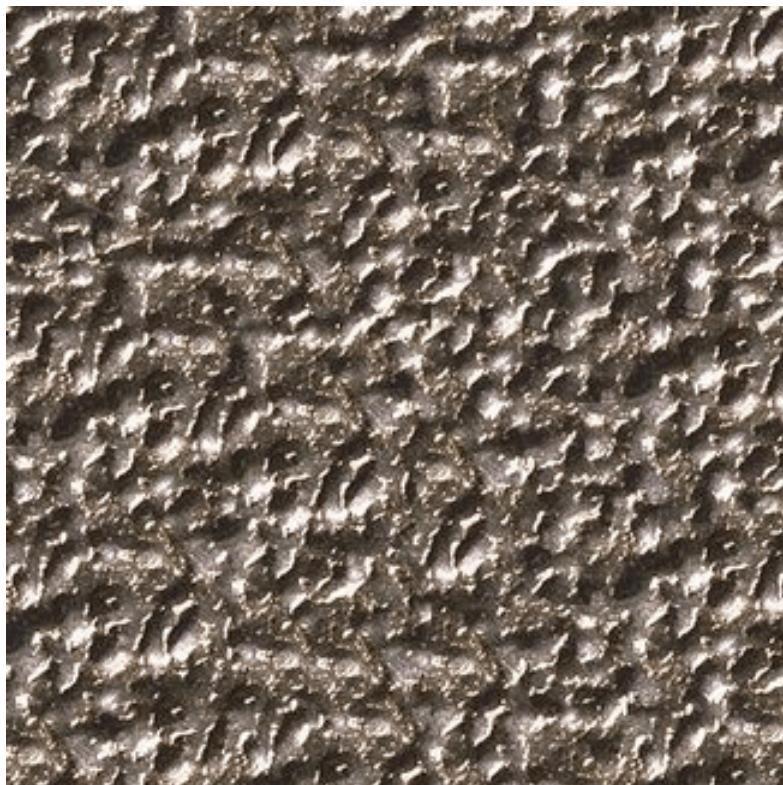
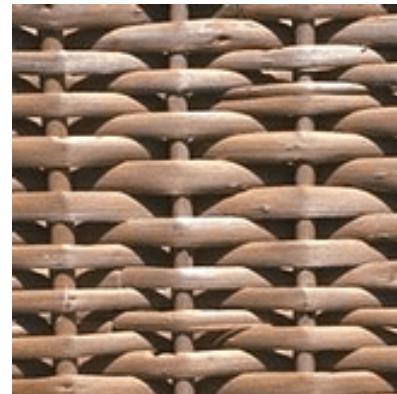


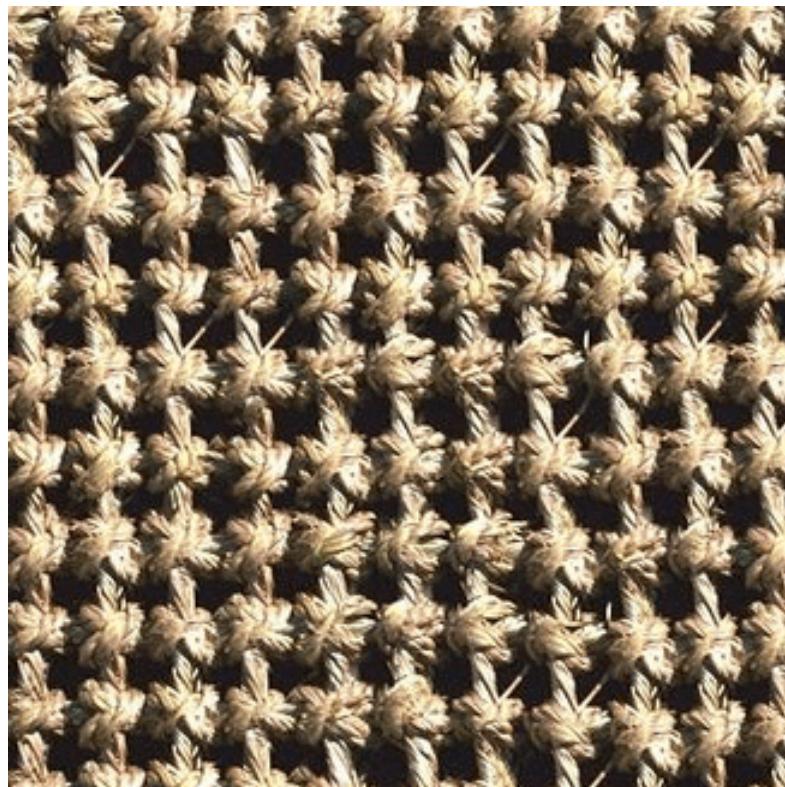
A diagram illustrating the calculation of overlap error. It shows two overlapping blocks from the first image. A bracket groups the two blocks, and a large number '2' is placed above the bracket. Below the bracket is a minus sign. To the right of the minus sign is an equals sign followed by a small image of a red jagged line on a black background, representing the error boundary.

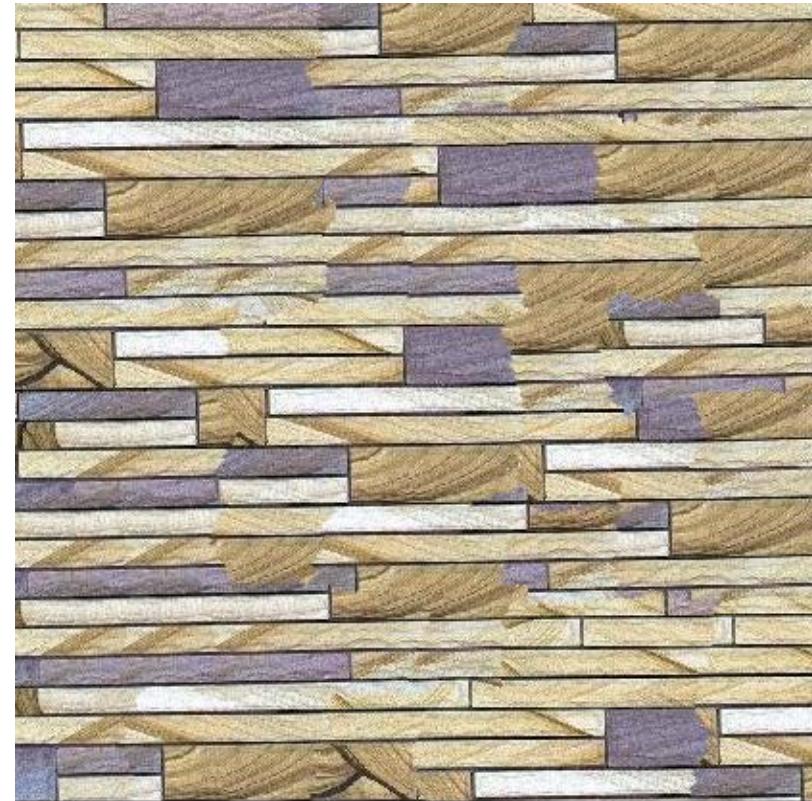
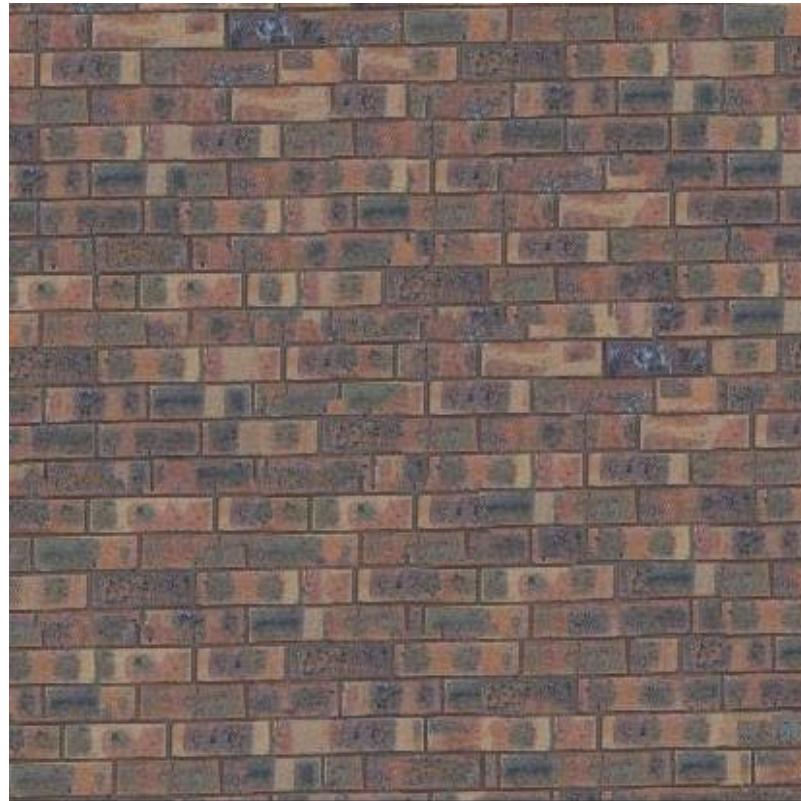
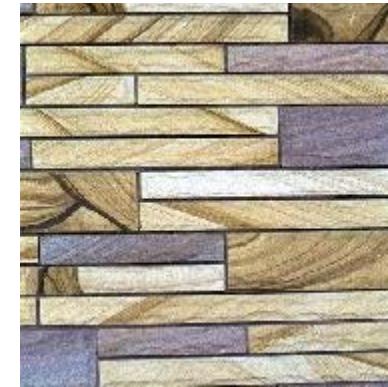
overlap error



min. error boundary

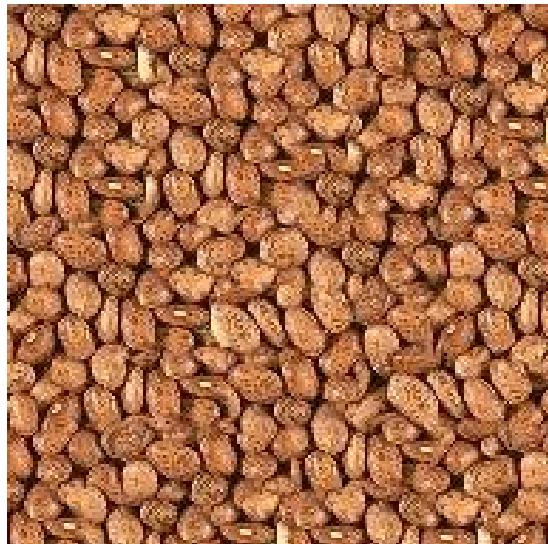
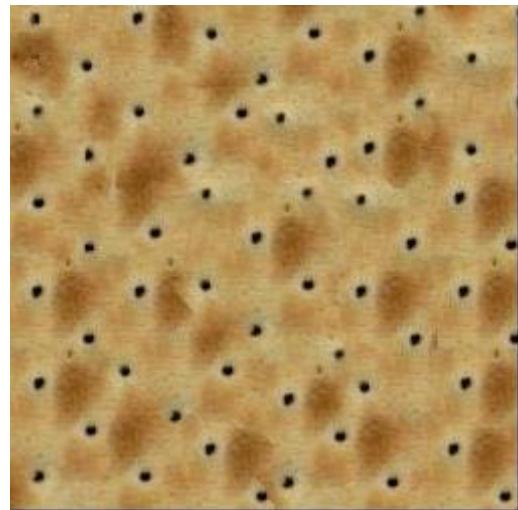
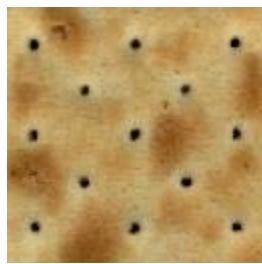








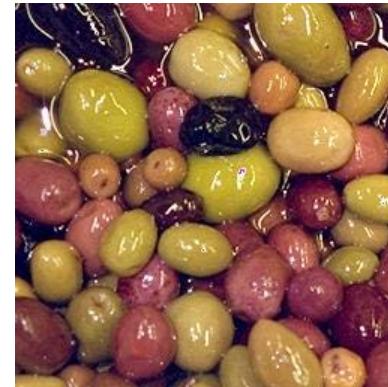






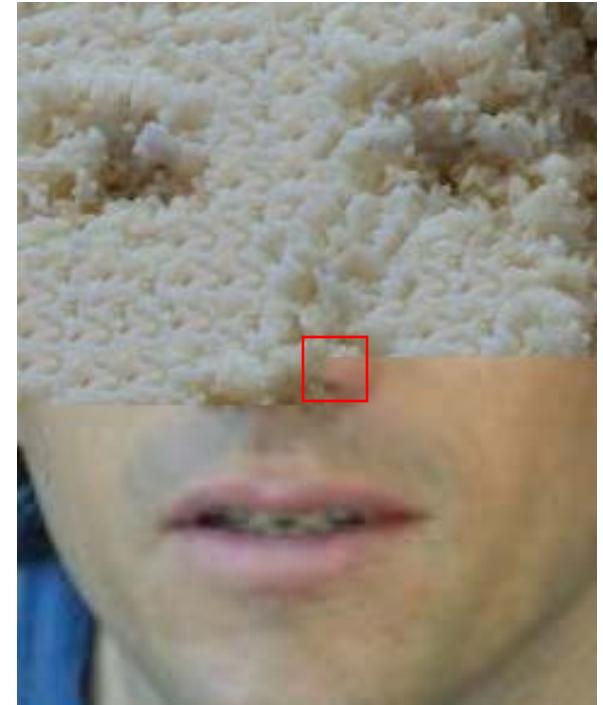
# Failures

(Chernobyl  
Harvest)



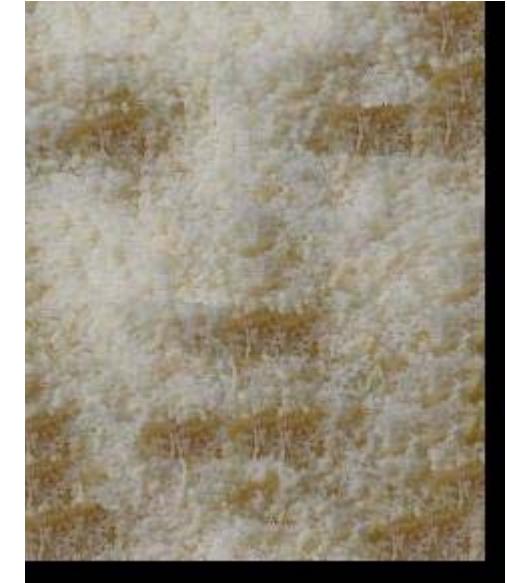
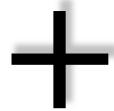
# Texture Transfer

- Take the texture from one object and “paint” it onto another object
  - This requires separating texture and shape
  - That’s HARD, but we can cheat
  - Assume we can capture shape by boundary and rough shading
- Then, just add another constraint when sampling: similarity to underlying image at that spot

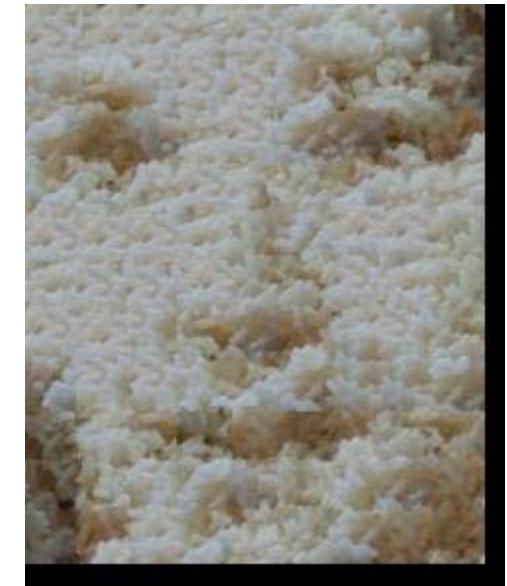




parmesan

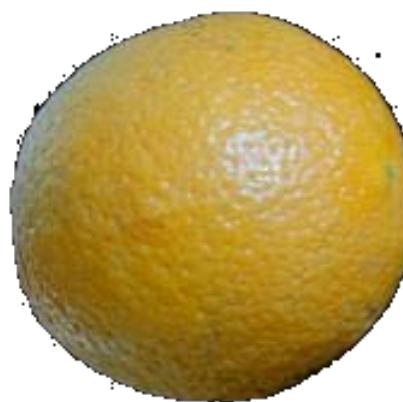


rice

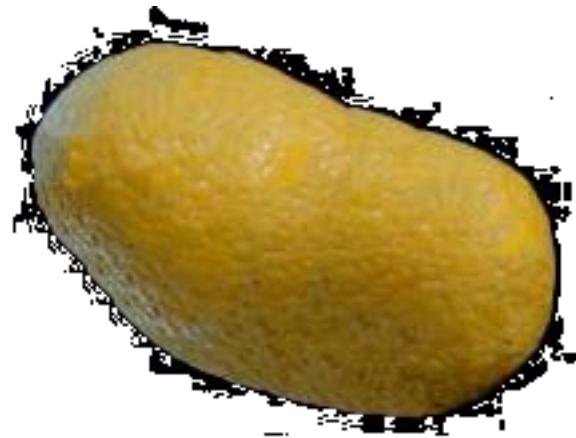


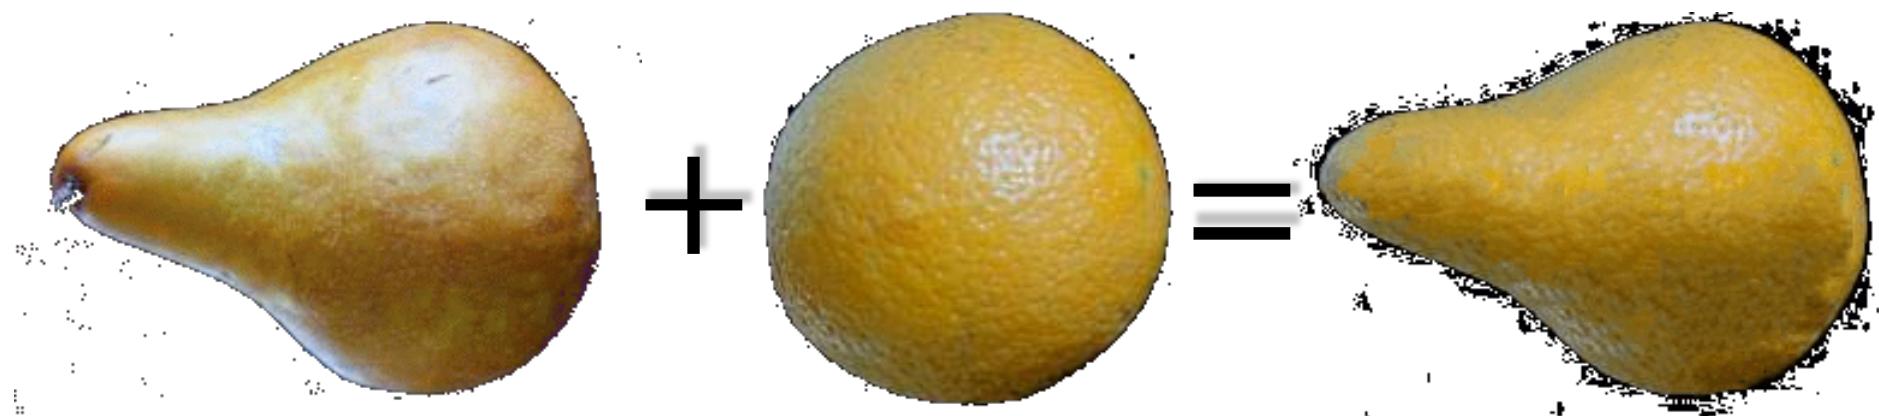


+



=





# (Manual) texture synthesis in the media



# (Manual) texture synthesis in the media



# WHATEVER IT TAKES



# Synthesizing textures when constructing 3d models of archaeological sites

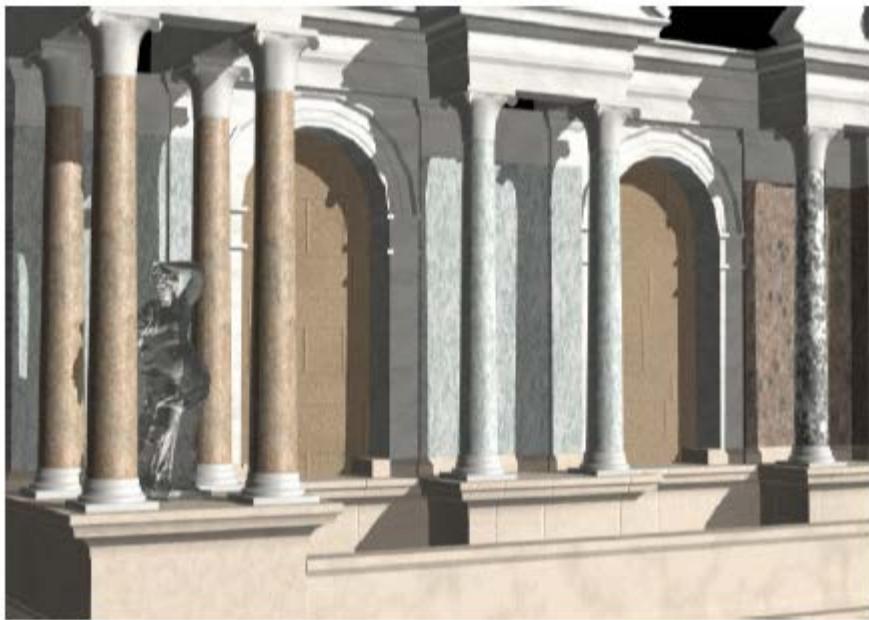


Figure 12. The Nymphaeum at the upper agora of Sagalassos with differently textured pillars. Overview of one half of the building (symmetric)



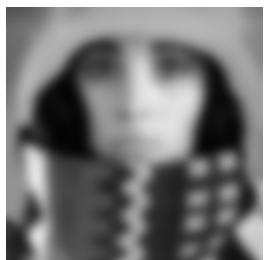
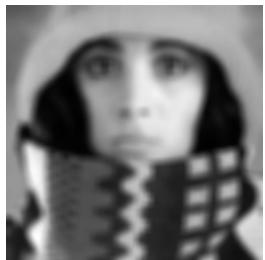
Figure 14. Nymphaeum pillars and back wall fragments in detail

A. Zalesny et al., Realistic Textures for Virtual Anastylosis

# Summary

- Texture is a useful property that is often indicative of materials, appearance cues
- **Texture representations** attempt to summarize repeating patterns of local structure
- **Filter banks** useful to measure redundant variety of structures in local neighborhood
  - Feature spaces can be multi-dimensional
- Neighborhood statistics can be exploited to “sample” or **synthesize** new texture regions
  - Example-based technique

# So far: features and filters



Transforming and  
describing images;  
textures, colors, edges

Kristen Grauman

