

# Computer Vision I - Algorithms and Applications: *Image Formation Process – Part 2*

Carsten Rother

17/11/2013

# Roadmap this lecture

- Geometric primitives and transformations (sec. 2.1.1-2.1.4)
- Geometric image formation (sec 2.1.5, 2.1.6)
  - Pinhole camera
  - Lens effects
- Photometric image formation process (sec 2.2)
- The Human eye
- Camera Types and Hardware (sec 2.3)
- Appearance based matching

# Reminder: Pinhole Camera (Geometry)

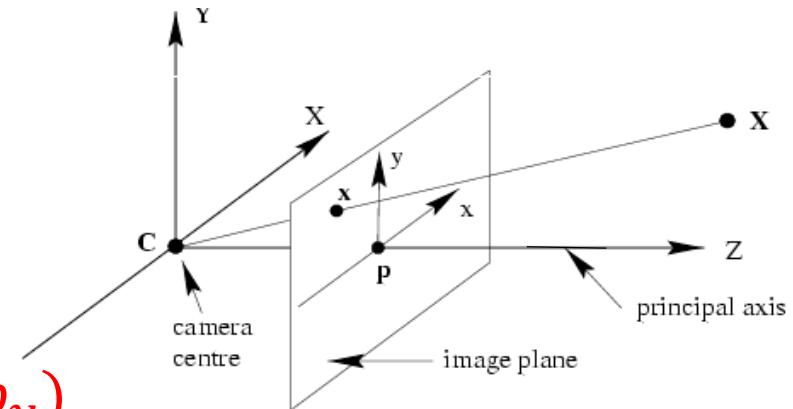
- Camera matrix  $P$  has 11 DoF

- Intrinsic parameters

- Principal point coordinates  $(p_x, p_y)$
- Focal length  $f$
- Pixel magnification factors  $m$
- Skew (non-rectangular pixels)  $s$

- Extrinsic parameters

- Rotation  $R$  (3DoF) and translation  $\tilde{C}$  (3DoF) relative to world coordinate system

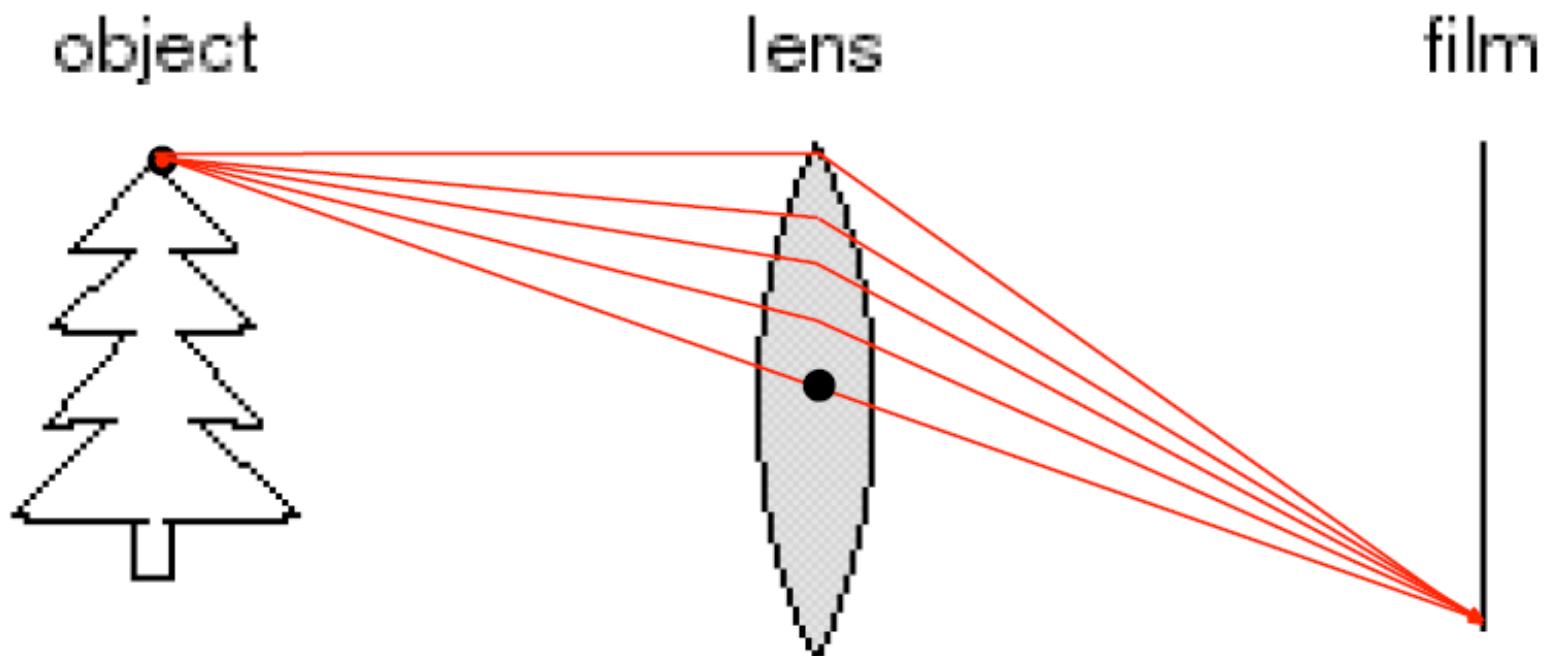


$$x = P X$$

$$x = K R (I_{3 \times 3} | - \tilde{C}) X$$

$$K = \begin{bmatrix} f & s & p_x \\ 0 & mf & p_y \\ 0 & 0 & 1 \end{bmatrix}$$

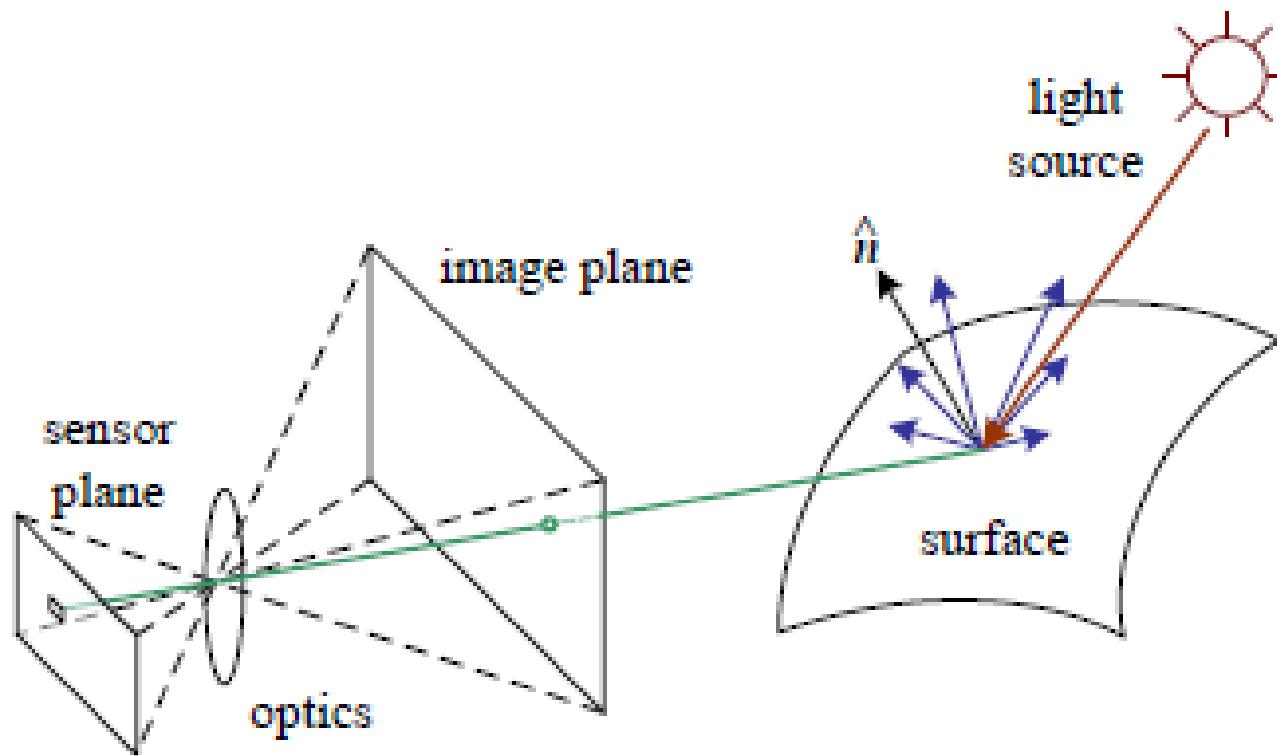
# Reminder: Lens effects



# Roadmap this lecture

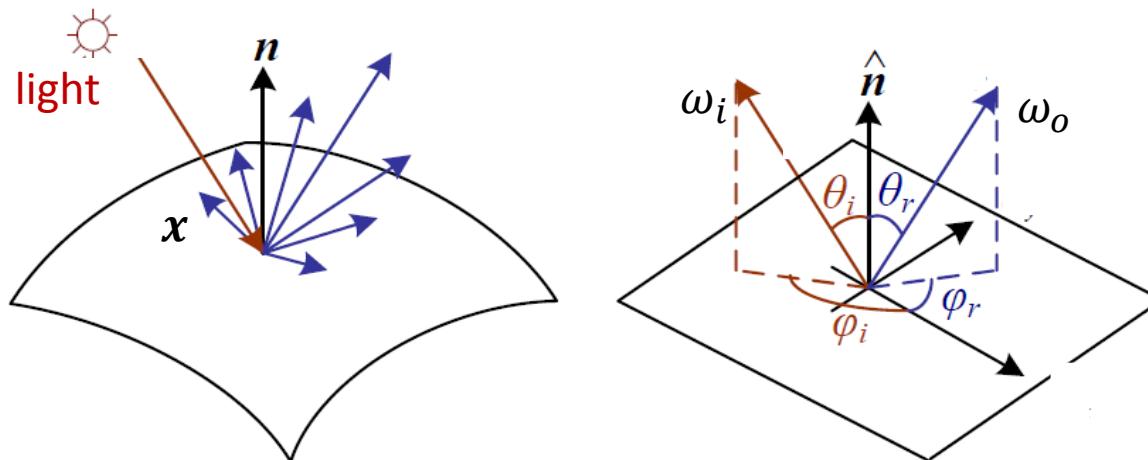
- Geometric primitives and transformations (sec. 2.1.1-2.1.4)
- Geometric image formation (sec 2.1.5, 2.1.6)
  - Pinhole camera
  - Lens effects
- Photometric image formation process (sec 2.2)
- The Human eye
- Camera Types and Hardware (sec 2.3)
- Appearance based matching

# Image formation Process



# Model of the Surface

BRDF (Bi-directional Reflectance Function)



Rendering equation:

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

Outgoing energy:

$t$  – time

$\mathbf{x}$  – position

$\lambda$  – wavelength

$\omega_i / \omega_o$  - incoming / outgoing direction

Incoming light

BRDF function

(4DoF only  $\omega_i, \omega_o$  considered)

Fore-shorting angle  
(always  $\geq 0$ )

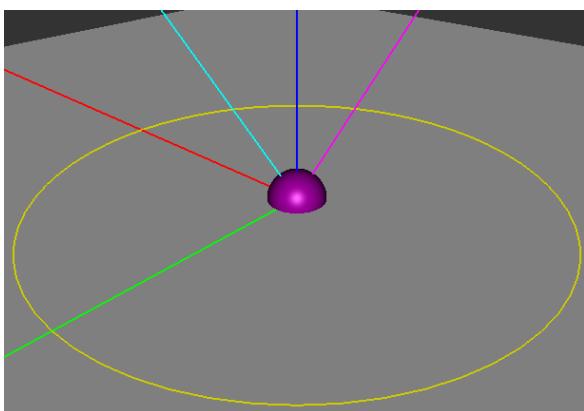
# Example: diffuse illumination

## Rendering equation

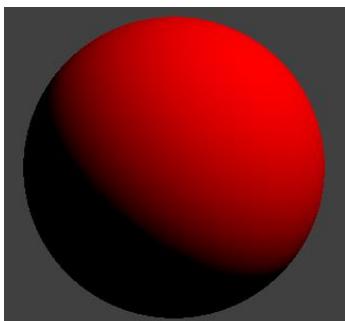
$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

0                    constant

## DIFFUSE



$f_r$  same for all  $\omega_o$



## Single Light source:

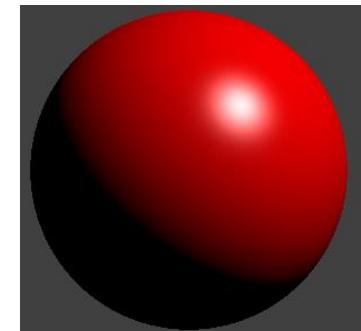
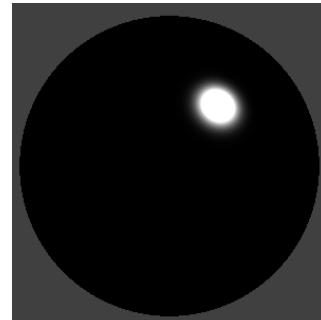
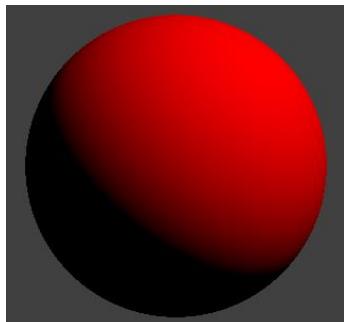
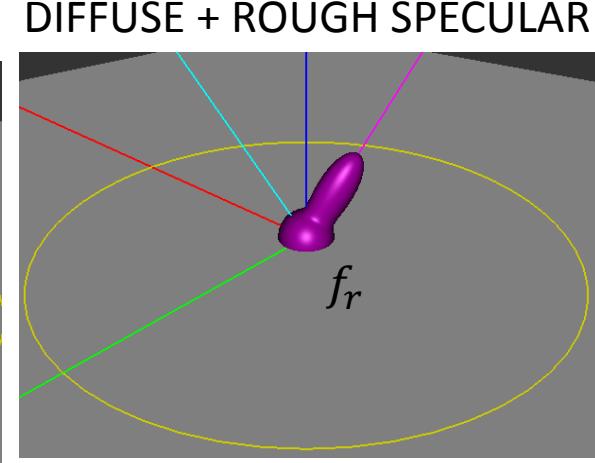
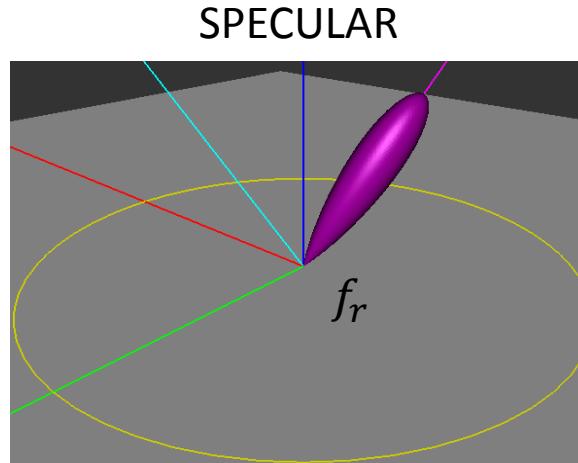
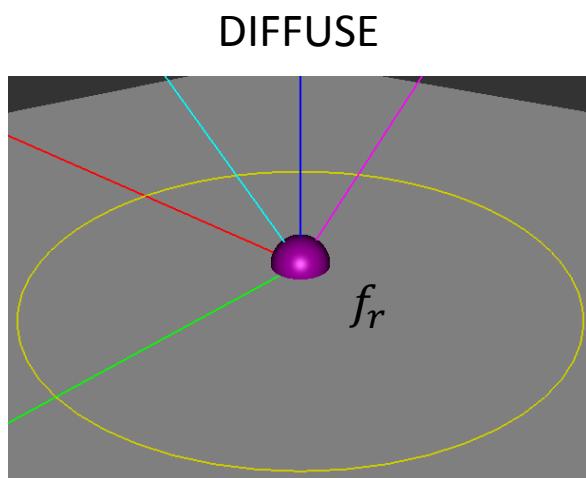
Shading effects come from:  $\max(0, w_i \cdot n)$

# Examples: BRDFs

Rendering equation (single light source)

$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \cancel{0} \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

constant  
(fixed  $\omega_i$ )



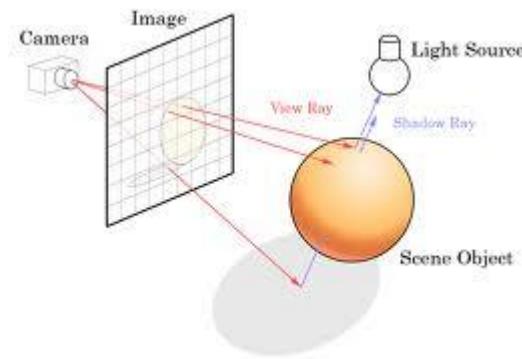
Single light source (top, right)

# Inside a graphics engine

- Rendering equation

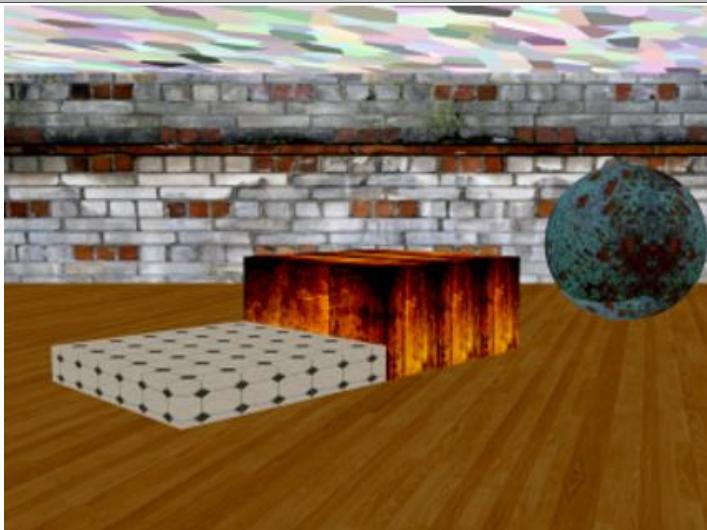
$$L_o(\mathbf{x}, \omega_o, \lambda, t) = L_e(\mathbf{x}, \omega_o, \lambda, t) + \int_{\Omega} f_r(\mathbf{x}, \omega_i, \omega_o, \lambda, t) L_i(\mathbf{x}, \omega_i, \lambda, t) (\omega_i \cdot \mathbf{n}) d\omega_i$$

- Ray tracing

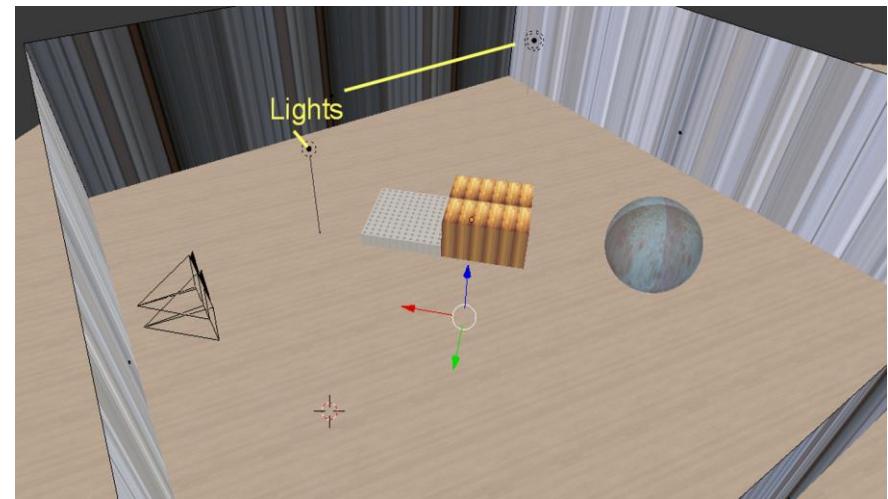


- Radiosity

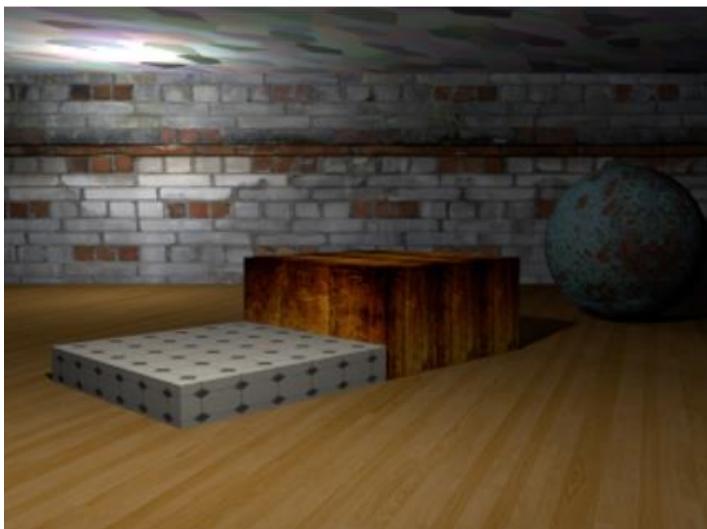
# Example



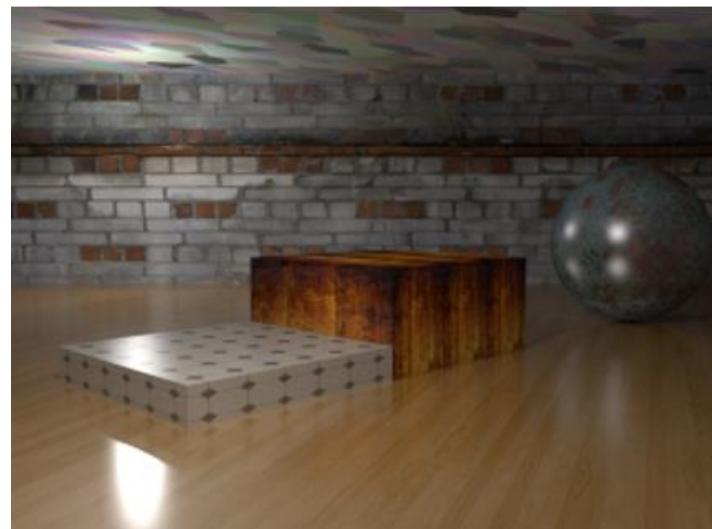
Texture only



3D scene in blender



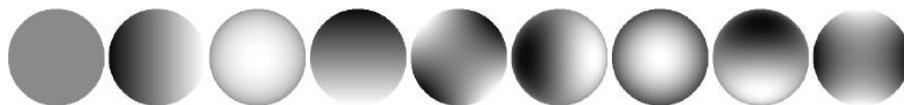
Direct light



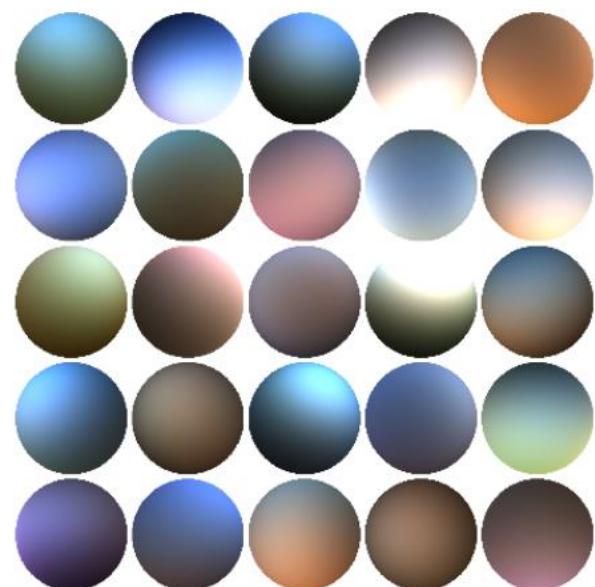
Global (direct + indirect) light

# How to model light?

- Point light source
- Area light source
- Non-local illumination situation can be approximated with spherical harmonics



First 9 spherical harmonics (gray scale)



They can produce this

# Capturing BRDF / BTF capture

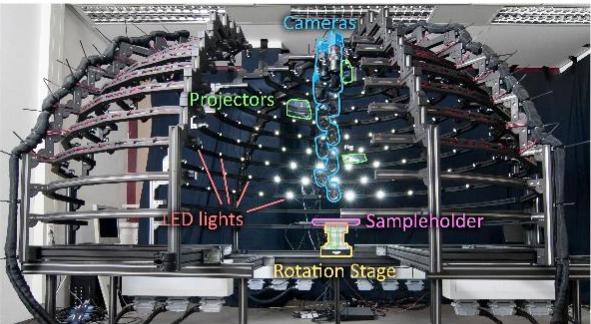


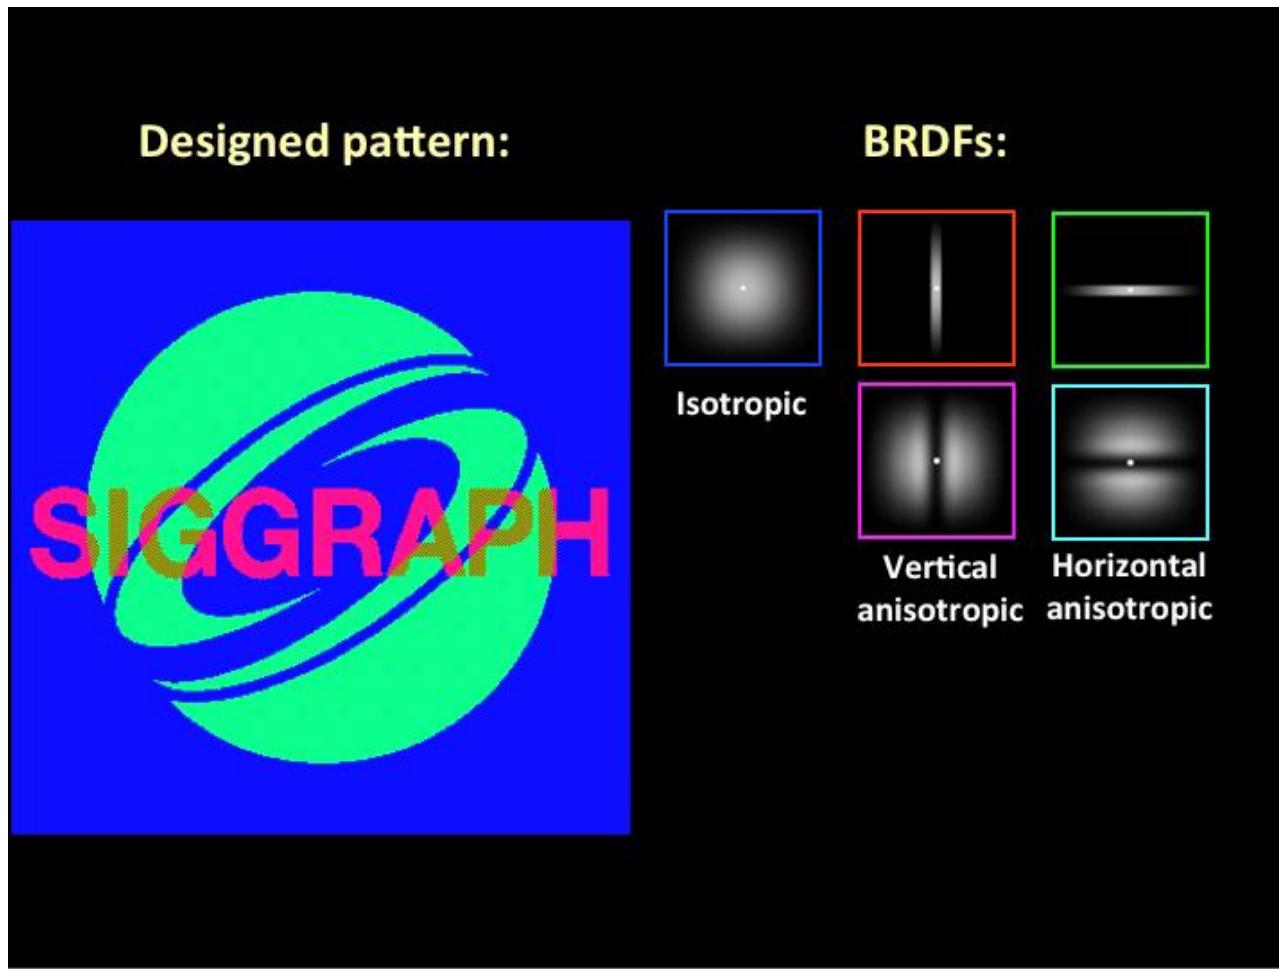
Figure 1: The DOME II BTF acquisition setup. One quarter has been slid open to expose the view on the inside.



BTF – Bidirectional Texture function. At every spatial location  $x$  a different BRDF

[Christopher Schwartz, Ralf Sarlette, Michael Weinmann and Reinhard Klein]

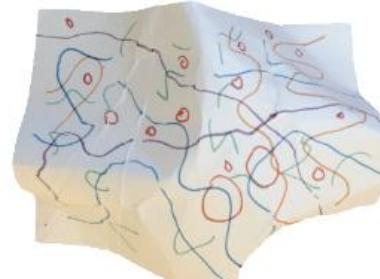
# Fabricating BRDFs



[Levin et al. Siggraph 2013]

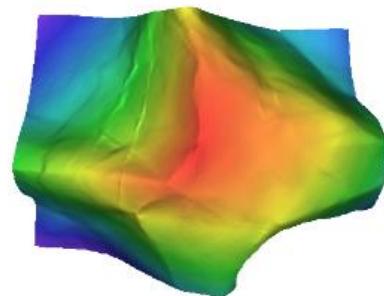
# Single Image

Input:

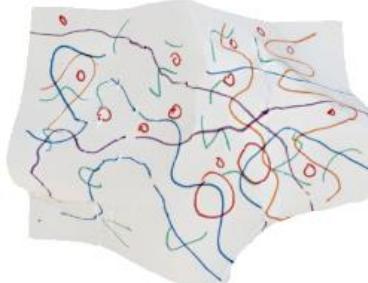


Image

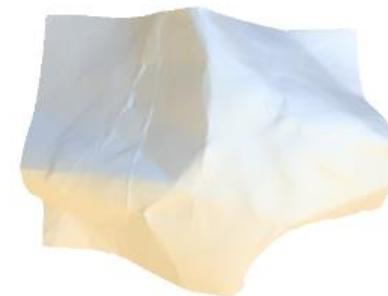
Output:



Shape



Reflectance



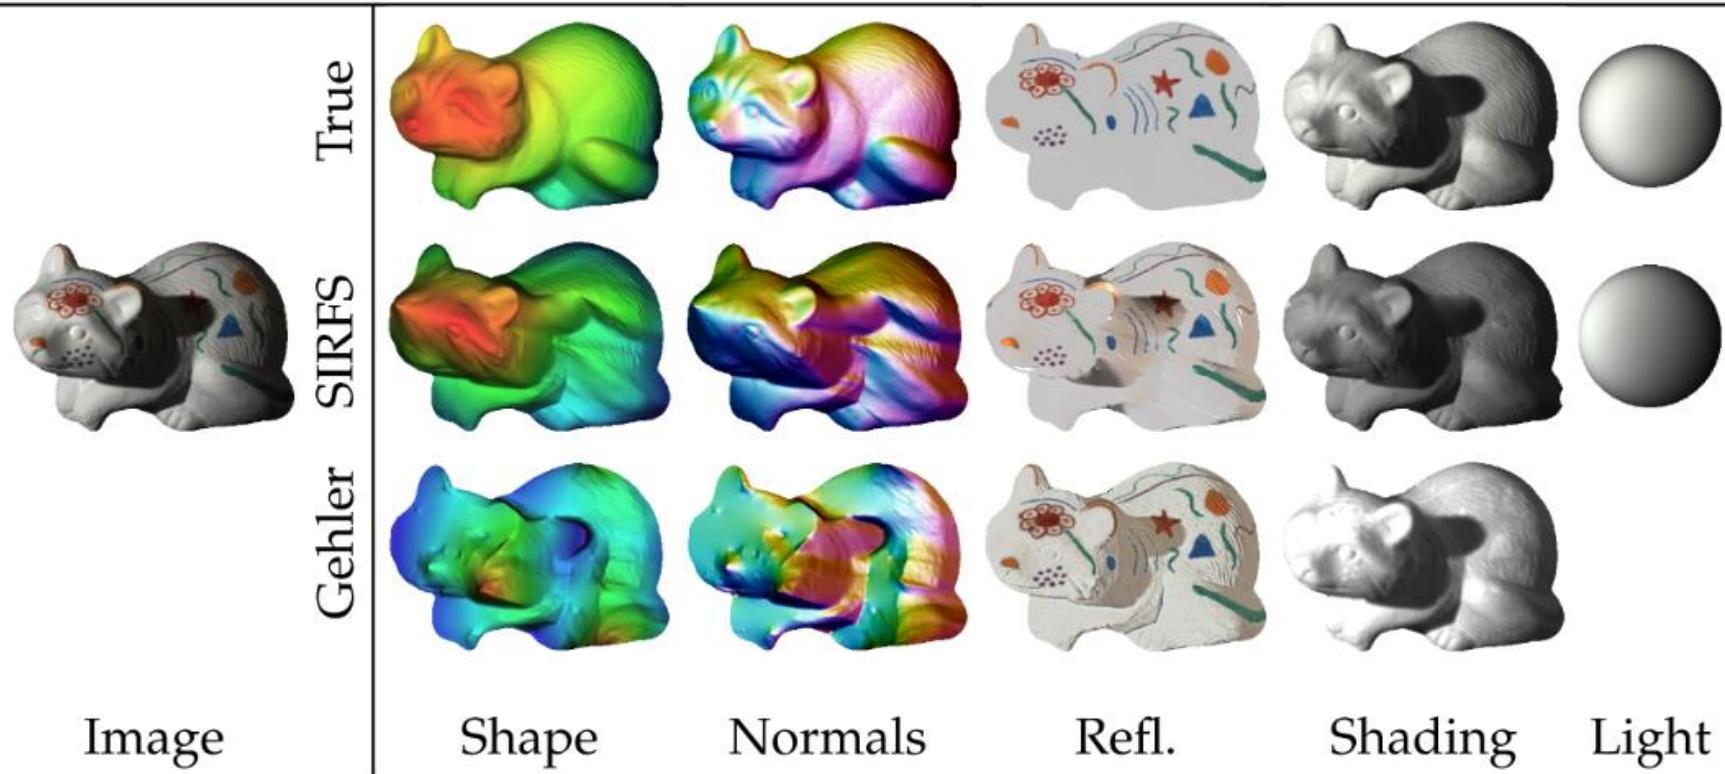
Shading



Illumination

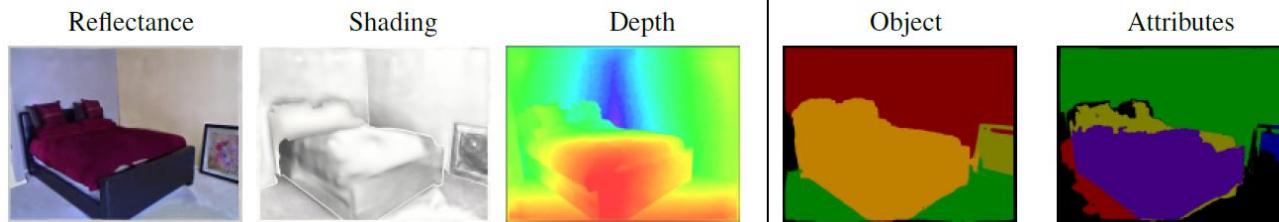
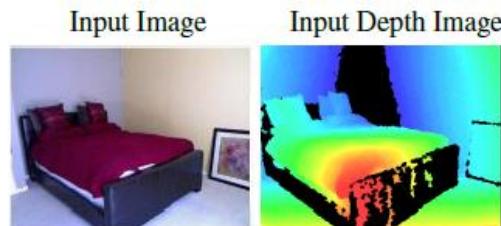
[Barron and Malik 2012]

# Single Image ... far from being solved



[Barron and Malik 2012]

# Reconstruction and recognition



Object-color coding

B-ground Wall Floor Picture Cabinet Chair Table Window Ceiling Lamp Counter Bed Blinds Bookshelf Curtain Monitor

Attribute-color coding

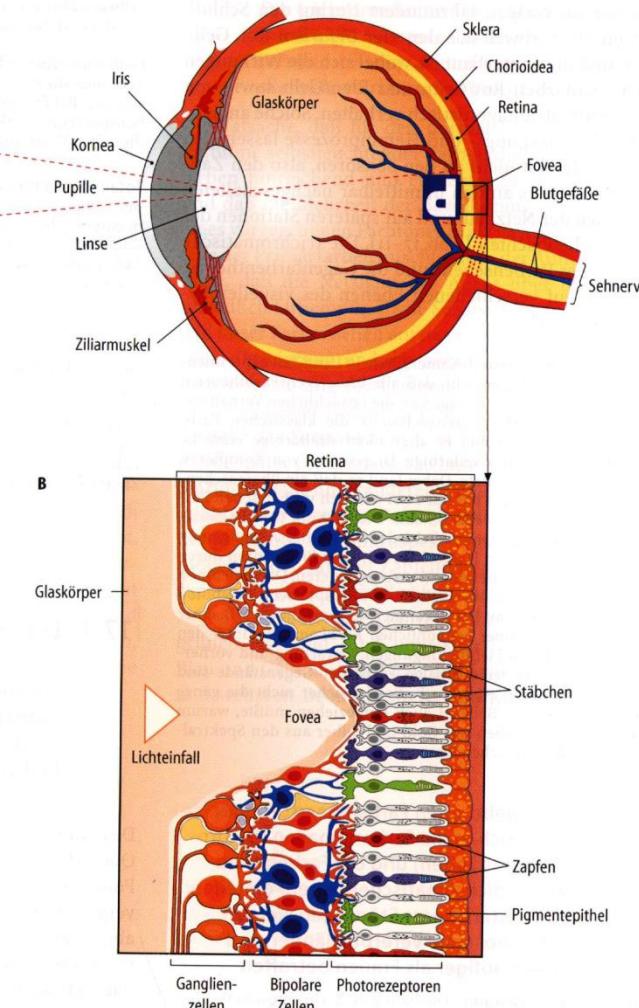
B-ground Wood Painted Cotton Glass Brick Plastic Shiny Dirty Skin Metal Cloth More than one attribute label

[Vinett, Rother, Torr, NIPS '13]

# Roadmap this lecture

- Geometric primitives and transformations (sec. 2.1.1-2.1.4)
- Geometric image formation (sec 2.1.5, 2.1.6)
  - Pinhole camera
  - Lens effects
- Photometric image formation process (sec 2.2)
- The Human eye
- Camera Types and Hardware (sec 2.3)
- Appearance based matching

# The Human eye



- The retina contains different types of sensors: cones “Zapfen” (colors, 6 million) and rods “Stäbchen” (gray levels, 120 million)
- The resolution is much higher In fovea centralis
- Light first passes through the layer of neurons before it reaches photo sensors (smoothing). Only in the “fovea centralis” the light hits directly the photo sensors
- Signal goes out the other way:  
Retina → Ganglion cells (1 million) → Optic nerve → 1 MPixel Camera ?

# Spatial resolution

## Spatial resolution

2MP Camera, far from the screen

# Spatial resolution

5MP Camera, close to the screen

# Spatial resolution (secrets)

Spatial resolution

Image Processing: Human seeing

5

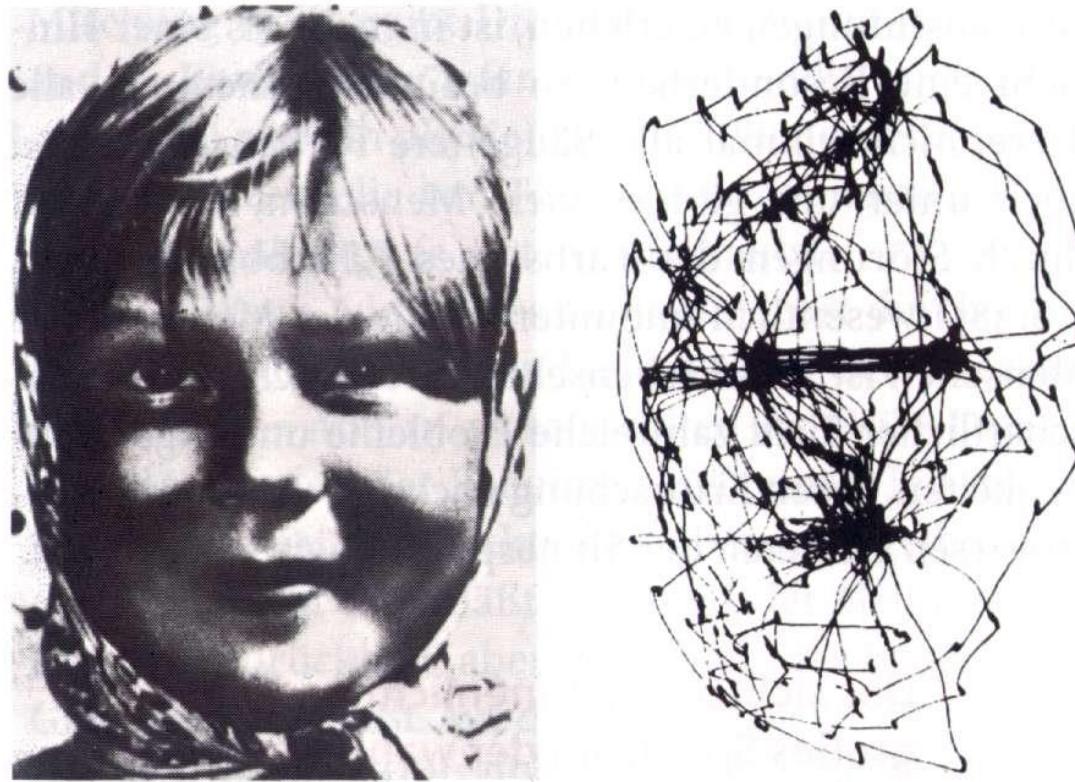
Spatial resolution

Image Processing: Human seeing

5

- The resolution is much higher In fovea centralis
- The Information is pre-processed by Ganglion cells  
(Compare:  $3072 \times 2304 = 7\text{MPixel}$ , 2.4 MB RGB JPEG lossless)
- No still image, but a „Video“ (super-resolution)
- Scanning technique – Saccades

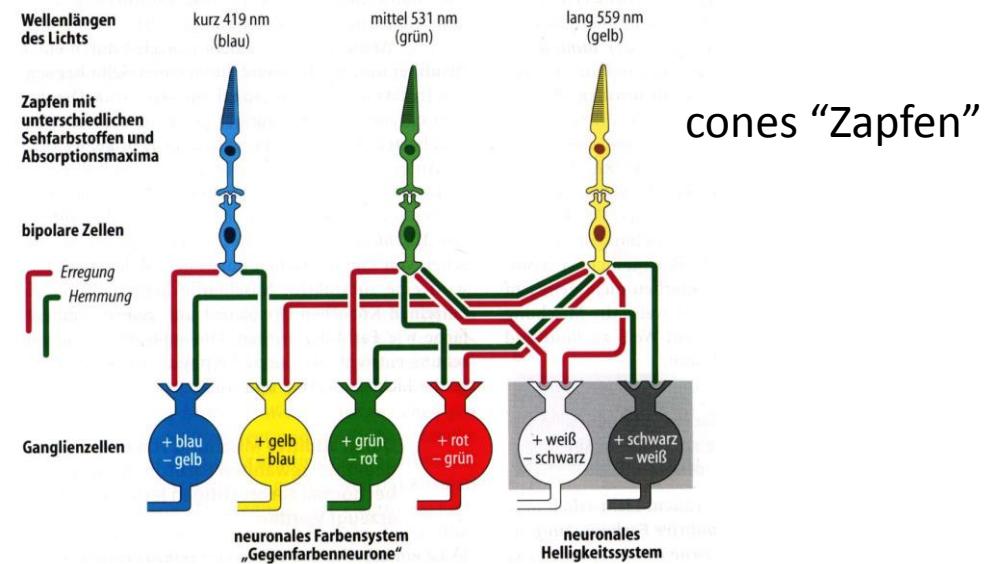
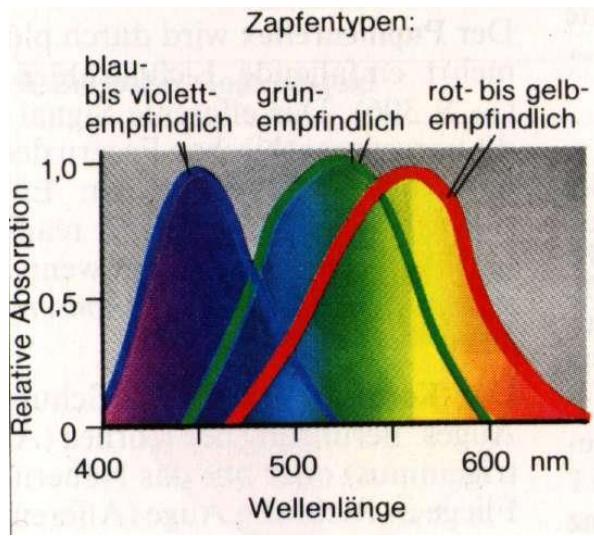
# Eye Saccades



Eyes never move uniformly, but jump in **saccades**  
(approximately 15-100 ms duration between fixation points)

Saccades are driven by the “importance” of the scene parts  
(eyes, mouth etc).

# The Human eye



What is light?

Spectrum, i.e. a function of the wavelength

Different colors can be computed by adding /subtracting signal of rods

Spectral resolution of the eye is relatively bad due to projection  $\infty \rightarrow 3$

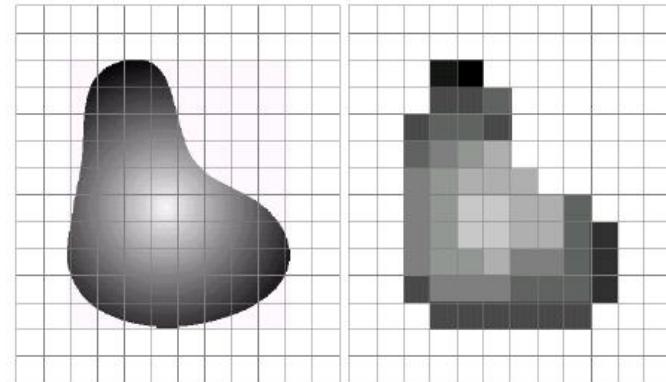
# Roadmap this lecture

- Geometric primitives and transformations (sec. 2.1.1-2.1.4)
- Geometric image formation (sec 2.1.5, 2.1.6)
  - Pinhole camera
  - Lens effects
- Photometric image formation process (sec 2.2)
- The Human eye
- Camera Types and Hardware (sec 2.3)
- Appearance based matching

# Camera types

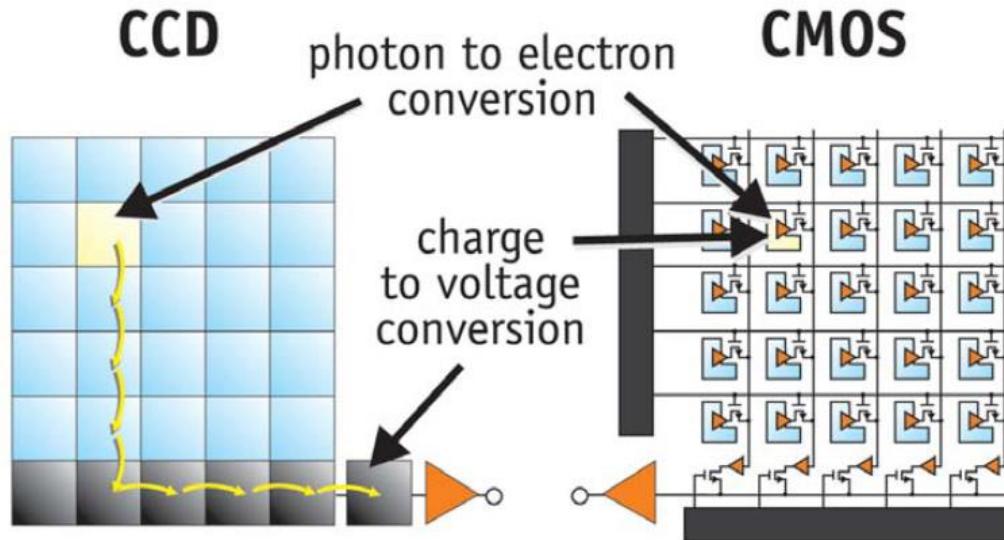
- RGB cameras
- Depth cameras:
  - Passive RGB stereo
  - Active Structured light
  - Active Time of Flight
- Lightfield cameras

# Digital RGB cameras



- A digital camera replaces film with a **sensor array**
- Each cell in the sensor array is **light-sensitive diode** that converts photons to electrons
- **Two common types:**
  - Charge Coupled Device (CCD)
  - Complementary metal oxide semiconductor (CMOS)

# CCD versus CMOS

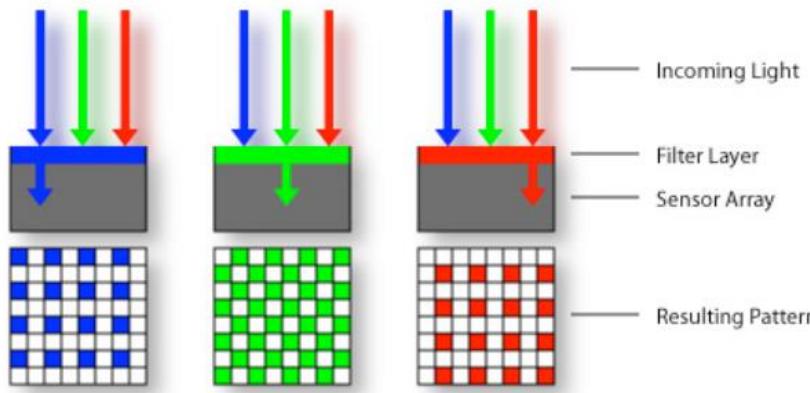
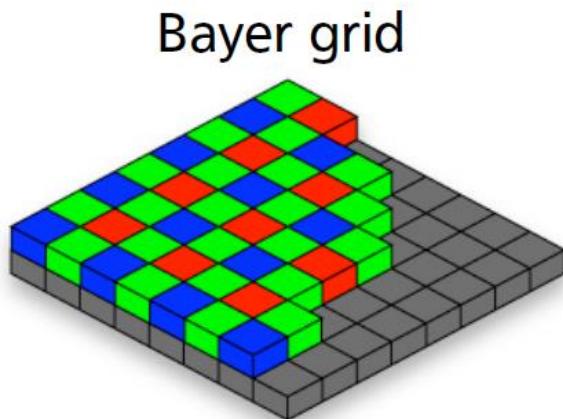


- CCD: move charge from pixel to pixel and then converts to voltage and digital signal
  - Negative: more expensive
- CMOS: converts to voltage inside the pixel
  - Negative: slower to read out image ("rolling shutter" effect)

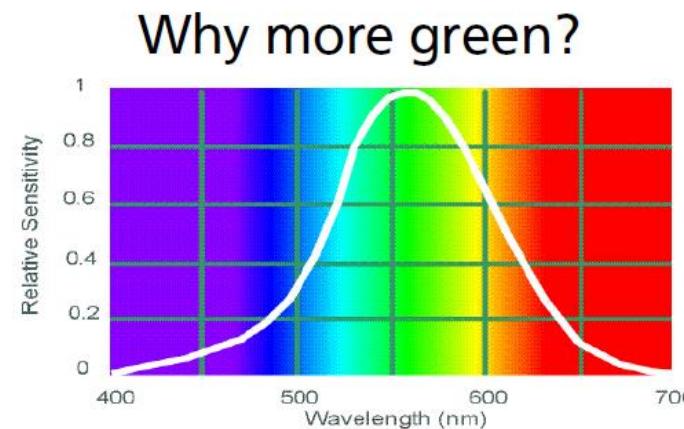
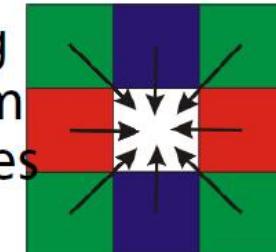


See details on: [http://www.dalsa.com/shared/content/pdfs/CCD\\_vs\\_CMOS\\_Litwiller\\_2005.pdf](http://www.dalsa.com/shared/content/pdfs/CCD_vs_CMOS_Litwiller_2005.pdf)

# Color - Filters



Estimate missing components from neighboring values (demosaicing)

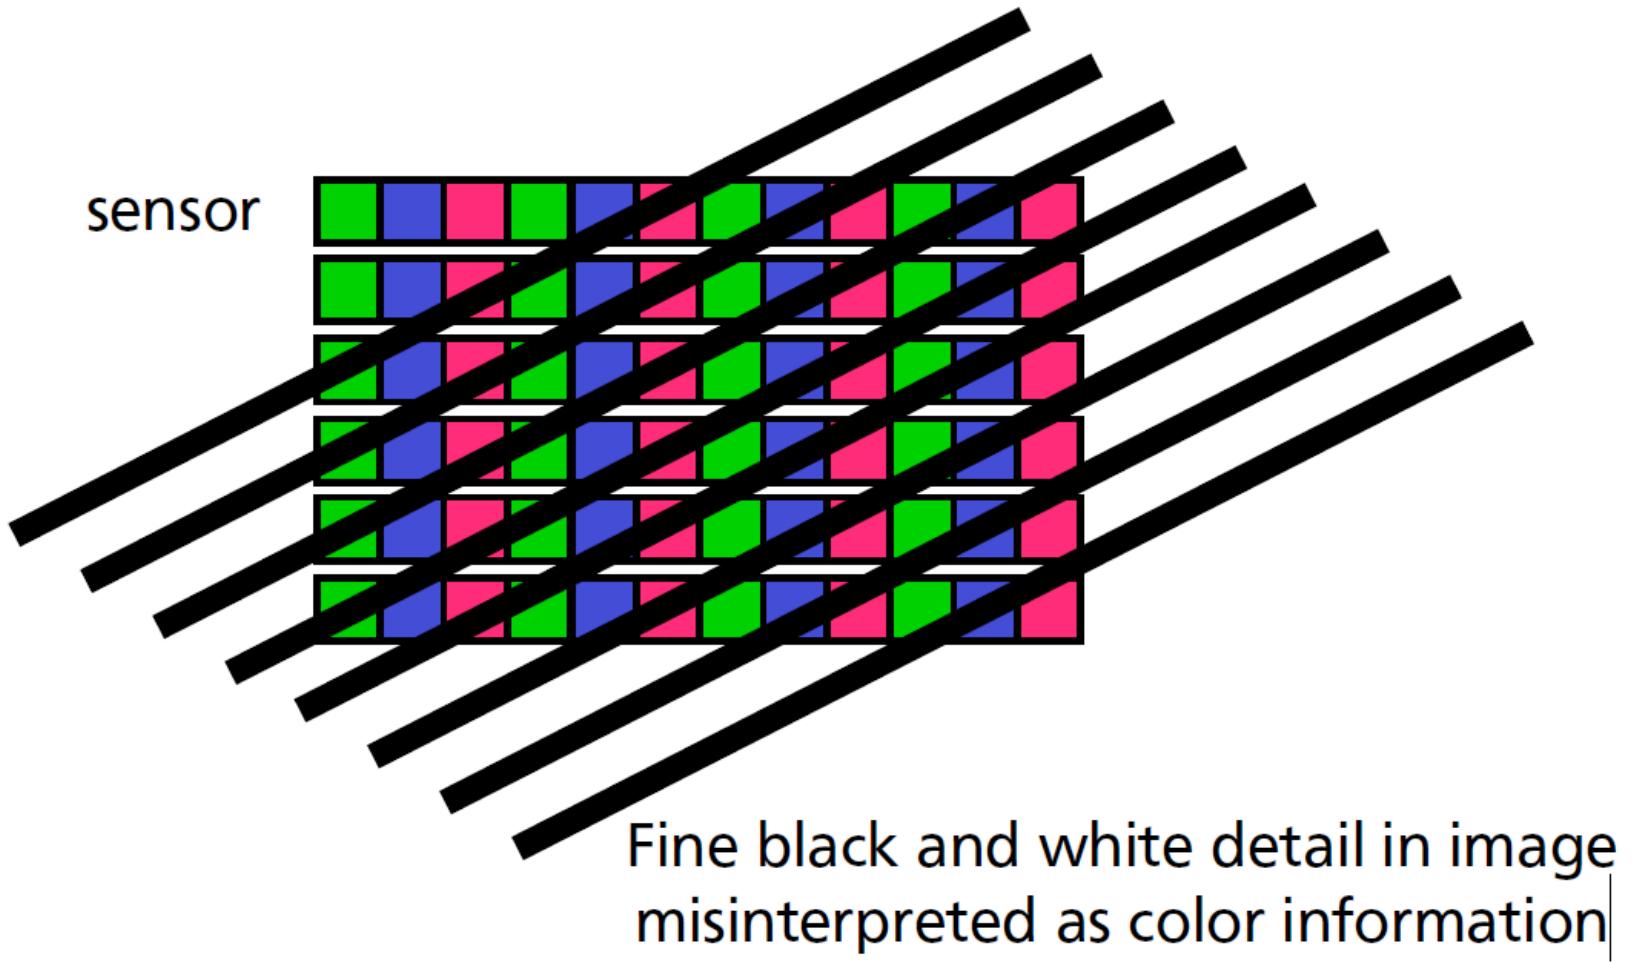


Human Luminance Sensitivity Function

# Problem with de-mosaicing: color moiré



# Cause of color moiré



# General problems with RGB cameras

- **Noise**
  - low light is where you most notice noise
  - light sensitivity (ISO) / noise tradeoff
  - stuck pixels
- **Resolution: Are more megapixels better?**
  - requires higher quality lens
  - noise issues
- **In-camera processing**
  - oversharpening can produce halos
- **RAW vs. compressed**
  - file size vs. quality tradeoff
- **Blooming**
  - charge overflowing into neighboring pixels
- **More info online:**
  - <http://electronics.howstuffworks.com/>
  - <http://www.dpreview.com/>
  - <http://www.dxomark.com/>

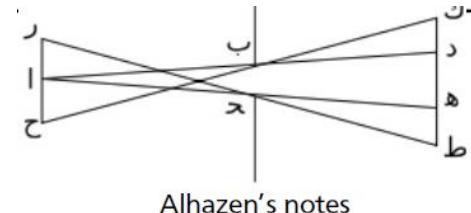


# In-camera processing



# Historical context of cameras

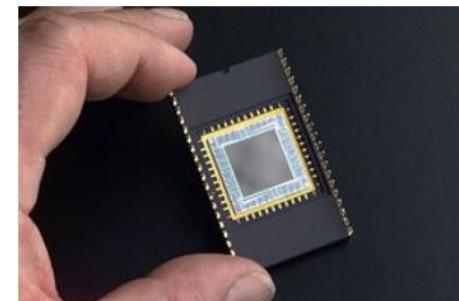
- Pinhole model: Mozi (470-390=1 BCE), Aristoteles (384-322 BCE)
- Principles of optics (including lenses): Alhazen (965-1039 CE)
- Camera obscura: Leonardo da Vinci (1452-1519), Johann Zahn (1631-1707)
- First photo: Joseph Nicéphore Niépce (1822)
- Daguerreotypes (1839)
- Photographic film (Eastman, 1889)
- Cinema (Lumière Brothers, 1895)
- Color Photography (Lumière Brothers, 1908)
- Television (Baird, Farnsworth, Zworykin, 1920s)
- First consumer camera with CCD: Sony Mavica (1981)
- First fully digital camera: Kodak DCS100 (1990)



Alhazen's notes



Niépce, "La Table Servie," 1822



# Camera types

- RGB cameras
- Depth cameras:
  - Passive RGB stereo
  - Active Structured light
  - Active Time of Flight
- Lightfield cameras

# Passive Depth Camera –RGB stereo



Easy to match

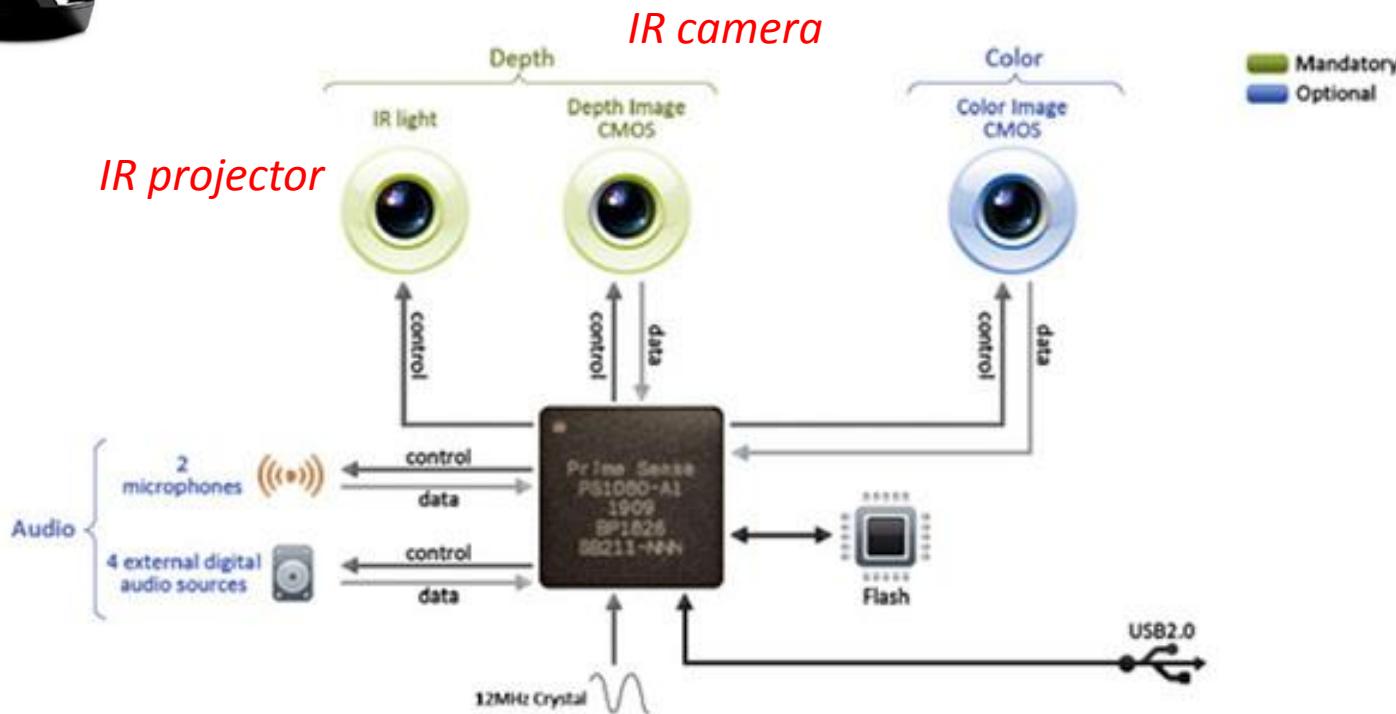


hard to match

# Active Depth Camera – structured light

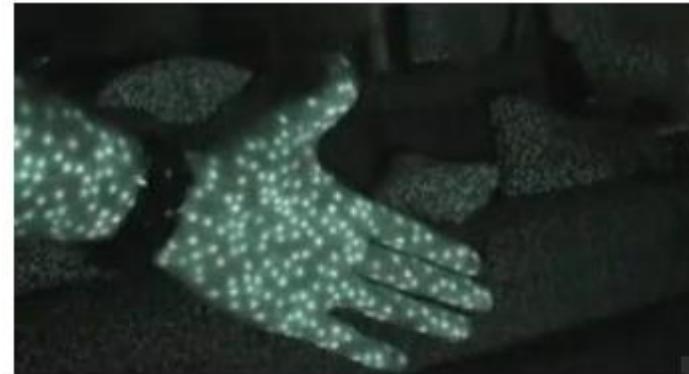


KINECT  
for XBOX360.



# Depth Camera – structured light

*Reference image from  
IR projector*



- We have to perform matching (as with passive stereo camera) but now we have a very textured image
- Only works well when external IR light is not too strong (not under sunlight)

# Depth Camera – structure Light



3 Minutes Break  
Question?

# Depth camera – time of flight



Intensity image



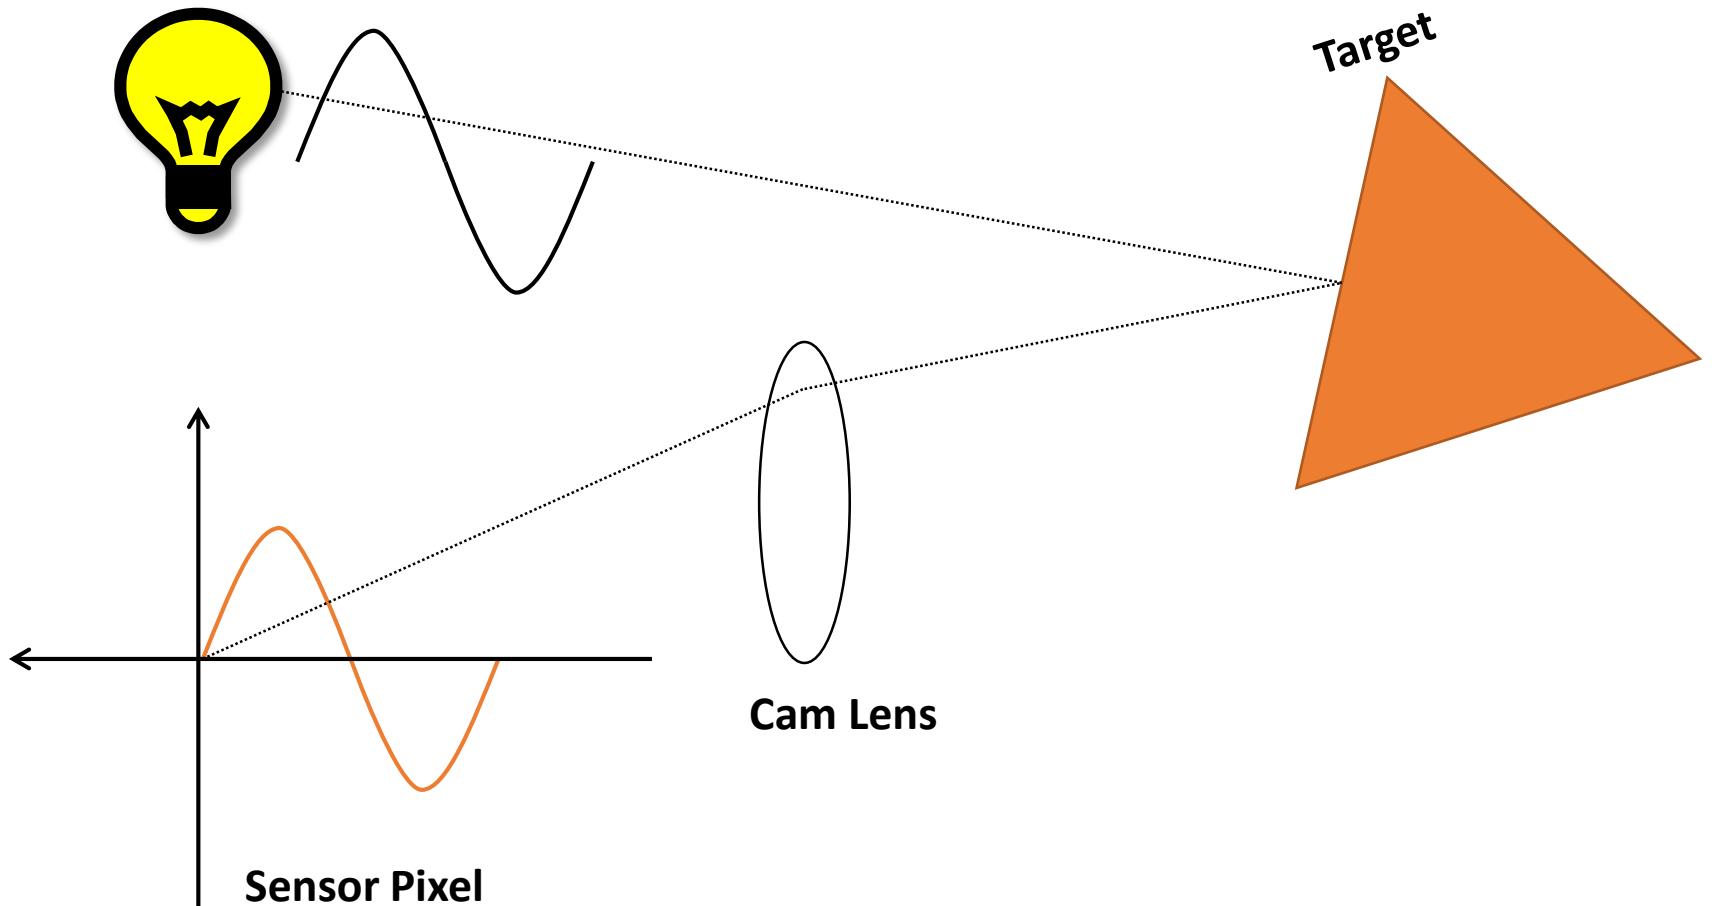
Depth Image



PMD Camera

# Principle Time of Flight

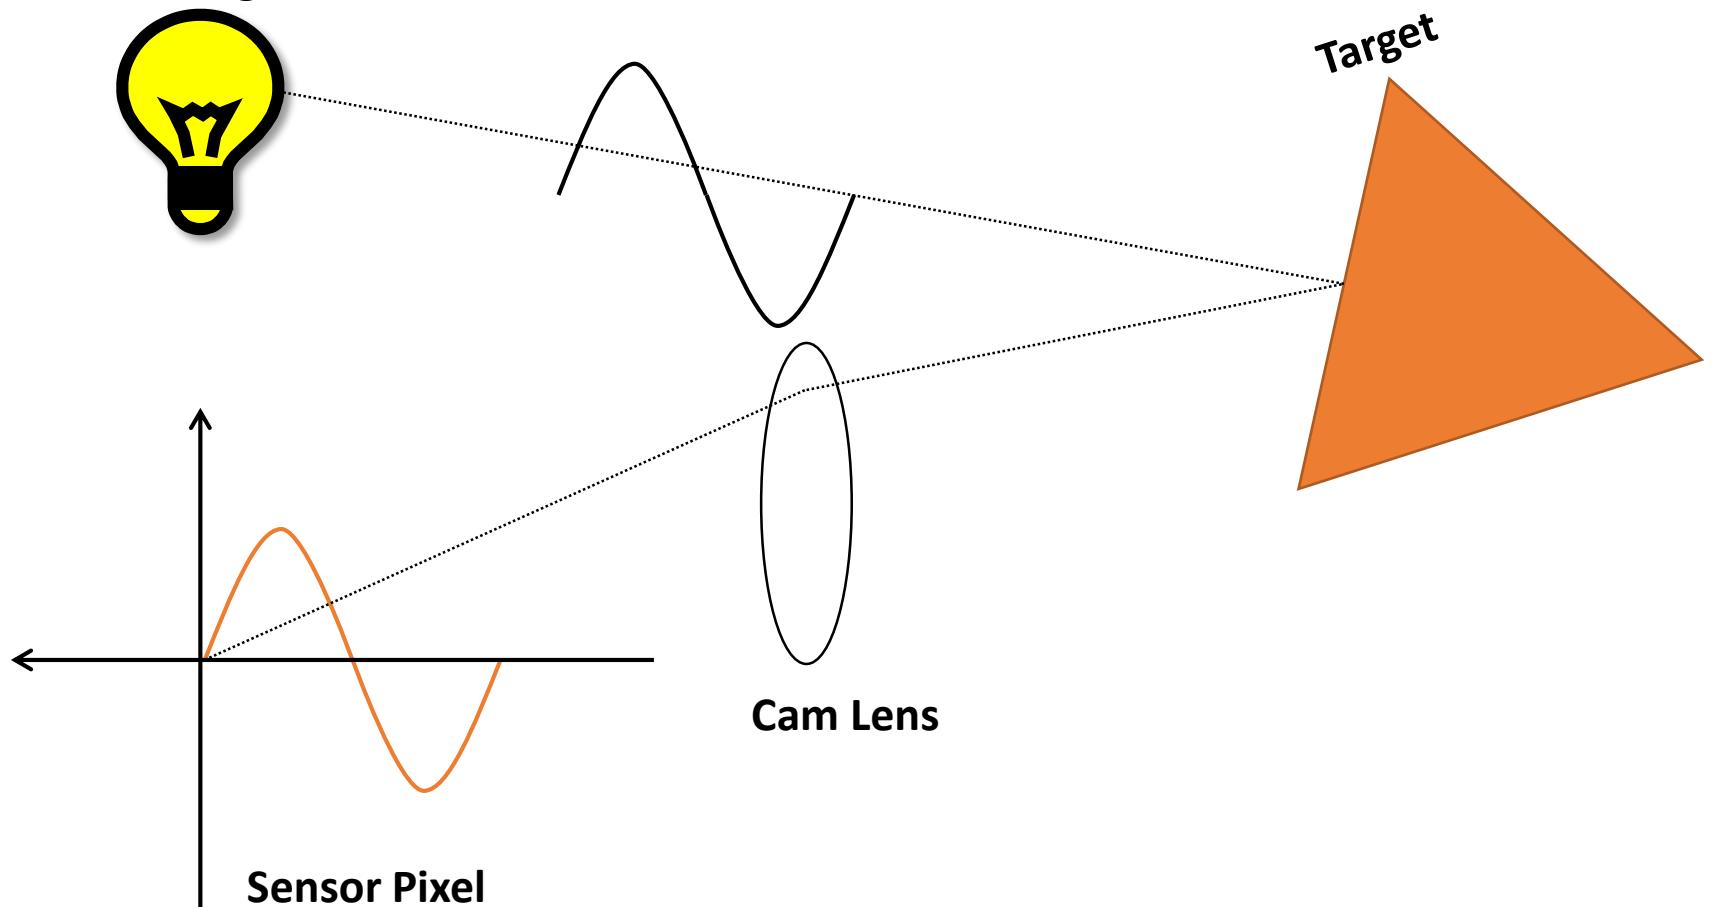
Modulated Light Source



[slide credits: Rahul Nair, Daniel Kondermann]

# Principle Time of Flight

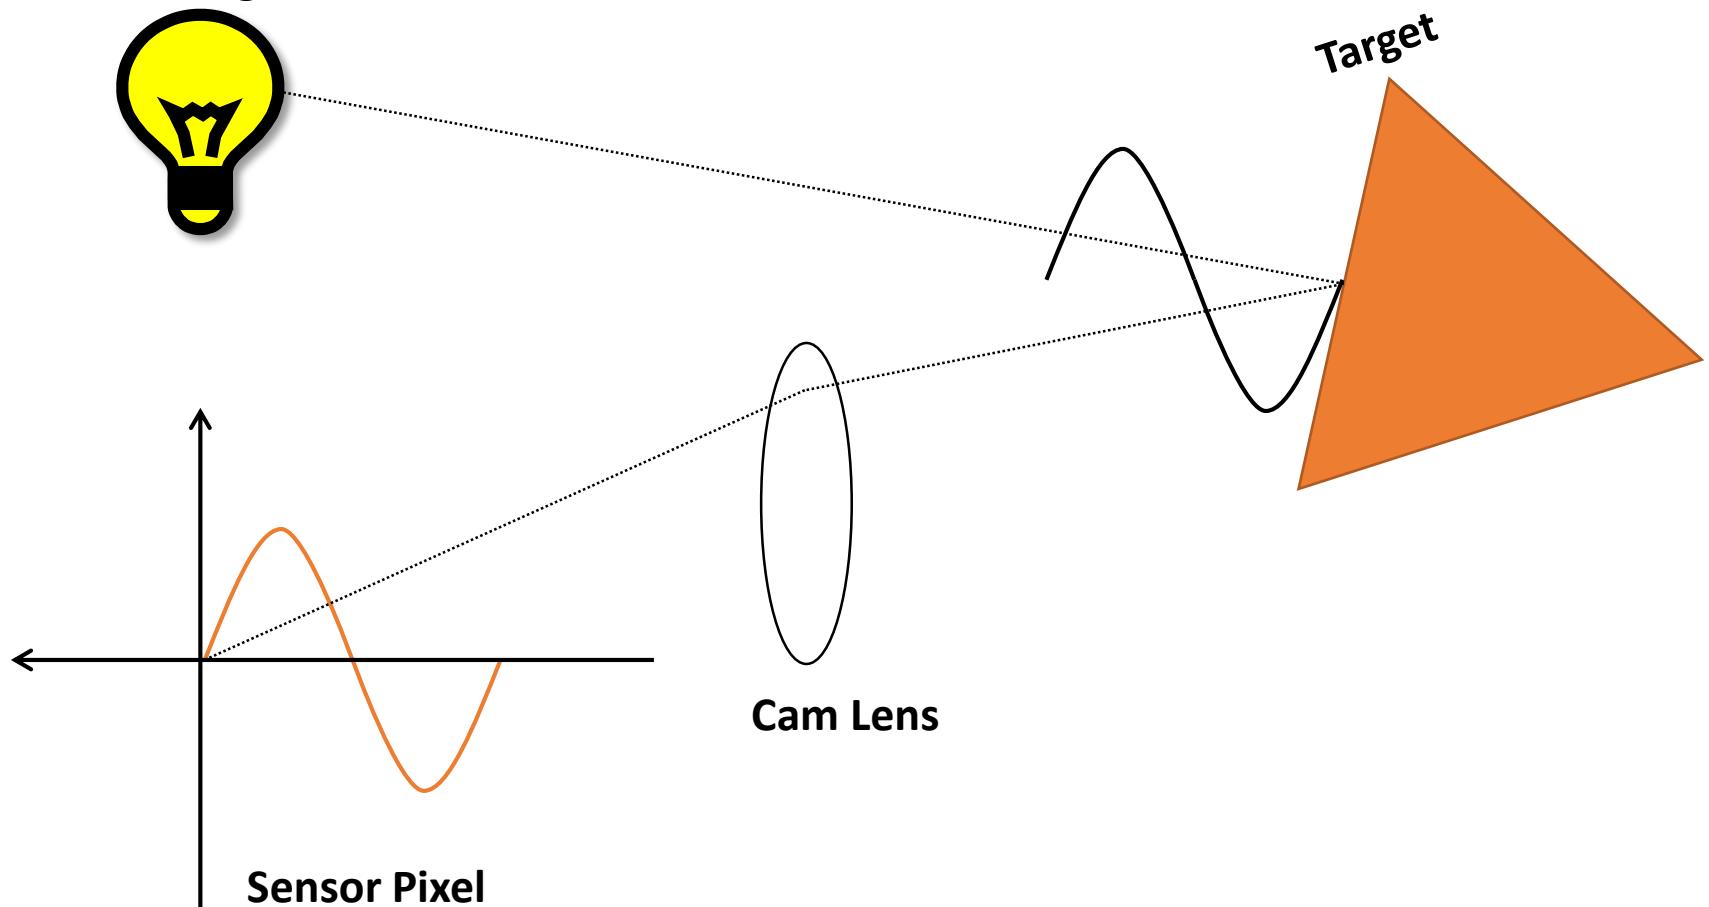
Modulated Light Source



[slide credits: Rahul Nair, Daniel Kondermann]

# Principle Time of Flight

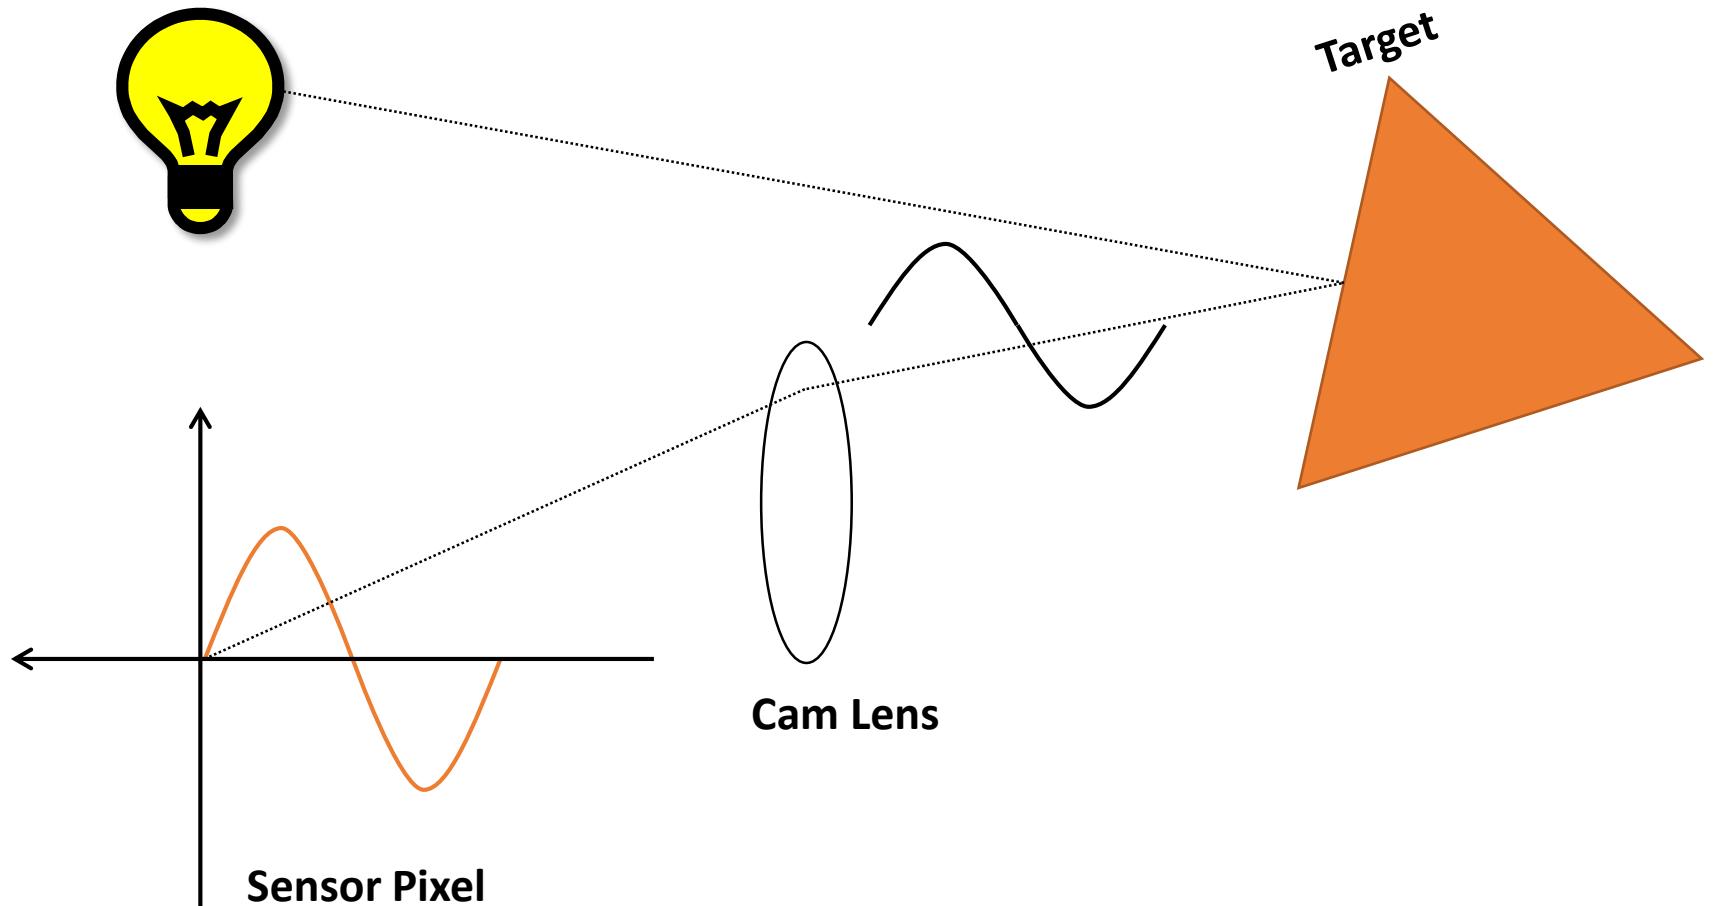
Modulated Light Source



[slide credits: Rahul Nair, Daniel Kondermann]

# Principle Time of Flight

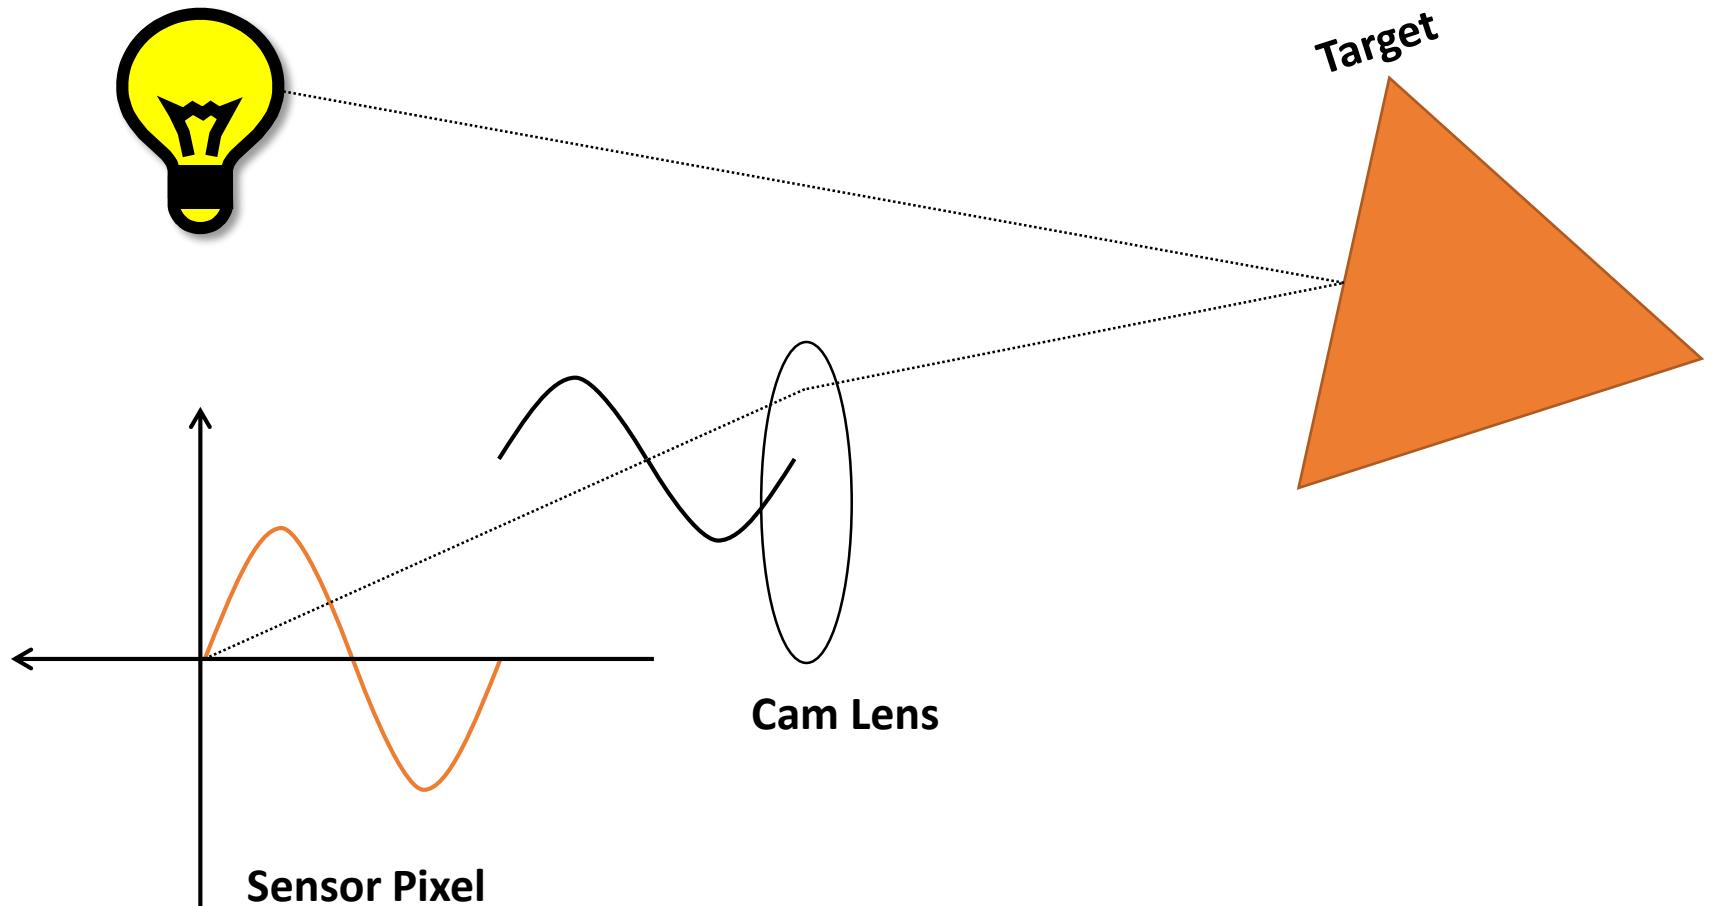
Modulated Light Source



[slide credits: Rahul Nair, Daniel Kondermann]

# Principle Time of Flight

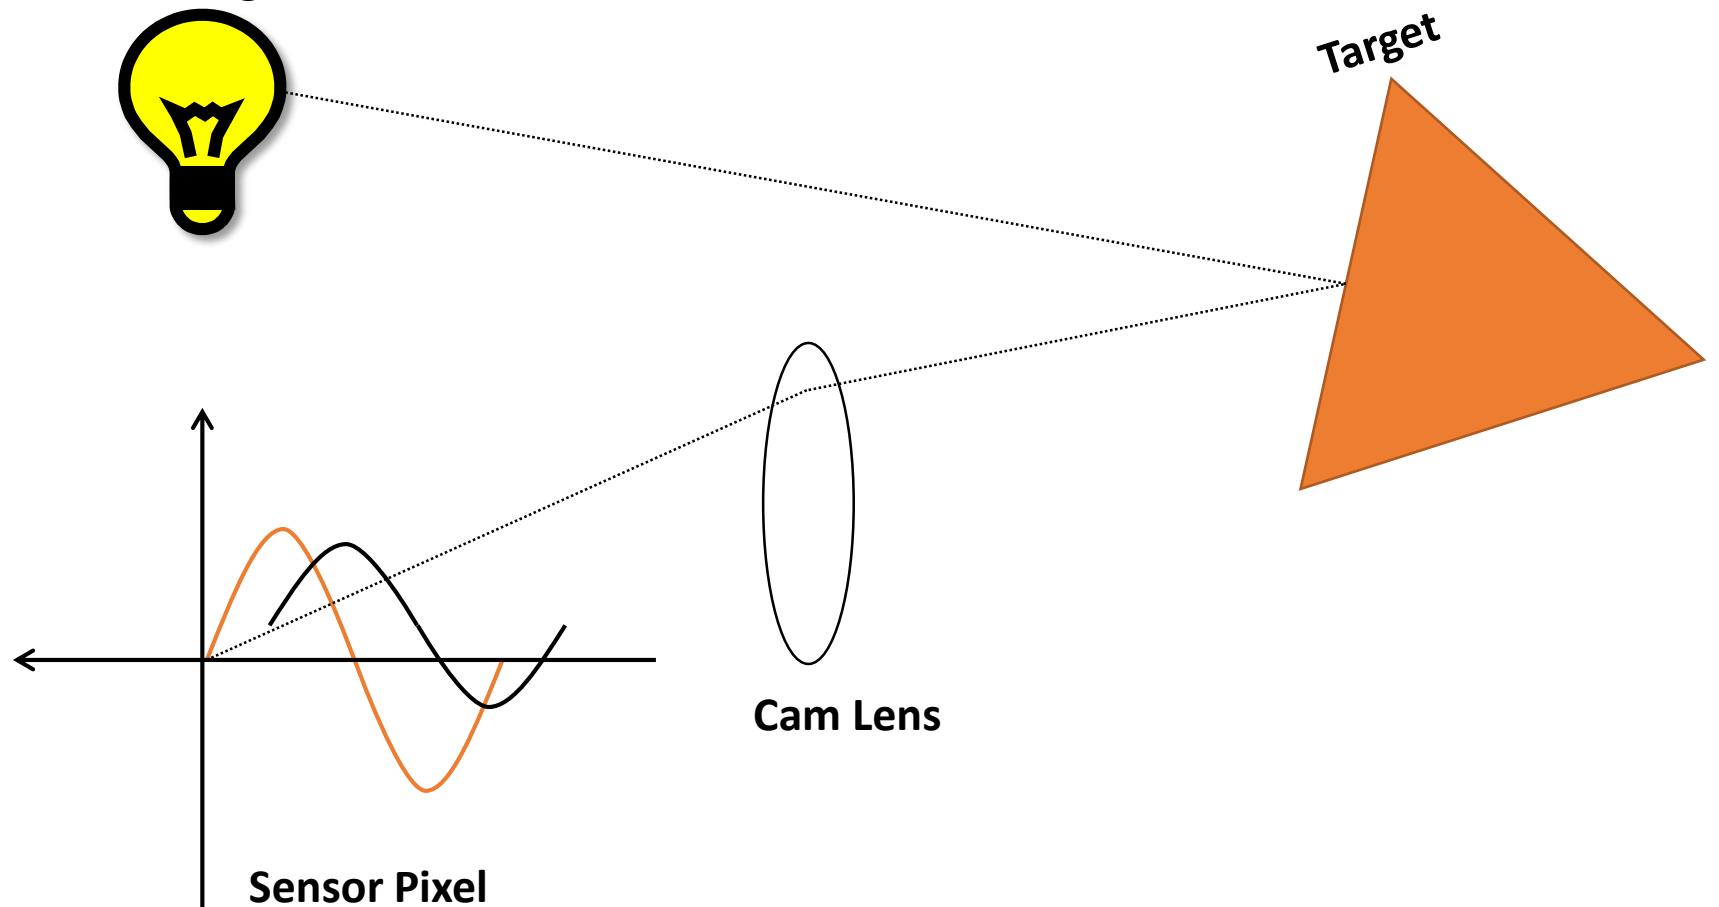
Modulated Light Source



[slide credits: Rahul Nair, Daniel Kondermann]

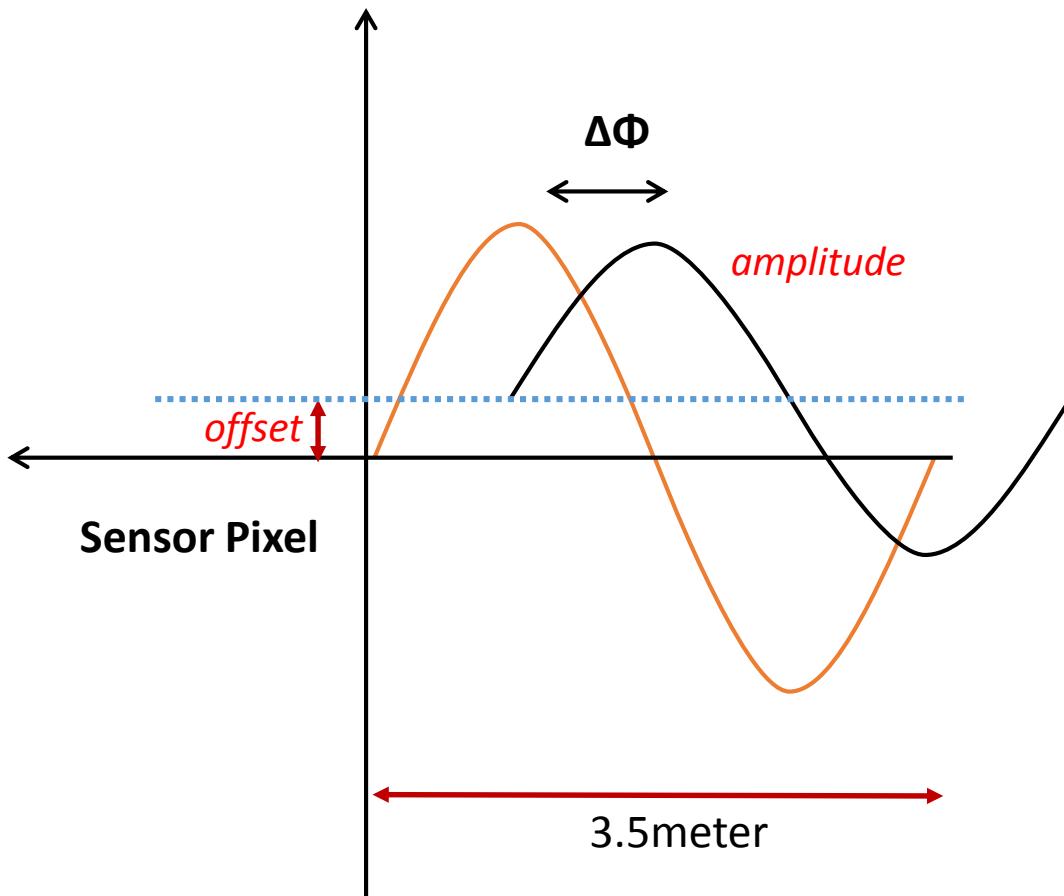
# Principle Time of Flight

Modulated Light Source



[slide credits: Rahul Nair, Daniel Kondermann]

# Principle Time of Flight



$\Delta\Phi$  means  $d = 0.3m$  or  $3.8m$  or  $7.3m$  or ...  
(take differently modulated frequencies)

## Measure:

- Phase shift
- Amplitude and offset

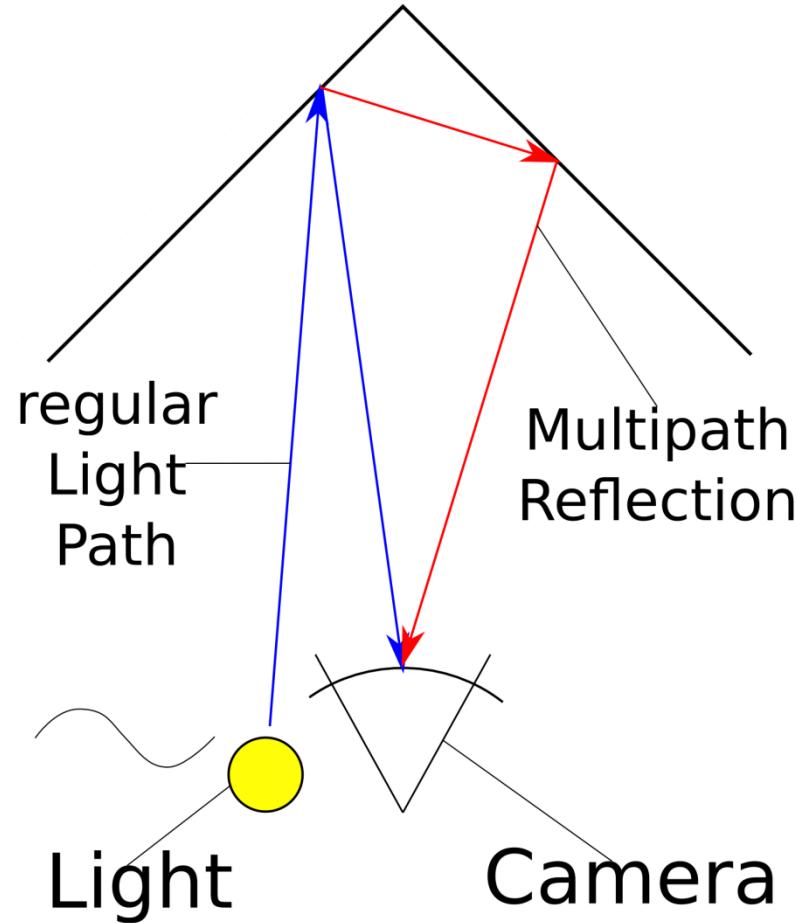
## Output:

- **Depth Image**  
(from phase shift)  
(wavelength determines the range of depth values; around 3.5 meter)
- **Intensity image** (from amplitude and offset)

[slide credits: Rahul Nair, Daniel Kondermann]

# Depth camera – time of flight

One of the biggest problems is multi-path



# Lightfield cameras

Capture all light:



<http://www.lytro.com/camera/>



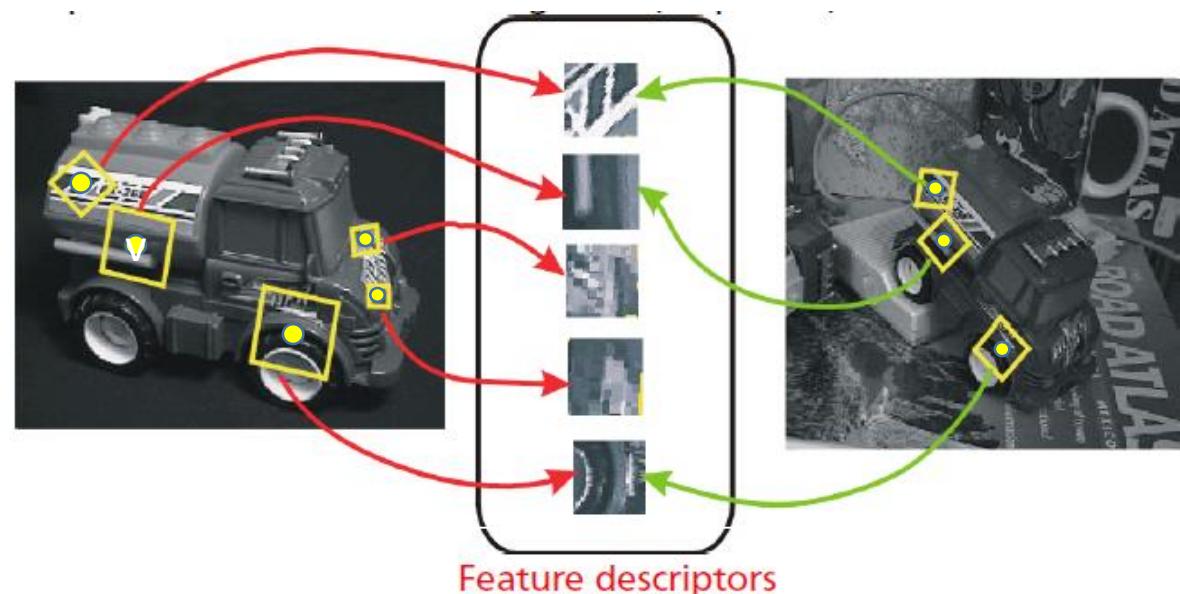
Refocus and change perspective with one image

# Roadmap this lecture

- Geometric primitives and transformations (sec. 2.1.1-2.1.4)
- Geometric image formation (sec 2.1.5, 2.1.6)
  - Pinhole camera
  - Lens effects
- Photometric image formation process (sec 2.2)
- The Human eye
- Camera Types and Hardware (sec 2.3)
- Appearance based matching

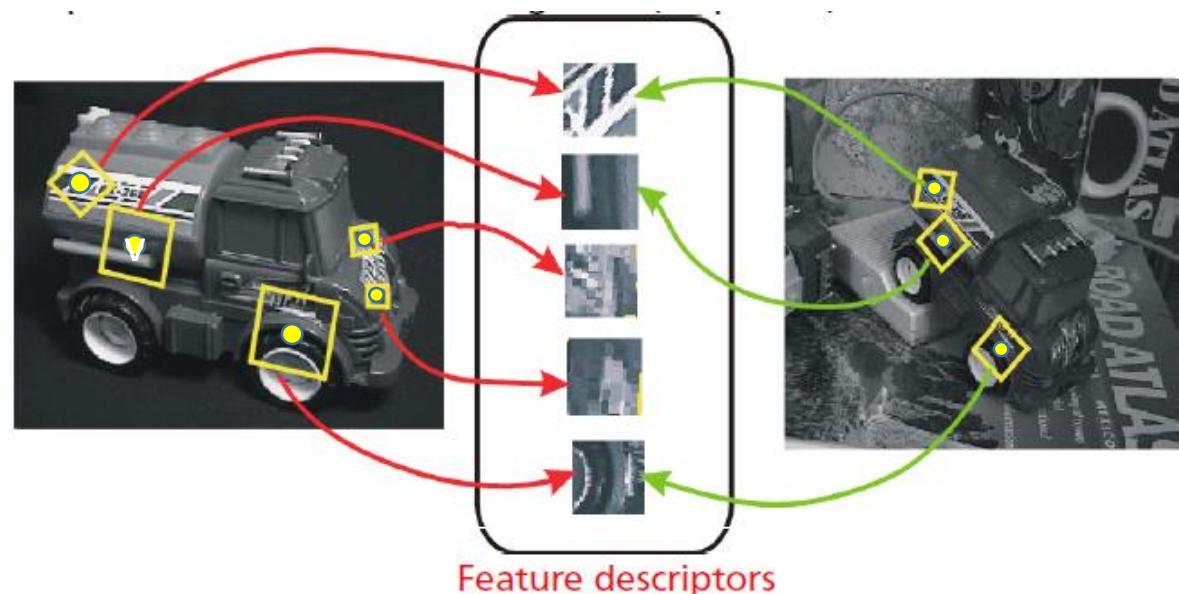
# Roadmap: matching 2 Images (appearance & geometry)

- Find interest points
- Find orientated patches around interest points to capture appearance
- Encode patch appearance in a descriptor
- Find matching patches according to appearance (similar descriptors)
- Verify matching patches according to geometry



# Roadmap: matching 2 Images (appearance & geometry)

- Find interest points
- Find orientated patches around interest points to capture appearance
- Encode patch appearance in a descriptor
- Find matching patches according to appearance (similar descriptors)
- Verify matching patches according to geometry



# Reminder Lecture 3: Harris Corner Detector



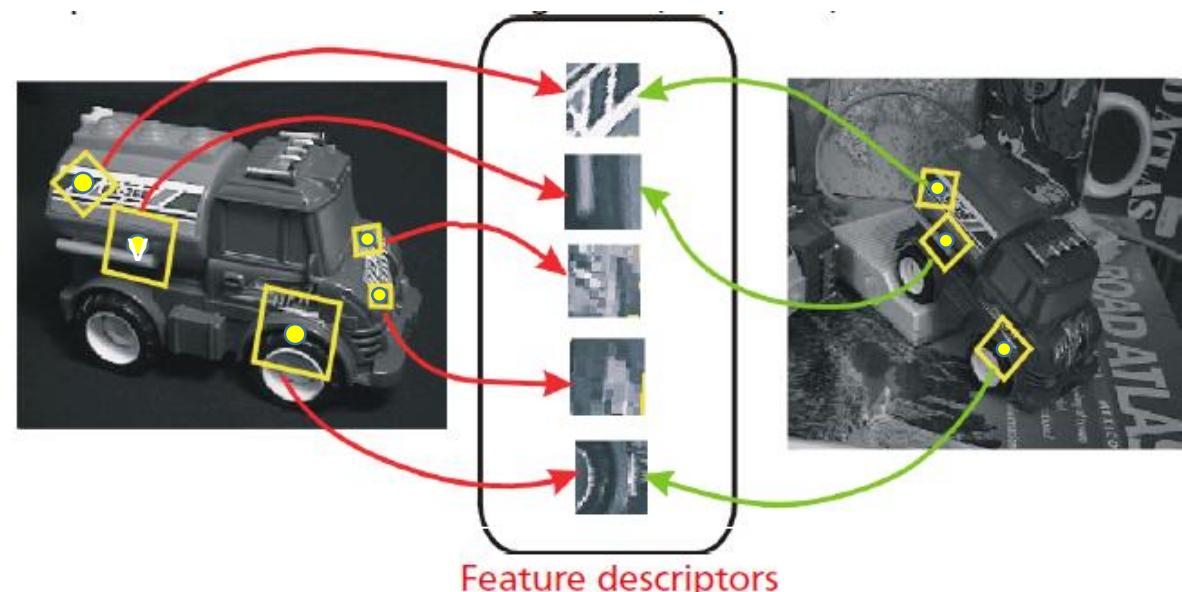
Auto-correlation function:  $c(x, y, \Delta x, \Delta y) \approx [\Delta x, \Delta y] Q(x, y) \begin{bmatrix} \Delta x \\ \Delta y \end{bmatrix}$

$$\lambda_1 \lambda_2 = \det Q(x, y) = AC - B^2, \quad \lambda_1 + \lambda_2 = \text{trace } Q(x, y) = A + C$$

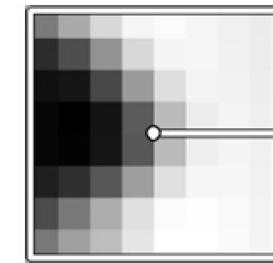
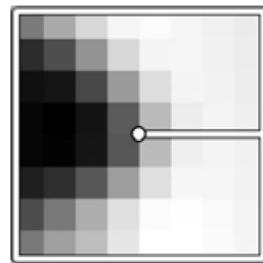
Harris measure:  $H = \lambda_1 \lambda_2 - 0.04(\lambda_1 + \lambda_2)^2$

# Roadmap: matching 2 Images (appearance & geometry)

- Find interest points
- **Find orientated patches around interest points to capture appearance**
- Encode patch appearance in a descriptor
- Find matching patches according to appearance (similar descriptors)
- Verify matching patches according to geometry



# How to deal with orientation

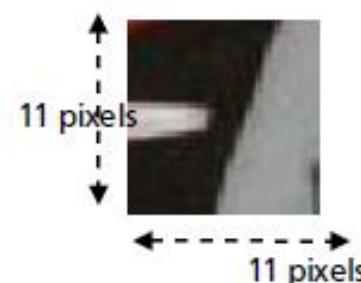
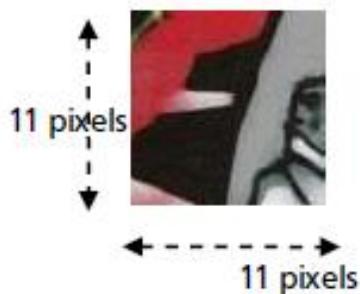
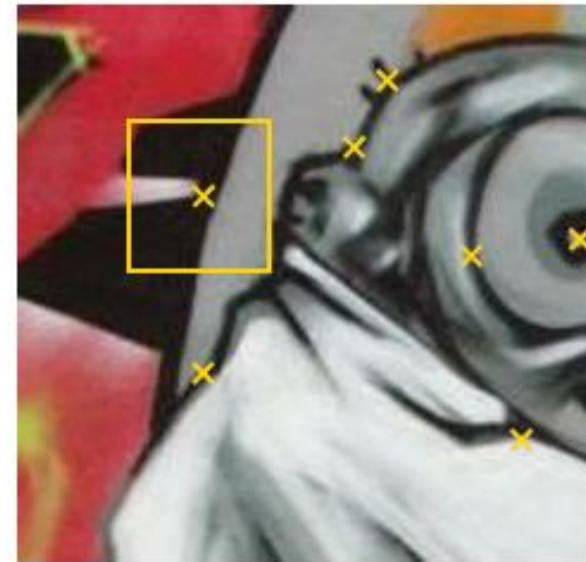


Orientate with image gradient:

$$\nabla I = (I_x, I_y) = \left( \frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right)$$

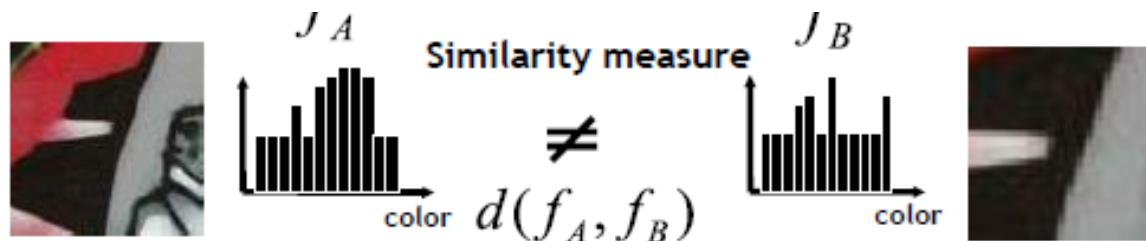
$$\theta = \text{atan}(I_y, I_x)$$

# Choose a patch around each point



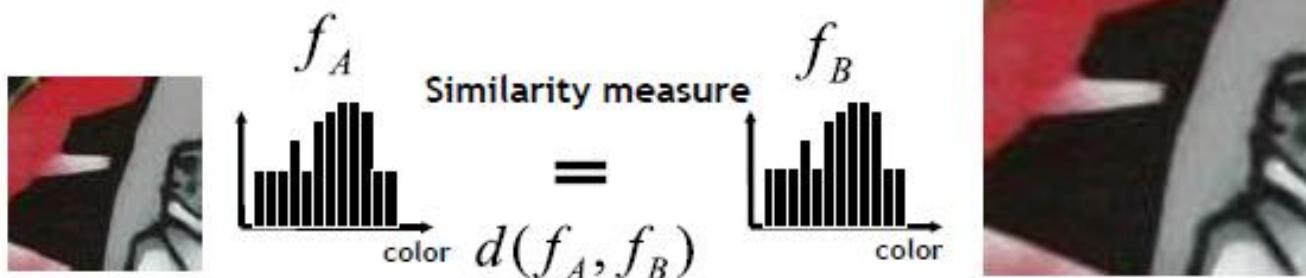
How to deal with scale?

# Choose a patch around each point



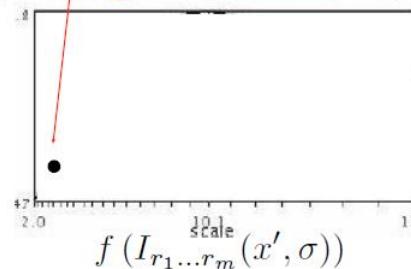
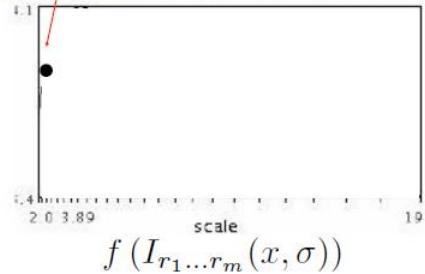
How to deal with scale?

# Choose a patch around each point



How to deal with scale?

# Scale selection (illustration)

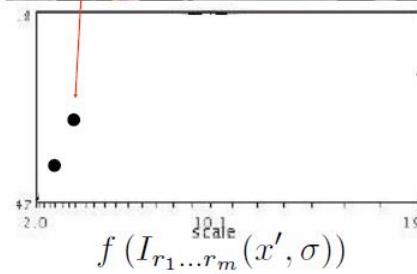
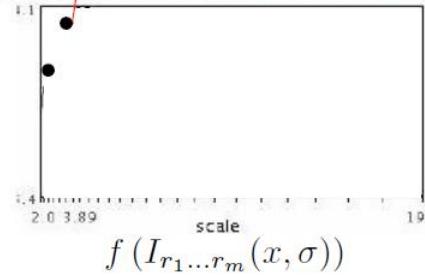


$f$  is Laplacian of Gaussian (LoG) operator.  
Measures an average edge-ness in all directions

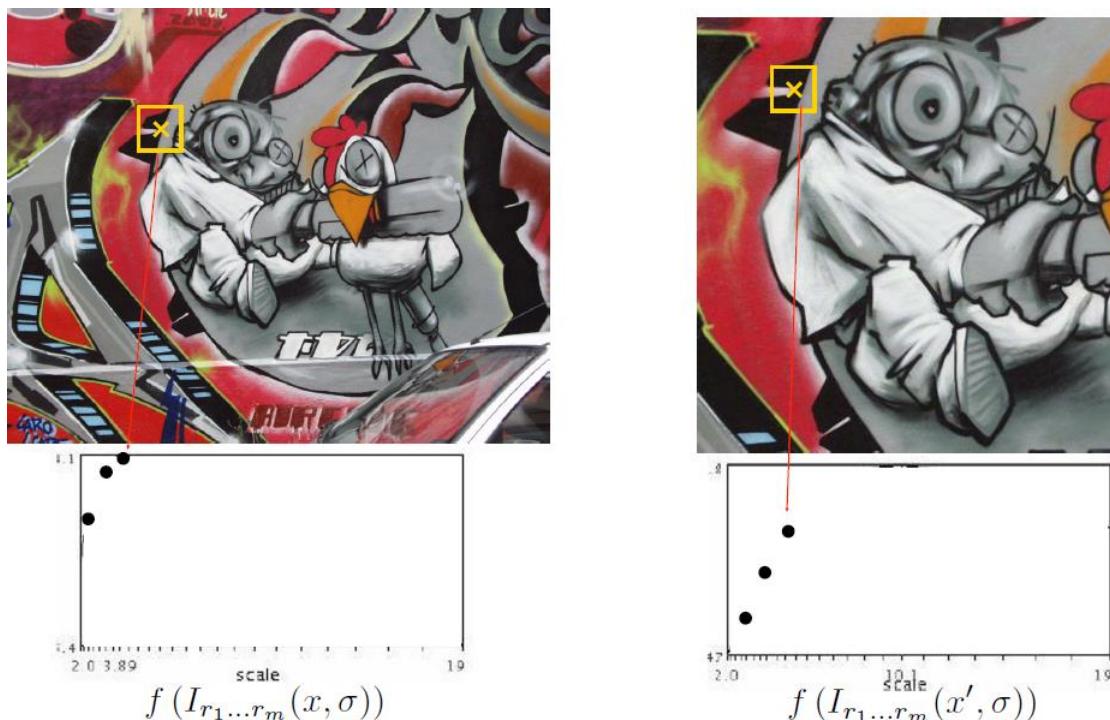
$$\nabla^2(G(\sigma) * I) = \frac{\partial^2 (G(\sigma) * I)}{\partial x^2} + \frac{\partial^2 (G(\sigma) * I)}{\partial y^2}$$

(details on page 191)

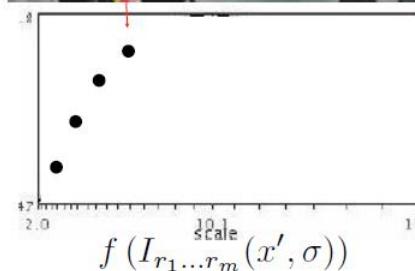
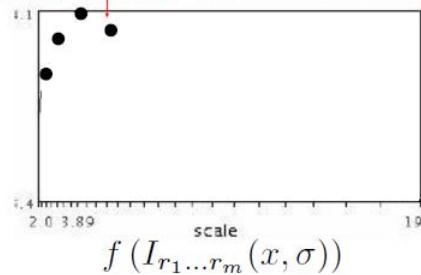
# Scale selection (illustration)



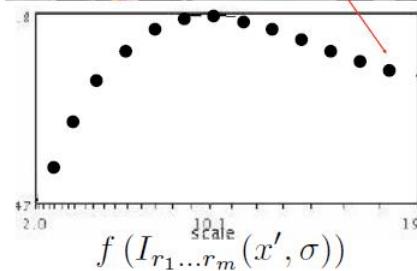
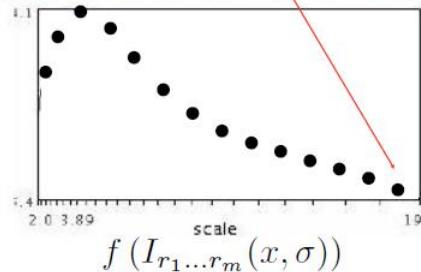
# Scale selection (illustration)



# Scale selection (illustration)

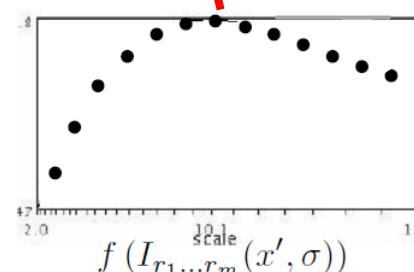
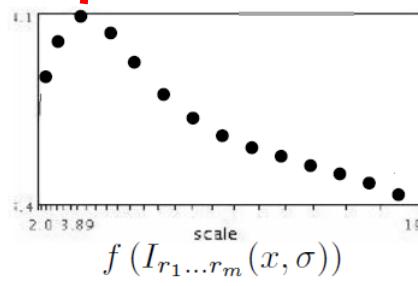


# Scale selection (illustration)



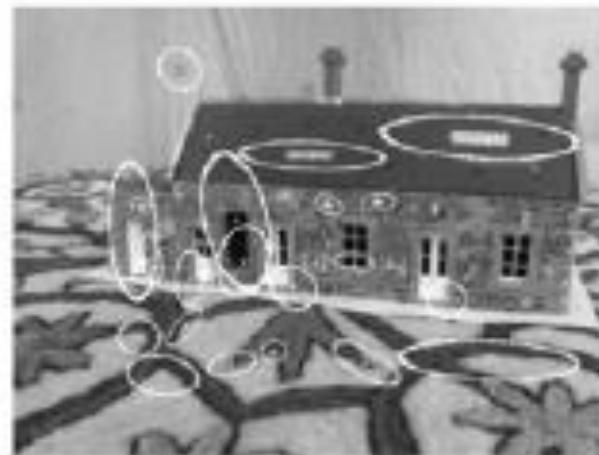
We could match-up these curves and find unique corresponding points

# Scale selection (illustration)



Simpler: Find maxima /minima in image

# Extensions: general affine transformations



$$\mathbf{x}_0 \rightarrow \mathbf{A}_0^{-1/2} \mathbf{x}'_0$$



$$\mathbf{x}'_0 \rightarrow \mathbf{R}\mathbf{x}'_1$$

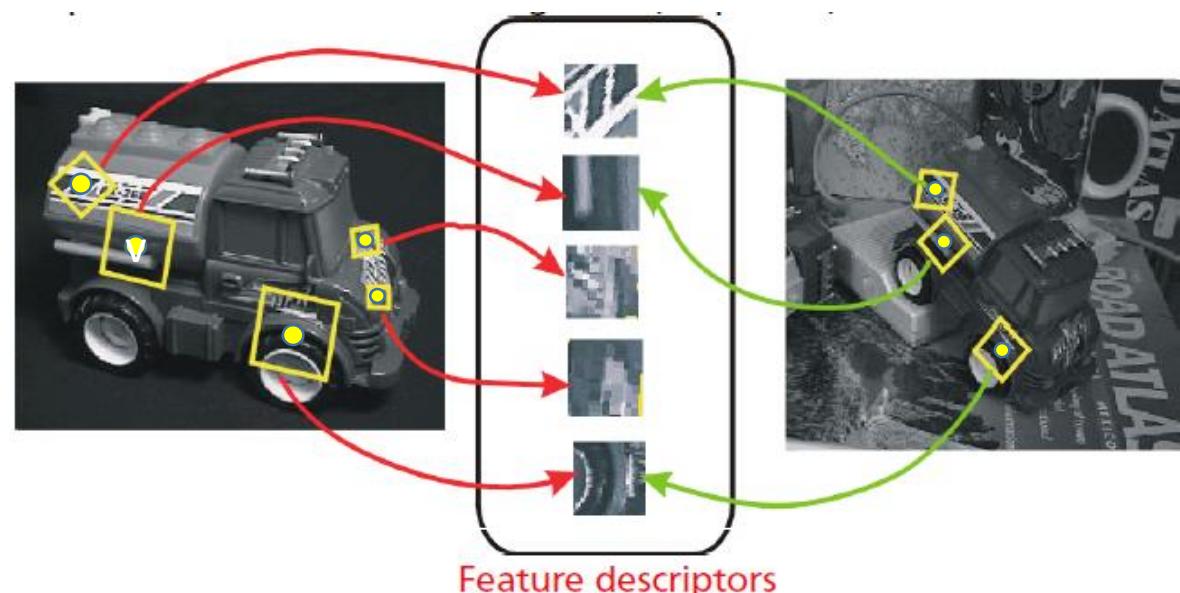


$$\mathbf{A}_1^{-1/2} \mathbf{x}'_1 \leftarrow \mathbf{x}_1$$

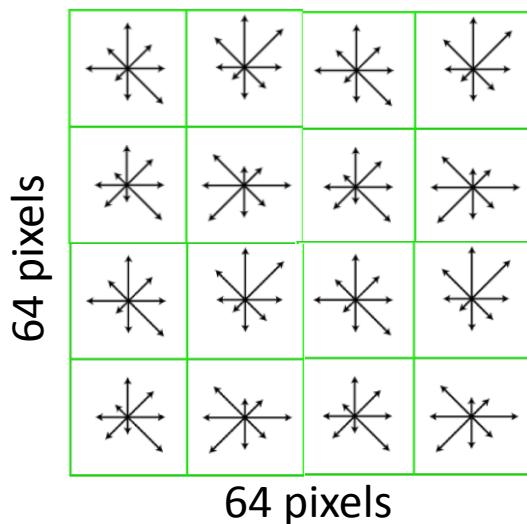


# Roadmap: matching 2 Images (appearance & geometry)

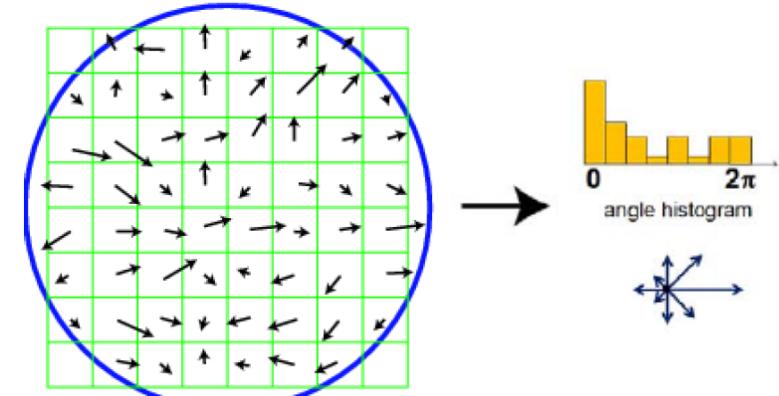
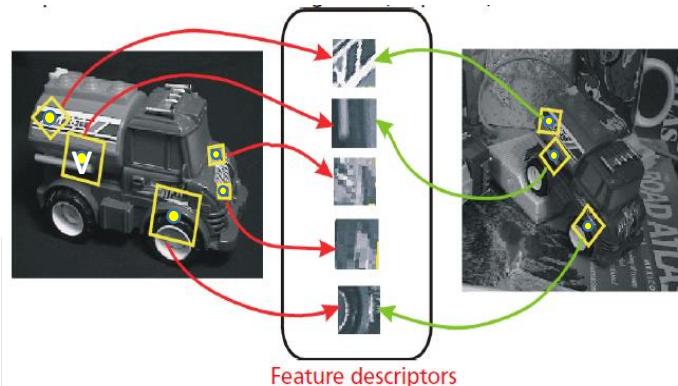
- Find interest points
- Find orientated patches around interest points to capture appearance
- **Encode patch appearance in a descriptor**
- Find matching patches according to appearance (similar descriptors)
- Verify matching patches according to geometry



# SIFT feature



- $4 \times 4 = 16$  cells
- Each cell has an 8 bin histogram (smoothed across cells)
- In total:  $16 \times 8$  values, i.e. **128D vector**



[Lowe 2004]

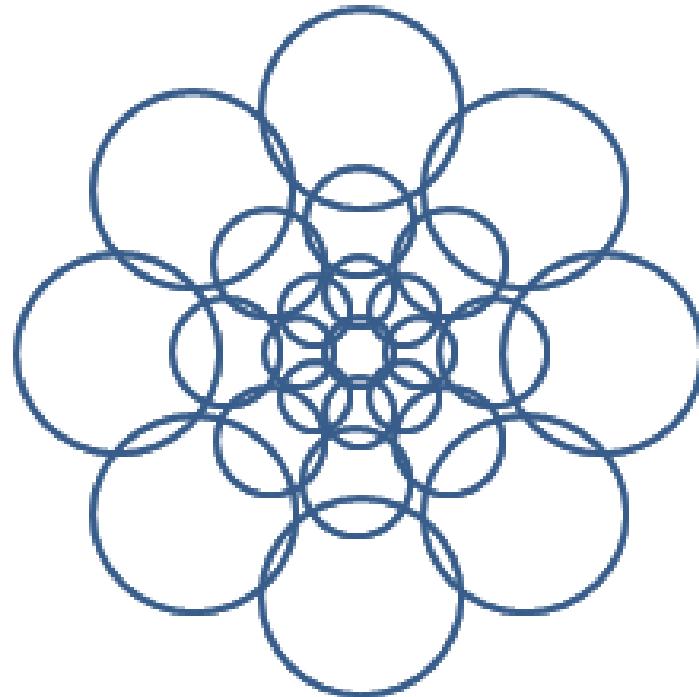
# SIFT feature is very popular

- Fast to compute
- Can handle large changes in viewpoint well (up to  $60^{\circ}$  out of plane rotation)
- Can handle photometric changes (even day versus night)



# Many other feature descriptors

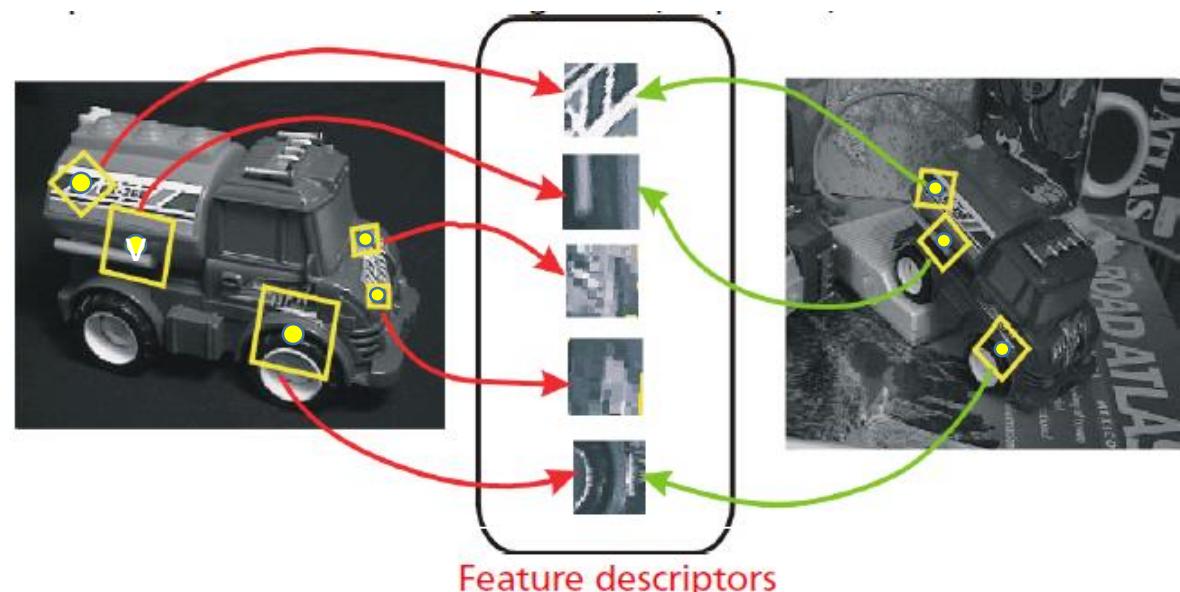
- MOPS [Brown, Szeliski and Winder 2005]
- SURF [Herbert Bay et al. 2006]
- DAISY [Tola, Lepetit, Fua 2010]
- Shape Context
- ....



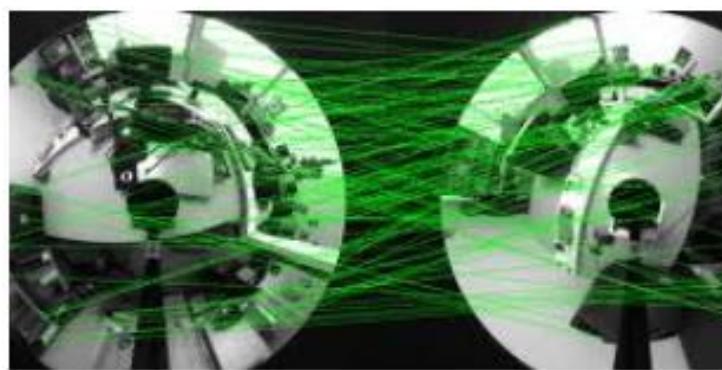
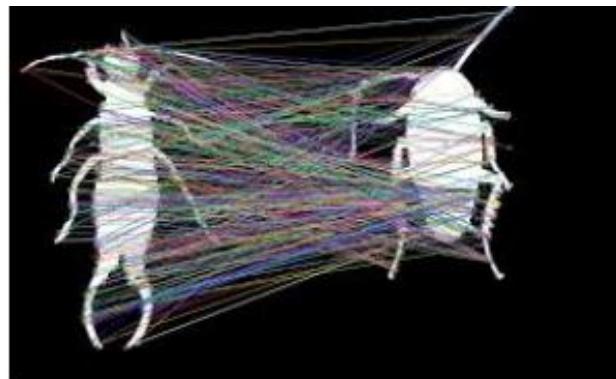
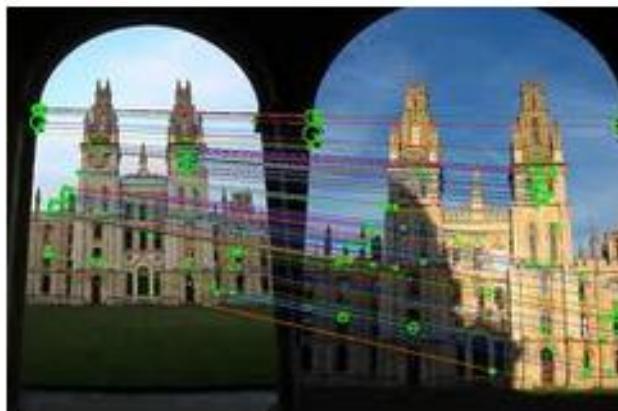
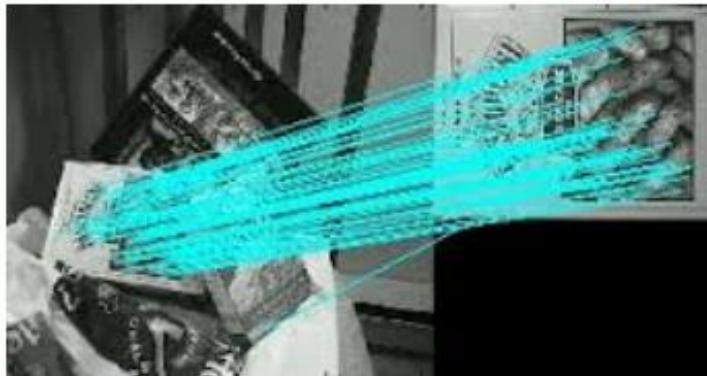
DAISY

# Roadmap: matching 2 Images (appearance & geometry)

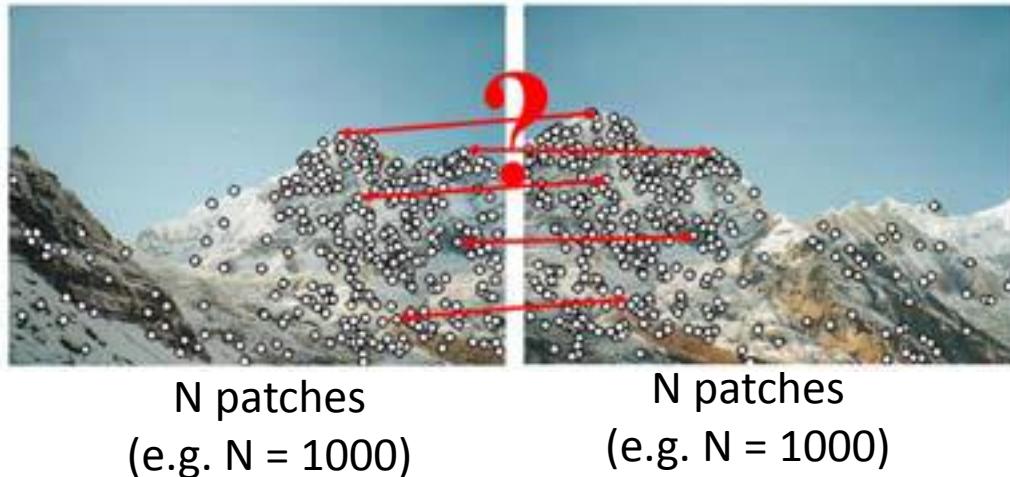
- Find interest points
- Find orientated patches around interest points to capture appearance
- Encode patch appearance in a descriptor
- Find matching patches according to appearance (similar descriptors)
- Verify matching patches according to geometry



# Appearance-based matching



# Appearance-based matching

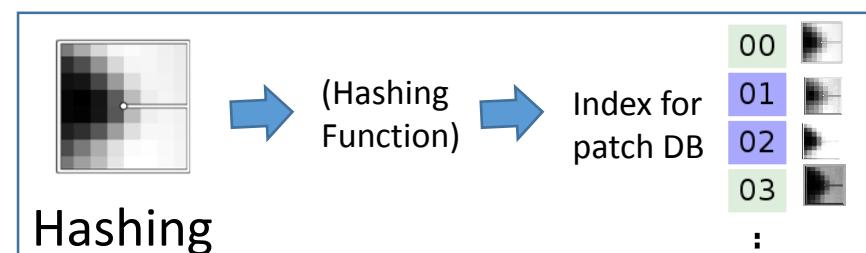


## Goal:

- 1) Find for each patch in left image the closest in right image
- 2) Accept all matches where descriptors are similar enough

## Methods:

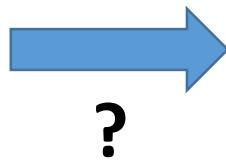
- Naïve:  $N^2$  tests (here 1 Million)
- Hashing (locality sensitive hashing)
- Kd-tree; on average  $N \log N$  tests (here 10,000)



# Subtask: Search for one patch



Query patch



Database image

# Nearest Neighbor Search

- Tracking in Video



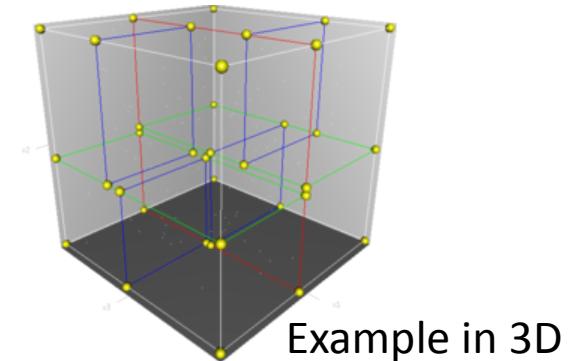
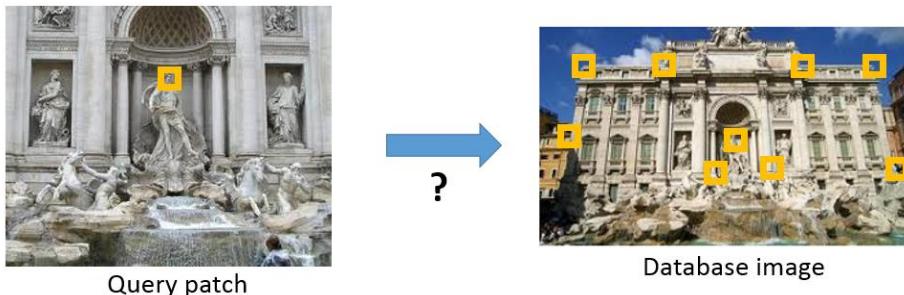
The video is the Database

- Image retrieval



- Whole image has one descriptor
- Database is an image collection

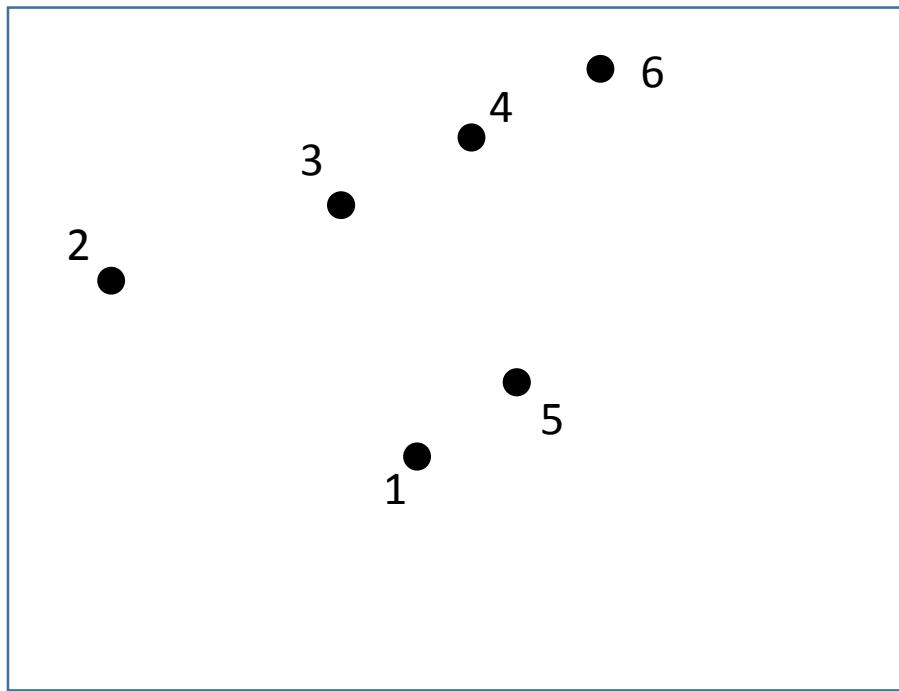
# Kd-tree (d stands for dimension)



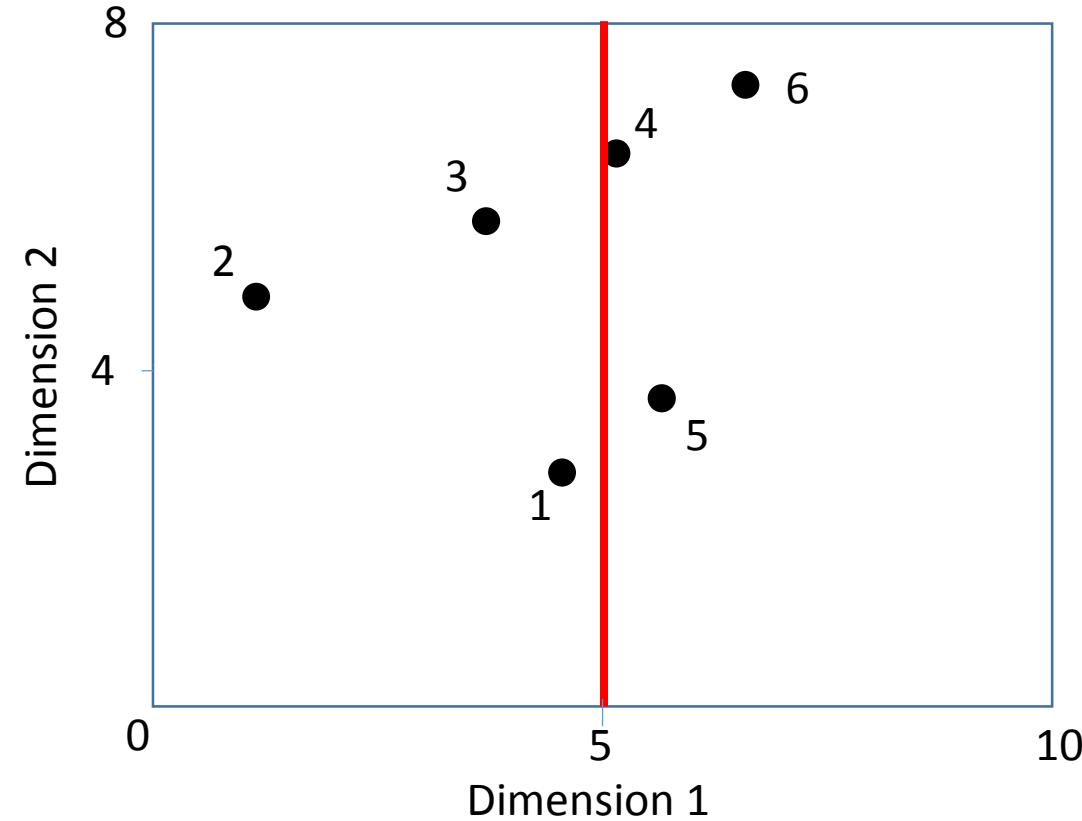
- Build the tree over database image:
  - 1) Cycle over dimensions: x,y,z,x,y,z,....
  - 2) Put in axis-aligned hyper-planes  
(split at median of point set)
- Result: balanced tree
- Nearest Neighbour search (to come)

[Invented by Jon Louis Bentley 1975]

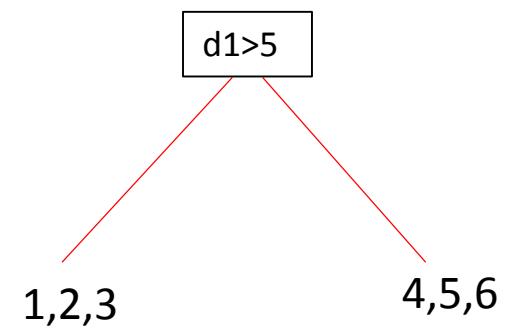
# Examples



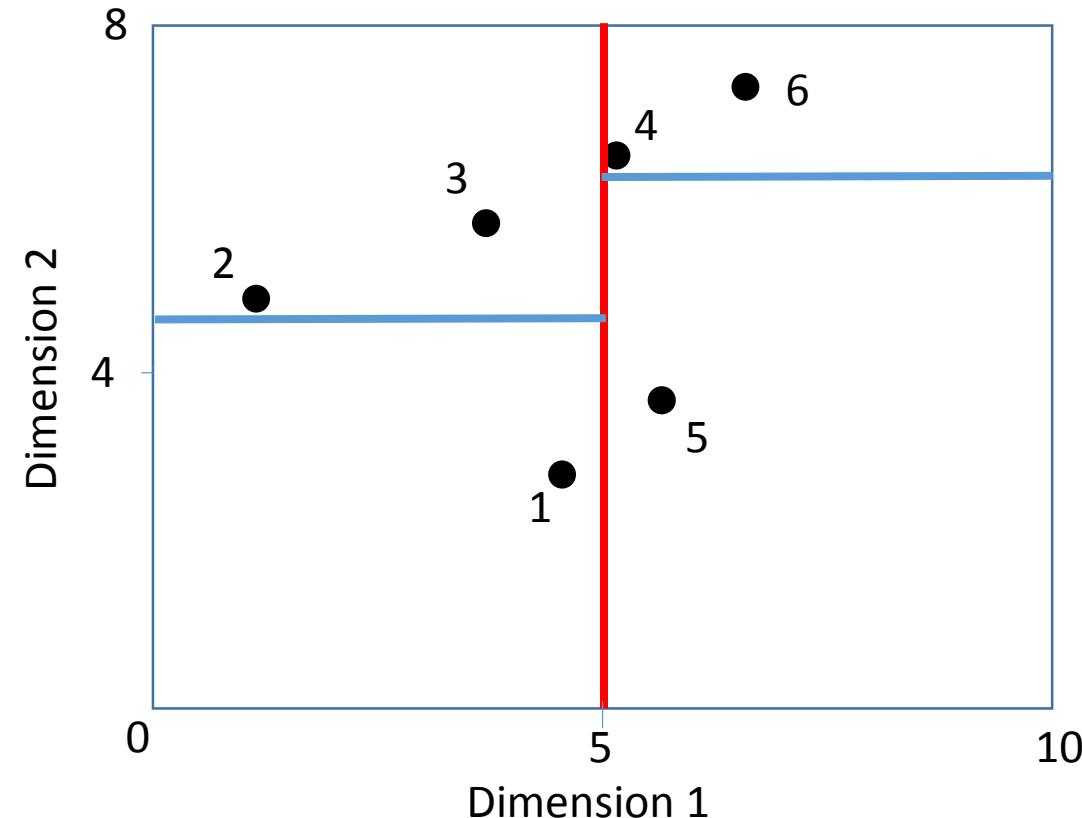
# Examples



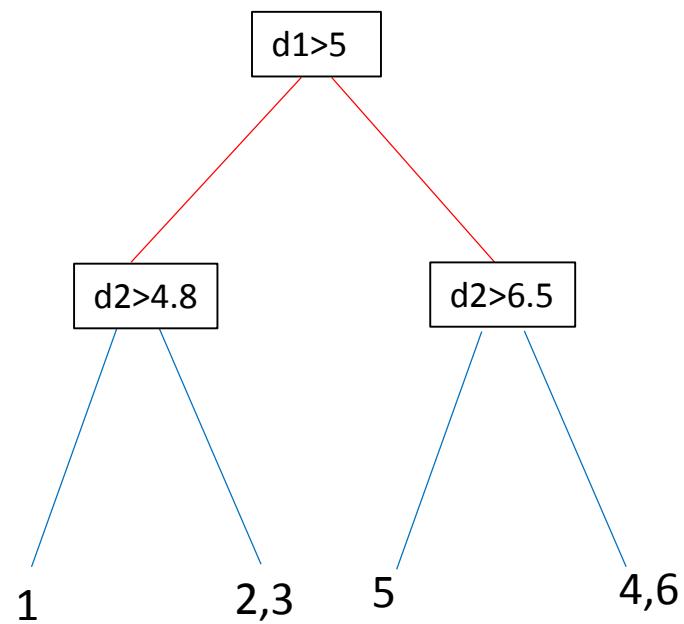
Kd-tree



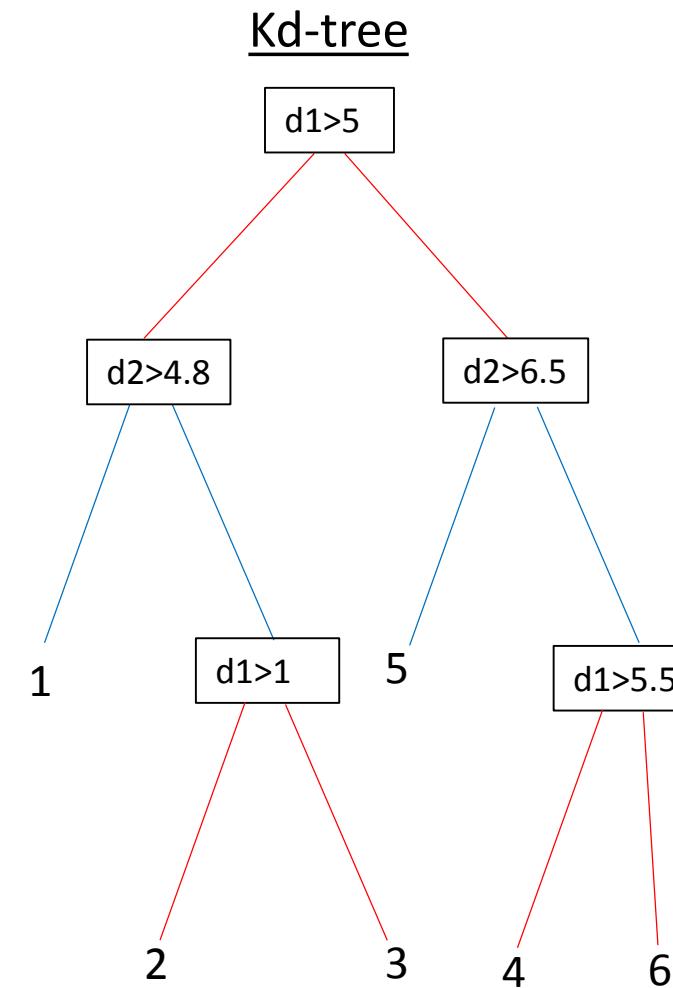
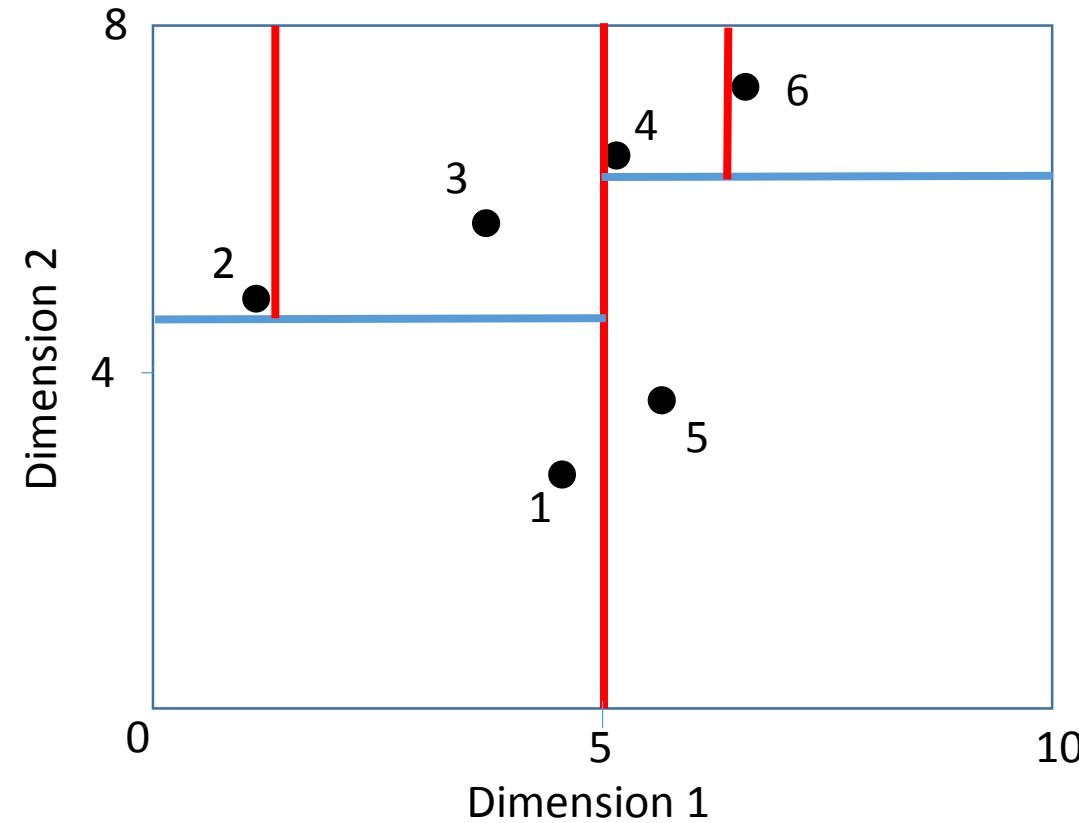
# Examples



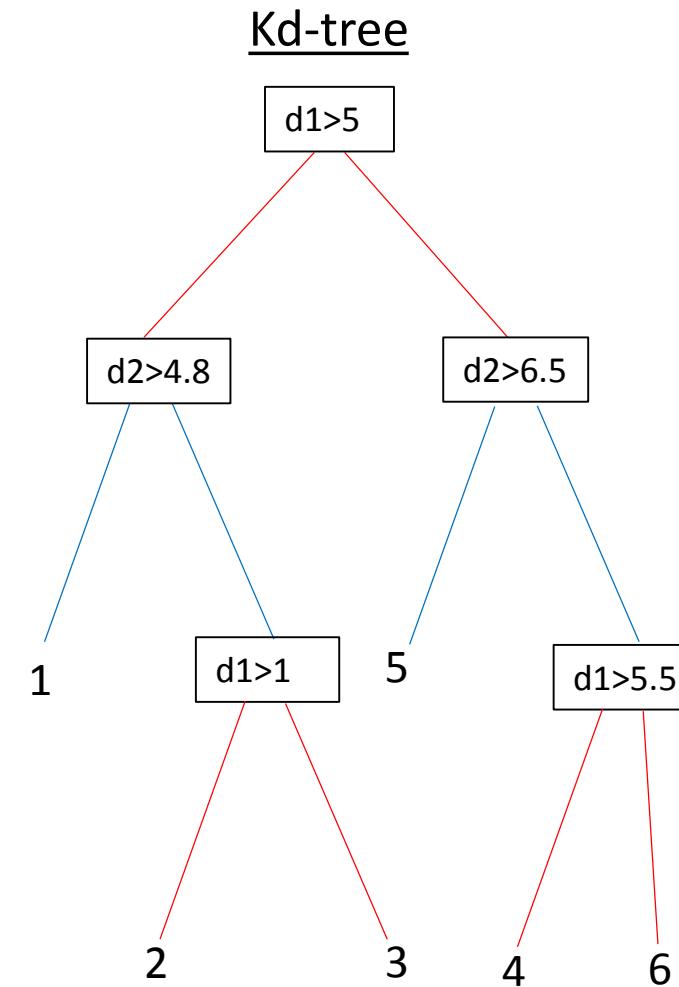
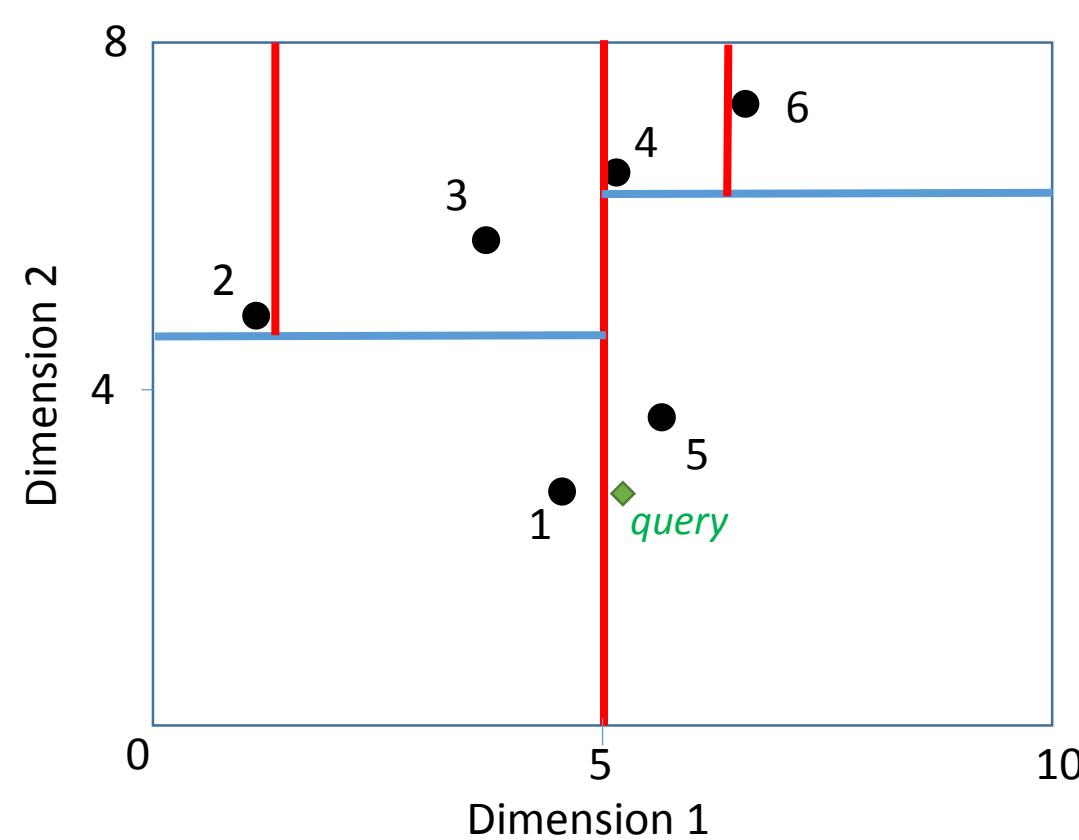
Kd-tree



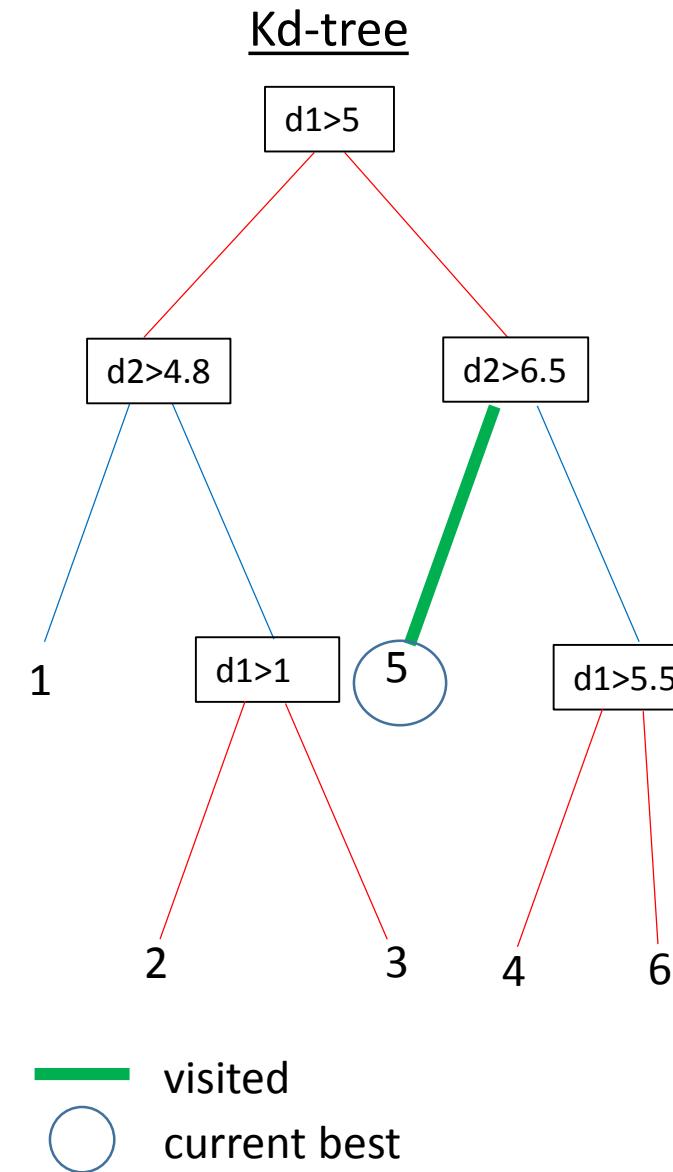
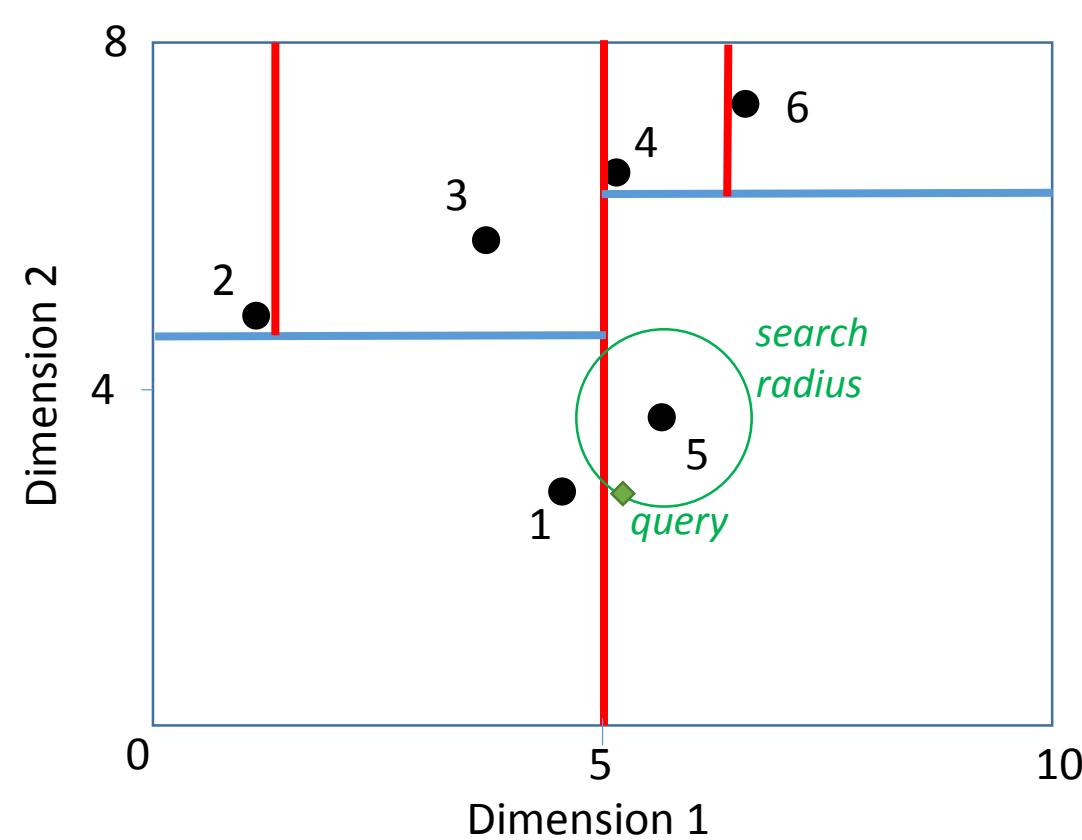
# Examples



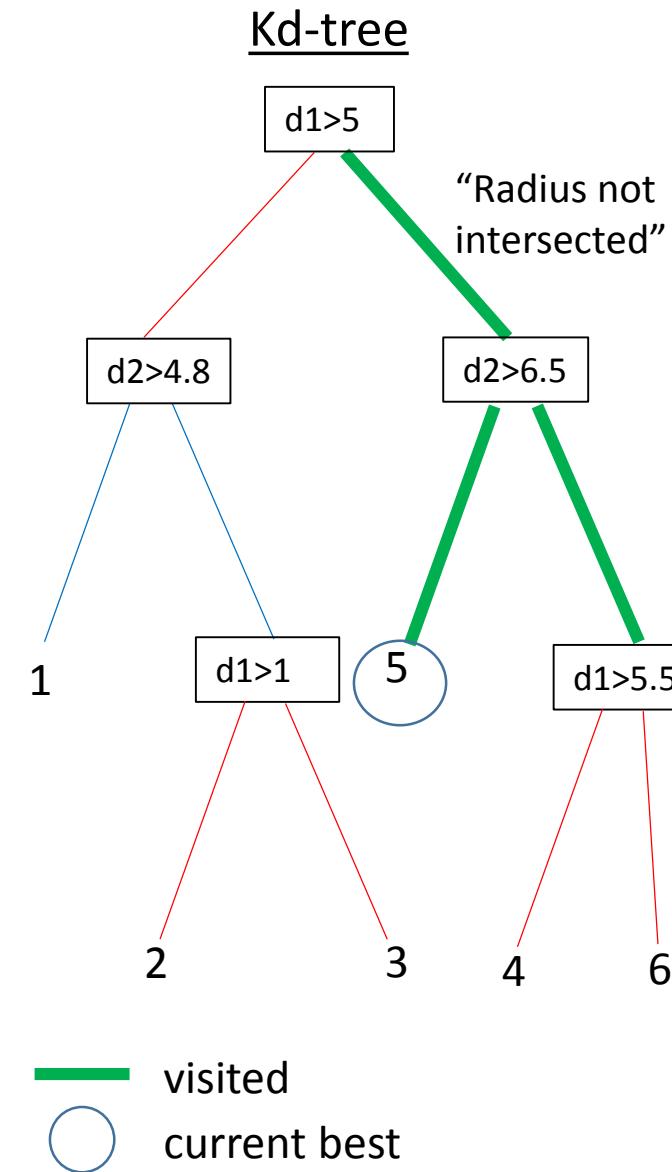
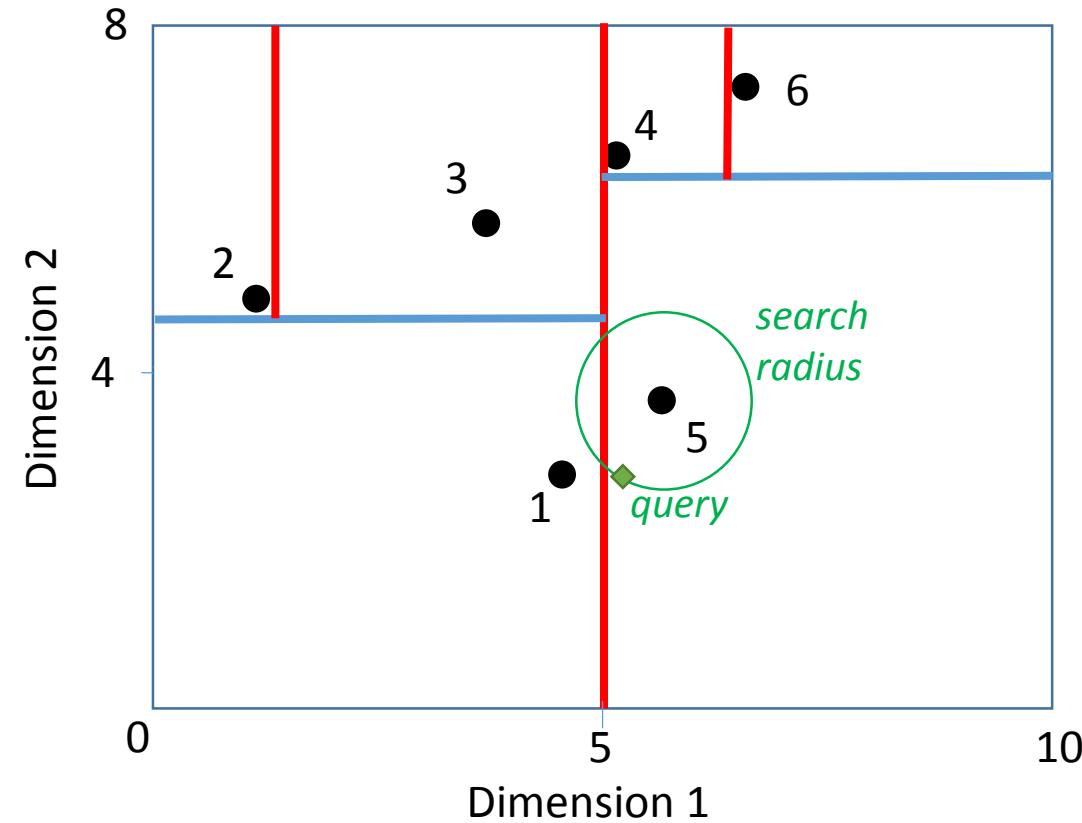
# Examples: nearest neighbor search



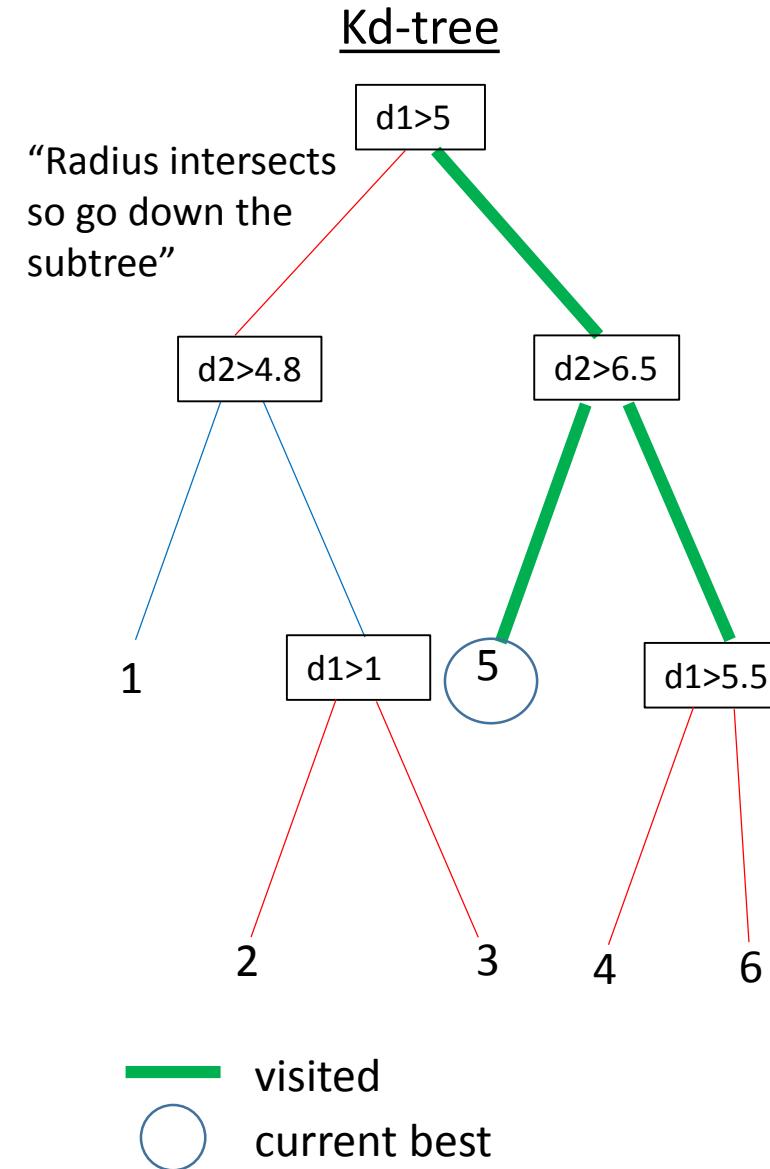
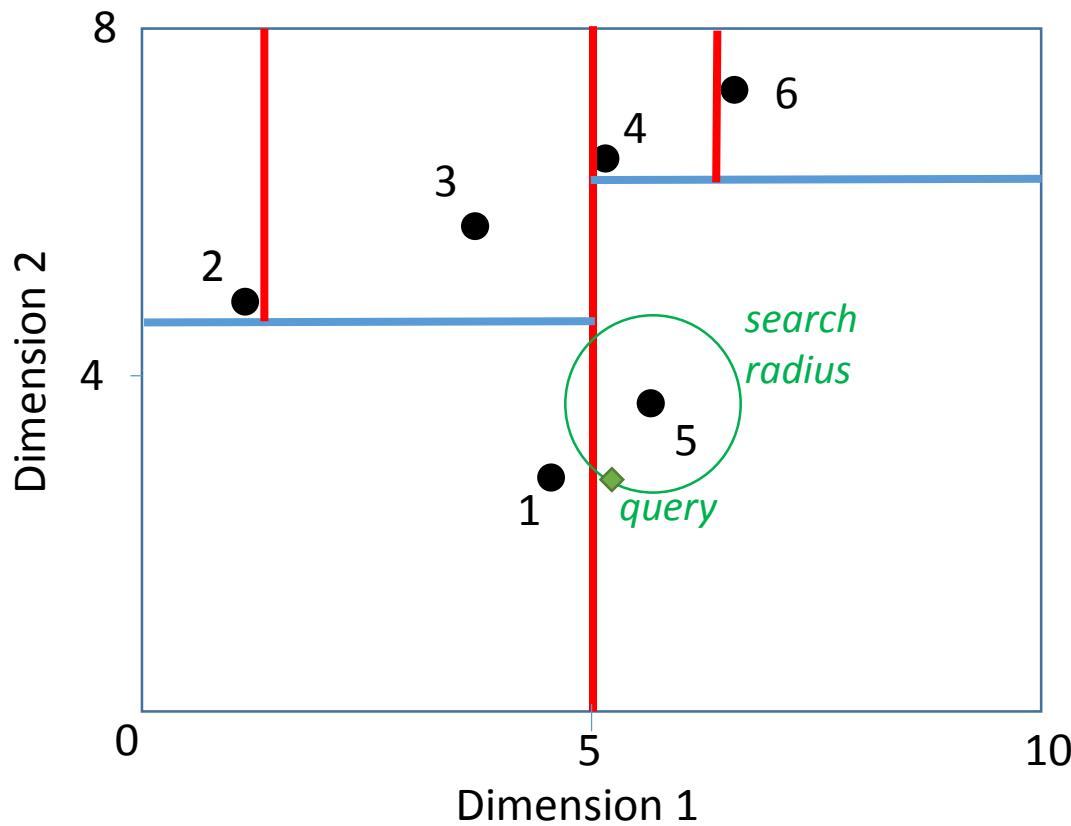
# Examples: nearest neighbor search



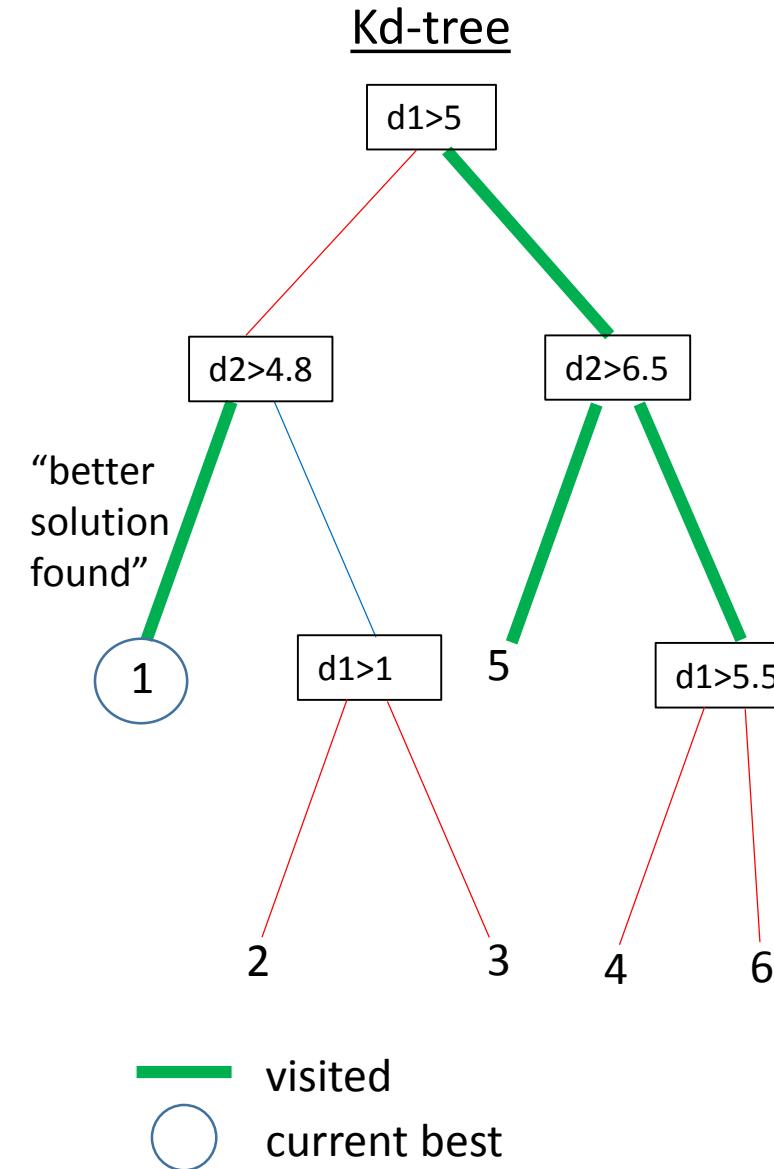
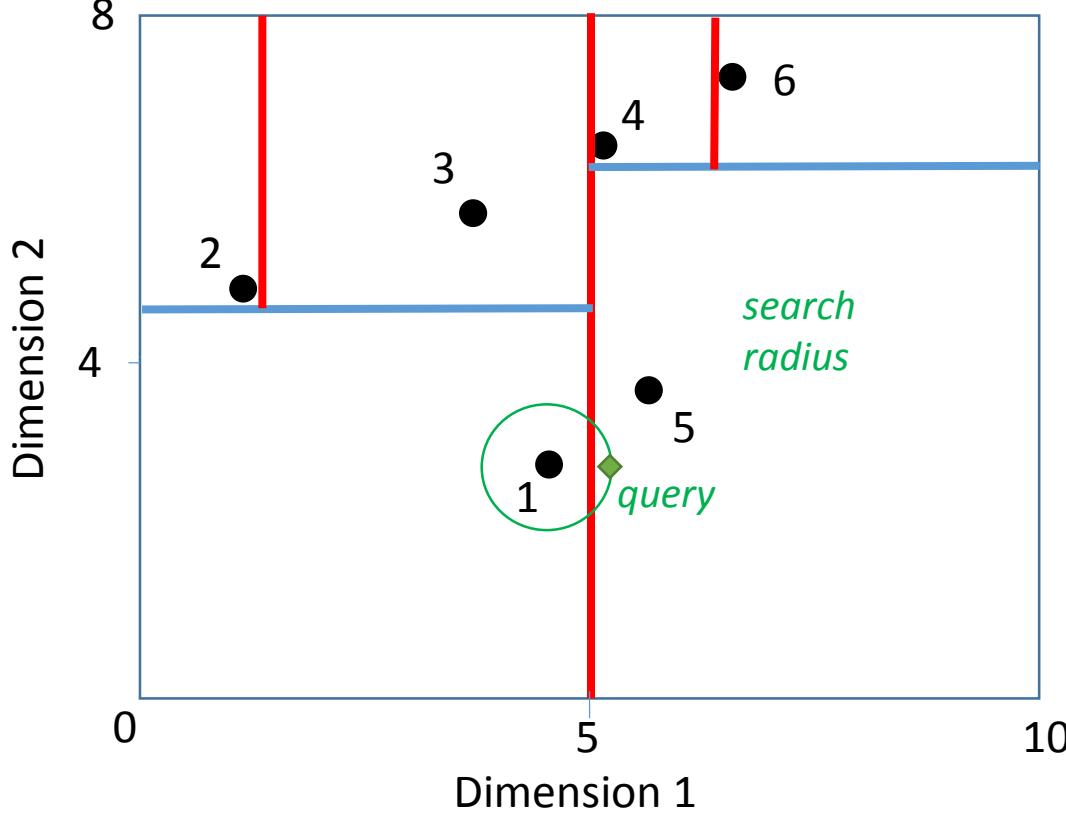
# Examples: nearest neighbor search



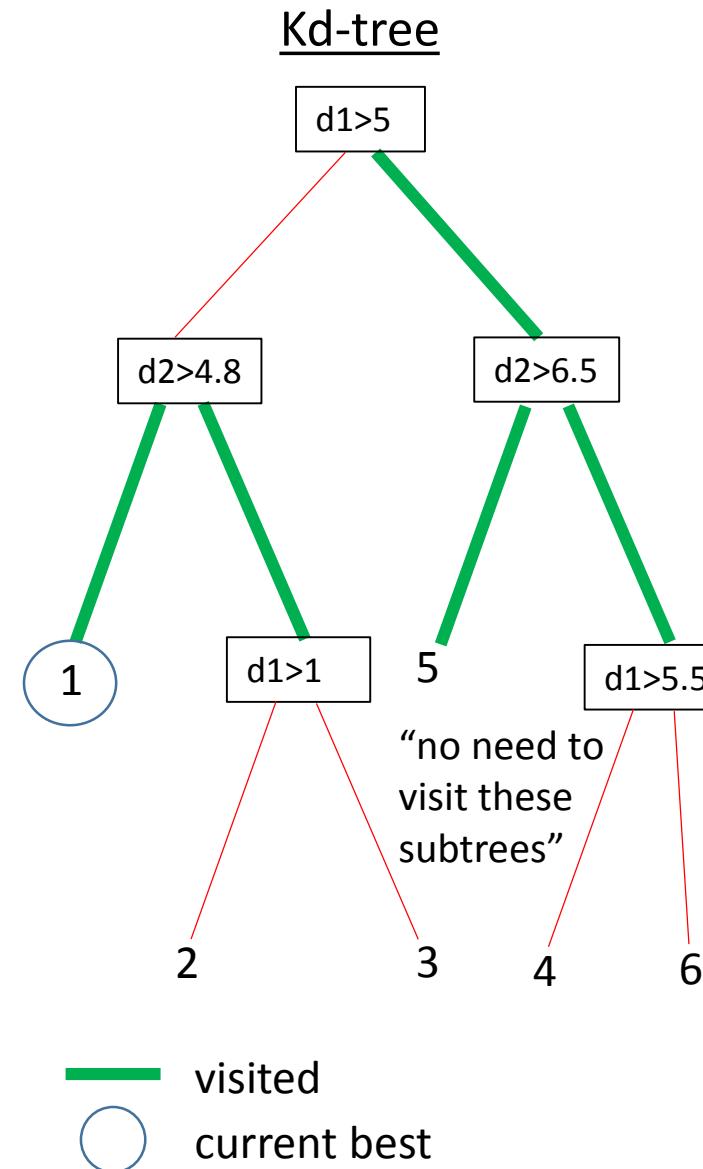
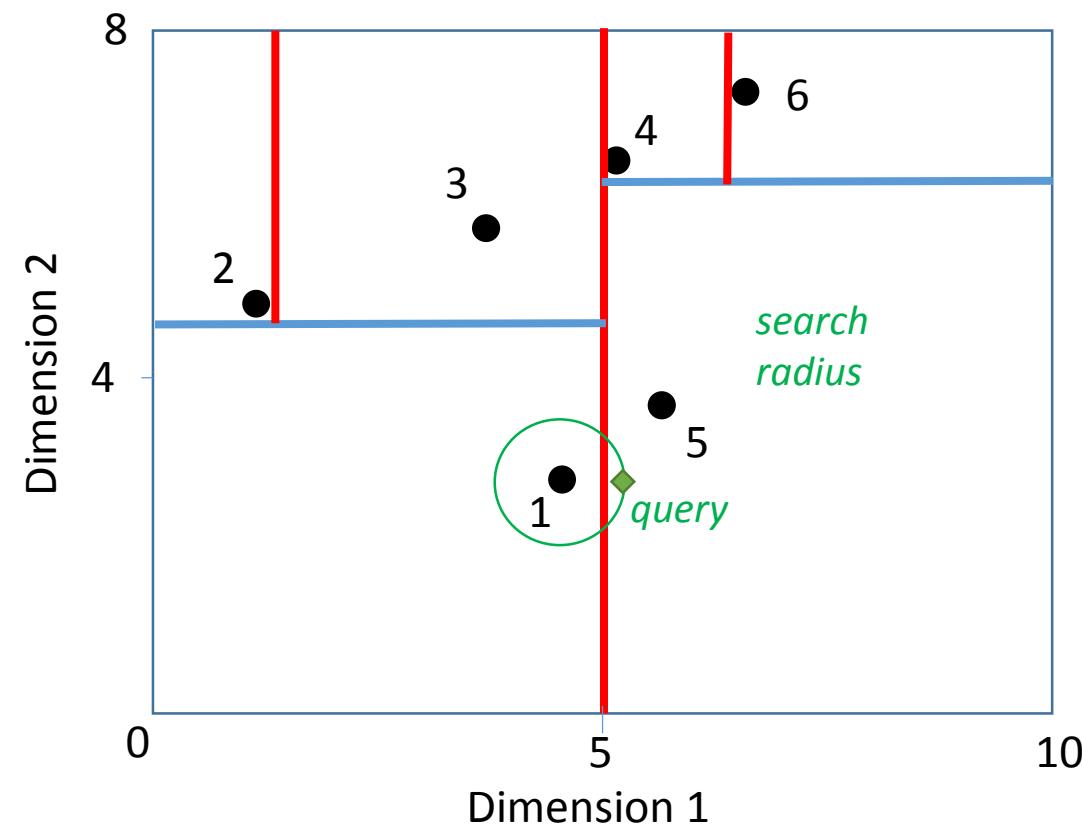
# Examples: nearest neighbor search



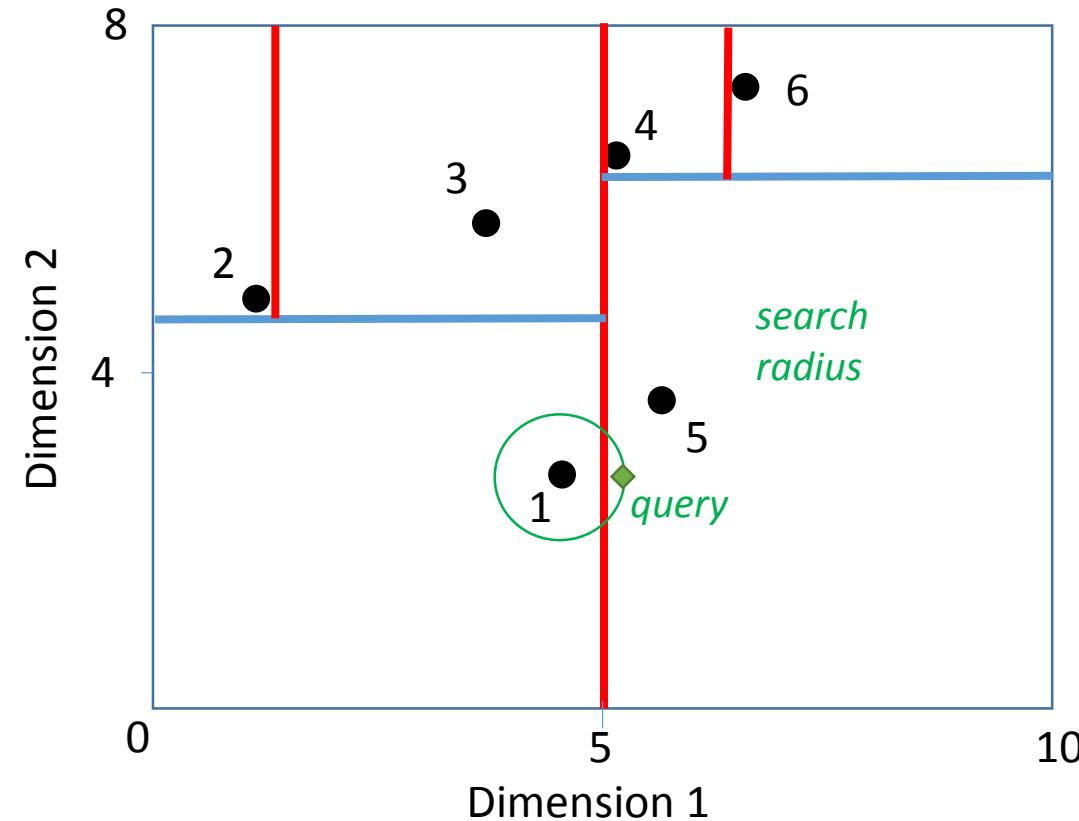
# Examples: nearest neighbor search



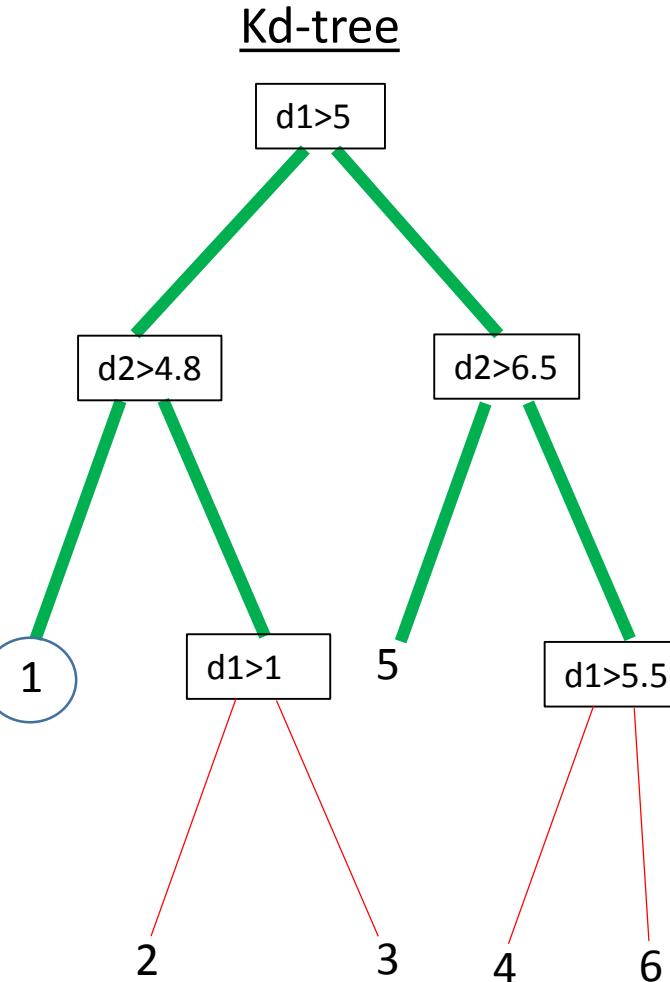
# Examples: nearest neighbor search



# Examples: nearest neighbor search



Done since root-node is marked in both ways



— visited  
○ current best

# Nearest neighbour search – pseudo code

Input: query point

Pseudo code

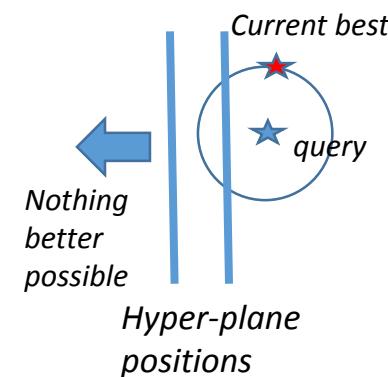
Step 1: Find leave node (bucket) with query point

Step 2: Make hyper-sphere with radius  
(current best and query point)

Step 3: go up the tree and see if hyper-plane intersects hyper-sphere

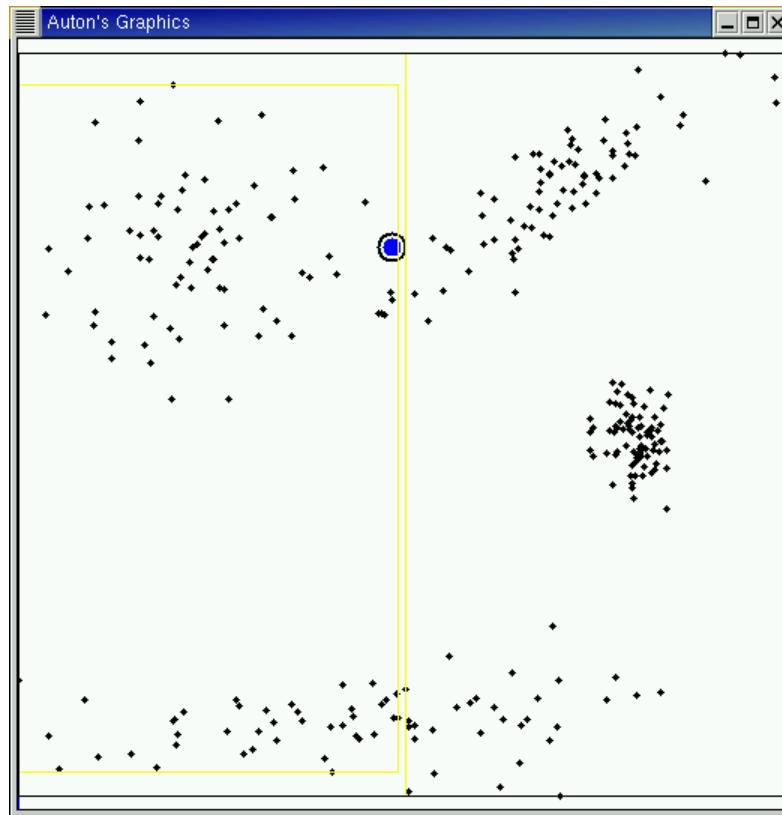
Step 3a: ***no intersection:*** mark tree branch as visited  
since no better point can be found there.  
If node is root node then stop.

Step 3b: ***intersection:*** go down the branch to find  
potentially a better point. If so, mark as  
current best and go to Step 2.



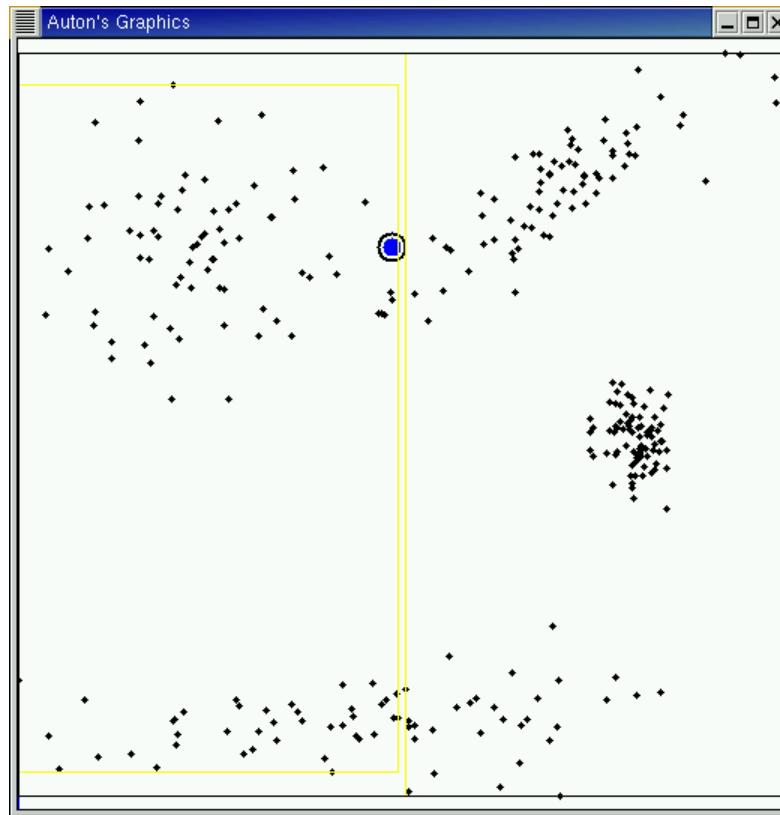
On average  $O(\log N)$

# Example with many points in 2D

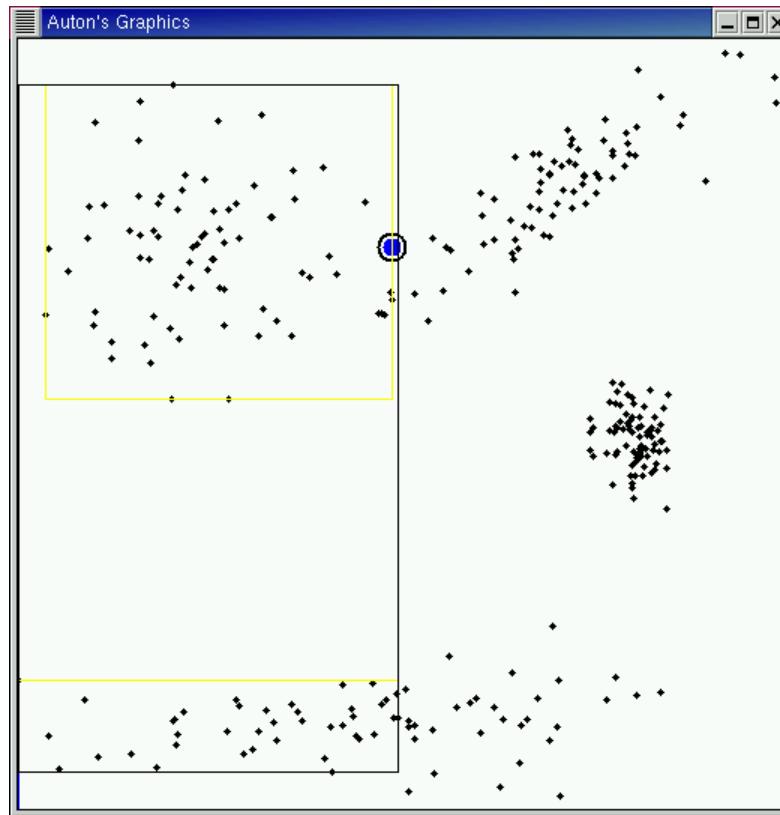


From Andrew Moore: <http://www.cs.cmu.edu/~awm/animations/kdtree/>

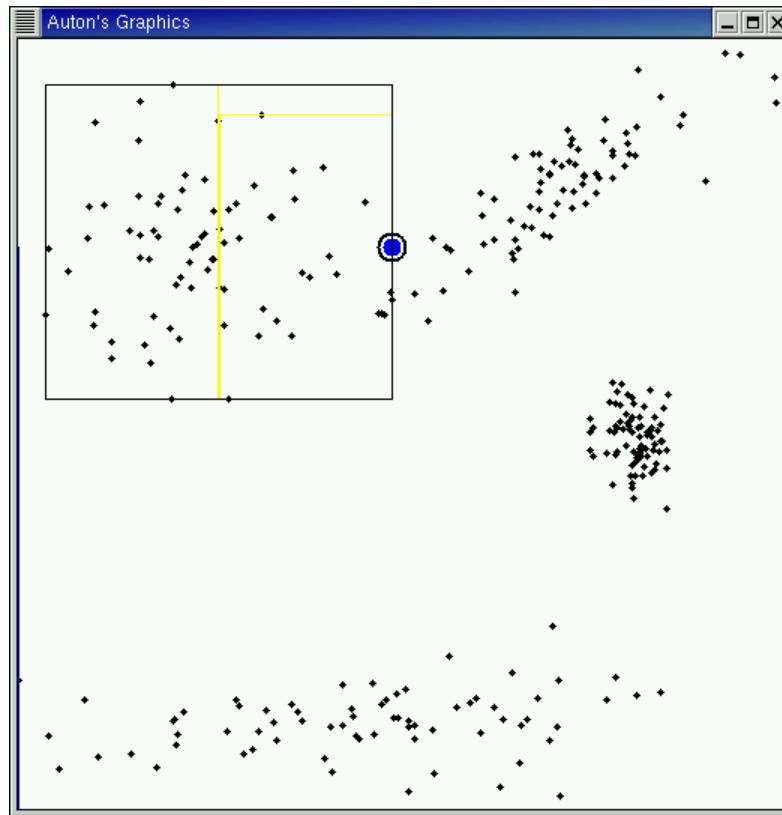
# Example with many points in 2D



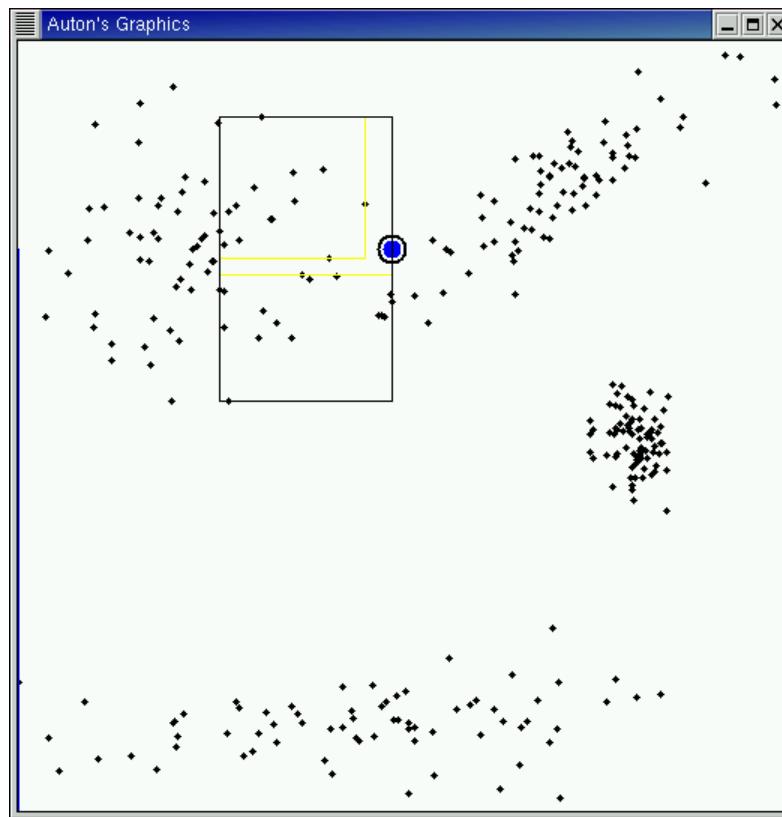
# Example with many points in 2D



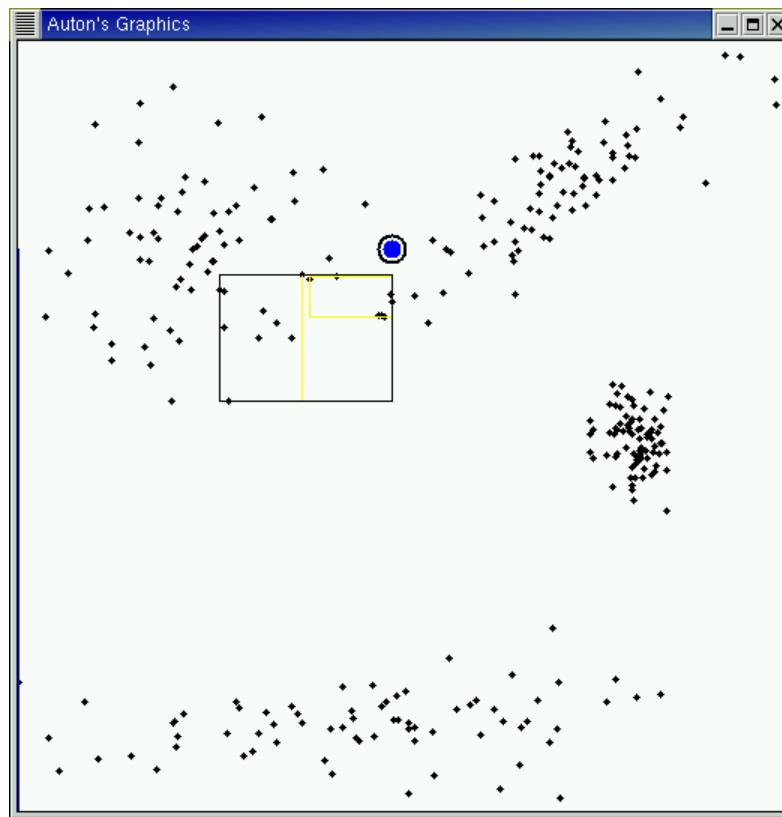
# Example with many points in 2D



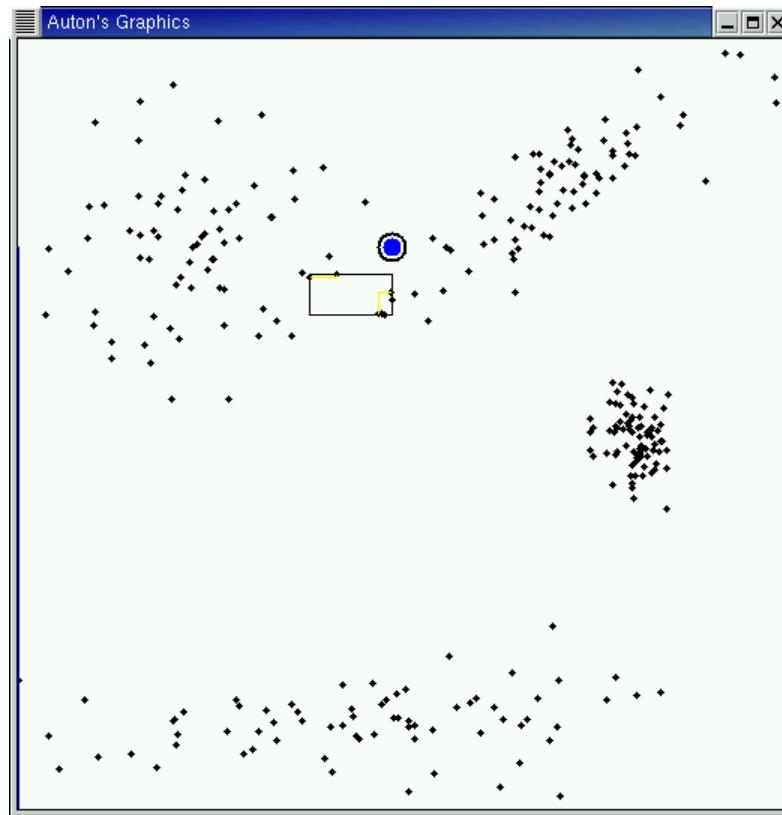
# Example with many points in 2D



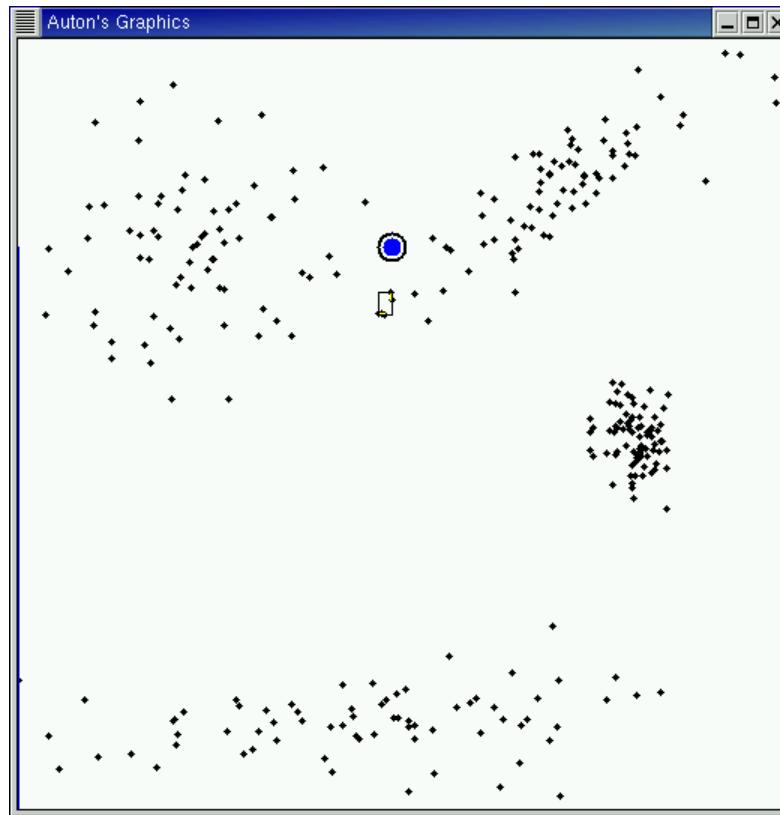
# Example with many points in 2D



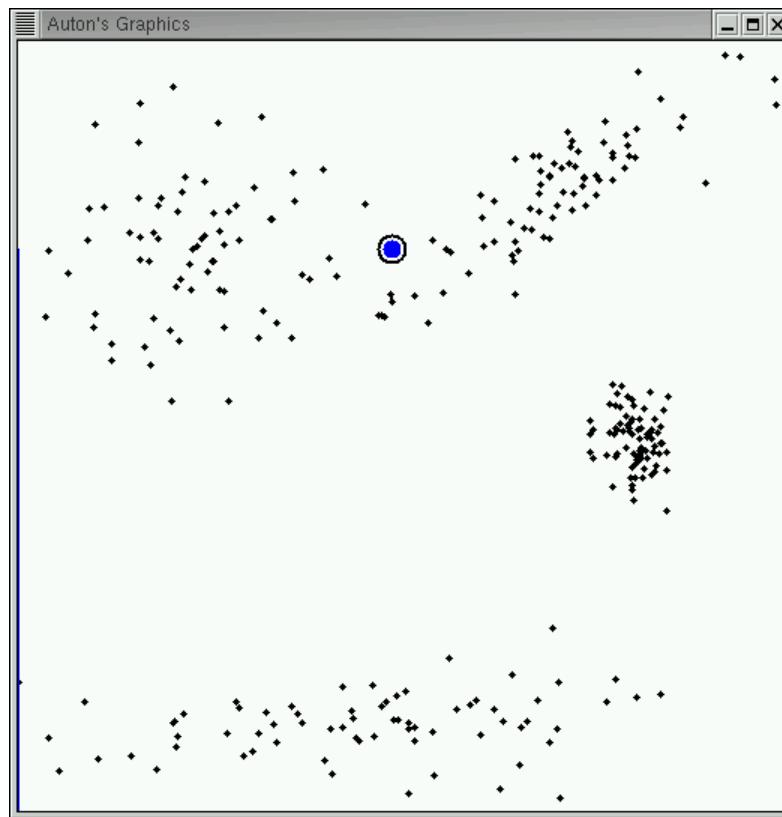
# Example with many points in 2D



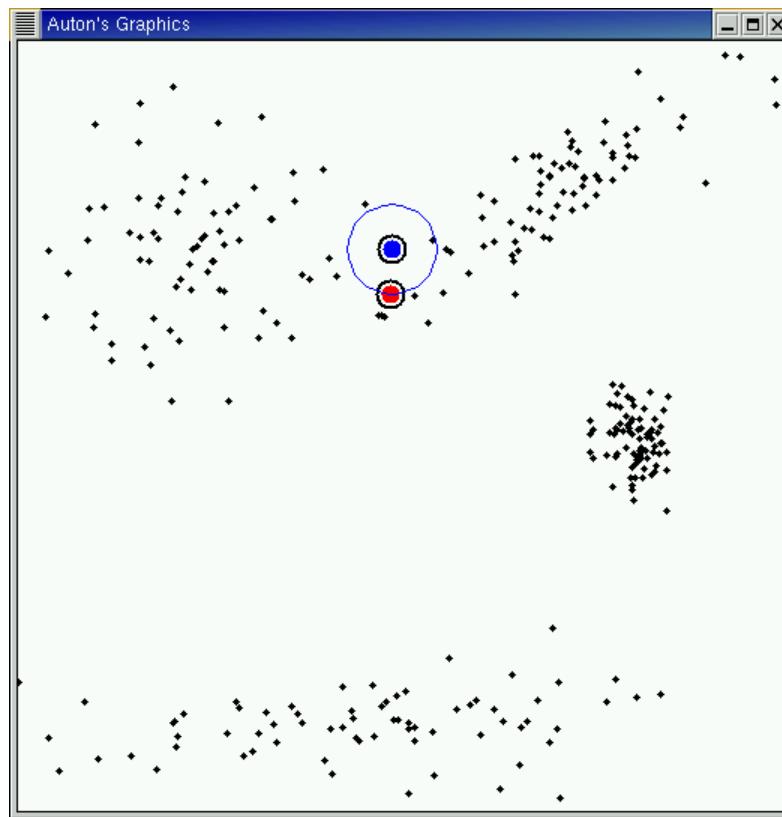
# Example with many points in 2D



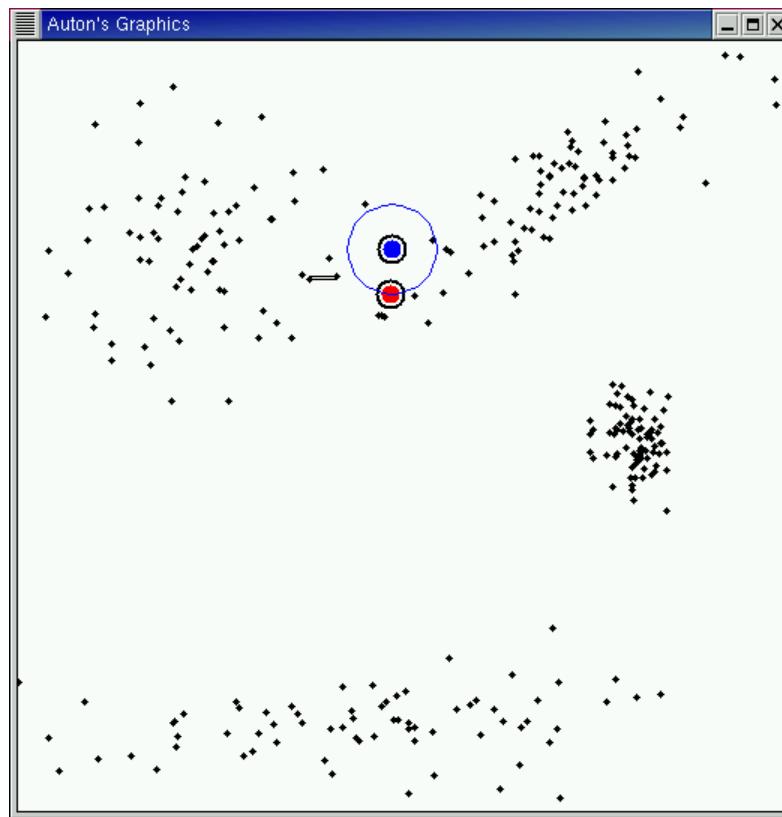
# Example with many points in 2D



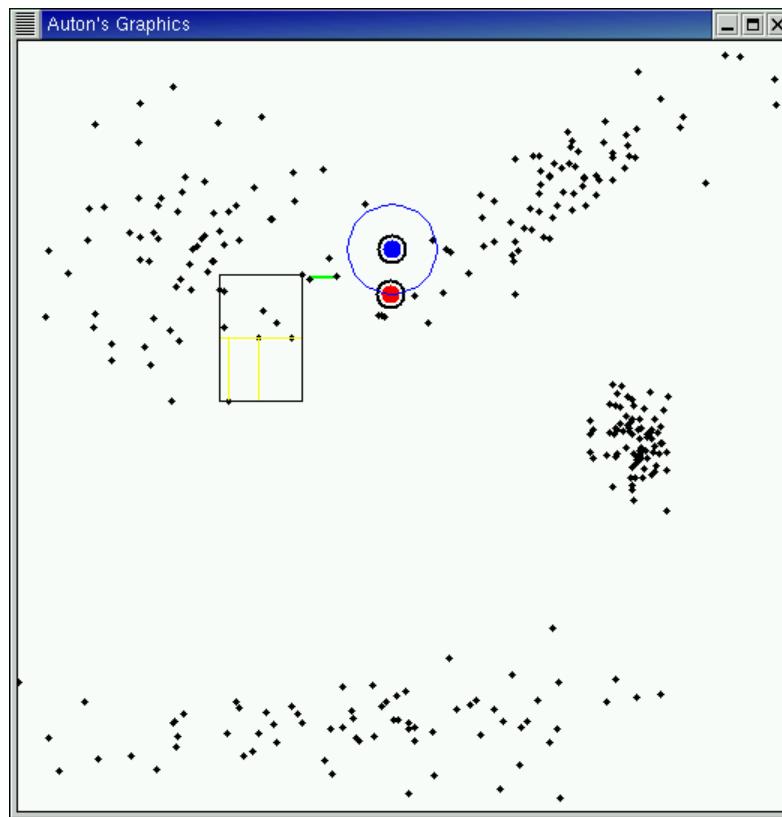
# Example with many points in 2D



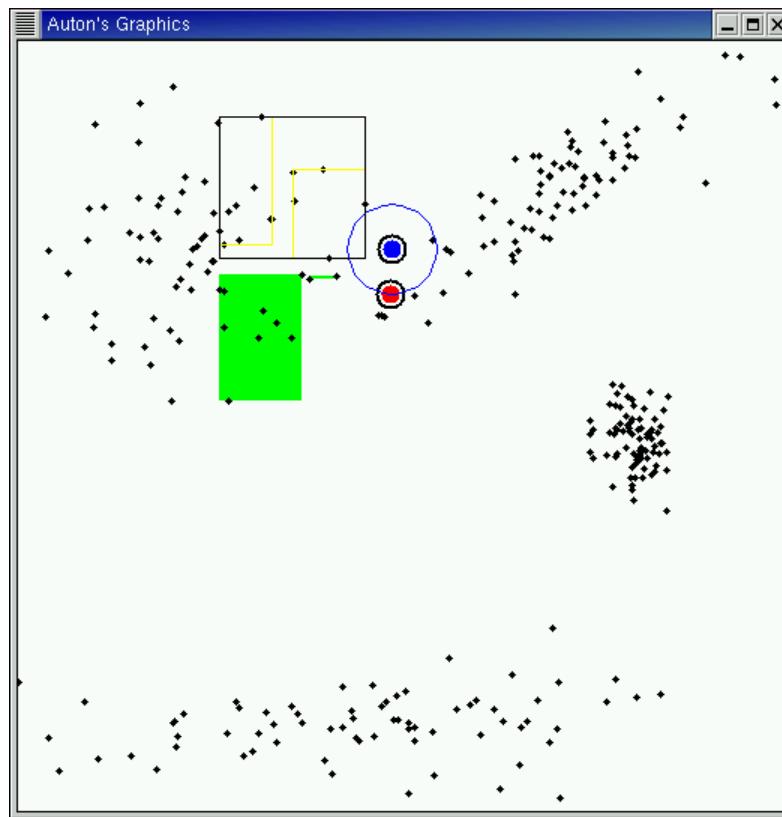
# Example with many points in 2D



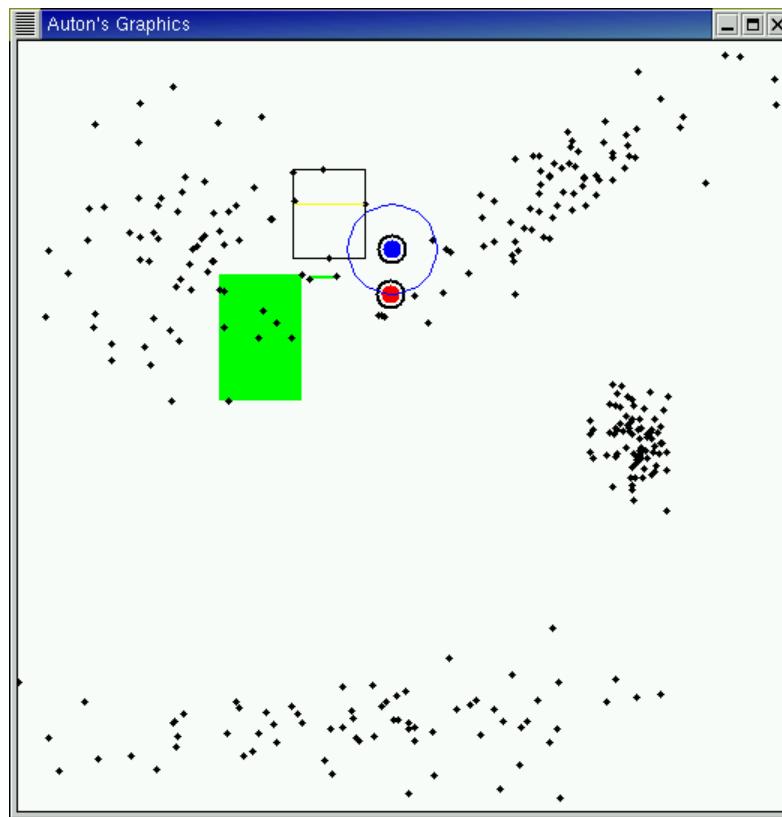
# Example with many points in 2D



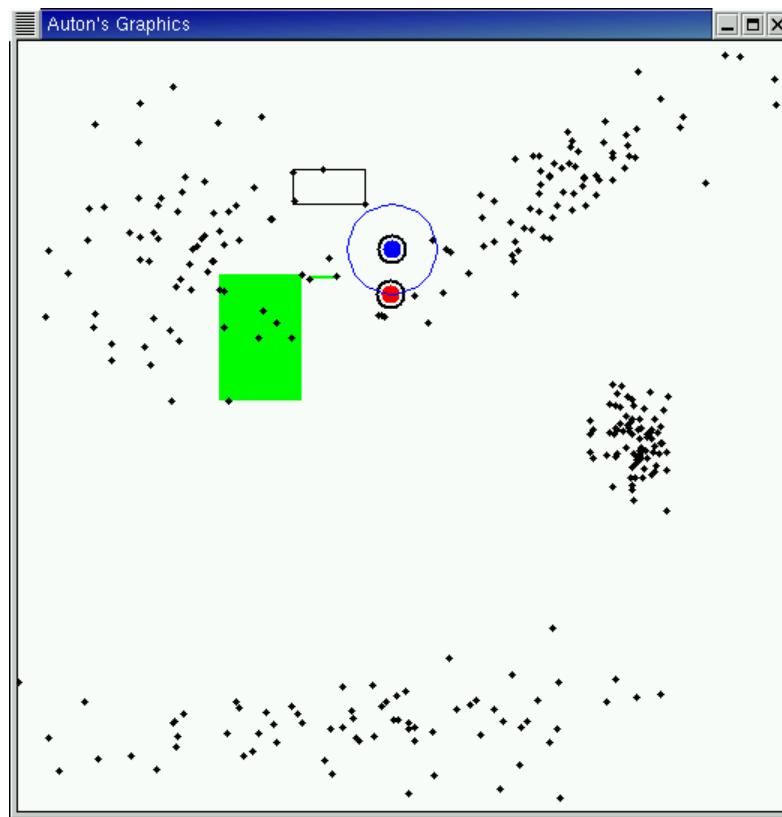
# Example with many points in 2D



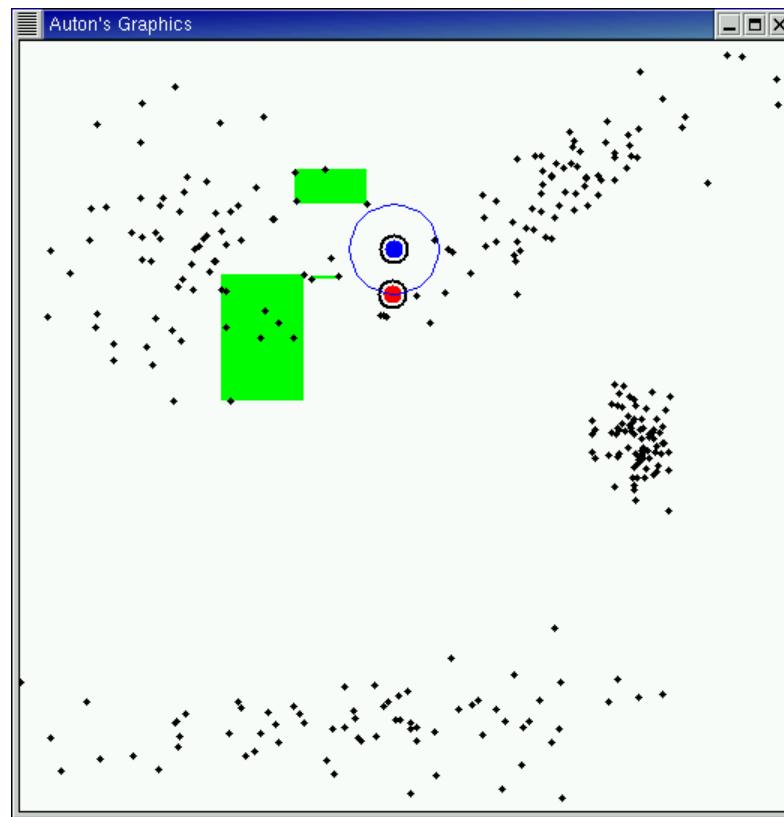
# Example with many points in 2D



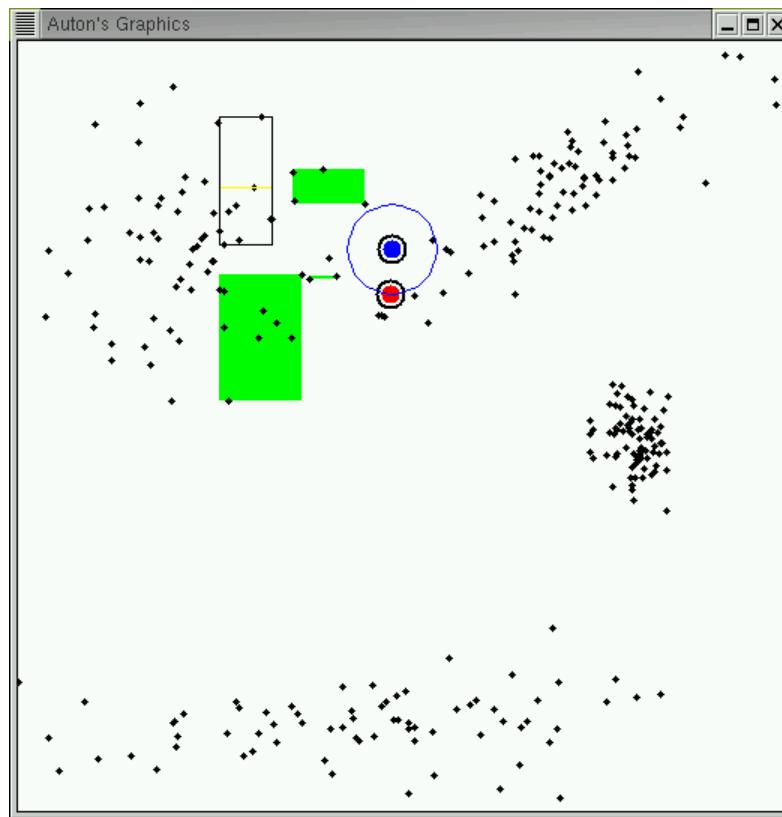
# Example with many points in 2D



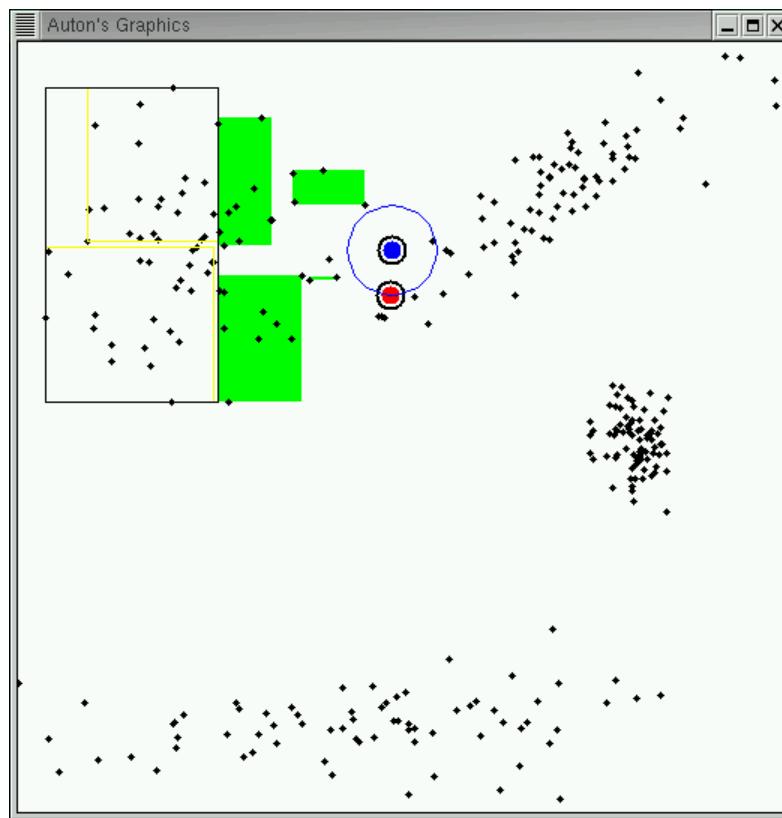
# Example with many points in 2D



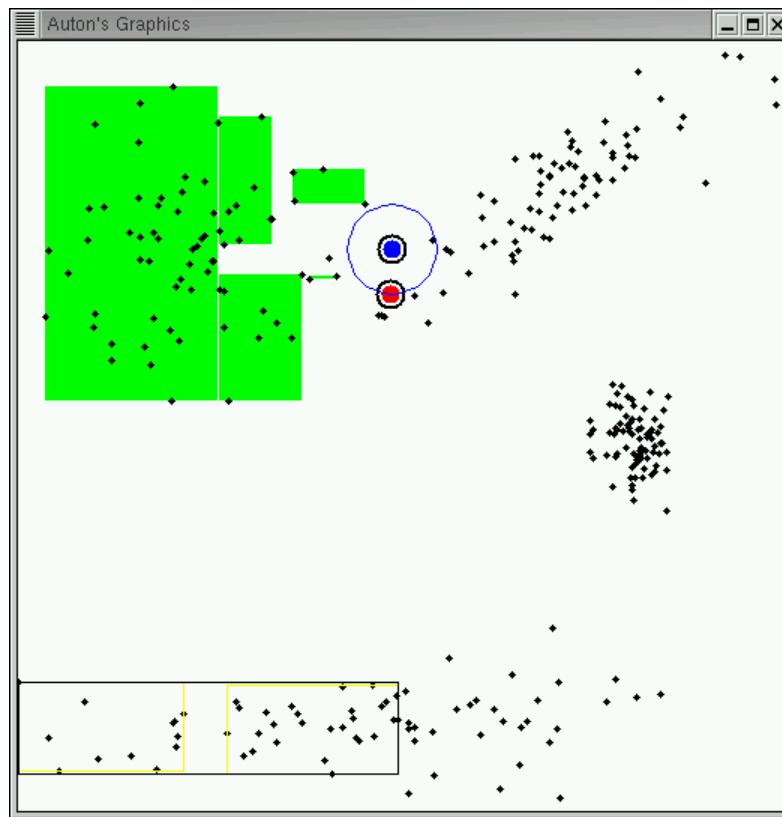
# Example with many points in 2D



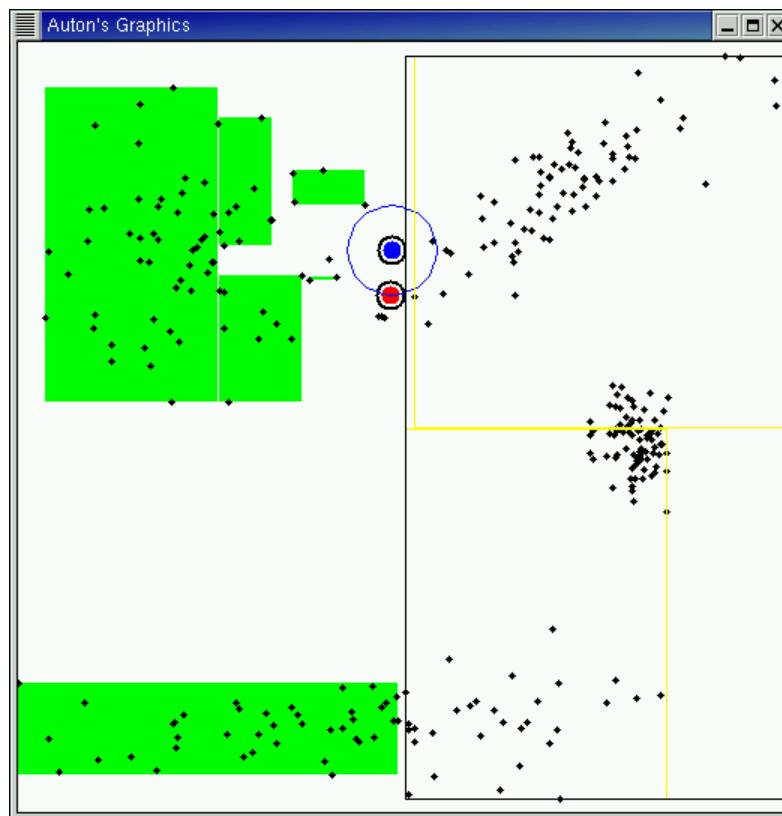
# Example with many points in 2D



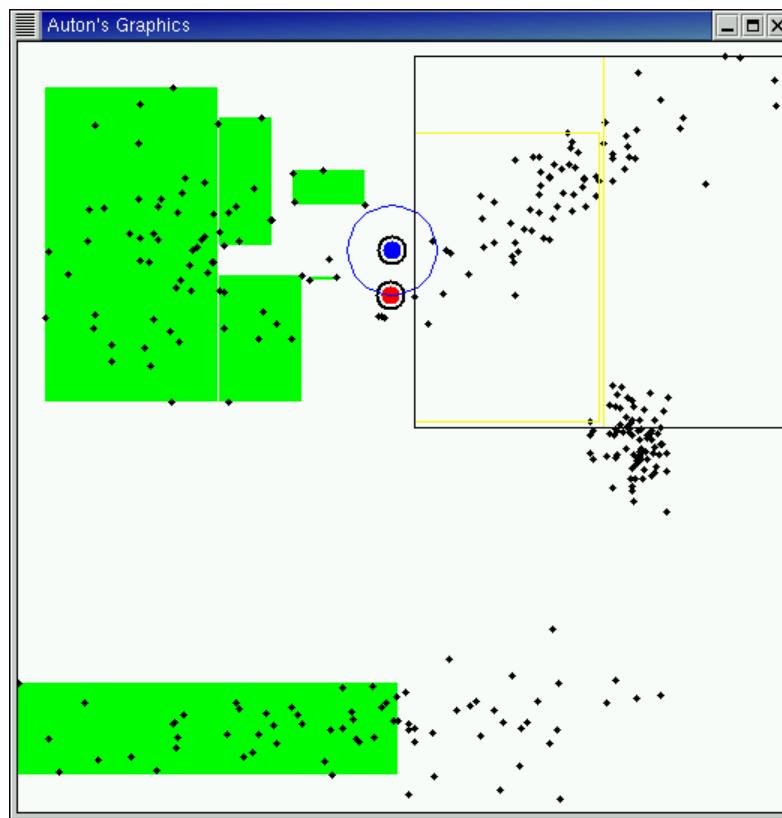
# Example with many points in 2D



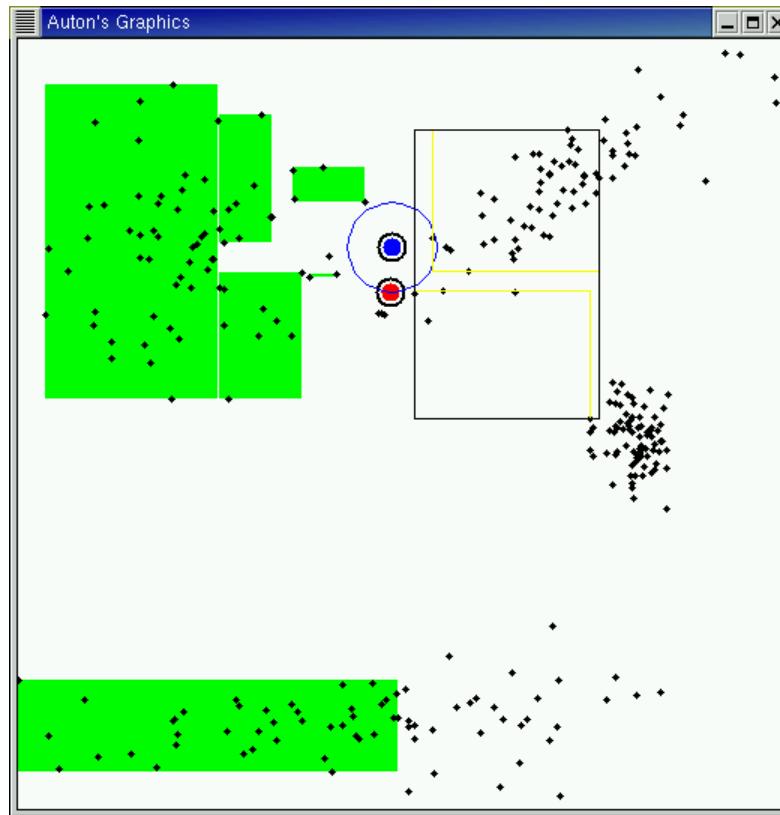
# Example with many points in 2D



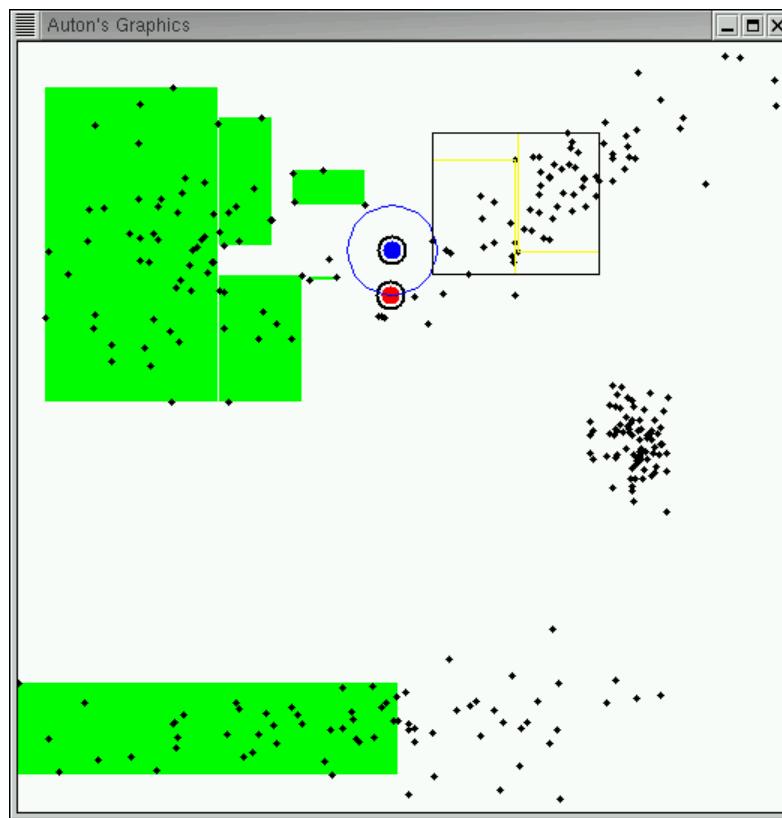
# Example with many points in 2D



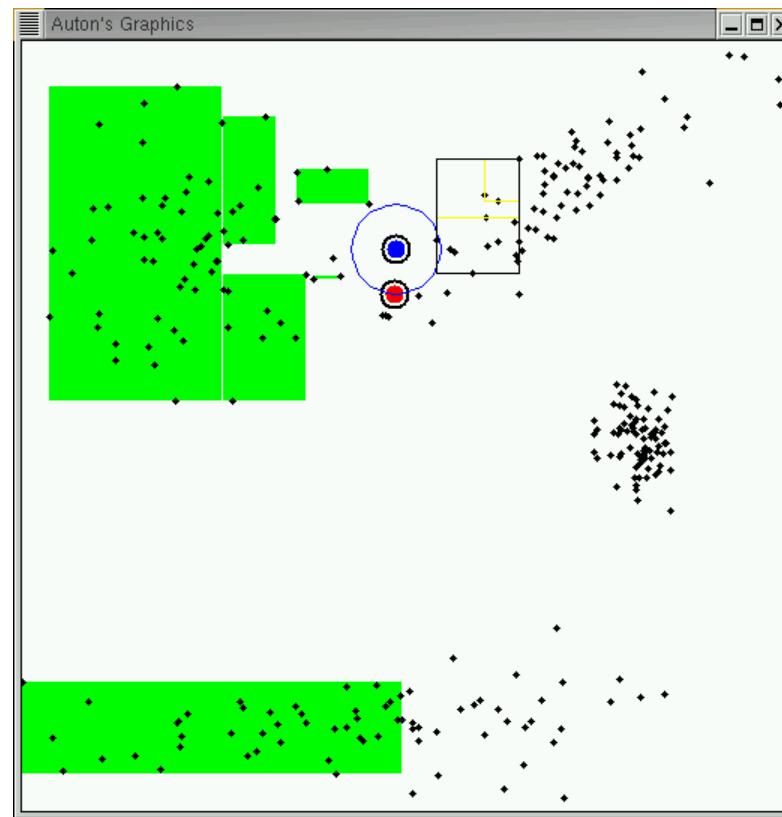
# Example with many points in 2D



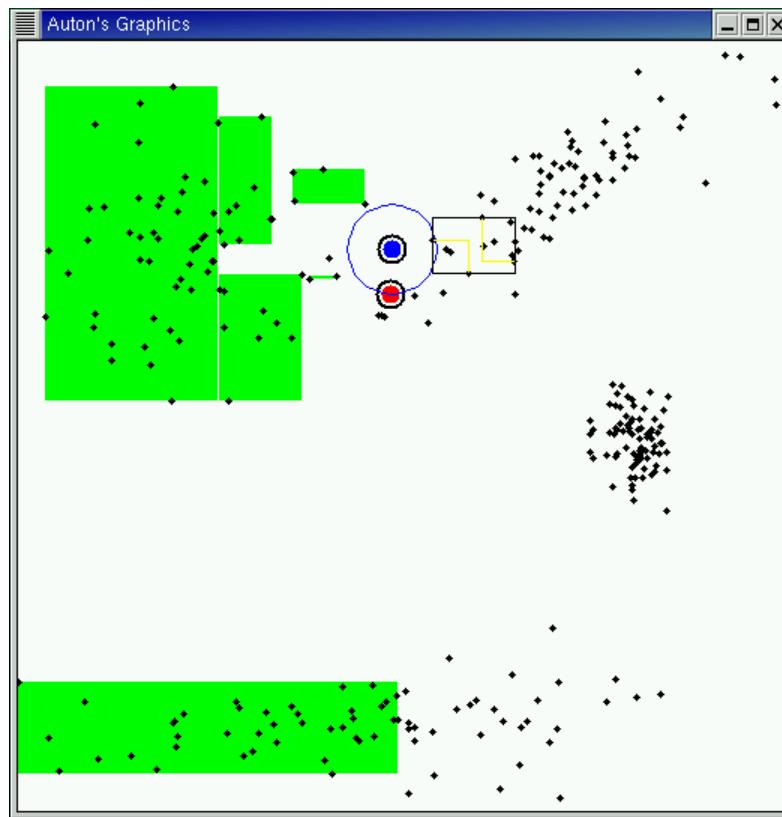
# Example with many points in 2D



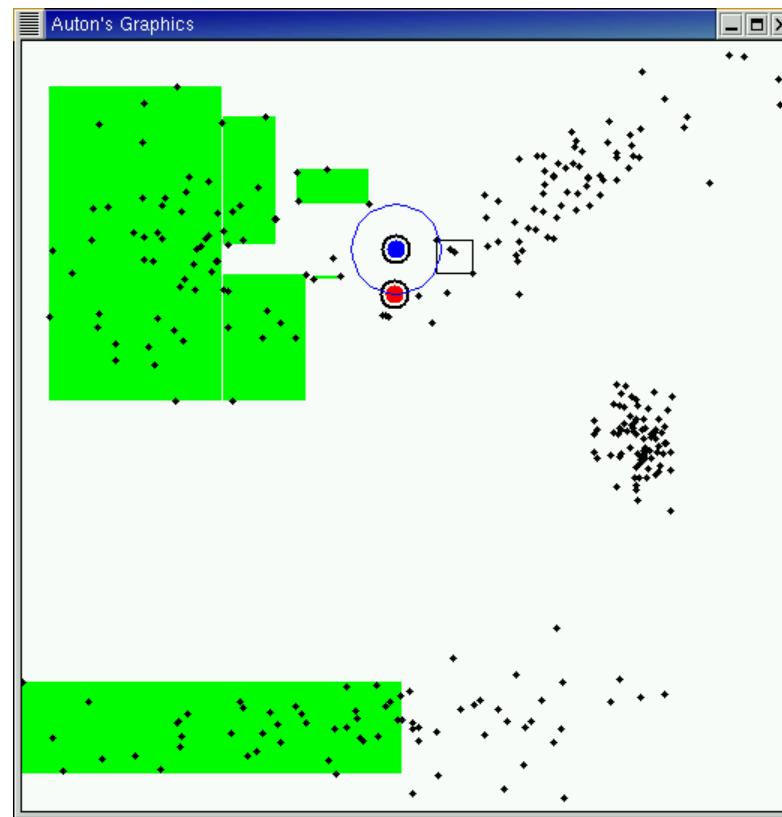
# Example with many points in 2D



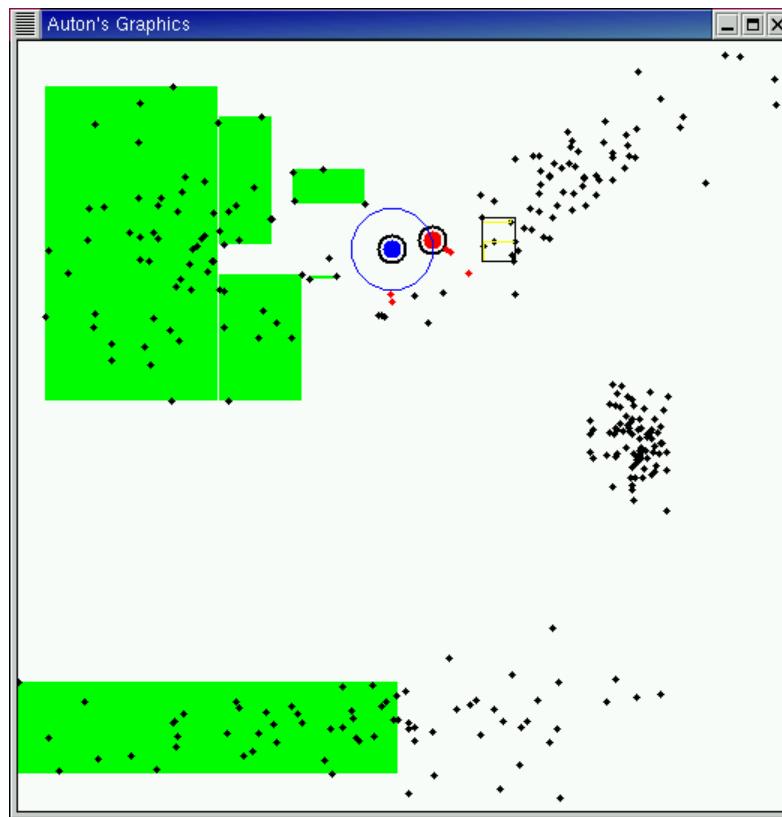
# Example with many points in 2D



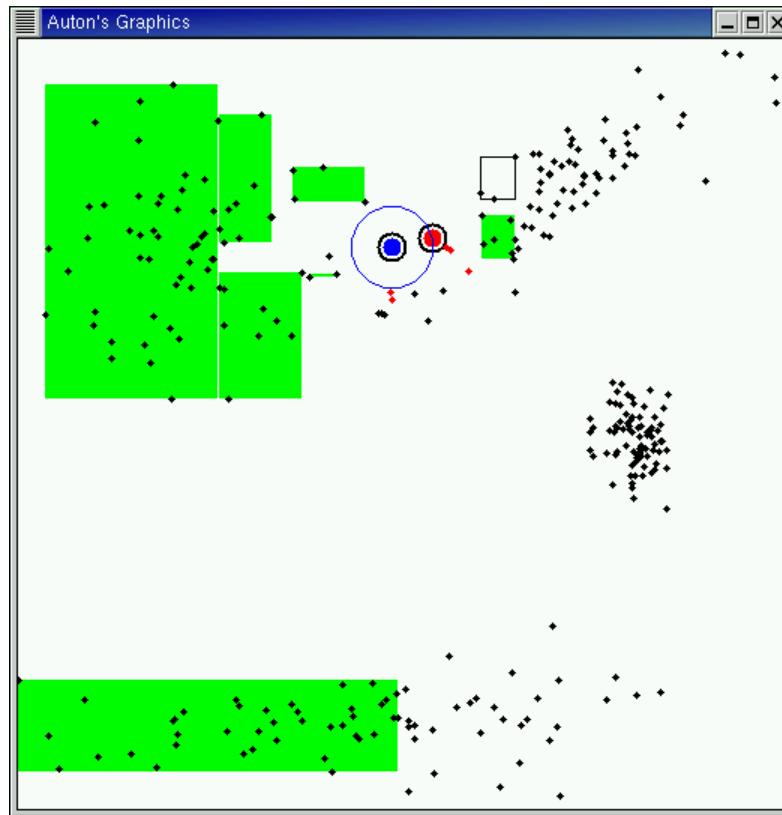
# Example with many points in 2D



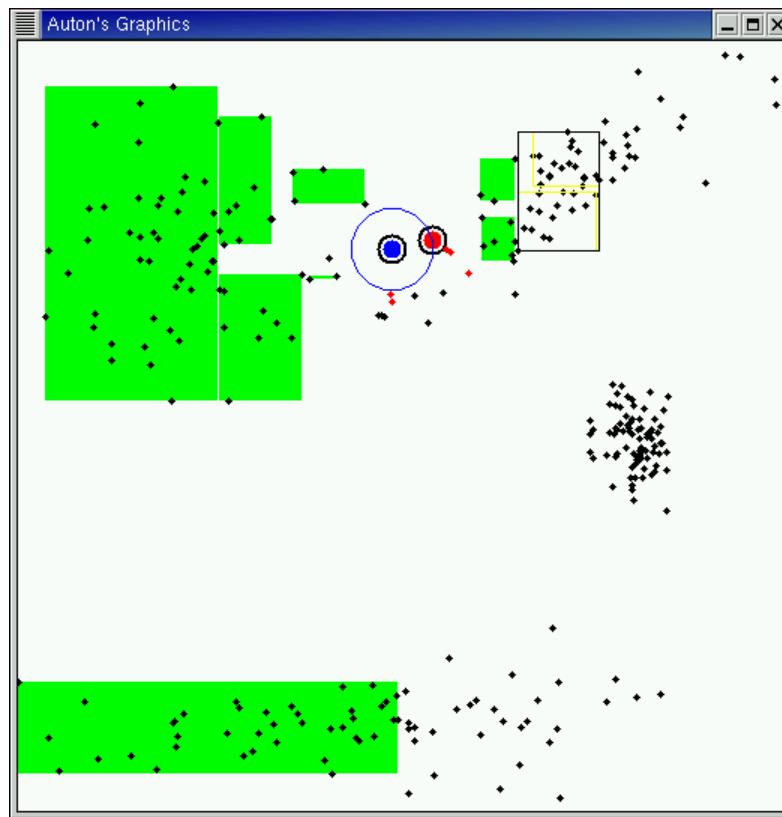
# Example with many points in 2D



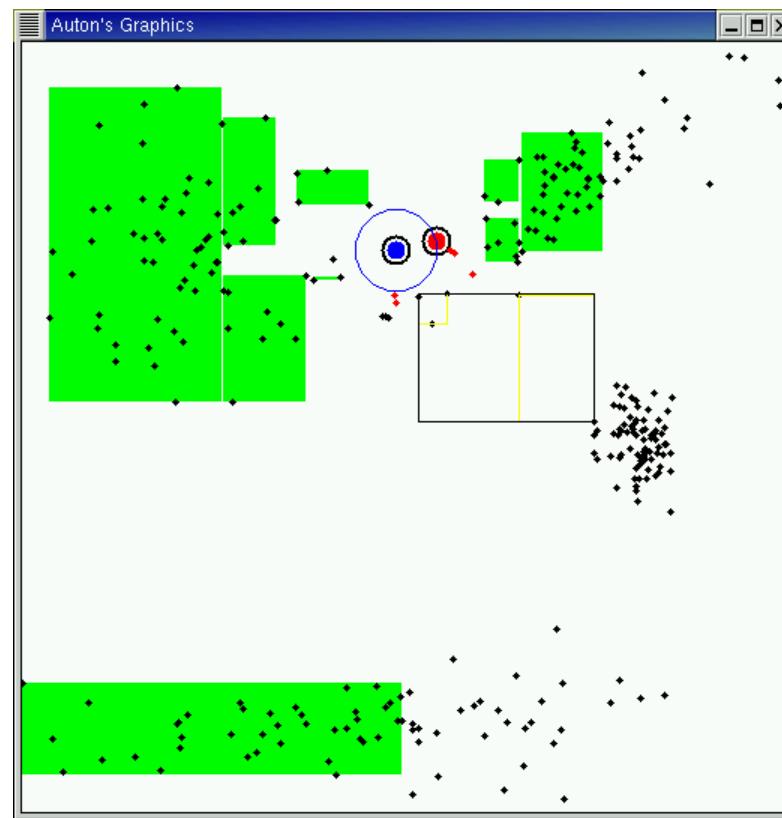
# Example with many points in 2D



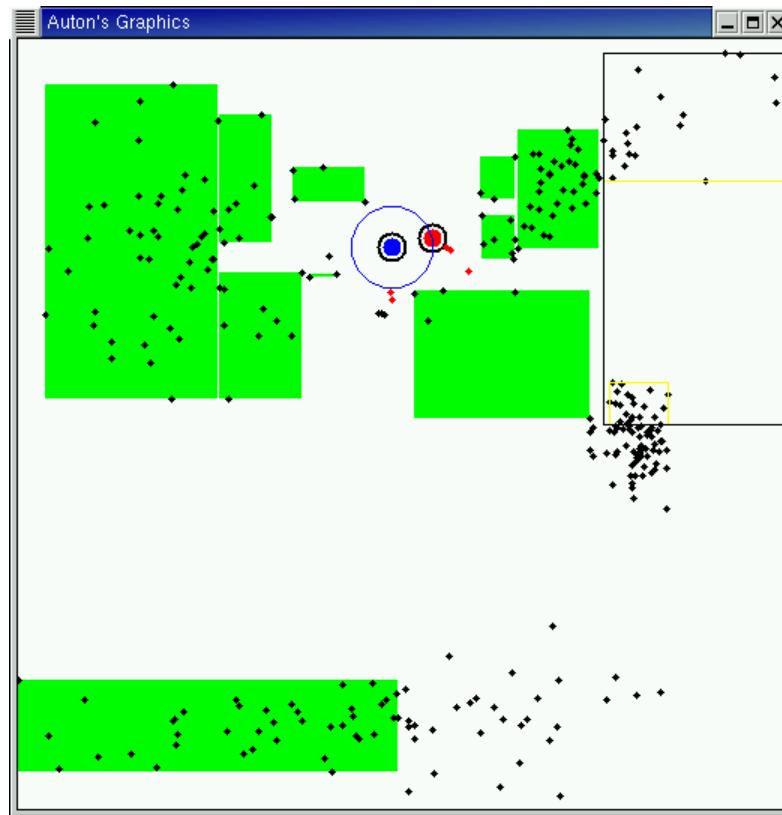
# Example with many points in 2D



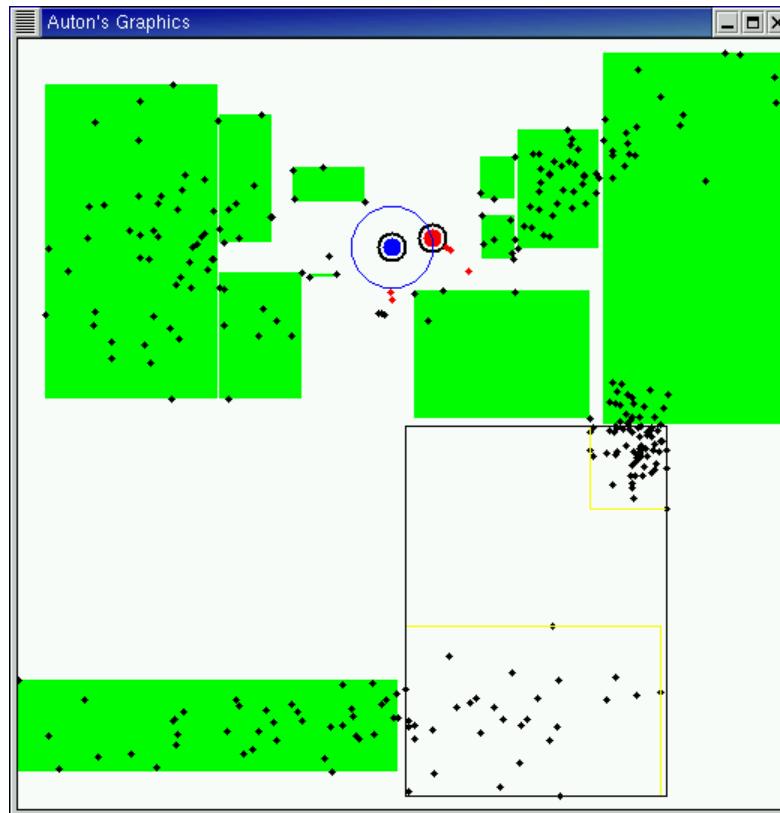
# Example with many points in 2D



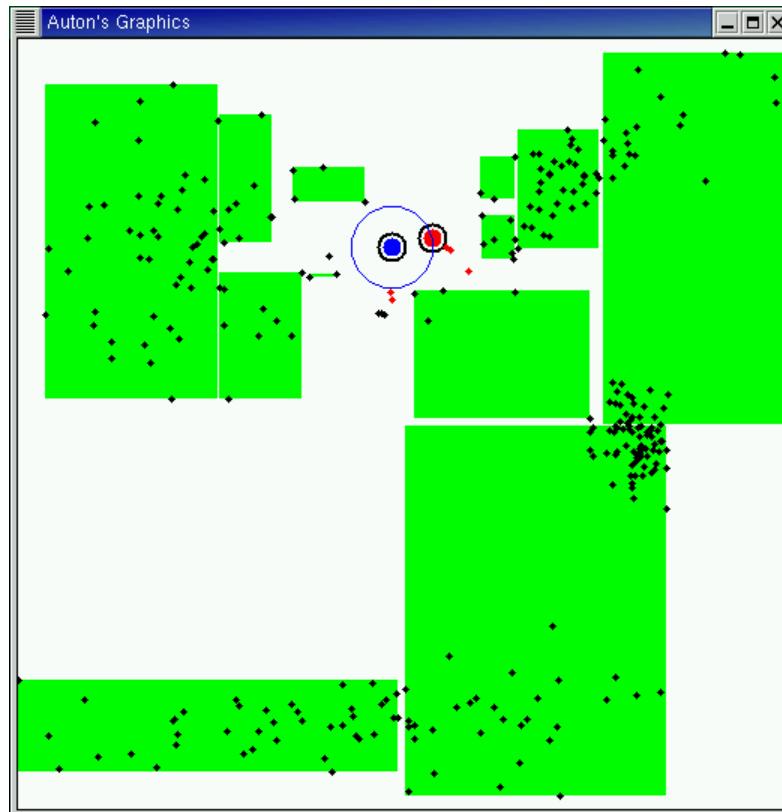
# Example with many points in 2D



# Example with many points in 2D

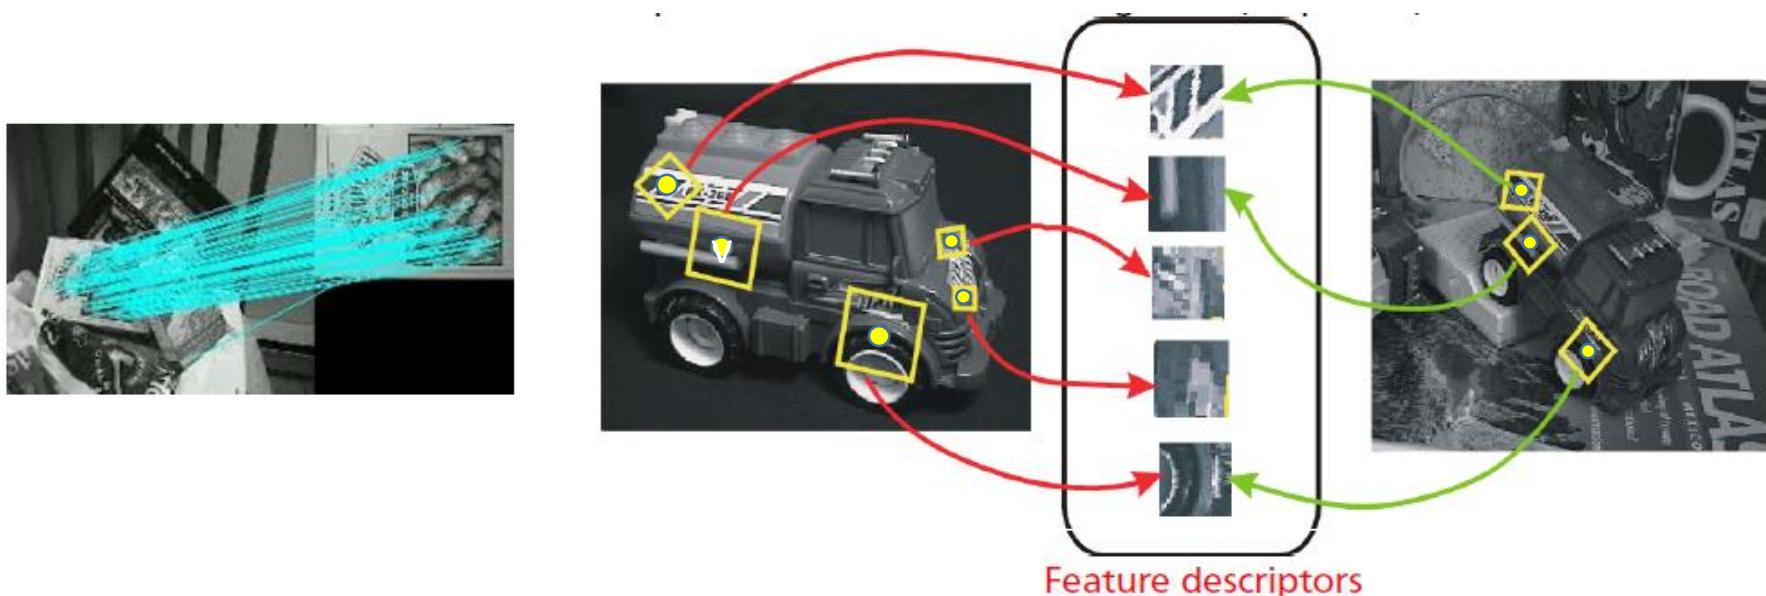


# Example with many points in 2D



# Roadmap: matching 2 Images (appearance & geometry)

- Find interest points (including different scales)
- Find orientated patches around interest points to capture appearance
- Encode patch in a descriptor
- Find matching patches according to appearance (similar descriptors)
- **Verify matching patches according to geometry**



# Reading for next class

## This lecture:

- Photometric image formation (sec 2.2)
- Camera Types and Hardware (sec 2.3)
- Appearance matching: (sec. 4.1.2-4.1.3)

## Next lecture:

- Two-view Geometry (Hartley Zissermann)